

Knowledge, Reality, and Value

A Mostly Common Sense Guide to Philosophy



Michael Huemer

Knowledge, Reality, and Value

A Mostly Common Sense Guide to Philosophy

Michael Huemer

Copyright © 2021 Michael Huemer

Text copyright ©2021 Michael Huemer, all rights reserved. Cover image: Iskra at Tampa Bay,
©2020 Michael Huemer, CC BY 4.0.

Contents

[Title Page](#)

[Copyright](#)

[Preface](#)

[Full Contents](#)

[Part I: Preliminaries](#)

[1. What Is Philosophy?](#)

[2. Logic](#)

[3. Critical Thinking, 1: Intellectual Virtue](#)

[4. Critical Thinking, 2: Fallacies](#)

[5. Absolute Truth](#)

[Part II: Epistemology](#)

[6. Skepticism About the External World](#)

[7. Global Skepticism vs. Foundationalism](#)

[8. Defining “Knowledge”](#)

[Part III: Metaphysics](#)

[9. Arguments for Theism](#)

[10. Arguments for Atheism](#)

[11. Free Will](#)

[12. Personal Identity](#)

[Part IV: Ethics](#)

[13. Metaethics](#)

[14. Ethical Theory, 1: Utilitarianism](#)

[15. Ethical Theory, 2: Deontology](#)

[16. Applied Ethics, 1: The Duty of Charity](#)

[17. Applied Ethics, 2: Animal Ethics](#)

[18. Concluding Thoughts](#)

[Appendix: A Guide to Writing](#)

[Glossary](#)

[Books By This Author](#)

Preface

Why Read This Book?

This is an introduction for students who would like a basic grasp of a wide variety of issues in the field of philosophy. There are many textbooks you could look to for this purpose, but this one is the best. Here is why:

i. The writing. It is written in a clear, simple style. It should be easier to read and won't put you to sleep as fast as other textbooks. (On the other hand, if you *want* to fall asleep quickly, I suggest checking out an academic journal.)

ii. The topics. I cover a broad array of big and interesting issues in philosophy, like free will, the existence of God, how we know about the world around us, and the existence of objective values. I don't spend too much time on the boring ones (which we won't even mention here).

iii. The price. I just checked the prices of some traditional textbooks. I won't mention them by name so as not to embarrass their publishers, but I see prices in the range of \$50, \$80 ... one is even listed at \$140. (You know why they do this, right? They know that students don't choose textbooks. *Professors* choose them, and students just have to buy them. The profs may not even know the prices, since they get *their* copies for free. This is also why most textbooks are written to please professors, not to please students. But I digress.) If I'd gone with a traditional textbook publisher, I'd have no control over the price, and it would probably wind up at \$80 or something ridiculous like that.

I also wouldn't be able to write it like this. They'd say the style was too informal and flippant and demand that I write more "professionally" and lethargically.

iv. The author. I'm smart, I know a lot, and I'm not confused – which means you can probably learn a lot from this book. You probably won't learn too many falsehoods, and you probably won't run into too many passages that don't make sense.

About the Author

I can hear you saying: “Oh sure, you *would* say that.” Okay, maybe you shouldn’t believe me yet, because you just met me, and maybe I’m biased. Maybe you want to know if I’m enough of an expert to write this textbook, especially since it hasn’t been certified by a big textbook publisher. So here is who I am, sticking just to objective facts:

I got my BA in philosophy from UC Berkeley. I got my PhD in philosophy from Rutgers University, which at the time was ranked #3 in philosophy in the United States (they’re now ranked #2).^[1] I am a tenured full professor at the University of Colorado at Boulder, where I have taught philosophy for over 20 years. As of this writing, I have published more than 70 academic articles in a variety of journals, including most of the top-ranked philosophy journals. (In philosophy, by the way, the good journals reject 90–95% of submissions.) My publications span a wide range of topics in different branches of philosophy, including many of the issues I introduce you to in the following pages.

I have written seven books before this one and edited an eighth, all published with traditional, academic publishers (which is why they’re so expensive). Here are my earlier books, in case you want to look up any of them:

Skepticism and the Veil of Perception

Epistemology: Contemporary Readings (edited volume)

Ethical Intuitionism

The Problem of Political Authority

Approaching Infinity

Paradox Lost

Dialogues on Ethical Vegetarianism

Justice Before the Law

My Approach in Writing This

That’s enough about me. Now here are some comments about my approach in writing this:

1. I have selected a few very prominent issues in each of the biggest areas of philosophy – issues that are commonly addressed in philosophy courses and that philosophy students like to know about, like the existence of God, free will, etc.
2. I give a basic presentation of each issue, including what I consider the most important and interesting arguments that can be

explained reasonably briefly. (In each case, there are of course many more complicated and nuanced views and arguments to be found in the literature.) By the way, when you read these arguments, don't just memorize them and move on (as students sometimes do). Spend some time thinking about whether you agree with them or not.

3. All of these are issues that people disagree about. In each case, my presentation aims to be (and I think is in fact) *fair*, but not *neutral*. That is:

1. I give each view a fair hearing, presenting its case as strongly as I can (given space constraints), in terms that I think are faithful to its proponents' intellectual motivations. I do not select evidence, distort people's words, or use any other tricks to try to skew the assessment of any of the philosophical theories. (Those sorts of tricks are unsuited to a philosopher.)

2. I do *not*, however, promise a *neutral* presentation – one that just reports other people's ideas without evaluation (which I consider terribly boring). I am going to tell you what I think, and I am going to defend it with logical arguments that try to show you why that view is right.

If you don't like that, this isn't the book for you. Go get another book, like Stuart Rachels' anthology or something.^[2]

Why Study Philosophy?

If you haven't studied philosophy, you probably don't know why you should. There are two main reasons to do it.

First, philosophical questions are inherently fascinating. At least, many of them are. I mentioned some of them above. If those *didn't* sound interesting to you, then philosophy probably isn't for you.

Second, studying philosophy helps you think better. Right now, you probably don't know what I mean by that, and I can't adequately explain it, but I will *inadequately* explain it presently. I can't prove it to you either, since appreciating the point requires, well, studying philosophy for a few years. So I'll just tell you my assessment based on my experience. I saw it happen to myself, and I have seen it happen to students over the years. I

came to the subject, at the beginning of college, in a state of confusion, but I did not then comprehend how confused I was. I had some sort of thoughts about great philosophical questions, but these thoughts very often, as I now believe, simply made no sense. It was not that they were mistaken, say, because I was missing some important piece of information. It was that I did not even really know *what* I was thinking. I used words but did not really know what I meant by them. I confused importantly different concepts with each other. I applied concepts to things that they logically cannot apply to. I might seemingly endorse a philosophical thesis at one moment, and in the next endorse a related but *incompatible* thesis, without noticing any problem.

I was not, I stress, an unusually confused student; I am sure I was much *less* confused than the average college student. It just happens that *virtually everyone* starts out *extremely confused*. That is our natural state. It takes effort and practice to learn to think clearly. Not even to get the right answers, mind you, just to think *clearly*. To know precisely what your ideas *are*, and not be constantly conflating them with completely different ideas.

By the way, it is not just studying in general or being educated in general that is important. The point I'm making is specifically about *philosophy*, and about a particular style of philosophy at that (what we in the biz call "analytic philosophy"). When I talk to academics from other fields, I often find them confused. That is a very common experience among philosophers. To be clear, academics in other fields, obviously, *know their subject* much better than people outside their field know that subject. That is, they know the facts that have been discovered, and the methods used to discover them, which outsiders, including philosophers, do not. But they're still confused when they think about big questions, including questions about the larger implications of the discoveries in their own fields. Whereas, when philosophers think about other fields, we tend to merely be ignorant, not *confused*.

Here is a metaphor (this doesn't *prove* my point; it just helps to *explain* what I'm saying): When we dream, we sometimes dream contradictory things, or things that conflict with basic, well-known features of reality, or things that just in general *make no sense*. You might, for instance, find yourself having a conversation with the color blue. (Okay, that is not a very typical dream. But that illustrates the idea of something that makes no sense in general.) And yet, almost always, we simply *do not notice*. We don't see

the contradictions. We don't have any problem with talking to the color blue. Nothing seems odd. It is only when we wake up that the dream seems strange. Only then do we see all the ways in which it was impossible. We were confused, but we did not know it.

That is how most people are when they think about philosophical questions, if they have not studied philosophy. By studying philosophy, one gradually wakes up and stops saying the things that make no sense. That doesn't guarantee that one knows the truth, of course. But at least one learns to say things that have definite meanings and are possible candidates for being true. This book won't get you all the way there; no book will. But it will get you started, and it will give you some interesting things to think about along the way.

Note: I've included a glossary at the end, which contains all the important terms that appear in boldface throughout the text.

Acknowledgements

I would like to thank Ari Armstrong and Jon Tresan (especially Ari) for their helpful comments on the manuscript, which helped to correct numerous shortcomings. I'd also like to thank Iskra for supporting anything and everything I do; God, if He exists, for creating the universe; and Satan for not maliciously inserting many more errors into the text. Naturally, none of these people are to blame for any errors that remain. All errors are entirely due to the failure of a time-traveling, editorial robot from Alpha Centauri to appear and correct all mistakes before I uploaded the final files. If such a robot had appeared, there wouldn't be any mistakes.

Full Contents

Preface

[Why Read This Book?](#)

[About the Author](#)

[My Approach in Writing This](#)

[Why Study Philosophy?](#)

Part I: Preliminaries

1. What Is Philosophy?

[1.1. The Ship of Theseus](#)

[1.2. What's the Definition of "Philosophy"?](#)

[1.3. Subject Matter & Branches](#)

[1.4. Methods](#)

[1.5. Myths About Philosophy](#)

2. Logic

[2.1. Why Logic?](#)

[2.2. Propositions](#)

[2.3. The Forms of Propositions](#)

[2.4. Characteristics of Propositions](#)

[2.5. Arguments](#)

[2.6. Kinds of Arguments](#)

[2.7. Characteristics of Arguments](#)

[2.8. Why I Hate These Definitions](#)

[2.9. Some Symbols](#)

[2.10. Some Confusions](#)

3. Critical Thinking, 1: Intellectual Virtue

[3.1. Rationality](#)

[3.1.1. What Is Rationality?](#)

[3.1.2. Why Be Rational?](#)

[3.1.3. Truth Is Good for You](#)

[3.1.4. Irrationality Is Immoral](#)

[3.1.5. Some Misunderstandings](#)

[3.2. Objectivity](#)

[3.2.1. The Virtue of Objectivity](#)

[3.2.2. Objectivity vs. Neutrality](#)

- [3.2.3. The Importance of Objectivity](#)
- [3.2.4. Attacks on Objectivity](#)
- [3.2.5. How to Be Objective](#)
- [3.2.6. Open-mindedness vs. Dogmatism](#)
- [3.3. Being a Good Philosophical Discussant](#)
- [3.3.1. Be Cooperative](#)
- [3.3.2. Be Modest](#)
- [3.3.3. Understand Others' Points of View](#)
- [4. Critical Thinking, 2: Fallacies](#)
- [4.1. Some Traditional Fallacies](#)
- [4.2. False Fallacies](#)
- [4.3. Fallacies You Need to Be Told About](#)
- [5. Absolute Truth](#)
- [5.1. What Is Relativism?](#)
- [5.1.1. Relative vs. Absolute](#)
- [5.1.2. Subjective vs. Objective](#)
- [5.1.3. Opinion vs. Fact](#)
- [5.2. Some Logical Points](#)
- [5.2.1. The Law of Non-Contradiction](#)
- [5.2.2. The Law of Excluded Middle](#)
- [5.2.3. What Questions Have Answers?](#)
- [5.3. Why Believe Relativism?](#)
- [5.3.1. The Argument from Disagreement](#)
- [5.3.2. The Argument from Tolerance](#)
- [5.4. Is Relativism Coherent?](#)
- [5.4.1. Conflicting Beliefs Can Be True?](#)
- [5.4.2. Is Relativism Relative?](#)
- [5.4.3. Meaningful Claims Exclude Alternatives](#)
- [5.4.4. Opposition to Ethnocentrism Is Ethnocentric](#)
- [5.5. What Is Truth?](#)
- [5.5.1. The Correspondence Theory](#)
- [5.5.2. Rival Theories](#)
- [5.5.3. Is Everything Relative?](#)
- [5.6. I Hate Relativism and You Should Too](#)
- [Part II: Epistemology](#)
- [6. Skepticism About the External World](#)
- [6.1. Defining Skepticism](#)

6.2. Skeptical Scenarios

6.2.1. The Dream Argument

6.2.2. The Brain-in-a-Vat Argument

6.2.3. The Deceiving God Argument

6.2.4. Certainty, Justification, and Crazyiness

6.3. Responses to Skepticism

6.3.1. Relevant Alternatives

6.3.2. Contextualism

6.3.3. Semantic Externalism

6.3.4. BIVH Is a Bad Theory.

6.3.5. Direct Realism

6.4. Conclusion

7. Global Skepticism vs. Foundationalism

7.1. The Infinite Regress Argument

7.2. The Reliability Argument

7.3. Self-Refutation

7.4. The Moorean Response

7.5. Foundationalism

7.5.1. The Foundationalist View

7.5.2. Arguments for Foundationalism

7.5.3. The Argument from Arbitrariness

7.5.4. Two Kinds of Reasons

7.5.5. A Foundationalist Reply to the Reliability Argument

7.6. Phenomenal Conservatism

7.6.1. The Thesis of Phenomenal Conservatism

7.6.2. The Self-Defeat Argument

7.6.3. PC Is a Good Theory.

7.7. Conclusion

8. Defining “Knowledge”

8.1. The Project of Analyzing “Knowledge”

8.2. The Traditional Analysis

8.3. Gettier Examples

8.4. Other Analyses

8.4.1. No False Lemmas

8.4.2. Reliabilism

8.4.3. Proper Function

8.4.4. Tracking

- [8.4.5. Defeasibility](#)
- [8.5. Lessons from the Failure of Analysis](#)
- [8.5.1. The Failure of Analysis](#)
- [8.5.2. A Lockean Theory of Concepts](#)
- [8.5.3. A Wittgensteinian View of Concepts](#)

[Part III: Metaphysics](#)

[9. Arguments for Theism](#)

- [9.1. Views About God](#)
- [9.2. The Ontological Argument](#)
 - [9.2.1. Anselm's Argument](#)
 - [9.2.2. Descartes' Version](#)
 - [9.2.3. The Perfect Pizza Objection](#)
 - [9.2.4. Existence Isn't a Property](#)
 - [9.2.5. Definitional Truths](#)
- [9.3. The Cosmological Argument](#)
 - [9.3.1. The Kalam Cosmological Argument](#)
 - [9.3.2. Reply: In Defense of Some Infinities](#)
 - [9.3.3. The Principle of Sufficient Reason](#)
 - [9.3.4. Reply: Against the PSR](#)
- [9.4. The Argument from Design](#)
 - [9.4.1. Design and Life](#)
 - [9.4.2. Fine Tuning](#)
 - [9.4.3. Bad Objections](#)
 - [9.4.4. The Multiverse Theory](#)
- [9.5. Pascal's Wager](#)
 - [9.5.1. Pascal's Argument](#)
 - [9.5.1. Objections](#)
- [9.6. Conclusion](#)

[10. Arguments for Atheism](#)

- [10.1. Cute Puzzles](#)
 - [10.2.1. Omnipotence and Immovable Stones](#)
 - [10.2.2. Omnipotence and Error](#)
 - [10.2.3. Omniscience and Free Will](#)
- [10.2. The Burden of Proof](#)
- [10.3. The Problem of Evil](#)
- [10.4. Theodicies and Defenses](#)
 - [10.4.1. How Do We Know What God Values?](#)

[10.4.2. How Would We Know What Goodness Is?](#)

[10.4.3. The Lord Works in Mysterious Ways](#)

[10.4.4. Satan Did It](#)

[10.4.5. God Will Fix It](#)

[10.4.6. Evil Is a Mere Absence](#)

[10.4.7. Evil Is a Product of Free Will](#)

[10.4.8. Evil Is Necessary for Virtue](#)

[10.4.9. God Creates All Good Worlds](#)

[10.4.10. There Is No Best World](#)

[10.4.11. The World Has Infinite Value](#)

[10.4.12. Weakening the Conception of God](#)

[10.4.13. The Case of the Serial Killer](#)

[10.5. Conclusion](#)

[11. Free Will](#)

[11.1. The Concept of Free Will](#)

[11.2. Opposition to Free Will](#)

[11.2.1. The Theory of Determinism](#)

[11.2.2. Evidence for Determinism?](#)

[11.2.2. No Free Will Either Way](#)

[11.3. Deterministic Free Will](#)

[11.3.1. Compatibilism](#)

[11.3.2. Analyses of Free Will](#)

[11.3.3. Freedom Requires Determinism](#)

[11.4. Libertarian Free Will](#)

[11.4.1. Incompatibilism](#)

[11.4.2. For Free Will: The Appeal to Introspection](#)

[11.4.3. Free Will and Other Common Sense Judgments](#)

[11.4.4. For Free Will: The Self-Defeat Argument](#)

[11.5. Other Reflections](#)

[11.5.1. How Does Libertarian Free Will Work?](#)

[11.5.2. Degrees of Freedom](#)

[12. Personal Identity](#)

[12.1. The Teletransporter](#)

[12.2. The Problem of Subject Identity](#)

[12.2.1. Basic Question](#)

[12.2.2. Persons and Subjects](#)

[12.2.3. Two Kinds of Identity](#)

[12.2.4. Identity over Time](#)

[12.3. Theories of Personal Identity](#)

[12.3.1. The Body Theory](#)

[12.3.2. The Brain Theory](#)

[12.3.3. The Naïve Memory Theory](#)

[12.3.4. Psychological Continuity](#)

[12.3.5. Spatiotemporal Continuity](#)

[12.3.6. The No-Branching Condition](#)

[12.3.7. The Closest-Continuer Theory](#)

[12.3.8. The Soul Theory](#)

[12.4. In Defense of the Soul](#)

[12.4.1. Objections to the Soul](#)

[12.4.2. Principles of Identity](#)

[12.4.3. Only the Soul Theory Satisfies the Principles of Subject Identity](#)

[12.4.4. Unanswered Questions](#)

[Part IV: Ethics](#)

[13. Metaethics](#)

[13.1. About Ethics and Metaethics](#)

[13.1.1. Ethics](#)

[13.1.2. Metaethics](#)

[13.1.3. Objectivity](#)

[13.1.4. Five Metaethical Theories](#)

[13.2. What's Wrong with Non-Cognitivism](#)

[13.2.1. The Non-Cognitivist View](#)

[13.2.2. The Linguistic Evidence](#)

[13.2.3. The Introspective Evidence](#)

[13.3. What's Wrong with Subjectivism](#)

[13.3.1. The Subjectivist View](#)

[13.3.2. Motives for Subjectivism, 1: Tolerance](#)

[13.3.3. Motives for Subjectivism, 2: Cultural Variation](#)

[13.3.4. The Nazi Objection](#)

[13.4. What's Wrong with Nihilism](#)

[13.4.1. The Nihilist View](#)

[13.4.2. Against Objective Values: The Humean Argument](#)

[13.4.3. Against Objective Values: The Argument from Weirdness](#)

[13.4.4. Nihilism Is Maximally Implausible](#)

13.5. What's Wrong with Ethical Naturalism

13.5.1. The Naturalist View

13.5.2. A Point About Meaning

13.5.3. Bad Theories

13.5.4. A Bad Analogy

13.6. Ethical Intuitionism

13.6.1. The Intuitionist View

13.6.2. Objection: Intuition Cannot Be Checked

13.6.3. Objection: Differing Intuitions

13.7. Conclusion

14. Ethical Theory, 1: Utilitarianism

14.1. An Ethical Puzzle

14.2. The Utilitarian View

14.3. Consequentialism

14.3.1. Objections to Consequentialism

14.3.2. For Consequentialism

14.4. Hedonism & Preferentism

14.4.1. For Hedonism or Preferentism

14.4.2. Against Hedonism & Preferentism

14.5. Impartialism

14.5.1. Partial vs. Impartial Ethical Theories

14.5.2. For Partiality

14.5.3. For Impartiality

14.6. Rule Utilitarianism

14.7. Conclusion

15. Ethical Theory, 2: Deontology

15.1. Absolute Deontology

15.1.1. Terminology

15.1.2. The Categorical Imperative, 1: Universalizability

15.1.3. The Categorical Imperative, 2: The End-in-Itself

15.1.3. The Doctrine of Double Effect

15.1.4. Rights

15.2. Objections to Absolutism

15.2.1. Extreme Consequences

15.2.2. Portions of a Life

15.2.3. Risks to Life

15.3. Moderate Deontology

- [15.4. Objections to Moderate Deontology](#)
 - [15.4.1. Arbitrary Cutoffs](#)
 - [15.4.2. The Aggregation Problem](#)
- [15.5. Conclusion](#)
- [16. Applied Ethics, 1: The Duty of Charity](#)
 - [16.1. The Shallow Pond Argument](#)
 - [16.2. Objections in Defense of Non-Giving](#)
 - [16.3. Poverty and Population](#)
 - [16.4. Effective Altruism](#)
 - [16.5. Government Policy](#)
 - [16.5.1. The Argument for Social Welfare Programs](#)
 - [16.5.2. The Charity Mugging Example](#)
 - [16.5.3. Other Problems with Government Programs](#)
 - [16.6. Conclusion](#)
- [17. Applied Ethics, 2: Animal Ethics](#)
 - [17.1. Arguments for Ethical Vegetarianism](#)
 - [17.1.1. Where Does Our Food Come From?](#)
 - [17.1.2. The Argument from Suffering](#)
 - [17.1.3. Arguments by Analogy](#)
 - [17.1.4. Animal Rights vs. Welfare](#)
 - [17.2. Arguments in Defense of Meat-Eating](#)
 - [17.3. Other Ethical Issues](#)
 - [17.3.1. The Importance of Factory Farm Meat](#)
 - [17.3.2. Other Animal Products](#)
 - [17.3.3. Humane Animal Products](#)
 - [17.3.4. Insentient Animals](#)
 - [17.3.5. Lab Grown Meat](#)
 - [17.3.6. Animal Experimentation](#)
 - [17.3.7. Responding to Other People's Immorality](#)
- [18. Concluding Thoughts](#)
 - [18.1. What Was this Book Good For?](#)
 - [18.2. How Good Philosophers Think](#)
 - [18.3. Further Reading](#)
- [Appendix: A Guide to Writing](#)

Part I: Preliminaries

1. What Is Philosophy?

1.1. The Ship of Theseus

Here is a classic philosophical problem. Note: I don't use this example because it's such an *important* problem; the reason I like it is that (a) it's easy to get people to quickly see the issue, (b) it is very clearly *not* either a scientific or a religious issue, or any other kind of issue besides philosophical. So it's good for illustrating what philosophy is.

Once there was a Greek hero named Theseus.^[3] He sailed around the Mediterranean Sea doing heroic things like capturing bulls, chopping heads off minotaurs, and abducting women. (Standards of heroism were a bit different back then.) As he was doing all this stuff, his ship suffered some wear and tear. When a particular plank of wood was damaged or rotted, he'd replace it with a new piece of wood. Just one at a time. And let's say that, after ten years of sailing, eventually every one of the original planks of wood had gotten replaced by a new one at one time or another.

Question: At the end of the ten years, did Theseus still have *the same ship* that he had at the beginning? Or was it a new ship?

Now for an amusing modification: Suppose there was someone following Theseus around all those years, collecting all the old pieces of wood as Theseus threw them aside. At the end of the ten years, this person reassembled all the original pieces of wood into a (tattered and ugly) ship. Was *this* ship the same as the original ship Theseus started with?

Notice how this is not a scientific question. It's not as if there's some kind of *experiment* you can do to figure out if it's the same ship. We could try getting a ship and swapping out its parts as in the story. But then what would we do? Observe the ship really closely? Weigh it carefully, observe it under microscopes, do a spectroscopic analysis? None of that would make the slightest difference. We already know the underlying facts of the case, because they're stipulated. We just don't know whether those facts *add up* to the ship being "the same ship" or not.

Notice also, though, that it's not as if there is nothing to say about the issue. You can see why one might think it was the same ship, and also why one might think it wasn't. It is odd to say that Theseus still had the same ship at the end of the story, since it has no parts at all in common with the

original. If anything, the ship reconstructed out of the old planks seems to have a better claim to being the original ship, since it has all the same parts, in the same configuration, as the original.

But if we say that the ship Theseus had at the end of the story *wasn't* the same ship as the original, then at what point did it cease to be the same? Which plank was it whose removal gave Theseus a new ship? To make the argument sharper, let S_0 be the original ship, S_1 be the ship after one plank has been replaced, S_2 the ship after a second plank has been replaced, and so on. Assume the ship has 1000 planks, so the series ends with S_{1000} . Now, presumably replacing just *one* plank of wood doesn't give you a different ship. Therefore, $S_0 = S_1$. But then, by the same reasoning, $S_1 = S_2$. And $S_2 = S_3$, and so on, all the way up to $S_{999} = S_{1000}$. But then it follows that $S_0 = S_{1000}$.

So you can see that one can construct seemingly logical arguments about this question. We're not going to try to resolve the question now. But that is the sort of question philosophers address. Most people in intellectual life – people in other fields – would just try to *avoid* that sort of question. Philosophers try to actually figure out the answer.

By the way, “the answer” need not be one of the answers that the question straightforwardly seems to call for – it doesn't have to be “Yes, it was the same” or “No, it wasn't the same.” (Among philosophers, by the way, almost everything is up for debate, including the terms of the debate and the question being debated.) The answer could be “It neither was nor wasn't the same” or “It's a semantic question” or “It was the same in one sense and different in another sense” or “There are degrees of sameness, and the degree of sameness decreased over time.” This situation is fairly typical of philosophical questions as well: Most questions in other fields of study are meant to be answered straightforwardly in the terms in which they are posed – you're generally not supposed to say the question contains a false presupposition, or has no answer, or needs to be rephrased, etc. But in philosophy, those sorts of responses are on the table.

1.2. What's the Definition of “Philosophy”?

Sorry, I'm not giving you a definition of “philosophy”. It's a field of study, but it does not have a generally accepted definition that differentiates it from all other fields of study. Fortunately, however, people normally do not acquire concepts by hearing definitions; we acquire concepts by seeing

examples. (For example: You acquire the concept “green” by seeing examples of green things, not by someone trying to *tell you* what green is.) That’s why I opened with an example of a philosophical issue. I’ll give some more examples below. I will also offer some generalizations about how philosophical thinking goes, to help distinguish it from, e.g., science and religion.

1.3. Subject Matter & Branches

Most fields of study are distinguished by a certain subject matter (what they study). Biology studies life, meteorology studies weather, UFOlogy studies space aliens, and so on. It’s hard to describe the subject matter of philosophy, because it is very wide-ranging. Here, I will just list the main branches (sub-fields) of philosophy, and what they each study. The first three (metaphysics, epistemology, ethics) are commonly considered the three central branches of philosophy.

*i.***Metaphysics:** Studies general questions about what exists and what sort of world this is. (Not *all* questions about what exists; only, well, the philosophical ones. Hereafter, I leave this qualifier implicit.) (Note: Terms in boldface, like that “metaphysics”, are important philosophical terms that appear in the glossary at the end of the book.)

Examples: Is there a God? Do we have free will, or is everything that happens predetermined (or random, or something else)? Is the future just as real as the present and the past? Do numbers and other abstract objects really exist? Is reality objective or subjective?

*ii.***Epistemology:** Studies the nature of knowledge, whether and how we know what we think we know, and whether and how our beliefs are justified.

Examples: What is the definition of “know”? How do we know that we can trust the five senses? How do you know that other people are conscious and not just mindless automata? Are all beliefs justified by observation, or are some things justified independent of observation?

*iii.***Ethics:** Studies right and wrong, good and bad.

Examples: Is pleasure the only good in life? Are we sometimes obligated to sacrifice our own interests for the good of others? What rights do people have? Is it ever permissible to violate

someone's rights for the good of society? Do non-human animals have rights?

iv.**Political Philosophy:** Studies good and bad social institutions, and how society ought to be arranged.

Examples: What gives the government authority over the rest of us? What is the proper function of government? What is the most just distribution of wealth in a society? When should the state restrict people's liberties for the good of society?

v.**Aesthetics:** Studies art, beauty, and related matters.

Examples: What is art? Is modern art really art? Is beauty objective, or is it in the eye of the beholder? In what way, if at all, can we learn about reality from reading fiction? Is a work artistically flawed if it expresses immoral values?

vi.**Logic:** Studies valid reasoning, certain general characteristics of propositions, and when propositions support or conflict with each other.

Examples: What are the rules for when an argument is valid? Must every proposition be either true or false? Can a proposition ever be both true *and* false?

vii.**Philosophy of Mind:** Studies the nature of the mind and consciousness.

Examples: Is the mind just the brain, or is it some kind of non-physical thing? Why is there consciousness; why do humans (and most animals) have experiences that *feel* like something, rather than just being complicated mechanisms with no experiences? How is it that we can have states that are "about" something or "represent" something?

viii.**Philosophy of Science:** Studies philosophical questions about how science works and the philosophical implications of scientific theories.

Examples: How do we know when a scientific theory is true? Why should we prefer simpler theories over more complex ones? Does quantum mechanics show that reality depends on observers? Does the theory of relativity show that the future is just as real as the past and present? Does the theory of evolution undermine belief in objective values?

Aside: You might have noticed that the above branches seem to overlap with each other in several ways. If you noticed that, you are correct. If you didn't notice that, pay more attention!

1.4. Methods

You might also have noticed that the above list of philosophical questions overlaps with some religious and scientific questions. So now I'm going to tell you some broad ways that philosophy differs from religion and science, even when they are studying similar questions. Those differences have to do with *methods*, i.e., philosophers use different ways of trying to reach conclusions.

Religions typically appeal to *authority* and (alleged) *supernatural sources* of knowledge. Note: This does not mean that religious figures never appeal to ordinary observations or reasoning. Of course, they often appeal to observation and reasoning. It's nevertheless true that appeals to authority and supernatural sources of knowledge play a crucial role in the world's established religions. In other words, in traditional religions, there are key claims that one is meant to accept because they come from a particular person, or institution, or because they appear in a particular book, or something like that. And one is supposed to trust that person or institution or book because it (or its author) had a form of supernatural access to the truth, something that goes beyond the ordinary ways of knowing that all of us have (such as reason and observation by the five senses). Thus, in Catholicism, one is meant to trust the Pope due to the Pope's special relationship to God. In Christianity more broadly, one is meant to trust the Bible because it is allegedly the inspired word of God. In Islam, one is meant to trust the Koran because it, again, allegedly derives from a divine revelation. Similarly for Judaism and the Torah. In Buddhism, one is meant to trust the Buddha's wisdom, because it allegedly derived from his attainment of Enlightenment, whereby he escaped the cycle of rebirth into Nirvana. (Aside: Buddhism is closer to the border between philosophy and religion than the other religions. In fact, some would call it a philosophy rather than a religion.)

Science, by contrast, does not appeal to supernatural knowledge sources to justify its theories. It appeals most prominently to observation, especially *specialized* observations. That is, it usually appeals to observations made by scientists that most people have not made but could make. These are usually observations that one has to collect by first setting up a very specific

experiment. Example: If you apply an electric voltage to a sample of water, you can observe bubbles forming at both electrodes. If you are very careful and very clever, you can verify that the water is turning into hydrogen and oxygen gas. That is part of how scientists know that water is H_2O . You've probably never observed this, but if you set up the experiment in the right way, you could.

Not all scientific evidence depends on an experimental manipulation of the environment. For instance, the main evidence showing that all the planets orbit the Sun comes from meticulous observations of the positions of planets in the night sky at different times, made by incredibly patient astronomers. You probably haven't made these observations, but, again, you could.

By the way, I am not saying any of this for the purpose of either attacking or defending religion, or attacking or defending science. That is not my concern. I am just factually describing how these pursuits work and are generally agreed to work. My point is to explain how they differ from philosophy.

Philosophy (at least modern, academic philosophy) appeals to (allegedly) logical arguments, where the premises of these arguments usually come from common experience, including well-known observations or common intuitions (that is, roughly, things that just seem to make sense when we think about them). It will generally not require supernatural access to the truth, nor will it generally require experiments or other highly specialized observations.^[4]

1.5. Myths About Philosophy

Now I'll address some things that people sometimes think about philosophy that are false.

Myth #1: Philosophers sit around all day arguing about the meaning of life and the nature of Truth.

Comment: Well, the meaning of life is a philosophical question, and philosophers argue about any philosophical question. But "What is the meaning of life?" happens not to be a very widely discussed philosophical question – very few philosophers have ever written anything about it. Similarly for the question, "What is truth?" There are some philosophers who work on theories of truth, but relatively few. The questions listed above (section 1.3) are more commonly discussed.

This myth isn't very bad, though, because it's just a matter of emphasis. The next myth is worse.

Myth #2: Philosophy never makes progress. Philosophers are still debating the same things they were debating 2000 years ago.

Comment: No, that's completely false.

a. On "debating the same questions": Here are some things that philosophers were not debating 2000 years ago: Criteria of ontological commitment. Modal realism. Reliabilism. Semantic externalism. Paraconsistent logic. Functionalism. Expressivist metaethics.

You probably don't know what any of those things are. But those are all well-known and important topics of contemporary debate which any philosophy professor will recognize, and none of them was discussed by Plato, or Aristotle, or any other ancient philosopher. Though Western philosophy has been around for 2000 years, *none* of those issues, to the best of my knowledge, was ever discussed by anyone more than 100 years ago. And having seen that list, any professional philosopher could now extend it with many more examples.

b. On progress: Here are some questions on which we've made progress:

i. Is slavery just? No joke! Aristotle, often considered history's greatest philosopher, thought slavery was just. No one thinks that anymore.

ii. Which is better: dictatorship or democracy? Seriously, Plato (also considered one of history's greatest philosophers) thought the answer was "dictatorship" (as long as the dictator is a philosopher!). No one thinks that anymore.

iii. Is homosexuality wrong? Historically, philosophers and non-philosophers alike have held different views on this question, with many thinking homosexuality was morally wrong, including such great philosophers as Thomas Aquinas and Immanuel Kant. Today, almost everyone agrees that homosexuality is obviously fine.

iv. Is nature teleological? Historically, many philosophers, following Aristotle, thought that inanimate objects and insentient life forms had natural goals built into them. Conscious beings had such goals too, and they didn't necessarily correspond to what those beings wanted. Today, hardly anyone thinks that. (The small number who

do are almost all Catholic philosophers, because that was what Catholicism's greatest philosopher, Thomas Aquinas, thought.)

v. What is knowledge? The orthodoxy in epistemology used to be that "knowledge" could be defined as "justified, true belief". Today, basically everyone agrees that that's wrong.

None of the above are minor cases. These are all significant changes on important issues. Granted, in some cases, philosophical progress consists in rejecting an old view about a question without achieving consensus on the correct view, as in case (v). But rejecting false views is an important kind of progress.

Some of the above examples might strike you as obvious, so you might be intellectually unimpressed. "Slavery is wrong. Well, duh", you might say. But in fact that was not at all obvious to people 2000 years ago, not even to the smartest and most educated people. And it is a super-important discovery. And by the way, it almost certainly *wouldn't* be obvious to *you*, if you hadn't been taught that slavery is wrong by other people in your society.

Have we found the answers to every question? Obviously not. But have we made important progress? Obviously so.

Myth #3: Doing philosophy is all about giving your opinion, or saying how you feel about things.

Comment: I don't know if many people think that, but it seems that some students think it. When you're doing philosophy – like when you're writing a paper, or talking in class, or talking with other philosophy buffs – no one wants to hear *mere* opinions. I mean by that, opinions that aren't supported by evidence or logical reasoning. We do not just express our feelings; if you're doing that, you're doing it wrong. Doing philosophy is about thinking things through carefully.

Myth #4: In philosophy, there are no answers.

Comment: Philosophers disagree about a lot of things, but one thing almost everyone in the field agrees on is that the above statement is false. By the way, it's also incoherent, since it is itself an alleged philosophical answer. No, there are answers. If you're wondering whether we ever get any closer to *finding* those answers, see Myth #2 above.

That will do for an initial explanation of philosophy. I hope the above remarks gave you some sense of what the field is like. You'll get a better sense from reading about the philosophical issues in the rest of the book.

2. Logic

2.1. Why Logic?

I know, you probably want to get to the good stuff – God, free will, good and evil, etc., etc. But before we do that, we have to learn some background about logic, and we have to draw some distinctions that philosophers use in debates about the good stuff. Otherwise, you won't be able to understand the arguments about the good stuff.

Logic tells us about what counts as good reasoning. There are well worked-out, precise systems of rules for classifying arguments as valid and invalid. But that isn't the main thing I want you to learn in this chapter, and I won't go through all those rules here. Usually, you don't need to look at the rules to see if an argument is valid; indeed, looking at the rules more often confuses students. (The rules are abstract and sort of mathematical. If someone doesn't get the logic of an argument intuitively, the best thing is usually to break down the argument into smaller steps, or to give some analogies, rather than to start appealing to abstract systems of rules.) So what do I want you to learn in this chapter?

I want you to learn some of the logical *concepts* and *distinctions* that philosophers use when we talk about arguments. I want you to know what I'm talking about if I say "that's not valid" or "that's a contingent proposition". I also want you to learn some symbolism that is useful later. In what follows, I will put important philosophical terms that you should learn in **boldface**, so you can easily find them later.

2.2. Propositions

Propositions are things that can be true or false – but wait, I need to distinguish three sorts of things that can be true or false.

1. *Sentences*. Sentences are sequences of words like what you're looking at right now. Not *all* sentences can be true or false; e.g., questions or commands cannot be. Only assertive sentences, or proposition-expressing sentences, can be true or false. For instance, "It is raining" is true or false; "Is it raining?" and "Make it rain!" are not.

2. *Beliefs.* Beliefs are a kind of mental state, a state of thinking something to be the case. They are typically expressed using assertive sentences. They need not actually be expressed, though; you could just think silently to yourself that it is raining. The thought must be either true or false. This contrasts with, e.g., emotions, desires, or sensations, which are neither true nor false.
3. *Propositions.* **Propositions** are the sort of things that beliefs and statements are *about*. When you have a belief, there is something that you believe to be the case; when you make an assertion, there is something you are asserting to be the case. That thing is a “proposition”. Propositions are sometimes thought of as ways the world could be (possible states of affairs), or ranges of possibilities.

A proposition should not be confused with a belief, since the proposition is the thing *that* one believes, not the belief itself. The same proposition can be believed by one person and doubted by another. One person may *believe* that we will colonize Mars, while another merely *hopes* that we will, a third *doubts* that we will, a fourth *is glad* that we will, and so on. In this case, these different attitudes would all be attitudes toward the same proposition. This proposition is what is denoted by the phrase “that we will colonize Mars”.

A proposition also should not be confused with a *sentence* or phrase in a particular language. The proposition is not the *phrase* “that we will colonize Mars”; it is the thing that that phrase *refers to*. (Compare: The Eiffel Tower is not to be confused with the *expression* “the Eiffel Tower”; the Tower is *the referent* of that expression.) The sentences “We will colonize Mars” and “Nous allons coloniser Mars” have something in common. (The second one is the French translation of the first.) They are obviously not the *same sentence*, but they do say the same thing – that is, they express the *same proposition*.

You have to know what a proposition is. If you’re not sure you understood the above, read it again. If you’ve already read it again, read some more about propositions below, and then read the above again.

2.3. The Forms of Propositions

Propositions have structures – that is, they have different kinds of components, which can be connected to each other in different ways. The

structure is often referred to as the “form” of the proposition.

The simplest kind of proposition has a simple subject-predicate form. That is, there is a thing the proposition is about (the “**subject**”), and there is a way that thing is said to be, or a property that is ascribed to the thing (the “**predicate**”). Example: [Donald is angry]. Note: I often use square brackets like that, to refer to propositions. This proposition is about Donald, and the way he is said to be (the property ascribed to him) is *angry*.^[5] So Donald is the subject, and the property of being angry is the predicate. Notice that neither of these things by itself is a proposition. Donald, by himself, cannot be true or false. Anger by itself cannot be true or false. The subject must be combined with a predicate to have a proposition.

Some propositions are compound, meaning that they have other propositions as components. Example: [If Donald is angry, then he is dangerous]. In this case, there are two simple propositions, [Donald is angry] and [Donald is dangerous], which are combined using an “if-then”. Sentences or propositions like this (using “if ... then”) are known as “**conditionals**”. The “if” part (in this case, [Donald is angry]) is known as the “**antecedent**” of the conditional. The “then” part (in this case, [Donald is dangerous]) is known as the “**consequent**” of the conditional.

Another type of compound proposition is a **conjunction** (an “and” statement), for example, [Donald is angry and dangerous]. The two parts are called the “**conjuncts**”. In this case, the first conjunct is [Donald is angry], and the second conjunct is [Donald is dangerous].

Another type is a **disjunction** (an “or” statement). The two parts are called **disjuncts**. So in the disjunction [Jesus is a liar, a lunatic, or the Lord], there are three disjuncts. The first disjunct is [Jesus is a liar], the second disjunct is [Jesus is a lunatic], and the third disjunct is [Jesus is the Lord].

We also sometimes talk about **negations**, which are propositions that deny another proposition. For instance, [Jesus is not a liar] is a negation; specifically, it is *the negation of* [Jesus is a liar].

2.4. Characteristics of Propositions

Here are some characteristics that propositions can have:

- i. They can be *true* or *false*. A proposition is true when it corresponds to (or just *is*) the way the world is. A proposition is false when it doesn’t correspond to (or isn’t) the way the world is. For instance, the proposition [Armadillos are furry] is true if and only if

armadillos are furry. In this case, the proposition is false, since armadillos aren't furry.^[6]

- ii. They can have *probabilities*. For instance, I am very unlikely to live past the age of 100 years. This is to say that the proposition [I will live past the age of 100] has a low probability.
- iii. They can be *possible* or *impossible*. A proposition is possible when it is or could have been true; impossible when it is not and could not have been true.

Note that philosophers usually (unless otherwise specified) use “possible” in an extremely broad sense, which we call “**metaphysical possibility**”. In this broad sense, things can be “possible” even if you know they are false, and even if they contradict the laws of nature. A good test for what is considered “metaphysically possible” is what it would make sense to entertain as part of a fictional story. Example: In the great television series *Star Trek*, people regularly travel faster than the speed of light, which contradicts the real-world laws of physics (Special Relativity in particular); nevertheless, there is nothing *nonsensical* about this. On the other hand, people also sometimes travel back in time, and when they do this, they very often take actions that would prevent them from ever being in the position to make the journey back in time in the first place. *Those* stories really are nonsensical. The self-undermining time travel aspect of the stories is metaphysically *impossible*, but the faster-than-light travel aspect is metaphysically *possible* – which is why savvy fans frequently object to the time travel but never object to the faster-than-light travel.

Some other examples: π could not have been 4. The sun could not have been completely green and at the same time completely orange. Time could not have been two dimensional. A coin could not have a $2/3$ probability of coming up heads and also a $2/3$ probability of coming up tails. All those things are impossible in the metaphysical sense.

However, there *could* have been eleven planets in the solar system. There could have been no Earth. Newtonian physics could have been true, rather than relativity and quantum mechanics. There could have been unicorns living on Mars. Those things are all metaphysically possible.

Other senses of “possible” include epistemic possibility (which applies to things that are consistent with everything we know) and physical possibility (which applies to things that are consistent with all the physical laws, such as the law of gravity, conservation of momentum, and so on).

iv. They can be *necessary* or *contingent*. A **necessary** proposition is one that could not have been false; that is, its negation is impossible. A **contingent** proposition is one that could have been true and could have been false; that is, it is neither necessary nor impossible. These terms are also usually used in the “metaphysical” sense (where a proposition is metaphysically necessary when its negation is metaphysically impossible), unless otherwise stated.

2.5. Arguments

An **argument** is a series of statements, some of which are supposed to provide reasons for others, where the whole series is meant to justify a particular conclusion. (We also sometimes speak of an argument as the series of *propositions* expressed by such statements.) These are the different parts of an argument:

i. *Premises*: The **premises** of an argument are the statements that are used to support the other statements. We reason *from* the premises. Premises are usually chosen to be things that are widely accepted or would seem obvious to most people. (If not, then you may need further arguments to support the premises.)

ii. *Conclusion*: The **conclusion** of an argument is the statement that the argument is meant to justify. We reason *to* the conclusion. The conclusion is usually something that is initially controversial or non-obvious. (Otherwise, we would not need the argument.)

iii. *Intermediary steps*: Sometimes an argument has intermediary steps. These are steps in between the premises and the conclusion that help you to see how the premises support the conclusion.

Philosophers talk a lot about arguments. We debate whether an argument is good or bad, which of two conflicting arguments is stronger, and things like that. In these discussions, it is crucial to distinguish the *premises* from the *conclusion* of someone’s reasoning; otherwise, you will totally fail to understand what’s going on. (Most common mistake about

this: Students confuse argument with mere assertion, and thus think that an author's premises are just paraphrases of the conclusion.)

Why are arguments so important? Because philosophers (also other people, but especially philosophers) think about a lot of issues that are difficult, and controversial, and it's hard to see the right answer. These are interesting and important issues, and we want to know the truth about them. So we try to find other things that are more obvious, less controversial, or in general easier to know, and use these other things to figure out the answers to the difficult philosophical issues. Arguments are not just for propagandistic purposes (to influence other people); they are, first of all, for helping *us* figure out what to believe. People with different views then tell each other their arguments, which explain why they hold the views they do, so that they can learn from each other.

Aside: I know, it often doesn't seem as if people are trying to *learn* from each other when they trade arguments! It might seem they are just trying to win some kind of personal contest, or score points for "their side". When that happens, you're doing it wrong. Now, maybe people almost always do it wrong. You should still try to do it right.

2.6. Kinds of Arguments

Here are some different kinds of arguments:

i. Deductive: A **deductive argument** is one in which the premises are supposed to support the conclusion conclusively, that is, in such a way that it would be *contradictory* for the premises to be true and the conclusion not to be true.

Now, why did I just say "supposed to"? Because a deductive argument could contain a mistake; it might appear that the conclusion logically follows from the premises, but on closer inspection, you might see that it doesn't really. In that case, it is still a "deductive argument"; it's just a *bad* (mistaken) deductive argument.

Example: "Socrates is a college student. All college students are sober all the time. Therefore, Socrates is sober all the time." That's a deductive argument.

ii. Non-deductive: A non-deductive argument is any argument that isn't deductive. This would be an argument in which the premises are supposed to provide *some* grounds for the conclusion, but not conclusive grounds. The premises are just supposed to show the

conclusion to be probable, or more probable than one would otherwise think. For examples, see below.

iii. *Inductive*: An **inductive argument** is *one type of* non-deductive argument, in which one generalizes from particular cases. Example: You check a lot of honey badgers, and they are all mean. You infer that, most likely, the next honey badger you run into will be mean.^[7]

iv. *Inference to the Best Explanation*: **Inference to the best explanation** is another kind of non-deductive argument, one in which the conclusion is supposed to be supported (probabilistically) because it provides the best explanation for the facts cited in the premises. Example: The east coast of South America is about the same shape as the west coast of Africa; also, various fossils found in eastern South America are similar to fossils found in west Africa. The best explanation for these things (and a bunch of other evidence), is that some time in the distant past, the two continents were joined, and then they drifted apart. So, that's probably what happened.

v. *Arguments by analogy*: **Argument by analogy** (or “analogical argument”) is a type of argument in which one tries to arrive at a conclusion about *x* by looking at *y*, which is thought to be *analogous* (that is, similar in the ways that matter) to *x*. This is used very often in philosophy, especially in ethics and political philosophy. It's hard to give a brief example, but here goes. Example: Say we're debating whether it is morally obligatory to donate some of our income to charity. Someone says: “What if you saw a child drowning in a shallow pond? You could easily wade in and pull the child out, at minimal cost to you. Surely it would be wrong *not* to save the child. This is similar to someone deciding whether to donate money to charity, where they could also save someone's life at small cost to themselves. So it's also wrong not to donate to charity.”

In this case, we say that the child in the pond is *analogous to* the poor people who could be helped if we gave money to charity. Refusing to give to charity is analogous to letting the child drown. The inconvenience of wading into the pond is analogous to the small sacrifice we have to make to donate to charity. Etc. To dispute the

argument, one would have to either disagree that saving the drowning child is obligatory, or claim that there is some *relevant difference* between the drowning child and the people who could be helped by charity.

- vi. *Reductio ad Absurdum*: Very commonly used in philosophy, a **reductio ad absurdum** is a kind of deductive argument in which one derives an absurd consequence from a theory, and thereby concludes that the theory must be false. Example: You're having a debate with an animal-rights advocate. You say, "It's okay to eat animals, because we are much more intelligent than they are." The animal-rights person points out that this implies that it would be okay to eat human babies, because they are also much less intelligent than we are. Since that's absurd, your theory must be mistaken (it's false that we may eat animals simply because we are much smarter than they are).

2.7. Characteristics of Arguments

Now here are some characteristics an argument can have:

- i. *Valid or invalid*: An argument is said to be **valid** (or "deductively valid" or "logically valid") when the premises entail the conclusion; that is, *it would be impossible (in the sense of contradictory) for all the premises to be true and the conclusion to be false*.

Note: This is not the ordinary English usage of "valid"; this is a special, technical usage among philosophers. Virtually all philosophers use the word this way, so you have to learn it. In this sense, "validity" does not require the premises of the argument to be correct, or reasonable, or even consistent. The only thing that is required is that it not be possible that the premises all be true *and* the conclusion be false.

Example: "Socrates is a fish. All fish live on Mars. Therefore, Socrates lives on Mars." That's valid, because it could not be that Socrates is a fish, and all fish live on Mars, *and* that Socrates doesn't live on Mars.

- ii. *Sound or unsound*: An argument is said to be **sound** when it is valid (in the sense given above) *and* all of its premises are true. (In this case, of course, the conclusion must also be true – you can see that if you understood the definition of "valid".) An argument is

unsound whenever it is invalid *or* has a false premise. Note: This is also a technical usage, not the ordinary English usage, and again, philosophers take the stated definition perfectly strictly and literally. Example: “The sky is blue. If the sky is blue, then it isn’t green. Therefore, the sky isn’t green.” That’s sound.

iii. *Circular or non-circular*: We say that an argument is **circular** or **begs the question** when the premises contain the conclusion, or they contain a statement that is so *similar* to the conclusion that you couldn’t believe the premises without already believing the conclusion, or they contain a statement whose justification *depends upon* the justification of the conclusion. Example: “God exists; therefore, God exists.”

Here is a more realistic example: “Everything the Bible says is true, since the Bible is the word of God. And we know the Bible is the word of God because the Bible says that the Bible is the word of God.” This argument is circular, because our (alleged) justification for believing that everything in the Bible is true is that the Bible is the word of God; but at the same time, our justification for believing that the Bible is the word of God depends upon our being justified in thinking that everything in the Bible is true.

By the way: The phrase “beg the question” *does not mean* “to raise the question”. It means “to give a circular argument”. It fills me with a burning rage when I see people abuse the phrase “beg the question”, so don’t do that.

iv. *Cogent or uncogent*: A **cogent** argument is one in which the premises make the conclusion more probable. An uncogent argument is one in which the premises don’t make the conclusion more probable. We need this concept for assessing non-deductive arguments. A deductive argument will be flawed if it is invalid, but a non-deductive argument is not necessarily flawed just because it isn’t valid (non-deductive arguments aren’t *trying* to be valid). Instead, a non-deductive argument is flawed if it is uncogent.

The above categories are used for assessing arguments, especially for discussing what might be wrong with a given argument. If you have a deductive argument, then the argument needs to be valid, sound, and non-circular. If it is invalid, or has a false premise, or begs the question, then it’s a bad argument. (It’s also bad if we merely lack justification for believing

one of the premises.) If you have a non-deductive argument, then it needs to be cogent, have true (and justified) premises, and not beg the question.

2.8. Why I Hate These Definitions

You don't have to read this section. It's very subversive and will get you kicked out of philosophy parties. (Note: Not really.) I'm about to tell you why I hate the way my fellow philosophers (and I!) use the words "valid" and "sound". Look at the following six cases, and try to say which of these are valid/sound arguments:

(A)(An argument with obviously false premises)

Socrates is a fish.

All fish live on Mars.

Therefore, Socrates lives on Mars.

According to the technical philosophers' usage, as discussed above, that's valid.

(B)(An argument with contradictory premises)

It is raining.

It is not raining.

Therefore, blue unicorns built the Taj Mahal.

Also valid. It is impossible for both of the premises to be true and the conclusion false. (Of course, that's because it's impossible for both premises to be true, *regardless* of what the conclusion does.) That satisfies the definition of "validity". So, if you think about it, having self-contradictory premises *guarantees that your argument is valid*. That's right. If you ask any philosopher, they will agree with this. It clearly follows directly from the definition, so what's your problem?

(C)(An argument with irrelevant premises)

The sky is blue.

Therefore, $7=7$.

It is impossible for the premise to be true and the conclusion false. (Of course, that's because it's impossible for the conclusion to be false, regardless of what the premise does.) So this counts as valid. Whenever the conclusion is a necessary truth, your argument is guaranteed to be valid, even if the premise has nothing at all to do with the conclusion. The above argument is also sound, since, in addition to its being valid, its premise is true.

(D)(A circular argument)

God exists.

Therefore, God exists.

It is impossible for the premise to be true and the conclusion to be false. That's because the premise is identical to the conclusion. So, if you think about it, giving a blatantly circular argument also guarantees that your argument is valid. If God in fact exists (whether or not anyone knows it), the above argument is also sound.

(E)(Arguments with unjustified premises)

(i)The number of stars in the universe is even.

If the number of stars in the universe is even, then the Theory of Evolution is true.

Therefore, the Theory of Evolution is true.

(ii)The number of stars in the universe is odd.

If the number of stars in the universe is odd, then the Theory of Evolution is true.

Therefore, the Theory of Evolution is true.

Both of those arguments are valid, and one of them, either (i) or (ii), is also sound. I don't know which one, because I don't know if the number of stars is even or odd. If it's even, then (i) is sound; if it's odd, then (ii) is sound. That's true even though it would be irrational to accept either argument, since both premises in both arguments are unjustified. One of them in fact has *true* premises, but even so, its premises are *unjustified* because we don't have any reason to think that either (i) or (ii) is more likely to be the correct one; hence, neither of these arguments could have anything to do with how anyone knows the Theory of Evolution.

(In case you're wondering why the second premise of one of these arguments would be true, it's because a conditional is considered to be true whenever the antecedent and the consequent are both true. If you don't like that, you can change the second premise in (i) to "The Theory of Evolution is true, or the number of stars isn't even" and in (ii) to "The Theory of Evolution is true, or the number of stars isn't odd".)

(F)A non-deductive argument:

In a random sample of 10,000 honey badgers drawn from the world population of honey badgers, it was found that 100% of them were mean.

So (most likely), the next honey badger we meet will be mean.

Here, it is *possible* (though highly unlikely) for the premise to be true and the conclusion false. So the argument is invalid, and hence also unsound – even though it is a perfectly reasonable argument to make (assuming we really made the observations described).

I hate the philosophical usage of “valid” and “sound” because in normal English, “valid” and “sound” both sound like they mean “perfectly okay”, or something like that. So a normal person would probably say arguments A-E are all “invalid” and “unsound”, while only F is “valid” and “sound”. But in philosophy-talk, A-E are all *valid*, while only F is *invalid*. Meanwhile, C, E, and possibly D are also *sound*. The philosophical usage just seems calculated to cause confusion. And yet, I still have to teach this annoying usage to you, because basically 100% of philosophers insist on using the words in this way. So if you use the words as in *ordinary English*, philosophers won’t understand you; and if you use the words as they are used in philosophy, *non-philosophers* won’t understand you. We’ve basically ruined these words.

And the annoyance of these definitions isn’t limited to their confusing students and impeding communication between philosophers and everybody else. Apart from what words we use for them, the categories established by these definitions are not even the most useful categories for evaluating arguments. If we’re introducing terms to refer to desirable properties of an argument, the different terms should refer to *separate* properties. E.g., there should be one term that denotes having premises that support the conclusion, and another term that denotes having true premises. The second term shouldn’t refer to what the first term refers to *conjoined* with some other thing. (As “sound”, in the philosophical usage, refers to the property of being valid *plus* having true premises.) There should also be a term for each major virtue that an argument should have. So there should be a term that refers to the property of having *justified* or *reasonable* premises. But there isn’t.

<End rant.> Now to be clear, for the rest of the book, if I use “valid” or “sound”, assume I’m using them in the philosophers’ technical sense, even though it’s terrible.

2.9. Some Symbols

The following symbols are commonly used in philosophy.



<u>Symbol</u>	<u>Meaning</u>	<u>Example</u>
A, B, C, ... $p, q, r,$...	Stands for some proposition. You can use any letter. But A, B, C, $p, q,$ and r are philosophers' favorites.	All cats are fluffy.
$\sim A$	It's not the case that A. Often read "not-A".	Not all cats are fluffy.
(A $\vee$ B)	Read "A or B". Note that the "or" is normally read inclusively, so it means "A or B or both."	All cats are fluffy or some hippos are fat.
(A & B)	A and B.	All cats are fluffy and some hippos are fat.
$\therefore$	Read "therefore". This is often placed before the conclusion of an argument.	

These symbols can be combined in various ways to represent propositions and arguments. This may help show the structure of an argument, and anyway it lets us write things concisely. Example: Consider the argument, "Either Skeezix or Ted ate the goldfish. Ted didn't do it. Therefore, Skeezix did it." This can be symbolized like so:

(S $\vee$ T)
 $\sim T$
 $\therefore S$

That, by the way, is valid.

If this were a logic book, I would give you a complicated series of rules for testing whether an argument is valid, based on how it is symbolized. But it isn't, and I won't.

2.10. Some Confusions

Almost done! Now despite how it might seem, I didn't write this chapter to torment you. I wrote it because there are certain concepts and distinctions that you simply have to understand in order to think clearly about philosophical issues (they're good for thinking about other things too). In my years of teaching philosophy, I have run into many a confused statement, and many a confused ... um, sequence of words. Most of them

are idiosyncratic – that is, they seem to be confusions of a particular individual, not generally shared. So I can't really anticipate them.

However, I have noticed some *types* of mistakes that recur frequently, and some *subject matters* that confusions seem to swarm around. So I'm going to say a little about that.

Category errors: A really common general type of mistake is known as a "category error". This is the mistake of applying a predicate to a type of thing (a "category") that that predicate just doesn't go with. For any predicate, there is generally a type of thing that it is supposed to apply to. For instance, the predicate "green" can be applied to material objects ("Kale is green") and light ("Green light was reflected off the surface"). It cannot be applied to thoughts, or numbers, or actions, etc. So if you say, "I just had a green idea", that's a category error.

Another example: Take the predicate "is the uncle of" (as in "John is the uncle of Barbara"). This refers to a relationship that two people might stand in. This predicate applies to *people*. You could also apply it to other animals that have sexual reproduction. You cannot apply it to plants, or amoebas, or tables, or concepts. So if you say, "The Earth is the moon's uncle", that is a category error.

Of course, people very rarely commit category errors about colors or uncles. But they very often commit category errors about *propositions*. That's why I've been blabbing on about propositions for so long. I'm trying to make you think about them long enough to form a settled category of *propositions* in your mind, as something distinct from statements, properties, arguments, etc.

Now let me give you some examples. Start by looking at the left column, and try to identify the problem in each case. Then look at the right column.

<u>Nonsense statement</u>	<u>Problem</u>
"That argument is false!"	Arguments cannot be true or false. Truth and falsity only apply to <i>propositions</i> . An argument can, however, be valid, invalid, sound, or unsound (and a proposition can't be any of those things).
"He tried to disprove rationality."	Proving and disproving can only be done to propositions. Rationality isn't a proposition; it's a character trait.

“Causation
contradicts
freedom.”

Contradiction is a relationship between propositions.
Causation and freedom aren't propositions.

“My
action was
caused by
determinism.”

Determinism is a proposition. But propositions don't
cause things. Events and states cause things. Your action
might be caused by your desires, or your upbringing, or
electro-chemical activity in your brain, or something like
that. It can't be caused by determinism.

Those are all category errors.

The problem with a category error isn't that you violated some arbitrary linguistic convention, one that just says you're not supposed to use this predicate in this certain way. The problem with a category error is that it doesn't express a coherent thought. If someone claims to have had a green idea, it's not just that they're violating a rule of language. It's that we don't know what they're trying to say, because ideas just can't have that property.

Now, in some of my examples, you might think that you *do* know what the person is trying to say; they just weren't completely correct and explicit in the way they expressed it. But in my experience, when you think you know what someone is trying to say, it very often turns out *not* to be what they are trying to say at all. That's why it's usually best to be explicit, and to say things correctly.

E.g., maybe you think you know what “That argument is false” means. Maybe you think it means, “That argument has a false conclusion.” But maybe it means “That argument contains at least one false statement” or “That arguments contains *all* false statements” or “That argument is invalid” or something else entirely. Don't make people guess what you're trying to say by misusing words.

Concepts vs. propositions: As the above examples illustrate, it is common for people to confuse concepts with proposition-expressing thoughts. E.g., there is a concept of rationality; this concept, by itself, doesn't assert anything. It doesn't assert that rationality is good, or that rationality is bad, or anything else. If you have the concept of rationality, you can believe or not believe whatever you want about rationality. So the concept of rationality should not be confused with a belief (or other propositional thought); similarly, the *term* “rationality” should not be confused with any *statement* about rationality.

Arguments vs. conditionals: These two things are sometimes conflated. An if-then statement is not an *argument*. It is just a *statement*. For instance, “If Skeezix looks guilty, then he ate the goldfish” is not an argument. It’s a single statement (specifically, a conditional statement). The following *is* an argument:

If Skeezix looks guilty, then he ate the goldfish.

Skeezix looks guilty.

Therefore, Skeezix ate the goldfish.

But a person who just says “*If* Skeezix looks guilty, then he ate the goldfish” is *not* committed to endorsing that argument – they might think Skeezix *doesn’t* look guilty and *didn’t* eat the goldfish. Or they might have no idea whether Skeezix looks guilty. So they may not be supporting any argument on the subject of the goldfish’s fate at all.

Concepts vs. things: The term “concept” also calls forth a great deal of confusion. For some reason, it is very common to commit category errors of the form “So-and-so is a concept.” As in, “Freedom is a concept.” Um, no, freedom isn’t a concept. Freedom is, roughly, an absence of constraints. That is an actual condition that you can be in (and it doesn’t consist of having some concept). Of course, *the concept of freedom* is a concept, but that’s trivial. It’s equally true that the concept of broccoli is a concept. But that doesn’t tell us anything interesting about broccoli (or the concept thereof). The fact that the concept of freedom is a concept also doesn’t tell us anything interesting about freedom or the concept thereof. The same goes for such nonsense sentences as “Time is a concept”, “Infinity is a concept”, and “God is a concept”. Those are all category errors.

My best guess is that when people make those statements, they are trying to say something like, “The concept of freedom is relatively abstract”, but they are also confusing concepts with the things that the concepts represent (confusing *x* with the concept of *x*).

Sometimes, people use the “So-and-so is a concept” locution for non-existent things. As in, “Unicorns are a concept”. No, they’re not. Unicorns are horses with horns on their heads – or at least, that’s what they *would* be, if there were any. Sadly, there are no unicorns; that is, the term “unicorn” fails to apply to anything. (That’s why I speak of what unicorns *would* be, *if* there were any.) There not being any doesn’t convert them into concepts. The fact that there’s nothing that the term “unicorn” correctly applies to

does not force the word to apply to an idea in someone's mind. An idea in someone's mind isn't a horse with a horn on its head.

My guess is that when people say things like this, they are again confusing ideas with the things the ideas purportedly represent, and in addition, they are assuming that all meaningful terms or ideas must in fact apply to something. Then, since we can't find anything in the physical world that "unicorn" applies to, they infer that it must apply to some mental thing.

3. Critical Thinking, 1: Intellectual Virtue

This chapter is about how we should think – especially when we’re forming philosophical beliefs, but also more generally. The short answer to that is that we should think *rationally* and *objectively*.

I’m going to try to explain what rationality and objectivity are, why they’re important, and how you can be more rational and objective.

3.1. Rationality

3.1.1. What Is Rationality?

What is it to think rationally? It is, in a certain sense, to think *correctly*. It is, for example, to accept the conclusions that one has good reason to accept, to avoid contradicting oneself, to avoid invalid deductive inferences or uncogent non-deductive inferences, to avoid reasoning circularly, and so on.

But now I have to clarify “thinking correctly”. I don’t mean necessarily having *true* beliefs (that is one sense of “to think correctly”, but not the one I meant) – it is possible to be fully rational but mistaken, and also possible to be highly irrational yet luckily arrive at the truth. Example:

Lucky is a poor college student who dreams of being a millionaire. As he dreams of all the things he could do with a million dollars, he gradually manages to convince himself that he *will* become a millionaire by the time he graduates from college. Lucky has no idea how this will happen and no plans to make it happen. He works at a local Taco Bell and can barely afford to pay his monthly rent. Nevertheless, he tells himself that somehow his fortune will soon be made, since he wants it so much. Lucky’s roommate, Thomas, tells Lucky that he is *not* going to become a millionaire, certainly not by the end of college. But a few months later, as chance would have it, Lucky finds a lottery ticket lying on the street. He picks it up and, on a whim, checks the numbers. It turns out to be a winning ticket, worth \$10 million. Lucky claims the prize and becomes a multi-millionaire.

In this example, Lucky forms an irrational belief – it is *irrational*, *unreasonable*, *unjustified*, and generally *doesn’t make sense* (these are all roughly synonymous) for Lucky to think he is going to become a

millionaire. Thomas is the rational one. Yet it turns out, improbably enough, that Lucky gets the money, so Lucky's belief is true and Thomas' belief false. This shows that one can think rationally yet be wrong, and one can think irrationally yet be right.

This *can* happen, but of course, it *usually* does not. One way of characterizing rational beliefs is to say that they are the beliefs that are *likely* to be true, given the experiences and information available to you at the time.^[8] That is what I mean by “thinking correctly”. At the time Lucky formed his belief, given everything he had to go on, it was very unlikely that he was going to become a millionaire. By contrast, given everything Thomas had to go on, it was very likely that Thomas' belief would be correct. Hence we call Lucky's belief irrational (unjustified, etc.) and Thomas' belief rational (justified, etc.).

Notice, by the way, that what is rational to believe is, as we say, “relative to an observer”: In other words, believing *p* could be rational for one person but not for another. This is because people might have different experiences and information available to them. If I somehow know about the lottery ticket that Lucky is about to find (maybe I planted it in his path), then it would be rational for *me* to believe that Lucky will soon become a millionaire, even though this belief is irrational for *Lucky*.

Philosophers have more complicated theories about rationality, and there are many more issues about rationality to explore. But the above is enough for us. The first thing I'm trying to do here is to get you to think about rationality, so that you will reflect on whether you're being rational and make an effort to be more rational, especially when you're forming philosophical beliefs.

3.1.2. *Why Be Rational?*

People sometimes wonder: Why should we be rational? Philosophers don't usually wonder this, but students first thinking about philosophy do.

Given my account of rationality above, this question is not really sensible. It is like the question, “Why are cats so feline?” or “What color was George Washington's white horse?”: If you understand what the question means, you don't need to ask it. To ask why one should be rational is to request a *reason* for being rational. But this is, in effect, to ask for a reason for responding to reasons – for example, a reason for believing what one has most reason to believe. If *p* is the rational thing for you to believe, then, by definition, it makes sense for you to believe that – because that's

what we *meant* by calling it “rational”. Once you have granted that *p* is rational, you’ve already granted that you have sufficient reason to believe *p*, so you don’t need an additional reason “to believe the rational things”.

Wait, that was a bit fast. Part of what I said above was that rationality is about thinking correctly, about accepting what one has most reason to accept, and stuff like that. The other part was that rational beliefs are *likely to be true*, given your information and experiences. So when someone asks, “Why be rational?”, perhaps they are asking: Why aim at having true beliefs? Let’s discuss this further in the next two subsections.

3.1.3. *Truth Is Good for You*

Truth is good for you. More precisely, knowing the truth is generally good for attaining your goals. For whatever goals you have in your life, it is almost always useful to have true beliefs. If you want to become rich, then you need to have a host of correct beliefs about financial matters, such as how much that job you are considering pays, how much an apartment you are going to rent costs, what is the expected yield on a stock market investment, and so on. If you want to become famous, you need to have correct beliefs about other things, such as what sorts of activities get people on TV and what things other people like to talk about. If you even have some very mundane goal such as getting a burrito, you need to have correct beliefs about such things as where burritos are, how store owners react when you try to take their burritos, how much they cost, and so on. It is *possible* to attain your goals while laboring under misconceptions, but it becomes drastically less likely as your misconceptions multiply.

What about philosophical beliefs in particular? They are important because they tend to be fundamental and have dramatic implications. For an obvious example, it would be good to form correct beliefs about whether there is a God, since if there is, you may need to follow that God’s commands in order to secure your place in the afterlife; if there isn’t, you may want to devote yourself to something else. Note that merely believing what you *want* to believe is not adequate: If you believe that there is a God, and there isn’t actually one, then you are likely to waste your life trying to please the nonexistent being and wind up never getting the reward you hoped for. Conversely, if you believe there is *no* God, and there *is* one, then you may be in big trouble when you finally meet Him. In both of these cases, it is bad to have the false belief, even if it is what you wanted to believe.

Ordinary errors cause you to make ordinary, small mistakes. E.g., being wrong about what stores sell burritos causes you to waste time and not get your burrito. Philosophical errors, on the other hand, cause you to make bigger mistakes, like wasting your life. That can happen because you form incorrect beliefs about what is worth doing in general.

This brings us close to our next point about why we should think rationally.

3.1.4. Irrationality Is Immoral

Some philosophers argue that it is *morally wrong* to adopt irrational beliefs. This view is sometimes labelled “**evidentialism**”.^[9] Why might this be true?

The first thing to note is that *if* one forms an irrational belief, and this has bad consequences, then one will be responsible for (one will be *to blame* for) those consequences. But if one has a fully rational belief, and it unluckily has bad consequences, one is not to blame for those consequences. Here is an example to illustrate; I call it the case of the Unlucky Samaritan (a “good Samaritan” is someone who goes out of their way to help someone they see in need):

Unlucky Samaritan: You have just witnessed a traffic accident in which a car has crashed into a tree. You approach the crashed car, whose driver is unconscious in the driver’s seat. You believe that the car is soon going to catch on fire, which will make it impossible to get the driver out and will likely result in the driver’s dying in the car. To prevent this, you unhook her seatbelt and pull the driver to safety. Later, it turns out that the driver had a spinal injury, which you exacerbated by moving her. Also, the car never did catch on fire; you were mistaken about that.

In this case, you acted out of a desire to save the driver. You did not know about the spinal injury, and you *thought* the car was going to catch on fire. *If* the car had been about to catch on fire, then pulling the driver out would have been the right thing to do. But *in fact* there was no fire, and the driver would have been better off if you had left her alone and waited for the paramedics. So are you blameworthy for what you did?

The answer is, “It depends.” It depends on *why* you thought the car was going to catch on fire. Consider two ways of filling in the story:

1. Suppose that you had no good reason to think the car was going to catch on fire. You just made that up because you’ve seen

lots of cool car explosions in movies, and you thought it would be awesome if you could save someone from one of those explosions just like the movie heroes. So you convinced yourself that this car was about to catch on fire and then explode. Furthermore, you have previously been told about spinal injuries and how one should avoid moving accident victims, but you just pushed that thought out of your mind as soon as it occurred to you. In this case, you are definitely to blame.

2. Now suppose, instead, that you had a very good reason to think the car was going to catch on fire. You could see gasoline leaking from the gas tank, and a pool of gasoline was moving underneath the car, toward the engine, which is very hot. Also, in this version of the story, you've never heard anything about why one shouldn't move accident victims. As it turned out, though, the car, luckily enough, never caught on fire. In this case, you are *not* to blame.

I'm assuming that you will agree with me about cases (a) and (b) – that you're to blame in (a) but not in (b). (If not, then the whole line of reasoning I'm doing here won't work on you.) What do we conclude from this? In cases (a) and (b), your action has the same consequences, in actual fact. Also, in both cases, your action makes sense *given* your beliefs; it is the right thing to do *given* that you think the car is going to catch on fire. So what's the difference between (a) and (b)? It looks like the answer is: In (a) your belief that the car will catch fire is *unjustified*, whereas in (b) your belief that the car will catch fire is *justified*. (What else could explain the difference?)

If that's right, then it looks like the lesson is: Forming irrational beliefs makes you *morally to blame* if you act on those beliefs and there is a bad outcome; by contrast, forming *rational* beliefs insulates you from that kind of blame. If you think rationally, and you do the thing that is right according to your rationally formed beliefs, then you are not morally to blame if things go wrong.

But what if you form irrational beliefs, and nothing bad happens? We might consider a third variant of the car accident story:

3. Things are just as in version (a) above, except that the accident victim has no spinal injury, so no harm is caused by your pulling her from the car. However, you had no way of knowing this at the time.

In this case, there is no actual harm to blame you for. Nevertheless, plausibly, your action is still immoral. It is immoral because, at the time you decided to pull the driver from the car, there was a significant *risk* that you would be harming the driver (that is, you had reason to suspect that you might be harming the driver, and you didn't do anything to rule that out).^[10] The degree to which an action is blameworthy should be determined by what was true of the agent (e.g., the agent's motives and what the agent had reason to believe) at the time of the action. It shouldn't be determined by *luck*, and the only difference between cases (a) and (c) is that the agent in (c) is *lucky*. Therefore, the action in (c) is equally blameworthy as the action in (a).

All this is supposed to establish that, if you form an irrational belief, and there is a *risk* of its causing something bad (given the information and experiences available to you at the time), then you are morally blameworthy.

So far, this doesn't show that you should *always* form rational beliefs. Maybe you only need to be rational in cases where there is a risk of causing bad outcomes. At this point, the most extreme evidentialists would argue that there is *always* a risk of causing bad outcomes. To make this plausible, notice that beliefs interact with each other in lots of complicated ways, which you can't really anticipate in advance. If you form one irrational belief, and you really believe that thing, then you will start inferring other conclusions from it. And then infer further things from those conclusions, and so on. For this reason, almost any belief can have practical consequences down the line.

Irrational beliefs can also have an impact on your belief-forming *methods*, causing you to adopt less rational methods of forming beliefs in the future. For instance, suppose you accept, purely on blind faith, that there is a God. This might lead to your adopting the more general belief that *blind faith is an acceptable way of forming beliefs*. (If you *don't* adopt that belief, it's going to be hard to explain to yourself why your blind-faith belief in God is acceptable.) But once you accept that, you are liable to form all

kinds of false beliefs, because there are so many false beliefs that could be adopted by blind faith.

Finally, notice that you are not in a good position to decide which beliefs do or do not carry a risk of causing bad outcomes, unless you are thinking rationally in the first place. That is, you have to start from some rational background beliefs in order to reason about what beliefs are likely to cause harm.

All this is leading to the conclusion that one is morally obligated to always form beliefs rationally. Here is a summary statement of the argument:

- 1.If one forms an irrational belief that causes a bad outcome, one is morally blameworthy.
- 2.Moral blameworthiness is determined by the risks assessable at the time of one's decision, not by what actually happens. Explanation for this:
 - a.Moral blameworthiness can't be a matter of luck.
 - b.If blameworthiness were determined by actual outcomes, rather than risks at the time of decision-making, then it would be a matter of luck.
- 3.Therefore, if one forms an irrational belief that has a *risk* of causing a bad outcome, one is morally blameworthy.
- 4.Irrational beliefs always have a risk of causing bad outcomes. Explanation:
 - a.They lead to further unreliable beliefs that can't be anticipated.
 - b.They risk worsening our general methods of forming beliefs.
- 5.Therefore, forming an irrational belief is *always* morally blameworthy.

(Quick exercise for the reader: What are the premises and conclusion there? What kind of argument is it, deductive or non-deductive? Is it valid or invalid? Cogent or uncogent? Sound or unsound? Circular or non-circular?^[11])

Now I'm going to comment on what I think of that argument. It might have sounded like I was endorsing it above, but don't assume that. I was just telling you what a strict evidentialist would say.

I think most of the argument is fine. The biggest gap I see is step 4. It does not seem to be true that irrational beliefs *always* have some non-negligible risk of causing bad outcomes. It's plausible that *as a general rule*

they carry such a risk, because of points 4a and 4b, but it's not as if there is some law of the universe that forces this to *always* be true. A person might have enough reliable, rational beliefs to know that some particular belief is relatively isolated, i.e., that he's not going to base any practical choices on it and he isn't going to infer a bunch of other beliefs in other areas from it.

So I don't think the argument succeeds in proving conclusion 5. However, I still think it's reasonable to conclude that forming irrational beliefs is *usually* morally blameworthy. Furthermore, I think it is likely to be blameworthy in *philosophy* in particular, which is what we're concerned with at the moment. The reason is that a person's philosophical beliefs are especially likely to have far-reaching implications – e.g., whether one believes in free will, or God, or objective values, or the authority of government, has far-reaching implications for one's belief system as a whole. So it is especially important to think rationally about that sort of thing.

3.1.5. *Some Misunderstandings*

When I say that irrationality is immoral, here are some things I am *not* saying. I'm not saying that everyone needs to be a completely unemotional (or even relatively unemotional), robotlike being. A rational thinker is not a person with no emotional reactions. It is, however, a person who strives (reasonably effectively) to base his beliefs on objective evidence, rather than on his emotions. Having feelings does not make you irrational. Believing that the world must be a certain way because of your feelings does.

I also am not saying that you are a bad person if you sometimes form biased beliefs. Nearly all human beings, perhaps all, sometimes form biased beliefs. This does not make all of us overall bad people. It obviously makes us imperfect in one respect. I am not, however, saying (nor do I think) that we are morally obligated to be perfect in that sense. What I think we are obligated to do (and what a good person does) is to make a reasonable effort to minimize the impact of bias on our beliefs.

3.2. Objectivity

3.2.1. *The Virtue of Objectivity*

Intellectual virtues are character traits that help us form beliefs well, particularly traits that tend to help you get to the truth and avoid error in normal circumstances. Rationality is the master intellectual virtue, the one that subsumes all the others. (So if you are rational, you are intellectually

virtuous, and vice versa.) But there is another intellectual virtue that is also extremely important, so much so that it also deserves a section of its own in this chapter. The virtue is *objectivity*.

Objectivity, like all other intellectual virtues, is part of rationality. The character trait of objectivity is a disposition to resist bias, and hence to base one's beliefs on the objective facts. The main *failures* of objectivity are cases where your beliefs are overly influenced by your personal interests, emotions, or desires, or by how the phenomenon in the world is related to *you*, as opposed to how the external world is independent of you.

For instance, when people hear about (alleged) bad behavior by politicians, their reactions are strongly influenced by which party the politician involved belongs to.^[12] When a politician from one's own party is accused of lying, or being inconsistent, or having sexual indiscretions, we tend to minimize the accusations. We might insist on a very high standard of proof, or try to think of excuses for the politician, or say that sexual indiscretions are not relevant to the job. But when a politician from another party is accused of similar bad behavior, we tend to just accept the accusations at face value, and perhaps even trumpet them as proof of how bad the other party is. That is a failure of objectivity: The way we evaluate the cases is determined, not by the relevant facts of the case (what the person did, what the evidence shows), but by whether we think of the politician involved as "on our side" or not.

When we are being biased (non-objective), we usually do not notice that we are doing this, nor do we actively decide to do it. It just happens automatically – e.g., we automatically, without even trying, find ourselves thinking of reasons why the person who is "on our side" shouldn't be blamed for what he did. On the other hand, when a person "on the other side" is accused of wrongdoing, no such excuses occur to us. This is to say that bias is usually unconscious (or only semi-conscious), and unintentional (or only semi-intentional). That is why it requires deliberate monitoring and effort to attain objectivity. You have to stop once in a while to *ask yourself* how you might be biased. If you don't, the bias will automatically happen.

Here is another example. You hear an argument about abortion. It's an argument against your position, whatever that is – if you're pro-abortion, imagine it's an anti-abortion argument; if you're anti-abortion, imagine it's a pro-abortion argument. If you have an opinion on this issue, you probably also have fairly strong feelings about it. Let's assume that's the case. When

you hear the argument against your view, you have a negative emotional reaction. Maybe it makes you angry. Maybe you feel dislike toward the person giving the argument. And those emotions affect your evaluation of the argument. They make you psychologically *averse* to thinking about the argument from your “opponent’s” point of view. (I am putting quotes around “opponent” because we often think of those we are arguing with as opponents, but we really shouldn’t do so. We *should* think of them as fellow truth-seekers.) You do not *want* to see the other person’s perspective, so you don’t.

What do you do instead? You might misinterpret the argument – in particular, you might interpret your “opponent” as saying the stupidest thing that they could possibly be interpreted as saying, and then respond to that. You might impose an impossibly high standard of proof for every premise used in the argument, using any doubt about any premise as an excuse for completely disregarding the argument. You might devote all your effort to thinking of ways your “opponent” could be wrong, while devoting none to thinking of ways that you yourself could be wrong.

Again, these are failures of objectivity: You let your treatment of ideas and arguments be determined by your personal feelings, rather than strictly by the rational merits of the ideas and arguments.

3.2.2. *Objectivity vs. Neutrality*

Objectivity is not to be confused with *neutrality*. Similarly, being partisan is not to be confused with being biased.

“**Neutrality**” is a matter of *not taking a stand* on a controversial issue. The neutral person may hold that both sides are equally good, or equally likely to be correct; or he may simply refuse to get involved in evaluating the issue. That is not what I’m talking about when I promote objectivity. I am *not* recommending that you refuse to take a side on controversial issues. It is generally false that both sides are equally good, and you should not refuse to evaluate issues.

What I am recommending is that, if you take a side, you nevertheless *treat the other side fairly*, even while defending your side. I am recommending that you treat intellectual debate as a mutual truth-seeking enterprise, rather than as a personal contest. This is an idea of crucial import, so it’s worth repeating. You can and should treat the other side fairly, even though you think they are wrong.

For example, when responding to opposing views, you should respond to the *most plausible* opposing views and address the *strongest arguments* for those views – that is, the views and arguments that have the greatest chance of being correct while being importantly different from your own view. When you explain what your “opponents” think, try to state their views in the way that they themselves would state them. If there is any ambiguity in your “opponents’” statements, choose the most reasonable interpretation of their words. Acknowledge the evidence that genuinely supports their side, and do not exaggerate the evidence for your side. All this is being objective.

Now, you might wonder: “If I do that, then how am I going to win debates?” If you have this concern, you’re thinking incorrectly about intellectual discussion. The purpose of intellectual discussion is promoting truth (for yourself and others). If your view can’t survive when you treat the opposing views fairly, then that pretty much means your view is wrong. As a rational thinker, you want your beliefs to be true, so you should welcome the opportunity to discover if your own current view is wrong; then you can eliminate a mistaken belief and move closer to the truth. If you are afraid to confront the strongest opposing views, represented in the fairest way possible, that means that you suspect that your own beliefs are not up to the challenge, which means you already suspect that your beliefs are false.

Nobody learns much from discussions in which two people unfairly caricature each other’s views, distort the evidence, and try to paper over the problems in their views. When two people with opposing beliefs argue things out while treating each other fairly and objectively, that is when people learn. If you’re talking to another person one-on-one, you will likely learn from each other and reach a more satisfying understanding, even if you don’t actually resolve the central disagreement. If you are having a public discussion (as in an internet discussion forum), the audience is also likely to be educated. Since you want to learn and promote others’ learning, you should try to have the kind of discussion filled with fair, objective treatments of one another, not the kind filled with distortions and evasions.

3.2.3. The Importance of Objectivity

Why is objectivity important? Because failures of objectivity are very common, and they often lead us very far astray. The main thing human beings need, to make progress on debates in philosophy (and religion, and politics), is more objectivity.

The human mind is not really designed for discovering abstract, philosophical truths. Our natural tendency is to try to advance our own interests or the interests of the group we identify with, and we tend to treat intellectual issues as a proxy battleground for that endeavor. Again, we don't expressly decide to do this; we do it *automatically* unless we are making a concerted, conscious effort *not* to. And naturally, when we do this, we form all sorts of false beliefs, because reality does not adjust itself to whatever is convenient for our particular social faction.

As I suggested above, one reason for treating opposing views fairly is that you yourself might be wrong – particularly if you are afraid to treat opposing views fairly. Another reason is that even if your central view is correct, you can often learn something from people with opposing views. It is very rare that some view (held by intelligent people – and I don't suggest having debates with stupid people, so I'll assume that we're talking about views held by intelligent people) has absolutely nothing to it, captures no important facts, responds to no relevant aspect of reality. Probably if someone (reasonably smart) disagrees with you, they know some relevant information that explains why they have their opposing view. Taking account of that information is likely to make your own view more sophisticated and accurate. At the very least, you can better understand how other people think.

Finally, by treating opposing views fairly, you are more likely to be persuasive. If you are arguing with another person, and you distort their views, or respond to only the weakest arguments for their views, then they won't be persuaded! To be persuaded, the other person has to feel that you understood what they were trying to say, and that you rebutted the *strongest* reasons that support their view. By the way, when you talk to other people about philosophy (or other abstract matters), the other people often do a less-than-ideal job of explaining their own view (often, they are still confused about what they want to say), so sometimes, you actually have to work to give their view a *better* presentation than they themselves gave it, before you rebut it.

You might have the chance to “score points” cheaply by attacking some misstatement of your “opponent”. But it is much better to actually persuade people than to score points.

3.2.4. *Attacks on Objectivity*

In modern intellectual life, you often actually hear attacks on the ideal of objectivity. This has been true for as long as I remember being in the world of philosophy (almost three decades). So I should probably say something about this.

One thing you might hear is that objectivity is impossible; everyone is biased. And you might be tempted to think that, if that's true, then there is no point to aiming at objectivity. It's senseless to aim at the impossible! But what is meant by claiming that objectivity is impossible?

Interpretation #1: "It is impossible for everyone to be 100% objective all the time."

This might be true. But this doesn't mean that we shouldn't promote objectivity. Falling short of perfection does not mean that one should not strive to be better. There are degrees of objectivity, and one can increase one's objectivity through effort.

Interpretation #2: "It is impossible for anyone to be at all objective, ever. We are all doomed to be maximally biased."

This is obviously false. People sometimes respond to the facts. Example: When President Nixon was first being investigated over the Watergate scandal, Republicans generally sided with Nixon, while Democrats sided against him, because Nixon was a Republican. So there was bias. But when the tapes came out with Nixon talking about bribing the Watergate burglars to keep quiet, pretty much everyone agreed that he was guilty. It's not as if Republicans are still defending Nixon today.

So people are not perfectly objective, but we are not perfectly biased either. We are a mixture. And that is exactly why we need to work at objectivity.

Here is a second thing you might hear: You might hear that the ideal of objectivity is a patriarchal, oppressive, Western value, or something like that. What does this mean? Sometimes, the concern seems to be that the ideal of objectivity fails to support the allegedly correct political conclusions (in particular, that certain left-wing ideologies won't thrive if we insist that people be objective). The problem here is that, if objective thinking leads people to reject your ideology, then probably your ideology is false. If you think your ideology would not survive objective examination, then you yourself probably already suspect that your ideology is false, without wanting to admit this. In that case, you should just admit

you're wrong and move on. Of course, anyone can hold onto their ideology (whichever ideology they have), simply by being sufficiently biased.

Here is a more reasonable concern. Often, the factors that make someone *biased* about a topic are also the factors that make them *knowledgeable* about it. Examples: Suppose your company is hiring a new employee, and one of the candidates is a friend of yours whom you have known for ten years. You would probably be more *knowledgeable* about the candidate than anyone else at your company, while at the same time being the most *biased*. Or suppose you are involved in a discussion about war, and you are a veteran of a past war. Again, you would probably be the most knowledgeable person present about what wars are like; but you would also likely have the most biases, because the experiences that gave you that knowledge also gave you strong feelings. Or, finally, suppose you are in a discussion about racism, and you are a member of a minority race. Then you are likely to be especially knowledgeable about what it is like to be a member of a minority, including how often such minorities experience discrimination. But the same experiences that gave you that knowledge are likely to have given you personal, emotional biases on the subject.

The people who criticize the ideal of objectivity are usually thinking of examples like the race example, but I have given those three examples to show that there are a variety of cases involving very different issues. Lesson: If we discount "non-objective" perspectives, that could mean throwing out the perspectives of the most knowledgeable people.

In response, we should start by acknowledging that the concern about objectivity rests on a genuine fact: It is true that the people with the most biases are also often the most knowledgeable. But the concern also rests on a misunderstanding of objectivity. Objectivity does not require that one *disregard* the testimony of anyone who is biased. Refusing to listen is not being objective. An objective person would listen to all relevant evidence, and try to weigh the various pieces of evidence fairly. In this process, one would take account of the possible biases of one's information sources, as well as how knowledgeable they are – and there really is no rational case for not doing that. It can't be denied that emotions can bias people, and that bias can lead to false conclusions. For instance, when you tell your company's manager that your friend of ten years is the ideal job candidate, the manager, even while acknowledging that you know the candidate far better than the manager does, obviously has to take account of the fact that

you might be biased because the candidate is your personal friend. That doesn't mean that he should ignore everything you have to say. But it probably means that he should give less weight to subjective judgments that you make (e.g., when you say the person is "great" or "likeable") than he would if you were talking about someone you were not already friends with.

For those with a left-leaning political perspective, it is worth pointing out that many of the problems liberals have fought against have been precisely failures of objectivity. For instance, traditional racism consists of privileging *one's own* race and discounting the interests and perspectives of people of other races – which is a paradigm failure of objectivity. Similar points apply to sexism, heterosexism, and other forms of prejudice. All the paradigm forms of prejudice are, centrally and obviously, failures to be objective.

If you hear someone attacking the ideals of objectivity or rationality, how should you react? First, I would suggest that if a person attacks rationality/objectivity, this is evidence that some key point of their ideology is false, and that they themselves know or suspect that. (Alternately, it could be that they don't understand what rationality and objectivity are.) If you were initially sympathetic to their views, you should greatly lower your confidence in those views, and in that person.

Here is an analogy: During the Watergate scandal, after investigators learned that President Nixon had taped all of his conversations in the White House, the investigators ordered Nixon to hand over the tapes, so they could see if Nixon had illegally conspired with the Watergate burglars. Some Nixon supporters were *happy* to learn of the tapes and were eager for Nixon to turn them over, because they assumed that the tapes would vindicate Nixon and the scandal would end. Nixon, however, fought tooth and nail against turning over the tapes. And that was when many people realized that he was guilty of something serious. If he were innocent, the tapes would vindicate him. The best explanation of his refusing to turn them over was that he knew the tapes would prove his guilt. (Which, of course, is what ultimately happened.)

Similarly, if your philosophical views are correct, then you should welcome an objective examination. The best explanation for someone's rejecting objectivity or rationality in philosophy is that, on some level, the person knows that an objective, rational examination would show his own

views to be false. (How a person can “hold views” that he knows to be false is an interesting question. But self-deception appears to be common in human beings.) If you can only maintain your beliefs by being biased or irrational, then your beliefs are almost certainly wrong.

A final point. Attacking rationality or objectivity is a short-sighted stratagem. If you manage to convince anyone to give up the ideals of rationality and objectivity, that does not mean that they will automatically come over to your side and support whatever *you* want. Irrationality and bias can support any ideology, including your opponents’. Nazis, Marxists, flat-Earthers, and partisans of any other crazy or evil view can base their beliefs on irrational biases, and there is no way to reason them out of it if you’ve rejected rationality and objectivity. So don’t attack objectivity and rationality. Unless you’re an asshole and you just want intellectual chaos.

3.2.5. How to Be Objective

How can we work to be more objective? There are three main steps that I recommend.

i. Identify your biases.

Just being aware of a bias makes that bias less influential. For instance, if you are an educator, and you believe that educators should be paid more money, acknowledge the fact that you could be biased because of your self-interest. If you are thinking about a controversial issue, and it makes you emotional, acknowledge the fact that your emotions could be clouding your judgment. E.g., if you feel angry when you think about abortion, then you might not be able to reason rationally and objectively about abortion, in which case, you probably should not be extremely confident that your opinions are correct.

By the way, in saying this, I am not saying that your emotions are inappropriate, nor that you should suppress them, nor that your political views are mistaken. I am only saying what I actually said: That your emotions could be clouding your judgment. It’s no defense to say, “But it’s appropriate to be emotional when babies are being murdered!” That’s just a completely different issue. It might well be appropriate, but that wouldn’t stop the emotions from clouding your judgment!

ii. Diversify your information sources.

When you learn about an issue, do not just learn from people you agree with. Gather information and ideas from people on different sides. For example, if you want to learn about gun control, collect information from both pro- and anti-gun sources.

Also, by the way, collect information from the *most sophisticated* sources, not (as most people do) the most *entertaining* sources. That usually means looking at academic sources, rather than popular media. (This is just a general point about reliability, not specifically about objectivity.)

iii. Challenge yourself.

When you think about a controversial issue, do not just try to think of reasons why other people are wrong. Try to think of reasons why your own views might be wrong. When you give an argument, ask yourself at each major step, “Is there anything that could be wrong with this step?” Spend some time trying to think of evidence against your own conclusions.

Note: If you hold a controversial view, and you haven’t thought of any objections to it, think more. If the only objections you can think of are really stupid, think more. Because the best explanation for your not knowing of any good objections to your view is not that you’re completely correct; the best explanation is that you’re too blinded by bias to have seen the problems with your view.

Note, again, that I am not calling for neutrality. I am not, for example, saying that all views are equally good. I am saying that if you hold a view that is *controversial*, then, *most* of the time, there is *some* reasonably strong evidence against the view – otherwise, there probably wouldn’t be disagreement about it. You might ultimately conclude that that evidence is misleading or simply outweighed by the evidence *for* the view. But if you can’t even think of what the evidence could be, that probably just means that you don’t know enough.

Example: Think about the abortion issue. If pro-lifers (or pro-choicers) make you angry, then you might be biased. If you personally have had an abortion, then you might be biased. Lastly, if you can’t think of any rational reason why anyone would think abortion was wrong, or you can’t think of any rational reason why anyone would think it *wasn’t* wrong, then you’re definitely biased. In that case, you should withhold judgment on that issue until you understand the rational reasons on the other side, if you ever do.

3.2.6. Open-mindedness vs. Dogmatism

Open-mindedness is the opposite of dogmatism (also called “closed-mindedness” – and please do not write “close-minded”; the opposite of “open” is “closed”, not “close”). Dogmatism is probably the most common kind of failure of objectivity. Dogmatic people have beliefs that are overly persistent and insufficiently receptive to disconfirmation. When given strong reasons for doubting their opinions, they don’t doubt; they confidently cling to those opinions.

From casual observation, it appears that the vast majority of people have this trait to some degree. It’s not *necessary* that people have this trait to any degree – you could imagine a person who, on the contrary, is too *quick* to abandon beliefs, so that they change their beliefs when given very slight reasons for doubt. But when we look around, it’s virtually impossible to find any people like that; so much so that we don’t even have a word for that vice. Most people err on the opposite side, while a few seem to be about right in their receptiveness to belief revision. I don’t know why this is.

Yet while the vast majority of people are dogmatic, no one *thinks* that they are. You, reader, are probably dogmatic, but you think you’re not. That’s partly because the word “dogmatic” sounds insulting, and hence it is unpleasant to entertain the hypothesis that one is dogmatic. To make it sound less bad, you can just replace it with the description, “systematically underestimates appropriate belief revision”. You probably systematically underestimate how much you should revise your beliefs when you acquire new information, because the vast majority of people do that, but you probably don’t realize that you do this.

The best counter to dogmatism is reflection: When reasoning about a controversial issue, ask yourself whether you are applying the same standards to “your side” as you do to the other side, or whether you are instead applying much stricter scrutiny to the other side. Ask yourself what, if anything, you would accept as proving yourself to be wrong. Collect evidence and arguments from the other side. Spend time thinking about what might be wrong with your own ideas and arguments, and how the other side might respond to your objections. Making an effort to be less dogmatic makes a difference.

3.3. Being a Good Philosophical Discussant

Here I'm going to give some general advice about how to be good at philosophical discussion. Some of this is redundant with the previous sections, but it won't hurt to repeat a little. These are points that are especially important in *philosophical* discussion. They apply to in-class discussion, one-on-one verbal discussion, discussion in online forums, etc. The person you are talking to/arguing with is called your "interlocutor".

3.3.1. *Be Cooperative*

First, a general principle of discussion: *Be cooperative*. In other words, when discussing philosophy, your aim is to make progress in the discussion, not to cause delays, "score points", prevent other people in the discussion from making their points, or sow chaos. This implies a number of more specific things (many of these points are overlapping):

i. Accept the hypo.

If someone gives a hypothetical example, do not raise objections to the example that will only make your interlocutor waste time thinking of other examples, or thinking of a series of increasingly elaborate modifications to the example. Do not start a debate about how realistic the example is or "what would really happen" in a situation of that kind; just accept the example as the other person intended it.

Example: Someone gives you a hypothetical example in which you have to choose between letting a train run over five people, and diverting the train to another track where it will only run over one person. They want to discuss what you should do in such a case. Do not start arguing about how realistic this scenario is, whether you could instead derail the train or move some of the people off the tracks, etc. Just accept that the scenario is as stated, and the available options are as stated.

The reason for this is that the example was given to illustrate some underlying philosophical issue, and you need to *address the central issue* that your interlocutor is trying to raise. If you start fighting the example, it diverts the discussion away from that issue and into irrelevant debates about the details of the particular example.

ii. Don't change the example.

Also, when someone gives an example, do not propose adding conditions to the example that would make it irrelevant to what you were talking about, or that turn it into an illustration of a

completely different issue. Do not “interpret” the example in ways that make it irrelevant to the topic of discussion.

For instance, in the above discussion, do not respond to the train example by saying, “What if one of the people on the track is baby Hitler?” Or, “What if the five people on the left are all really old and are about to die anyway, but the one on the right is a baby?” You should not say these things, because they are changing the example from what it was intended to be into an illustration of completely different issues.

iii. Don't raise extraneous controversies.

Modifying other people's hypothetical examples is one way of raising extraneous issues. Another way is just making controversial statements needlessly. For instance, in the discussion of the train example, you might just throw in your opinion that the current President is an asshole. Another way is unnecessarily asking broader questions. For instance, in a discussion of the train example, you say, “Well, what are ‘right’ and ‘wrong’, really?”, inviting general debate about a much larger issue.

You should not do that, because that prevents you from making progress on the original topic of discussion. No progress can be made if you keep changing the subject. Especially if you change it to some huge other controversy.

Similarly, when you give examples, do not give examples that presuppose opinions that you have about something else that the other person might not agree with. For instance, if you want to give an example of an immoral action, pick something like “murder”, not something like “voting Republican”.

iv. Be charitable.

If anything is ambiguous in what the other person is saying, do not search for the stupidest thing that they might be saying. Search for the most reasonable thing they might be saying. This is known as being charitable. Do not ascribe to the other person much stronger claims than necessary. If you are not sure what the other person is saying, ask them. If it sounds as if a person is saying something ridiculous, try asking them if that is what they meant, before going ahead and assuming that it is.

For instance, suppose you're debating affirmative action. The other person says, "I think there are fewer women mathematicians because women aren't as good at math as men are." Is the person saying *no* woman is as good as *any* man? That would be a very strong (and stupid) thesis, so if you assume that's what they are saying, then you're being uncharitable. Before assuming this, you could try asking them, "Do you mean that no woman is as good at math as any man?" (They will say, "No.") A more charitable interpretation would be that they meant, "The average mathematical ability of women is lower than the average for men."

v. *Don't quibble.*

This imperative requires skill and judgment to follow. Basically, I mean that you shouldn't make people spend time talking about objections that don't go to the heart of the matter, or objections that could be met by making minor modifications to what the other person said. Instead, you should spend time talking about *core* objections. Speak to the spirit of what the other person has said, not merely its letter.

vi. *Try to see the point.*

Do not just focus on making your own points, and don't try to stop the other person from making their point. Try to see the main point the other person is getting at, and let it be directly, centrally addressed. The above points (i)-(v) are really all parts of this.

3.3.2. *Be Modest*

Not everything that seems obvious to you is right. In fact, when it comes to abstract, philosophical questions, probably most of the thoughts that occur to you, even the ones that seem obviously right to you, are wrong. (The way I know this is that the things that seem obvious to people very often conflict with each other, so the percentage of things that seem obvious that are actually true must be pretty low.) Thinking well in philosophy requires being *much, much* more careful than people are naturally inclined to be. Here are some more specific suggestions:

i. Use weak, widely-shared premises.

A "strong" claim is one that says a lot; a "weak" claim says not very much. For instance, "All politicians are liars" is a strong claim; by contrast, "*Some* politicians are liars" is a much weaker claim. In general, the more controversial claims you make, and the stronger the

claims are, the more likely that your argument is wrong. So try to build arguments that use the weakest, least controversial premises possible. If you can argue for your desired conclusion using only the premise that *some* politicians are liars, do not go overboard and claim that *all* politicians are liars (even if you believe this). Don't claim more than you have to.

ii. Look for multiple supports.

Almost any argument that you make might be wrong. So, even if you find one argument convincing, still look for other arguments. A theory with multiple independent reasons supporting it is better than one that rests on a single reason.

iii. Limit the topic.

If the main thing you want to know about is X, don't try to address any other issues that don't need to be resolved in order to address X. Just focus on X.

3.3.3. *Understand Others' Points of View*

Misunderstandings are common in philosophical discourse. Sometimes they go on for a long time without anyone noticing. To avoid them:

i. Don't assume.

If something another person has said is ambiguous or unclear, don't assume that you have the right interpretation; usually, you don't. Ask for clarification.

Also, do not assume that other people are trying to *imply* something beyond what they actually *said*. If someone says, "Sue's argument is stronger than John's", do not assume that they agree with Sue's argument. They didn't say that. They just said that one argument is *stronger than* the other; that's compatible with both arguments being crappy. Similarly, if someone says, "I don't agree with that argument", do not assume that they disagree with the argument's *conclusion*. It could be that they accept the conclusion, but they just don't think that the particular argument gives a good *reason* for it.

ii. Don't imply.

This is the flip side of point (i): Don't count on your audience to understand what you're trying to imply. It is very common that they don't. State your view as explicitly as possible.

iii. Know when to use the same word.

In philosophy, we often need to rely on subtle distinctions, which we mark with slightly different words. (Example: Doing something “by accident” and doing something “by mistake” sound the same. But they are actually different.^[13] In some contexts, that difference would be important to an argument.) Because this happens a lot, you have to be careful about using words. If you use one word to talk about something, and then you shift to using a different word (that sounds to you like a synonym) to talk about that same thing, this can cause confusion – readers/listeners might think that you’re trying to make one of these subtle distinctions. Therefore, if you’re talking about one thing, keep using the same word for it unless you’re making a subtle distinction.

iv. Be charitable.

As discussed in section 3.3.2.

4. Critical Thinking, 2: Fallacies

4.1. Some Traditional Fallacies

A fallacy is a type of inference that superficially appears good (at least to some observers) but in fact is a mistake. “Fallacy” can also refer to a rhetorical trick that tends to mislead audiences. If you get a book on critical thinking (or “informal logic”, as it’s sometimes called), the book will have a list of fallacies (or alleged fallacies), which will include all or most of the following. I’m going to start by describing these in roughly the way they are traditionally described; in the next section, though, I’ll have objections to some of these descriptions. So take the following with a grain of salt for now.

Affirming the Consequent: The error of arguing, “If A then B. B. Therefore, A”. E.g., suppose you hear that if a person walks on the moon without a space suit, they die (which is true!). You also hear that Uncle Joe has recently died. You infer that Uncle Joe walked on the moon without a space suit. That’s fallacious!

Appeal to Authority: This is where you accept an idea because of good characteristics of the person advancing it, particularly that the person has some sort of expertise in some other area. Example: Thinking that we should get rid of nuclear weapons because Albert Einstein said so (Einstein was an expert in physics, but not on nuclear arms policy).

Argumentum ad Hominem (Latin for “argument against the man”): This is the mistake of rejecting an idea because of irrelevant bad characteristics of the person advancing the idea. E.g., suppose you reject Christianity because Christians are too preachy and annoying. This is fallacious since their annoying preachiness doesn’t show that their belief isn’t true. “Argument ad hominem” also covers the mistake of rejecting a theory because the person advancing it previously said or did something that conflicts with it (which would show that there is something defective about that person, but would not show that the theory isn’t true). *Note:* “Ad hominem” does not simply mean “insult”. Insulting a person is not

“committing an *ad hominem*”, unless that insult is used to draw some unwarranted conclusion. (It’s still jerky behavior, though!)

Argumentum ad Ignorantiam (appeal to ignorance): Concluding that something is the case merely because we don’t know anything to the contrary. Authors sometimes try to lure you into this mistake by writing things like, “There is no reason why X would be true” (hoping that you’ll infer that X isn’t true) or “There is no reason to doubt X” (hoping you’ll infer that X is true).

Argumentum ad Populum (appeal to the people): Inferring that something is true from the fact that it is popularly believed.

Attacking a Straw Man (a.k.a. “straw-manning” your opponent): Attacking a position that your interlocutor does not hold because it’s much easier to refute than their actual position. Usually, this mistake consists in attributing to someone a view that is more extreme, more simplistic, or otherwise just dumber than what they actually think. (People often attack straw men without realizing it.)

Begging the Question: Circular reasoning; reasoning in which one of the premises contains the conclusion, or presupposes the conclusion, or depends for its justification on the conclusion.

Complex Question: This is a question that contains an unstated presupposition, which makes the question unanswerable if one doesn’t accept that presupposition. E.g., “Have you stopped voting for degenerate bastards who want to ruin the country?”

Denying the Antecedent: The error of arguing, “If A then B. ~A. Therefore, ~B”. E.g., “If a person walks on the moon without a space suit, they die. Uncle Joe has *not* walked on the moon without a spacesuit. Therefore, Uncle Joe has *not* died.” That’s fallacious!

Emotional Appeals: Attempts to provoke emotions in the audience that will cause the audience to form beliefs based on those feelings. E.g., a lawyer might try to get his client acquitted by making the jury feel sorry for the client.

Equivocation: A type of argument in which a word or expression is used in two different senses, but they are treated as the same. Example: “All jackasses have long ears. Carl is a jackass. Therefore, Carl has long ears.”^[14]

False Analogy: An argument by analogy that’s no good, because the two things being compared are not really comparable. E.g., “The

government should be able to exclude foreigners, just as I can exclude strangers from my house.” The house might not be analogous to (not a fair comparison to) the whole country (perhaps because the government does not own the whole country in the same way an individual owns a house).

False Dilemma: This is where an interlocutor tries to make you choose between two alternatives, presupposing that these are the only alternatives, when in reality they are not. E.g., someone asks you, “Do you think abortion is murder, or do you think it’s a woman’s right to choose?” This is a false dilemma, since there are other possibilities (e.g., perhaps abortion is wrong, but not as bad as murder; perhaps it’s not wrong, but it’s still not one’s *right*; etc.)

Genetic Fallacy: Confusing a thing’s origins with its current characteristics. E.g., inferring that all governments are (currently) evil, because governments first originated in gangs of exploiters and conquerors.

Guilt by Association: The mistake of rejecting an idea because it is associated with some undesirable person or idea (but it doesn’t actually entail the bad idea). E.g., inferring that eugenics is bad because Adolf Hitler believed in it, and Hitler was terrible. Or arguing that drug prohibition is bad because some of the early drug laws were motivated by racism.

Hasty Generalization: Drawing a generalization from a small amount of evidence. E.g., concluding that all Canadians are rude because the first two you met were rude.

Non Sequitur (Latin for “it does not follow”): A sort of catch-all for cases in which an argument’s premises don’t at all support the conclusion, but it doesn’t fit under one of the other named fallacies.

Persuasive Definition: An attempt to make people accept your conclusion by building it into a “definition”. E.g., a socialist might try to define “capitalism” as “a system of oppression in which greedy businessmen exploit the poor.” The problem: Whether the system is exploitative or oppressive needs to be established by argument. A definition is supposed to explain the meaning of a word, not summarize one’s own personal opinions about the thing the word refers to.

Poisoning the Well: This is a rhetorical strategy of trying to undermine an interlocutor by warning the audience that he can't be trusted for some reason. This is supposed to make it impossible for the interlocutor to defend himself, since the audience won't listen to what he might say in his own defense. (Only works on naive audiences!)

Post Hoc Ergo Propter Hoc (Latin for "after this; therefore, because of this"): The mistake of assuming that because B follows A, A must cause B. E.g., many people die shortly after being rushed to the hospital. But it's not the case that being rushed to the hospital causes death.

Red Herring: Red herrings are issues that are irrelevant to the topic of conversation, or at least are not necessary to resolve, and that serve to distract people from the main issue. E.g., if you start out debating about the morality of abortion, you might get sidetracked into talking about the legality of abortion. Then someone might say, "I think the legal issue is a red herring."

Tu Quoque: Responding to a criticism by saying that your accuser is guilty of the same failing. Example: Sue tells Jack that he should stop eating meat. Jack responds by saying that Sue has bought some animal products. This is irrelevant, since the other person's being guilty of a failing doesn't show that you are innocent. Unfortunately, this tactic often succeeds in distracting people. (See also "ad hominem" and "red herring".)

Weak Manning: Attacking the weakest opponent you can find (or one of the weaker ones), rather than the strongest. This is not the same as straw-manning, because there really are people who hold the position you're attacking; it's just that they are among the least reasonable opponents to your view. This can mislead an audience (and yourself) into thinking that your own position is better supported than it is. The opposite of this is "**steel-manning**" – seeking out the strongest opponent to your view.

4.2. False Fallacies

So I've just produced another list of fallacies of the sort that you find in traditional critical thinking books and classes. But I don't much like these lists (not even my own). I have two reasons for disliking them.

First, I think they misdirect attention. They direct attention to some problems that occur rarely, while neglecting much more common errors. (Not *all* the traditional fallacies are rare, of course, but several of them are quite rare.) I'm not sure I've *ever* seen someone affirm the consequent or deny the antecedent. To the extent that the list identifies genuine errors, most of them are pretty dumb, so you probably don't need much discussion of them. For some more common and useful-to-discuss errors, see section 4.3 below.

Second, I think the fallacy lists lure people into thinking that some *perfectly good* inferences are wrong, because these perfectly good inferences sound like what the fallacy definitions are talking about. I refer to this as *the "fallacy" fallacy*: The fallacy of rejecting a good inference because it has been superficially labelled as a "fallacy". Let me explain with some examples.

Ad Hominem

Students who learn about the "ad hominem" fallacy are liable to draw the lesson that one should never reject an idea or argument because of who says it. But in fact, negative information about an individual is often very relevant to whether you should believe what they say.

Example: You see a television ad for "clean coal". The ad contains some evidence and arguments for the claim that your country should rely more on "clean coal" for its energy needs. Now, suppose you find out that this ad was produced by a coal company that would stand to profit if people accept the ad's message. The particular company in question is an especially immoral one that has been in trouble with the law on several occasions for safety and environmental violations.^[15] Now, how should you treat this information? Ignore it, because the bad traits of the company are irrelevant to the truth of its message?

That might be what you would think after reading about the "ad hominem fallacy" in your critical thinking book. But of course, that would be wrong. The bias and the immoral qualities of the company make it very likely that the ad is going to be misleading or outright wrong. If the ad-makers are any good at their job, you (without extensive expertise in the area) probably wouldn't be able to identify exactly how it is misleading. Therefore, you should apply a heavy skepticism to the ad and all of its content.

In this case, you would be rejecting ideas and arguments because of the immorality of the party putting them forward. This sounds like exactly what people are calling the “ad hominem fallacy”. But it’s not fallacious; it’s smart.

We could avoid this problem by just *defining* “ad hominem argument” so as to make it automatically fallacious – e.g., defining it as the mistake of rejecting an idea because of *irrelevant* negative information about the idea’s proponent. But then it must be said that ad hominem arguments are (i) rarer and (ii) harder to recognize than you might think. And the standard accounts of the fallacy aren’t very helpful. In order to know whether someone has given an ad hominem argument, you’d have to first figure out whether their argument was good or bad.

Ad Populum

This is the “fallacy” of believing something because most people believe it. But what exactly is supposed to be wrong with that? Here are three interpretations:

(i) Maybe the idea is that most people believing *p* is *irrelevant* to whether *p* is true. I.e., if most people believe it, that doesn’t mean it is more likely to be correct. Problem: This is obviously wrong. If most people believe something, that obviously *does* make it more likely to be correct than if most people don’t believe it. If most of our beliefs weren’t true, the human species would die out pretty much immediately.

Sometimes, people elaborate on this “fallacy” by citing examples of beliefs that were once widely held but were false – e.g., that the sun orbits the Earth. So let me now just mention a few typical examples of beliefs that are widely held:

Dogs exist.

It’s generally lighter in the daytime than at night.

The sky is blue, not red, green, or yellow.

There are more than three human beings in existence.

Human beings commonly have beliefs and desires.

Putting your hand in a fire hurts.

Six is more than two.

The Earth has existed for more than five minutes.

When you drop rocks near the surface of the Earth, they generally fall.

No objects are completely red and simultaneously completely green.

...

Once you get the hang of it, I'm sure you can extend that list for a long time. Now, which would you say there are more of: Widely-held beliefs that are true, or ones that are false? If you don't think most of those items are true, there's something seriously wrong.

(ii) Maybe the idea is just that most people believing p does not *conclusively prove* that p is true. That's true, of course. But it's also a frivolous point to make. Of course it isn't conclusive proof; so what? Who was expecting *conclusive proof*? You may as well complain that it hasn't been conclusively proved that the Earth orbits the sun (this is true – it's merely *overwhelmingly likely* that the Earth orbits the sun!), and thence conclude that modern astronomy rests on a "fallacy".

(iii) Maybe the idea is simply that people often put *too much weight* on popular opinion. The fact that many people believe P is indeed evidence for P , but it is not *as* probative as people think. This is indeed very plausible in many cases. It's easy to overgeneralize this point, though. So bear in mind that people don't always overestimate the reliability of popular opinion. E.g., consider the examples of popular beliefs listed under (i) above: those beliefs are just as reliable as people generally take them to be.

Appeal to Authority

Students who read about the "appeal to authority fallacy" may conclude that one should never believe something because of who says it. But often one *should*. Especially if the person who says p is very smart and reasonable, then p is likely to be true. This doesn't *guarantee* that p is true, but it often makes it likely.

This might extend to Einstein on nuclear policy. Einstein was smart and reasonable, so his views are likely to be correct. I don't mean to imply that there is no problem here, though. Because Einstein was a popular celebrity scientist, people are liable to attach *more weight* to his views than they deserve. But that isn't what the textbooks imply; they seem to suggest that Einstein's opinion on nuclear policy should be given *no* weight.

We could define “appeal to authority” so as to make it automatically fallacious – e.g., define it as the mistake of attaching *too much* weight to an authority. But again, it’s not clear how often this actually happens, and the textbook presentations are not generally very helpful for recognizing when it does.

Begging the Question

The concept of “begging the question” is often misused by philosophers (one of the few confusions that is distinctive of philosophers!). The misuse comes about something like this: The philosopher starts with the idea that an argument begs the question (and therefore is fallacious) when *someone who rejects the conclusion wouldn’t* (or shouldn’t, or couldn’t reasonably be expected to) *accept all the premises*. That italicized phrase is treated as something like a definition of the fallacy. The philosopher then looks at some particular deductive argument. He notices that if you start by assuming the conclusion of the argument is false, you can deduce that one of the premises is false. Usually, the philosopher identifies a specific premise that is least obvious and says that, if the argument’s conclusion is false, then that specific premise is false. He concludes that someone who rejected the argument’s conclusion would also reject that premise. Therefore, to assert that premise is to beg the question.

People who fall for this mistake fail to notice that it represents a rejection of *all valid deductive reasoning*. In a valid deductive argument, by definition, if all the premises are true, the conclusion must be true. That is logically equivalent to the following: If the conclusion is false, then one of the premises must be false. So if you start by assuming the conclusion is false, and the argument was valid, you can *always* deduce that (at least) one of the premises is false. Example: Take the argument, “Miley Cyrus is a person. All people are mortal. Therefore, Miley Cyrus is mortal.” This could be said to beg the question because, if you don’t think Miley is mortal, then you should not accept the premise that all people are mortal. Given the obvious fact that Miley is a person, to assert that all people are mortal just “assumes” that Miley is mortal. Or so you might claim.

Presumably, it’s false that all valid arguments are fallacious. So something went wrong there. The problem is the definition of “begging the question”. Bad definition: “You beg the question when

someone who rejected your conclusion would reject one of your premises.” Better definition: “You beg the question when the *justification* of one of the premises *depends upon* the justification of the conclusion.” In the “Miley” inference mentioned above, it’s true that someone who insists that Miley is immortal would presumably also deny that all people are mortal. But it’s false that the *justification* for “All people are mortal” must *depend upon* the justification for “Miley is mortal.” Rather, “All people are mortal” could be justified, say, by an inductive inference (so far, all people who have ever lived have died within 125 years of their birth).

Post Hoc

When A is followed by B, that *is evidence* that A causes B, provided that you don’t know anything to the contrary. Of course, it is not *conclusive* evidence, and in most cases, you need more information to form a justified belief. But talk of the post hoc “fallacy” is facile and unhelpful. It tempts students to think either (i) that the fact that A is followed by B is evidentially *irrelevant* to the causal claim (which is wrong), or (ii) that an inference is only good if the premise *conclusively proves* the conclusion (also wrong).

A related slogan is “Correlation doesn’t imply causation.” The saying means that just because A and B go together regularly does not mean that one causes the other. Students learn the slogan in college and think it’s sophisticated, but it’s kind of simplistic. Granted, if there is a reliable correlation between A and B, that does not *guarantee* that there is a causal connection. It *could* just be a coincidence. But if the correlation is well established, it becomes vanishingly improbable that it’s just a coincidence. There will be *some causal explanation*. Maybe A causes B, or B causes A, or some third factor, C, causes both A and B.

All of that is to help inoculate you against false charges of fallaciousness. Sometimes, a “fallacy” is not a fallacy.

4.3. Fallacies You Need to Be Told About

Now I’m going to tell you about some more interesting errors that human beings are prone to. If you’re like most people, you probably actually need to be told about these things.

Anecdotal Evidence

Often, people try to support generalizations by citing a single case, or a few cases that support the generalization. Scientists call this “anecdotal evidence”. Example: You try to show that immigrants are dangerous by citing a few examples of immigrants who committed crimes.

Anecdotal evidence has two problems. First, usually, when people do this, they don’t pick a case randomly; they search for a case that supports their conclusion while ignoring cases that don’t. (See: cherry picking.) Second, random variation: Even if you picked the cases randomly, it can easily happen just by chance that you picked a few atypical cases. In the immigration example, what you should actually do is look up the statistics on crime rates for immigrants compared with native-born citizens.

Assumptions

One of the major ways we go wrong is that we simply assume things that we don’t know. Unfortunately, when you assume things, you go wrong *a lot* more often than you expect. (You should assume that most of your assumptions are wrong!) It is hard to combat this, because we often don’t notice what we’re assuming, and it doesn’t even occur to us to question it.

Here are a couple of examples. Suppose you hear a statistic about how common intimate partner violence is in the United States (this is where someone physically abuses their girlfriend, boyfriend, or spouse). You naturally assume that the vast majority of these cases are men beating up women, and you might just go on reasoning from that implicit assumption. In reality, though, survey evidence suggests that men and women suffer this kind of abuse about equally often.^[16]

Or suppose you hear a statistic stating that most murder victims are killed by a family member or someone they knew. You naturally assume that most murders result from domestic disagreements, and that the murders are committed by ordinary people who lost control during an argument with a family member, or something like that. In fact, it turns out that almost everyone who commits a murder has a prior criminal record. Also, the vast majority of the *victims* are also criminals. (The category “a family member or someone they knew” includes such people as the victim’s drug dealer, the victim’s criminal partner, the victim’s fellow gang members, and so on.) You just

assumed that these were ordinary people, but the original statistic didn't say that.

I can't really properly convey to you just how often assuming things leads you astray – you need to experience being wrong over and over again, in order to appreciate the point. Unfortunately, most people never come to appreciate the point, because they never check on their assumptions to find out how many are wrong.

Base Rate Neglect

A “base rate” is the frequency with which some type of phenomenon happens in general. E.g., the base rate for heart disease is the percentage of people in the general population who have heart disease. The base rate for war is the percentage of the time that a country is at war. Etc.

When you want to know whether some kind of event is going to happen (or has happened, etc.), the best place to *start* is with the base rate. If you want to know whether you have a certain disease, first find out how common the disease is in general. If 1% of the population has it, then a good *initial* estimate is that you have a 1% chance of having it. From there, you should adjust that estimate up or down according to any special risk factors (or low-risk factors) that you have.

Most people don't do this; people commonly ignore base rates. Example: Suppose there is a rare disease that afflicts 1 in a million people. There is a test for the disease that's 90% accurate. Suppose you took the test, and you tested positive (the test says you have the disease). Question: Given all this information, what is the probability that you have the disease?

Many people think it is 90%. Even doctors sometimes get this wrong (which is disturbing). The correct answer is about 0.0009% (less than one in a hundred thousand). Explanation: Say there are 300 million people in the country. Of these, 300 (one millionth) have the disease, and 299,999,700 don't. The test is 90% accurate, so 270 of the 300 people who *have* the disease would test positive (that's 90%), and 29,999,970 of the 299,999,700 who *don't* have the disease would also test positive (that's 10%). So, out of all the people who test positive, the proportion who actually have the disease is $270/(270+29,999,970) \approx 0.000009$.

Cherry Picking

“Cherry picking” refers to the practice of sifting through evidence and selecting out only the bits that support a particular conclusion, ignoring the rest. Simple example: I have a bag of marbles. I want to convince you that most of the marbles in the bag are black. I look inside the bag, which is full of many colors of marbles – black, red, teal, chartreuse, and so on. I pick out five black ones, show them to you, and say, “See, these marbles came from this bag.” I don’t show you any of the other colored marbles that were in the bag. You might be misled into concluding that the bag is full of black marbles.

That’s like what people do in political debate. If I want to convince you, say, that affirmative action is bad, I might search for cases where affirmative action was tried and it didn’t work or it had harmful effects. If I want to convince you that it’s good, I search for cases where it really helped someone. *Of course* both kinds of cases exist – it’s a big society, full of millions of people! Almost any policy is going to benefit some people and harm others. Because of this, you should be suspicious when someone tells you stories designed to support a conclusion – always ask yourself whether they have a bias that might have caused them to cherry pick the data.

Confirmation Bias

When asked to evaluate a theory, people have a systematic tendency to look for evidence *supporting* the theory and not look for evidence *against* it. (This happens especially for theories that we already believe, but can also happen for theories we initially have no opinion about.) E.g., if asked whether liberal politicians are more corrupt than conservative politicians, a conservative would search through his memory for any cases of a liberal doing something corrupt, and he would *not* search through his memory for cases of conservatives being corrupt. A liberal, on the other hand, would do the reverse. Each just looks for cases that support his existing belief, and does not look for evidence against it. This is called “confirmation bias”.

To combat this, it is necessary to make a conscious effort to think of exceptions to the generalizations that you accept, and to look for evidence against your existing beliefs. Whenever you feel inclined to cite some examples supporting belief A, stop and ask yourself whether you can also think of similar examples supporting $\sim A$.

Credulity

Humans are born credulous – we instinctively believe what people tell us, even with no corroboration. We are especially credulous about statistics or other information that sounds like objective facts. Unfortunately, we are not so scrupulous when it comes to accurately and *non-misleadingly* reporting facts. There is an enormous amount of disinformation in the world, particularly about politics and other matters of public interest. If the public is interested in it, there is bullshit about it.

I have noticed that this bullshit tends to fall into three main categories. First, *ideological propaganda*. If you “learn” about an issue from a partisan source – for instance, you read about gun control on a gun control advocacy website, or you hear the day’s news from a conservative radio show – you will get pretty much 100% propaganda. Facts will be exaggerated, cherry picked, deceptively phrased, or otherwise misleading. Normally, you will have no way of guessing the specific way in which the information is deceptive, making the information essentially worthless for drawing inferences.

Second, *sensationalism*. Mainstream news sources make money by getting as many people as possible to watch their shows, read their articles, and so on. To do that, they try to make everything sound as scary, exciting, outrageous, or otherwise *dramatic* as possible.

Third, *laziness*. Most people who write for public consumption are lazy and lack expertise about the things they write about. If a story has some technical aspect (e.g., science news), journalists probably won’t understand it, and they may get basic facts backwards. Also, they often just talk to one or a few sources and print whatever those sources say, even if the sources have obvious biases.

I can’t give you adequate evidence for all that right now. But here’s an anecdote that illustrates what I mean. I once heard a story on NPR (National Public Radio, a left-leaning radio news source). It was about a man on death row who was about to be executed. From the story, it appeared that the man was innocent. New evidence had emerged after the trial, several of the witnesses had recanted their testimony, yet the courts had refused to grant a new trial. The only remaining hope was for the governor to grant a stay of execution. There was an online petition that listeners could sign.

Usually, I just accept news stories and then go on with my day. But on that occasion, I decided to look into the story before signing the petition. With a little googling, I found the court decision from the convict's most recent appeal, which had been denied. I read the decision, which contained a summary of the facts of the case and an explanation of the judges' decision.

What it revealed was that the NPR story was bullshit. What NPR said was basically just what *the defendant's lawyer* had claimed. The court carefully explained why each of those claims was bogus and provided no basis for an appeal. The most striking claim (which had initially made me think the defendant was probably innocent) was that multiple witnesses had "recanted" their testimony. What had actually happened was this: The defense lawyer went back to the witnesses *many years* after the original trial and questioned them on details of the case. Several of them either couldn't remember the details, or reported details slightly differently (e.g., what color shirt someone was wearing). The lawyer described this as "recanting their testimony". But none of them had changed their mind about the defendant being guilty.

The NPR journalists had apparently just credulously reported what the lawyer told them, without bothering to look up the court documents from the case. Why would they do that? Three reasons: (i) *Ideological bias*: The story painted the death penalty in a bad light, which a left-leaning news outlet would like. (ii) *Sensationalism*: The story of an innocent man about to be executed grabbed the audience's attention and inflamed their passions. (iii) *Laziness*: Checking on the story would have required work. Why put in that work when you know that almost all of your audience will just accept whatever you say? Long experience has led me to think that that case was not unusual; this is the way news media work.

Lesson: Popular media stories are untrustworthy. (By the way, it's no good checking them against *other* popular news sources, because they basically all copy from each other.) That also goes for, e.g., most bloggers, your next door neighbor, and other casual information sources. For relatively reliable information, look at academic books and articles and government reports (e.g., Census Bureau reports, FBI crime reports).

Dogmatism and Overconfidence

People who study rationality have a notion called “calibration”. Your beliefs are said to be well-calibrated when your level of confidence matches the probability of your being correct. For example, for all the beliefs that you hold with 90% confidence, about 90% of them should be true. When you’re 100% confident of things, they should be true 100% of the time. Etc.

Most people are badly calibrated. In fact, almost everyone errs in a particular direction: Almost everyone’s beliefs are *too confident*. People say they are “100% certain” of a bunch of things, but then it turns out that only, say, 85% of those things are actually true. (There are psychological studies of this.^[17]) This is the problem of overconfidence. Almost everyone has it, and almost no one has the opposite problem (underconfidence), so you should assume that you are probably overconfident too. You should therefore try to reduce your confidence in your beliefs, particularly about controversial things, and particularly for speculative and subjective claims.

Ideological “Cause” Judgments

Back in 2008–2009, America suffered a severe economic recession. A lot of people lost money, lost their jobs, and were generally unhappy. What set it off was problems in real estate. Home prices had gotten very high, then they dropped, a lot of people started defaulting on (not repaying) their home loans, banks were in a lot of trouble, and other investors and financial companies were in trouble because they’d made investments that depended on home prices staying high and home loans getting repaid.

In the wake of the crisis, many people tried to explain why it had all happened. This included people with opposing ideologies. Roughly, there were people with pro-government and people with anti-government ideologies, and both tried to explain the crisis. Can you guess what the two sides said? The pro-government people said the recession happened “because” there wasn’t enough regulation – and they listed regulations that, if they had been in place, would probably have prevented the crisis. The anti-government people said the recession happened “because” there was *too much* government intervention – and they listed existing government policies that, if they *hadn’t* been in place, the crisis probably wouldn’t have happened.

Notice that the basic factual claims of both sides are perfectly consistent: It's perfectly possible that there were some actions the government took such that, if the government *hadn't* taken them, the crisis wouldn't have happened, *and also* there were some actions the government *failed* to take such that, if it *had* taken them, the crisis wouldn't have happened. It's perfectly plausible that the crisis could have been averted in more than one way: either by adding certain government interventions, or by removing some other government interventions. Which alternative you focus on depends on your initial ideology.

Both sides took the episode to further support their ideology: "We have too much government" or "We need more government." These conclusions were supported by their respective causal interpretations: "The recession was caused by government interventions" or "The recession was caused by government failure to intervene."

Who was right? Assume the facts are as stated (that some additional interventions would have prevented the recession *and* the repeal of some other interventions would have prevented the recession). In that case, we should either accept both causal claims or reject both causal claims, depending on what we mean by "cause". If we mean "sole cause", then we should reject both causal claims (i.e., we should say the recession was *not* caused either by government intervention or by failure to intervene). If we just mean "factor such that, if it were changed, the effect wouldn't have happened", then we should accept both causal claims (the recession was caused by intervention *and* by failure to intervene).

It's okay to say that *x* was caused by *y*, provided that you also recognize all the other things that caused *x* in the same sense. If there are many different causes, then you need additional evidence or arguments to establish *which* one of those causes is the best one to change. In the recession case, we would need independent arguments to establish which cause of the recession (intervention or failure to intervene) it would have been better to change.

Oversimplification

People very often oversimplify philosophical issues. Say you're thinking about the morality of abortion. A tempting simplification would be to say that there are two positions: pro-choice and pro-life

(or pro- and anti-abortion). Either fetuses are people and killing them is murder, or fetuses aren't people and killing them is perfectly fine.

But this overlooks the possibility that late-term fetuses are people but early-term fetuses are not; or maybe personhood comes in degrees and fetuses become progressively more personlike as they develop; or maybe fetuses are persons in some senses but non-persons in other senses. So there is a range of possible positions, not just two.

Viewing things in black-and-white terms is a common oversimplification. We look at two simple positions rather than considering a spectrum of possibilities. The problem is that often, the truth is a more subtle position that doesn't clearly fall under either of the two simplest categories of view.

***p*-hacking**

Similar to cherry picking, “*p*-hacking” or “data mining” sometimes happens in science. A scientist has a large amount of statistical data, with different variables. Even if all the data is completely random, any complex set of data is going to show some patterns that look significant. Essentially, one can take the data and use it to test many different possible hypotheses. Even if all the hypotheses are false, eventually, just by chance (due to random variations in the data), one of the hypotheses will pass a test for “statistical significance”. This is one reason why many published research results, especially in medicine, psychology, and social science, are false. E.g., a study will find that some food increases the risk of cancer for non-smoking, middle-aged men; but then someone tries to replicate it, and they don't get the same result, because the original result was just due to chance.

[\[18\]](#)

Speculation

Speculative claims are essentially guesses about things that we lack the evidence to establish as yet. Claims about the future, or claims about what would have happened in hypothetical alternative possibilities, are good examples of speculative claims.

Example: You're arguing about whether it's good for government to try to stimulate the economy by spending money. You say this is good because, e.g., if the government *hadn't* stimulated the economy back in 2009, the recession would have continued much longer. This is speculative – we don't *know* what would have happened, because in

fact the government *did* pass a stimulus plan, and we can't now go back in time and change that to see what would have happened if they hadn't.

The problem with speculative claims is that people with different philosophical (or political, religious, etc.) beliefs tend to find very different speculations plausible. E.g., people who are suspicious of government will find it more plausible that, without government stimulus, the recession would have been *shorter*. So arguments that start from speculative premises are typically not rationally persuasive.

Advice: If you want to rationally persuade people of something, try to avoid speculation.

Subjective Claims

Roughly, a “subjective” claim is one that requires a judgment call, so it can't just be straightforwardly and decisively established. For example, the judgment that political candidate A is “unqualified” for the office; the judgment that it's worse to be unjustly imprisoned for 5 years than to be prevented from migrating to the country one wants to live in; the judgment that Louis CK's jokes are “offensive”; etc. (This differs from speculative claims, because in the case of speculation, there might be ways that the claim *could* in principle be decisively verified; it just hasn't *in fact* been verified.)

Note: I am not saying that there is “no fact” or “no answer” as to whether these things are the case, or that they are dependent on people's “opinions”. What I am saying is that there are not clear, established criteria for these claims, so it is difficult to verify them. Maybe it's true that Louis is offensive, but if someone doesn't find him offensive, there is no decisive way of proving that he is.

People often rely on subjective premises when arguing about controversial issues. The problem with this is that subjective claims are more open to bias than relatively objective (that's the opposite of “subjective”) claims. So people with different philosophical (or political, or religious) views will tend to disagree a lot about subjective claims. And for that reason, they are ill suited to serve as premises in philosophical, political, or religious arguments. Advice: Try to base your arguments, as much as possible, on relatively objective claims.

Treatment Effects vs. Selection Effects

Let's say you have created a new educational program for pre-school children. You want to know whether it improves learning or not. What you would do is look at kids after they've had your program, and compare them to kids of the same age who *didn't* have your program, and see if the first group perform better on tests. Let's say kids who had your special program perform 10% better on later tests, on average. Then you'd probably conclude that your program works.

But wait. Here is another possibility. Suppose (as would usually be the case) that the kids who entered your special educational program were the kids whose *parents chose* to enroll them in that program. The rest were kids whose parents did not decide to enroll them. Furthermore, maybe the parents who enroll their kids in special programs are on average *smarter* and *value education more* than the parents who don't do that. Furthermore, maybe intelligence and value placed on learning are partly genetic, and so these parents passed those traits on to their kids. So the children who went into your program were *already*, on average, smarter and more interested in learning than the children who didn't go into the program. And maybe *that* explains why they did 10% better on tests after the program. Maybe your program has no effect at all; it's just that you got the smart kids in it, and that made the program look good.

That is an example of a "selection effect" – a case where it looks like A causes B, but it's actually just that the instances of A that you tested were already more likely to be B's for other reasons. Selection effects are contrasted with "treatment effects" – cases where the thing you're testing really causes the effect that it's thought to cause. In the education example, academic success is correlated with taking the special program. This could be due to a treatment effect (meaning the program *causes* kids to learn more), or due to a selection effect (meaning the program *selects* students who are already good at learning).

Selection effects are very often mistaken for treatment effects. Another example: You want to know if some vitamin improves people's health. So you look at people who take supplements of that vitamin regularly, and you find that they are healthier than the people who don't take it. You think this shows that the vitamin supplements are good for people ... but actually, it's more likely a selection effect:

People who take vitamins are more likely to also be exercising, eating healthy foods, and so on, which is why they would be healthier than average, even if the vitamins did absolutely nothing.

Whataboutism

Similar to *tu quoque* (see section 4.1 above), whataboutism occurs when someone criticizes something bad, and you respond with, “What about x?”, where x is some other bad thing. Example: Someone complains that the current President’s proposed budget has a very high deficit. You say, “What about the *previous* President? He had high deficits too!” Or: Someone complains that the President just murdered a child. You respond that some other political figure, from an opposing party, also murdered children. “What about that?” you demand.

The reason people engage in whataboutism is that, rather than being interested in practical issues about what should be done in our current situation, they instead see political discussion as a kind of tribal contest, a competition between “their side” and “the other side”, where whoever makes their side look better wins. So they don’t want attention focused on any flaws in one of their side’s people (e.g., a politician from their own political party). So they try to divert attention to something that’s bad about someone on the *other* side.

The problem is that this practice *systematically prevents evils from being addressed*. For *any* evil in the world (unless it’s literally the worst thing in the world), one can always identify some other, even worse evil, and say “What about that?” For any evil done by any political leader, it will virtually always be true that some other leader from another party has some time committed a similar evil (and also that members of that person’s party didn’t do anything about it). If your response when you hear about any evil currently happening is to deflect attention to some *past* evil committed by another person or group, that means that evils never get addressed. Attention always gets deflected away by whataboutism. The next time someone else is doing something evil, that won’t be addressed either, because people will say “what about” the *present* evil that wasn’t properly addressed.

5. Absolute Truth

Beginning philosophy students sometimes want to know whether there is “absolute truth” or “objective reality”. These questions are not much discussed in contemporary, academic philosophy because there is not much disagreement about them among philosophy professors. Still, we need to discuss them here because students wonder about them, and how one thinks about them can affect one’s thinking about the rest of philosophy.

5.1. What Is Relativism?

5.1.1. *Relative vs. Absolute*

In philosophical contexts, to say that a thing is “**relative**” is to say that it varies from one person to another, or from one society to another (or perhaps from one species to another, etc.). To be more explicit, we sometimes say a thing is “relative to an observer”, “relative to a society”, and so on. By contrast, to say a thing is **absolute** is to say that it does *not* vary from one person to another (or one society to another, etc.); it is constant.

(By the way, notice how the definition of “absolute” exactly matches the definition of “relative”, except with a “not” inserted. This is deliberate. In philosophy, we commonly define two terms such that one simply covers everything that *isn’t* covered by the other term. That’s because we want to be sure that we’ve covered all the possibilities.)

For example, a proposition can be *certain* for one person but *uncertain* for another. If I’m in Paris and I see and feel rain falling on me, then for me it is certain that it is raining in Paris. On the other hand, if you are in New York at the time, and you cannot observe the weather in Paris, then for you it is uncertain whether it is raining in Paris. Thus, we can say that the level of certainty of propositions is “relative to an observer”.

Another example: Suppose you have some homework problems to do for your math class. It may be difficult for you to complete the problems, yet easy for the professor to complete those same problems. Thus, we can say that the *difficulty* of a task is “relative to an individual”.

Relativism about truth (a.k.a. “truth relativism”) holds that *truth* is relative to an individual. That is, the same proposition can be true for one person but not true for someone else. (What does it mean to be “true for” a

person? More on that below.) Absolutism holds that truth is not relative: Propositions are simply true or false, not true *for* a person.

5.1.2. *Subjective vs. Objective*

Relativists also often say that “reality is **subjective**”. What this means is that the world (“reality”) is dependent on observers. That is, it depends on there being some people (or other beings with minds) to be aware of it. The contrast to “subjective” is “objective”. **Objective** phenomena exist on their own, independent of observers.

It is fairly uncontroversial that *some* things are subjective in this sense. For example, consider the property of being *funny*. A plausible analysis is that for a joke to be “funny” is for it to have a tendency to make ordinary humans who hear the joke laugh, feel amused, etc. – or something like that. Funniness isn’t an intrinsic property of funny things; it is in the ear of the observer. The funniness just *consists of* the tendency to provoke amusement in us.

Note: This is a *different sense* of “subjective” than the sense used in section 4.3 above. There, “subjective” was used for claims that require judgment and lack a decisive method of verification. Here, “subjective” is used for phenomena that constitutively depend on observers. Many words (within philosophy and outside it) have multiple senses, depending on the context. Get used to it.

Almost everyone regards some things as objective. For instance, for an object to be *square*, it is not necessary that anyone observe the object, feel any way about it, or have any other reaction to it; the squareness is just a matter of the spatial arrangement of the object’s parts, independent of us. The great majority of things in the world seem to be objective in this sense. Relativists, however, are known to deny this sort of thing, claiming instead that everything is in some way dependent on the mind.

5.1.3. *Opinion vs. Fact*

One way of understanding relativism is that it is the view that “everything is a matter of opinion”. But what does this mean? American high school students are frequently taught a distinction between *facts* and *opinions*; unfortunately, they are often taught a confused account that presupposes controversial views, and incorrectly taught it as if *it* were a matter of fact.

There are a few different distinctions in the vicinity. E.g., the distinction might be between things that are believed to be true and things that *are* true;

or between our beliefs and the aspects of the world that our beliefs are *about*; or between propositions that are conclusively verified and those that have not been (or cannot be) verified; or between propositions that are true and those that are false; or between propositions that are true and those that are neither true nor false (if there are any of those?); or between objective things and subjective things.

Notice that those are *six different* distinctions. Unfortunately, “fact” vs. “opinion” (or “matter of fact” vs. “matter of opinion”) appears to be a jumble of all these different distinctions. For this reason, I shall avoid talking about “facts” versus “opinions” in the rest of this discussion.

5.2. Some Logical Points

5.2.1. The Law of Non-Contradiction

As a preliminary matter – and this is really good background for any philosophical discussion – it’s worth reviewing some basic logical points ... starting with the most famous, basic principle of logic, the **law of non-contradiction**. This is, basically, the principle that contradictions are always false. Or: For any proposition A, $\sim(A \ \& \ \sim A)$.

A proposition of the form $(A \ \& \ \sim A)$ (read “A and it’s not the case that A”) is known as an explicit contradiction. (We also sometimes use “contradiction” to cover statements that are not already of the form $(A \ \& \ \sim A)$ but *entail* something of the form $(A \ \& \ \sim A)$; these would be implicit contradictions, not explicit.) Why is it that contradictions are never true?

The answer is basically “because of the meaning of the word ‘not’”. A proposition, A, has a certain range of possibilities in which it counts as true. (In some cases, the “range” might be empty, i.e., it never counts as true.) The negation of A (represented “ $\sim A$ ”), *by definition*, just refers to all the *other* possibilities. If you think you can imagine a situation in which both A and $\sim A$ are true, then you haven’t understood how the symbol “ $\sim$ ” is used (or how the English word “not” is used). If A obtains in a certain situation, then $\sim A$, by definition, doesn’t. That’s just what “ $\sim A$ ” means.

Another way to put the point: If a person asserts A, and then asserts $\sim A$, then they are basically telling you that they themselves are wrong. That is, the second half of what they said was that the first half was wrong; therefore, overall, they’re guaranteed to be wrong. That’s the problem with contradicting yourself.

5.2.2. The Law of Excluded Middle

Now for the second most famous principle of logic, the **law of excluded middle**: For any proposition, either that proposition or its negation obtains; there is no third alternative. That is, for any A , $(A \vee \sim A)$. Why is this true?

Again, the answer is “because of the meaning of ‘not’”. We noted above that the proposition $\sim A$ is just defined as excluding all the cases in which A obtains. It is also defined as *including* all the *other* cases. If you think you’re imagining a case in which neither A nor $\sim A$ obtains, then you’re confused about the use of “ $\sim$ ”. If A doesn’t obtain in a certain situation, then $\sim A$, by definition, *does*. That’s just what “ $\sim A$ ” means.

Another way to put the point: Suppose someone tells you that neither A nor $\sim A$ obtain. In that case, *one* of the things they are saying is that A doesn’t obtain. The *other* thing they are saying is that $\sim A$ doesn’t obtain. But “ $\sim A$ ” just means that A doesn’t obtain. So what they are saying is: $\sim A$, but also, $\sim(\sim A)$. But that’s an explicit contradiction.^[19]

Caveat: The preceding points apply only when “ A ” picks out a definite proposition. If you have a sentence that does not have a clear enough meaning to assert any determinate proposition, then neither that sentence nor its negation will be true. Thus, “All blugs are torf” is not true, and neither is “Not all blugs are torf”, since “blug” and “torf” do not have definite meanings. For another example, suppose I announce, out of the blue, “He has arrived”. You ask whom I’m talking about, and where that person arrived, and I reply that I didn’t really have any particular person or place in mind. In that case, my sentence is neither true nor false. “He has arrived” isn’t true, and neither is “He has not arrived.”

5.2.3. What Questions Have Answers?

It is sometimes said that philosophical questions “have no answers”. (Almost no philosopher would agree with that statement, but often students and lay people say it.) What should we think about this view? On the face of it, it is hard to make sense of the idea.

Take the question of whether God exists, which is a good example of a philosophical question. Suppose someone says that this question “has no answer”. Now, it appears that the possible answers to the question would be “Yes, God exists” and “No, God doesn’t exist.” If either of those is correct, then the question has an answer. So to say the question has no answer must be to claim that neither of those answers is correct: It is neither the case that God exists, nor the case that God doesn’t exist. But that is just to say that

it's not the case that God exists, and it's also not the case that it's not the case that God exists. An explicit contradiction.

It doesn't matter what question we pick. You can substitute the question "Do animals have rights?" To say this question has no answer must be to claim, at least, that it's not the case that animals have rights, and it's also not the case that animals don't have rights. Again, an explicit contradiction.

But wait; there *are* ways that a question could lack an answer. One way is if the question is not sufficiently meaningful (compare the caveat about the law of excluded middle above). "Is the moon torf?" has no answer since "torf" has no meaning. "When is 14?" likewise lacks an answer since it doesn't make sense. Also, a question might be said to have no answer (or maybe just no *appropriate* answer) if it contains a false presupposition. Suppose someone asks me, "Have you stopped stealing kittens?" If I have never stolen a kitten, then I can't say "Yes, I've stopped", but it wouldn't really be appropriate to say, "No, I haven't stopped" either.

However, neither of these things apply to typical philosophical questions. "Is there a God?" isn't meaningless, and it doesn't contain a false presupposition. So it remains unclear in what sense it could fail to have an answer.

Perhaps the idea is just that philosophical questions lack answers *that can be decisively verified*. If this is what is meant, then "Philosophical questions have no answers" is a simple misstatement. Compare: If you don't know who stole your cookies, you should *not* say, "There was *no* thief"; you should just say "The thief is unknown." Similarly, if we don't know the answer to a philosophical question, we should not say "There *is* no answer"; we should just say "The answer is unknown."

All this is related to the question of truth relativism, because relativists often say that philosophical questions have no answers (or maybe *no* question has an answer?), and this seems to be intended as closely related to the idea that there are no "absolute truths".

5.3. Why Believe Relativism?

5.3.1. The Argument from Disagreement

The most popular "argument for relativism"^[20] begins by observing that there is a great deal of variation in people's beliefs across cultures. Some cultures believe that when we die, we go to heaven; others, that we are reincarnated in this world; others, that we are simply gone forever. Some believe that polygamy is wrong; others, that it is perfectly cool. And so on.

(Anthropologists like to go on and on about the variation among cultures.) Therefore, it is said, you can see that truth is relative to a culture. (Or, if you want to say truth is relative to an *individual*, start by going on about the variation in beliefs among individuals.) The argument appears to go like this:

P1. Beliefs vary from culture to culture. (Premise)

C. Therefore, truth varies from culture to culture. (Conclusion)

Is this argument sound? Certainly the premise is true; no one doubts that. Is the inference valid? Does it follow, just from the fact that beliefs vary, that truth varies?

No, it does not. It could be that beliefs vary across cultures, and yet there is only one truth; it might just be that most (possibly all) of these cultural beliefs are false. To make the argument valid, we would have to add a premise to it, something like this:

P1. Beliefs vary from culture to culture.

P2. All beliefs are true.

C. Therefore, truth varies from culture to culture.

Now *that* is valid. C clearly follows from P1 and P2. But now the problem is that P2 is obviously false. Not all beliefs are true! I bet you can think of some times that you had a false belief.

We could try weakening the second premise to “All beliefs that vary from culture to culture are true” to make it slightly less ridiculous, but it would still be obviously false, or at best unjustified. We would need an argument that all these cultural beliefs are true.

In fact, the argument has a bigger problem than merely a false or unjustified premise: The first premise *logically contradicts* the second one. For in saying that beliefs *vary* from culture to culture, what is of course meant is that different cultures have *conflicting* beliefs – this is borne out by the standard examples. For instance, as noted, some cultures think that when we die, we go to heaven; others, that we are reincarnated in this world. Those two possibilities are incompatible with each other; we couldn’t be in both places at once. Some cultures think polygamy is wrong; others, that it is *not* wrong. Again, those are mutually inconsistent views. The fact that they are inconsistent just *means* that they can’t both be true. So P1, understood in the sense that it is intended, just directly entails that P2 is false.

5.3.2. *The Argument from Tolerance*

Why has anyone ever been a relativist? The original motivation appears to have been sort of political: Relativists think that *toleration* is an important virtue. We should not try to impose our practices or beliefs on other cultures or other individuals. It was thought that being a relativist was a way of expressing tolerance and open-mindedness. If you are an absolutist, after all, then you must think that other people and other cultures, when they disagree with you, are *wrong*. This sounds closed-minded and intolerant. It might be offensive to people from other cultures. It could even lead to your trying to force the other people to conform to your beliefs. In the past, for example, people who were convinced that they knew the one true religion would try to forcibly convert others – this led to wars, inquisitions, torture, and lots of awful stuff like that. The best way to prevent that sort of thing, the relativists think, is to give up on thinking that there is any one truth.

Notice a peculiar feature of this argument: It is not actually an argument that relativism is true. It just says that it would be *socially beneficial* if people were to believe relativism. That's compatible with the theory being factually false. We could agree that toleration is good, and that being a relativist makes people tolerant, but also hold that relativism is false.

The other problem with the argument is that it overlooks other ways of promoting tolerance. Here is one way: We could adopt the view that tolerance is good. Maybe even *objectively* good. Wouldn't that be the most logical approach, if we're trying to promote tolerance? We don't have to go through any logical contortions trying to figure out how conflicting propositions can be simultaneously true. In fact, the people who accept relativism on the basis of the value of toleration have *already* accepted that toleration is good. They could have just stopped there.

Here is another, closely related possibility: We could hold that people have *rights*. Including, say, a right not to be coerced as long as they are not violating anyone else's rights. Philosophers have had a good deal of discussion and debate about exactly what rights we have, but we don't need to work out the details here. For present purposes, it suffices to say that, on pretty much anyone's conception of rights (among people who believe in rights at all), forcing people from other cultures to adopt your cultural practices or beliefs would normally count as a rights violation. We don't have to say that their cultural beliefs are all true; even if someone has false

beliefs, you still can't use force against them without provocation. People with mistaken beliefs still have rights not to be coerced.

Notice how this is perfectly consistent with absolutism. The staunchest absolutist could (and most of them do) embrace the idea of individual rights and toleration. In fact, holding that individual rights are *objective* would presumably make one more inclined to respect them – and therefore, to be more consistently tolerant than people who don't accept any objective truths.

5.4. Is Relativism Coherent?

5.4.1. *Conflicting Beliefs Can Be True?*

Among professional philosophers, truth relativism is often seen as incoherent or otherwise absurd. For this reason, the view is rarely discussed in academic books or journal articles, unless it is to object to some *other* theory by accusing the theory of leading to relativism.^[21] Why are academic philosophers so anti-relativist?

Mostly because philosophers don't like inconsistency. Logic is kind of our thing. And the core drive of relativism seems to be to somehow embrace inconsistencies. We see a bunch of conflicting beliefs, especially beliefs of different cultures that contradict one another – e.g., some think polygamy is wrong, others think it's fine. The relativists see this, and they want to somehow *let everyone be right*. That motivation, just on its face, seems like a desire to embrace contradictions. The fact that two beliefs contradict each other *just means* that they can't both be right. If one belief says that *x* is wrong and another says that *x* is *not* wrong, then just by definition, the two beliefs can't both be correct (because of the meaning of “not”, as discussed in section 5.2).

It also seems as though relativists are allowing their politics (specifically, their desire to avoid offending people from other cultures) to override logic, as discussed in section 5.3.2.

That said, relativists try to avoid actual inconsistency precisely by holding that truth is relative. If you and I have conflicting beliefs, it would of course be contradictory to say that both our beliefs are simply *true*. So instead, the relativist says that the one belief is *true for me*, and the other belief is *true for you*. Of course, they're not both true for the same person, nor are they both true absolutely.

This formally avoids inconsistency. But it only helps if it's possible to say what expressions like “true for me” *mean*. Otherwise, we've just traded

a contradictory statement for a meaningless statement. Unfortunately, relativists rarely have anything to say about what “true for me” means, which arouses suspicion that they don’t actually mean anything by the phrase.

Sometimes, it sounds as though “ p is true for me” just means “I believe p ”. But then all the relativist is saying is this: When two people have conflicting beliefs, each belief is believed by that person. E.g., if I believe p and you believe $\sim p$, then p is believed-by-me, and $\sim p$ is believed-by-you. But this would *trivialize* relativism.

Note

A thesis is said to be “**trivial**” when it is so obvious that it is not worth saying (especially if it is just defined to be true). For example, the thesis that *all tall people are tall* is trivial. To “trivialize” a statement is to interpret words in such a way that the statement would be trivial. Philosophers generally *reject* trivializing interpretations of our statements, because we want to be saying something that’s worth saying.

To put the point in more technical terms: Well, duh. *Obviously* each belief is believed by the person who has it. What’s the point of saying that? How does that help with the fact that the two beliefs contradict each other? It certainly doesn’t do anything to show how both beliefs could be in any sense correct. (See section 5.5 below for more on the meaning of “true for me”.)

5.4.2. Is Relativism Relative?

Perhaps the most popular objection to relativism is that relativism, if true, could only be *relatively* true, not absolutely true. If we say relativism is absolutely true, we contradict ourselves.

The relativist might respond: “Yep, the truth of relativism is relative! What’s wrong with that?”

Maybe the objection assumes that to call something *relative* or *relatively true* implies that it is *not really true*. In that case, a theorist could not hold their own theory to be relative. But the relativist would presumably deny that “relatively true” implies “not really true”; they would say that relative truth just is truth. So so far, the objection doesn’t show anything.

Here’s another try. If the truth of relativism is relative, that means it is only true *for relativists*. For the rest of us (i.e., for absolutists), absolutism is true. But it is very difficult to understand this. So, in the relativist’s view, it is true *relative* to absolutists that absolutism is true *absolutely* (and not just

relative to them). Huh? I don't know what it means for something to be true absolutely, relative to someone. That just sounds incoherent.

Whatever this might mean, if it means anything, it would not satisfy the aim of relativism to promote tolerance. For now the absolutists get to hold on to their absolutist view (it's true for them!), which means they can go on oppressing everyone else (*if* indeed that was a consequence of absolutism in the first place). Just as relativism is supposed to stop us from saying that other cultures are wrong, it must also stop the relativist from saying that absolutism is wrong. But then, if they're not rejecting absolutism, there seems to be no point.

5.4.3. *Meaningful Claims Exclude Alternatives*

To make a meaningful, informative claim is to exclude some alternatives. We can think of the range of possible ways the world might be, metaphorically, as a space, the "space of possibilities". Making an informative statement (a statement intended to communicate some information to the audience) is drawing a line around a region in that space and saying "The actual world is in here." If you then add, "But I'm not excluding the possibility that it might be outside this region", then you rob your own statement of all content; now you're telling your audience nothing. E.g., if you say, "The sky is blue", you are conveying information about the color of the sky, which excludes the possibility that it's green, or red, or yellow, etc. But if you then add, "... or it's some other color, or no color, or maybe the sky doesn't exist", then you defeat the point of your own statement; now you've told us nothing. ^[22]

The same point applies to philosophical beliefs. If I say that God created the world, I am excluding the possibility that the world always existed, or that the world was created by someone other than God, or that it was created by entirely natural forces. So if someone else believes one of those other possibilities, I am necessarily denying what they believe. If I say that I'm *not* ruling out any of those alternatives (nor any other alternatives), then I am essentially not saying anything about how the world did or didn't come about.

What the relativist wants is to have his cake and eat it too: He wants everyone to be able to hold on to their own beliefs, but at the same time to not have to reject anyone else's beliefs. That only makes sense if we have beliefs that don't exclude any alternatives. That is, our beliefs must be meaningless. Since the relativist wants *everyone* to refrain from rejecting

each other's beliefs, what the relativist really wants is for all beliefs to be meaningless (including relativism itself).

5.4.4. *Opposition to Ethnocentrism Is Ethnocentric*

Ethnocentrism is the habit of regarding one's own culture as superior to other cultures. Relativists, and especially cultural anthropologists, are famously opposed to ethnocentrism, which they associate with intolerance. They hold that toleration and belief in relativism are better than intolerance and ethnocentrism.

Now here is an interesting fact: Virtually all other human cultures have been intolerant and ethnocentric. People in other societies consider their own ways to be right and superior to those of other cultures. Attempts to subordinate other societies by force are extremely common in human history, all over the world. In fact, the belief in tolerance is a recent feature of our own culture, much more so than traditional cultures.

So if tolerance is better than intolerance and ethnocentrism, then tolerant cultures like our own must be better than intolerant, ethnocentric cultures (like almost all other cultures). From the premise that ethnocentrism is bad, we can infer that our culture is better than other cultures ... but that conclusion is itself ethnocentric! We seem to have arrived at incoherence.

The problem is the blanket assumption "ethnocentrism is wrong". The correct insight in this area is this: You cannot assume, merely because some practice is the practice of your own culture, that it is the best. Your culture is not necessarily the best just because it's your own. But here is the flip side: You also cannot assume, merely because some practice is the practice of your own culture, that it *isn't* the best. Being part of your own culture does not automatically make a belief correct, but *nor does it make it not correct*. Ideas have to stand or fall on their own merits, regardless of what society or person they come from or don't come from.

5.5. What Is Truth?

I don't know how we've gotten this far without talking about the meaning of "truth". To assess whether truth might be relative, surely we should say something about what truth *is*. Let's get to that now.

5.5.1. *The Correspondence Theory*

The traditional account of truth is known as the **correspondence theory of truth**. It says that truth is correspondence with reality. That is, truth is understood as a certain relationship, a kind of match, between a sentence or

a belief and the world: A sentence says that things are a certain way, or a person thinks that things are a certain way, and things *are* indeed that way. When that happens, you have a “true” sentence or belief.

Here is the most famous explanation of truth, which comes from Aristotle: “To say of what is that it is not, or of what is not that it is, is false, while to say of what is that it is, and of what is not that it is not, is true.”^[23]

5.5.2. *Rival Theories*

There have been other theories of truth. According to the pragmatic theory, truth is just whatever it is good to believe.^[24] (It could be good for a variety of reasons, including that it makes you feel good. But it has to be good overall, in the long run.) According to the coherence theory, truth is what coheres (fits together) with our belief system. According to the verificationist theory, truth is that which can in principle be verified.

These theories make room for relativism, because they suggest a coherent interpretation of such phrases as “true for me”: Perhaps a proposition is true for me when it is good for *me* to believe it, or when it coheres with *my* belief system, or when *I* could verify it. Notice that the same proposition might not be good for you to believe, might not cohere with your belief system, or might not be verifiable by you. So the relativist could use these theories of truth to argue that truth is relative.

The only problem is that all these theories of truth are wrong. (Yes, some smart people believed them, and some still do. Smart people believe a lot of false things.) What? How do I know that? Because I understand the use of the word “true” in English. Here are two things you should accept if you understand the word “true” in standard English:^[25]

1. “It’s true that P” entails “P”.
2. “P” entails “It’s true that P”.

For example, if it’s true that cats eat mice, then cats eat mice. Also, if cats eat mice, then it’s true that they eat mice. These aren’t profound or controversial points that I’m making here; these are just the most basic, trivial points about how the word “true” works. If some philosopher doesn’t agree with these things, then that philosopher must be using some different concept, not the concept of truth as used in ordinary English.

But the above three theories of truth all conflict with these trivial principles. Take the pragmatic theory: Truth is that which is useful to believe. This implies: (Necessarily) it’s true that cats eat mice if and only if it is useful to believe that cats eat mice. But as we’ve already said,

(necessarily) it's true that cats eat mice if and only if *cats eat mice*. If we combine these two claims, we can infer: (Necessarily) cats eat mice if and only if it's useful to believe that cats eat mice. That's obviously false. Cats, alas, don't care about us – they're not going to hold off on eating mice depending on whether it's useful for us to believe that they do it.

Maybe we're lucky: Maybe it turns out that for literally *everything* in the universe that happens it somehow is useful for us to believe that that thing happens. Even this amazing coincidence wouldn't really save the pragmatic theory. Because what the pragmatist is actually committed to (provided he accepts 1 and 2 above) is this: For every proposition P, "P" entails "It's useful to believe that P", and "It's useful to believe that P" entails "P". But both of those entailment claims are uncontroversially false. It is not logically impossible for cats to eat mice and yet for it not to be useful for us to believe that, or vice versa. (For another example: Suppose God rewards everyone who believes in Santa Claus with eternal life. Then it would be useful to believe in Santa. But this wouldn't make Santa pop into existence.)

Essentially the same point applies to the other two mentioned theories. The coherence theory requires us to accept that "Cats eat mice" entails "The belief that cats eat mice coheres with our belief system", and vice versa. The verificationist theory requires us to accept that "Cats eat mice" entails "We can verify that cats eat mice", and vice versa. Both false. (Granted, "We can verify that P" actually does entail "P", but not vice versa.)

Conclusion: We still have no good way of understanding the notion of relative truth.

5.5.3. *Is Everything Relative?*

Relativists hold that truth is relative. Given our above principles 1 and 2, that means that they would have to say *everything* is relative. For instance, if the *truth* of "Cats eat mice" is relative, then *cats eating mice* must be relative: Cats can't eat mice absolutely, they can only eat mice *relative to* some individual or culture.

If you're having trouble understanding what that means, join the club. I have no idea what it would mean for cats to eat mice *relative to* a person or culture. But we'd have to somehow make sense of that, to make sense of relativism.

In general, the relativist view (given principles 1 and 2 above) would have to be that *no* sentence in any language refers to a state of affairs

existing in the external world, apart from us; rather, *every* sentence refers to a relationship between a person or culture and something else. “Cats eat mice” would have to refer to a relationship between a person/culture and ... (something to do with cats and mice). “ $2+2=4$ ” would have to refer to a relationship between a person/culture and ... (something to do with numbers). Etc. This would have to be the case, again, because the relativist would have to think that cats can only eat mice relative to a person/culture, that $2+2$ can only equal 4 relative to a person/culture, and so on.

So, is that true?

Um, no. *Some* expressions in our language refer to relationships to people. For instance, “difficult” refers to a relationship that a task bears to a person (as in, “Handstands are difficult for me”). “Useful” refers to another relationship that a thing can bear to a person (as in, “Popsockets are useful to me”). That’s why we have no trouble understanding statements like, “Handstands are difficult for me, but not for Jo” and “Popsockets are useful for me but not for my cat.” But obviously not *every* damn predicate in the language refers to a relationship to a person. “Square” does not refer to a relationship to a person; that’s why “This table is square for me but not for Sue” draws a blank; there is no clear meaning of that.

5.6. I Hate Relativism and You Should Too

Philosophy professors, at least those from major research universities, tend to hate truth relativism. (Sometimes, we wonder where students learned relativism and what can be done about it. It wasn’t from us! Maybe they learned it in high school?) Why should we hate relativism?

Part of the reason is that truth relativism is an extremely unjustified view, for reasons explained above. It seems to straddle the fence between being contradictory and lacking any clear meaning. The central motivations for the theory appear to be ideologically propagandistic (a desire to promote tolerance), rather than stemming from anything that on its face would appear to be evidence for the theory. It’s more than just that the theory isn’t true or justified (after all, nearly all philosophical theories are false, but we don’t hate them). It’s that the theory doesn’t even seem to be *trying* to be true or justified. Philosophers tend to place a high value on rationality and truth, so we tend to take a dim view of philosophical positions that do not seem to aim at rationally identifying any truths.

But it’s more than that. Truth relativism does not just fail to be true, and it does not just fail to aim at truth; truth relativism *actively discourages* the

pursuit of truth. How so? The relativist essentially holds that all beliefs are equally good. But if that's the case, then there is no point to engaging in philosophical reasoning. We might as well just believe whatever we want, since our beliefs will be just as good either way. But this undermines essentially everything that we're trying to do. When we teach philosophy, we're trying to teach students to think carefully, and rationally, and objectively about the big philosophical questions (which hopefully will help you think well about other stuff too). When we do research in philosophy, we try to uncover more of the truth about these questions, so that we can all better understand our place in the world. All of that is undermined if we decide that it doesn't matter what we think since all beliefs are equally good.

Officially, relativism is a theory about the logical structure of the concept truth (that this concept is relational and always contains an implicit reference to an observer or group of observers); unofficially, however, it is an attack on the concepts of truth and objectivity, which are perhaps the two most important concepts for all intellectual inquiry. Inquiry (including philosophy, science, and all other forms of investigation) is about trying to bring our beliefs into line with reality. The world is a certain way, apart from us, and we need to try to make our minds accurately represent it. That kind of correspondence is known as "truth"; that is, this is what the standard English word "truth" refers to. By proposing that there is no absolute truth but only different "truths" relative to different people, the relativist is erasing the whole bit about matching reality. Which is to say, erasing the actual point of intellectual inquiry. They might then propose some *other* purpose of inquiry, but they have no room for what the rest of us thought was the point of it.

Traditionally, it was thought that relativism promotes tolerance and open-mindedness, so at least it would have good effects on people. But that might not even be true; it might in fact do the opposite. First, relativism might have the effect of *closing* people's minds, for the reason just discussed: It takes away the point of inquiry, thus potentially leading people to stop asking questions, stop trying to figure things out. That is the opposite of opening the mind.

Second, relativism might have the effect of promoting intolerance. For remember, the theory says that there are no objective/absolute/observer-independent truths. Whatever you believe is "true for you", and it doesn't

make sense to question whether your beliefs are *really* true, because, on this view, relative truth is all there is. Therefore, you may as well stick dogmatically to your current beliefs. Furthermore, if *you believe* that you should oppress other people and force them to adopt your practices, then *that* belief, too, will be “true for you”. So why not oppress others? There would be no basis for saying that you shouldn’t really do that, because the theory has removed objectivity from the picture.

The only response I can see to this last problem would be if the relativist declares that people should not act on the basis of what is “true for them”, because it isn’t objectively true. But if *that’s* what they say, then they’d also have to say that *no one should act on anything* – that is, we should all be completely apathetic – because, remember, the theory says there *is nothing* other than relative truth. If relative truth isn’t a basis for action, then there is no basis for action, on the theory. Thus, truth relativism potentially has very serious negative consequences, both intellectually and practically.

Part II: Epistemology

6. Skepticism About the External World

6.1. Defining Skepticism

In philosophy, “**skepticism**” basically refers to any view that implies that we can’t know a lot of the things we normally think we know, or that a lot of the beliefs we normally think of as justified are unjustified. There are more and less extreme kinds of skepticism, and there are many different things philosophers have been skeptical about. E.g., some people are skeptical just about morality (claiming that there is no moral knowledge); others are skeptical about inductive reasoning; others about the entire external world. A few people are skeptical about everything.

In this chapter, we’re going to discuss **external world skepticism**. External world skeptics think that we can’t know (or justifiedly believe) any contingent truths about the external world. That’s the world outside our own minds, so this view would say that you don’t know whether tables exist, whether there are any other people, whether you actually have two hands, etc. (Note: But these skeptics generally do *not* object to knowledge of *necessary* truths, like $[2+2=4]$, $[\sim(A \ \& \ \sim A)]$, and so on. Nor do they object to knowledge of one’s own mind – e.g., you can know what thoughts, feelings, and sensations you are experiencing. Hereafter, I’ll drop the tedious qualifier about “contingent” truths.)

6.2. Skeptical Scenarios

Skeptical scenarios are possible situations in which everything would appear to you as it presently does, but your beliefs would be radically mistaken. Skeptics use these to try to convince you that you don’t know anything about the world around you.

6.2.1. The Dream Argument

Have you ever thought you were awake when you were actually dreaming? (If not, you’re pretty weird, because this has happened to almost everyone.) In the dream, things might *seem* perfectly real to you, yet none of what you seem to see, hear, or otherwise perceive is real. Given that, how can you know that you’re not dreaming *right now*?

If we’re just thinking about normal dreams, of the sort that we all remember experiencing many times, there might be ways of telling that you’re not dreaming. Maybe you could try pinching yourself; if you feel

pain, then you're not dreaming. Or you could try to remember how you arrived at your present location; if you're awake, you should be able to remember roughly what you've done today, from the time you got up till you arrived at wherever you are. Or you could pick up something written (like this book) and just try reading it – if you're dreaming, you won't be able to read the book, because your unconscious mind does not in fact have the information that's contained in a real book.

Those sorts of things are all well and good. But now consider the hypothesis that maybe *all* of your life that you remember has been one huge dream. What you think of as your past waking experiences were just part of the same really long dream, and what you think of as your past *dreams* were actually *dreams within dreams*. (Some people, by the way, have actually had dreams within dreams: I have had dreams in which I dreamt that I was dreaming *and that I woke up*. But in reality, I was still dreaming.) So, all the rules you've learned about how you can tell when you're dreaming are actually just rules for how to tell when you're having a dream *within* the larger dream, as opposed to merely being in the larger dream. Maybe, in the larger dream, you actually *can* experience pain – i.e., you can dream pain – and so on.

When you think about it, it seems impossible to refute this kind of hypothesis. Any evidence you cite, any experience you have, the skeptic can just explain as part of the dream.

This is all leading up to the following skeptical argument:

1. You can have knowledge of the external world only if you can know that you're not dreaming.
2. You can't know that you're not dreaming.
3. Therefore, you cannot have knowledge of the external world.

You might want to pause and think about that argument. Is it right? What might be wrong with it?

Interlude: About René Descartes

Pretty much everyone who takes an introductory philosophy class has to learn about Descartes. He's sometimes called the founder of modern philosophy (where by "modern", we mean "the last 400 years". Seriously, that's how philosophers talk.) Here's what you need to know.

He was a French philosopher of the 1600s. He invented analytic geometry; he said "I think; therefore, I am"; and he wrote a very famous and great book called *Meditations on First Philosophy* ("the *Meditations*",

for those in the know), which philosophy professors commonly use to befuddle beginning college students.

In the *Meditations*, he starts by trying to doubt everything. He entertains scenarios like “Maybe I’m dreaming”, “Maybe God is deceiving me”, and “Maybe my cognitive faculties are unreliable.” He wants to find something that one cannot have any reason at all to doubt, so that he can build the rest of his belief system on that unshakable foundation. He first decides that nothing about the physical world is certain, given the skeptical scenarios. He then decides that his own existence, and the facts about his own present, conscious mental states, are impossible to doubt. So they should be the foundation for the rest of his belief system. So far, so good.

He then tries to prove that God exists, starting just from his idea of God. (This is where most people think the *Meditations* goes off the rails. But the arguments are too long and weird to detail here. But see §9.2.2 for one of them.) Then he argues that since God is by definition a perfect being, God cannot be a deceiver. Therefore, God would not have given Descartes inherently unreliable faculties; so, as long as Descartes uses his faculties properly, he can trust them. And therefore, the physical world around him must really exist.

Most importantly, if you want to avoid sounding like a yokel: His last name is pronounced like “day-kart”, not “dess-karts” as some freshmen are wont to say.

6.2.2. *The Brain-in-a-Vat Argument*

Here’s something else that could happen (maybe?). Let’s say scientists in the year 3000 have perfected technology for keeping a brain alive, floating in a vat of liquid. They can also attach lots of tiny wires to the brain, so that they are able to feed the brain exactly the same pattern of electrical stimulation that a normal brain receives when it is in a normal human body, moving around the world. (Electrical signals from your nerve endings are in fact what causes your sensory experiences.) They can also attach tiny sensors to the brain to detect how the brain thinks it is moving its body, so they can modify the brain’s experience accordingly – e.g., the brain “sees” its arm go up when it tries to raise its arm, and so on. They have a sophisticated computer programmed to give the brain the exact pattern of stimulation to perfectly simulate a normal life in a normal body.

This is an odd scenario, but nothing about it seems absurd or impossible. As far as we know, this could in principle be done. And if so,

the scientists could program, let's say, a simulation of an unremarkable life in the twenty-first century. They might even include in the simulation a funny bit where the brain has the experience of reading a silly story about a brain in a vat. The scientists have a good laugh when the brain thinks to itself, "That's silly; of course I'm not a brain in a vat."

So now, how do you know that *you* are not a brain in a vat right now? Again, it seems impossible to refute the scenario, because for any evidence that you try to cite, the skeptic can just explain that as part of the BIV (brain-in-a-vat) simulation. The logic of the skeptic's argument here is basically the same as that of the dream argument:

1. You can have knowledge of the external world only if you can know that you're not a BIV.
2. You can't know that you're not a BIV.
3. Therefore, you cannot have knowledge of the external world.

This is actually the most discussed argument for skepticism. Epistemologists have spent a lot of time trying to figure out what's wrong with it, since the conclusion seems pretty crazy to most of us. A small number of philosophers have endorsed the argument and become external-world skeptics.

6.2.3. *The Deceiving God Argument*

You can probably already guess how this one goes from the title. The skeptic asks you to consider the hypothesis that there might be an all-powerful being, similar to God, except that he wants to deceive you. This being can give you hallucinatory sensory experiences, false memories, and so on. There's no way of proving that there isn't such a being, because any evidence you cite could have just been produced by the deceiving god to trick you. You're then supposed to infer that you can't know anything about the external world, since everything you believe about the world *could* be a result of this being's deception.

6.2.4. *Certainty, Justification, and Crazyiness*

Note that the skeptic is not *completely* crazy. The skeptic isn't saying that any of these scenarios are actually *true* or even *likely*. That's not the issue. The issue isn't, e.g., whether there is in fact a desk in front of me now. The issue is whether I *know* that there is. Most skeptics say that the mere *possibility* that I'm dreaming, or that I'm a BIV, or that there's a deceiving god, means that I don't *know* that the table I see is real. And so the skeptic only has to claim that the skeptical scenarios are *possible*.

Now to be more precise. There are two kinds of external-world skeptic: **certainty skeptics**, and **justification skeptics**. The former say that you lack knowledge (of the external world) because it is not *absolutely certain* that your beliefs are true. The latter say that you lack knowledge because your beliefs are not even *justified*. What does this mean? Basically, a **justified** belief is one that makes sense to hold, that a reasonable person would hold, that represents what a person (rationally) *ought* to think in your situation.

Certainty skepticism is more common, but justification skepticism is much more interesting. It's interesting because if it turns out that our beliefs are not justified, then we should presumably change them. On the other hand, if our beliefs are merely uncertain but still justified, then we don't need to change anything important – we should keep holding our ordinary beliefs, keep using them to navigate the world and so on, but merely stop *calling* them “knowledge”. Who cares about that?

Now, how would the argument for justification skepticism go? Pretty much the same as the arguments above.

1. Your beliefs about the external world are justified only if you have some justification for believing that you're not a BIV.
2. You have no justification for believing that you're not a BIV.
3. Therefore, your beliefs about the external world are not justified.

Premise 1 here still seems true, just as much as in the BIV argument (“You can have *knowledge* of the external world only if you can *know* that you're not a BIV”).

Premise 2 is maybe less obvious ... but there's still a pretty obvious case for it, on its face. To have justification for denying that you're a BIV, it seems that you would need to have *at least some evidence* that you're not a BIV. But, as discussed above, it doesn't seem that you can have any such evidence. So it's not just that you can't be *absolutely certain* that you're not a BIV; it looks like you have *no reason at all* to believe that you're not a BIV.

Of course, you also have no evidence at all that you *are* a BIV. But the skeptic isn't claiming that you know you are a BIV, so there doesn't need to be any evidence that you are one. The skeptic is claiming that you *don't know whether* you are one or not. That's perfectly consistent with the fact that there is no evidence either way.

6.3. Responses to Skepticism

6.3.1. Relevant Alternatives

Now I'm going to start telling you about things that philosophers have said to try to avoid skepticism. I find many of them unsatisfying, but that's par for the course when you're dealing with a big philosophical problem.

So here's the first response: The skeptic has misunderstood something about language, about how the word "know" is used in English. (This is the certainty skeptic we're talking about now.) To illustrate the idea, here is an analogy. Imagine that you work at a warehouse that stores merchandise. All the merchandise in the warehouse was supposed to be moved out this morning, and it was in fact moved out, which you observed. Now your boss calls you up on the phone and asks, "Is the warehouse empty now?"

Clearly the right answer to give is "Yes." You should not answer "No" on the grounds that there is still some dust on the floor, or a spider web in one of the corners, or light bulbs in the ceiling fixtures. You shouldn't do that, because those are obviously not the kind of thing that the boss was concerned about. When he asked if the warehouse was empty, he meant "Has all the *merchandise* been moved out?", not "Is it a hard vacuum inside?" This example leads to the idea that "empty", in English, does not normally mean "having nothing whatsoever inside"; it means something like, "having nothing *of the relevant kind* inside", where the relevant kind is determined by the context of the conversation.

The skeptic is like the person who says the warehouse isn't empty because there's dust on the floor – except the skeptic is misunderstanding the word "know" rather than the word "empty". The skeptic thinks that to know a proposition, one must rule out every possible alternative whatsoever. In fact, though, this isn't how the word "know" works in English. To know something, in the standard English sense, it is only necessary to be able to rule out every **relevant alternative**. Or so the anti-skeptic would argue.

What are the relevant alternatives? Well, this is going to have something to do with what alternatives are close to reality, so that they really had a chance of being true, in some objective sense.^[26] Let's not worry about the details of what determines "relevance". The important points are that the relevant alternatives are a proper subset of the logically possible alternatives, and that the skeptical scenarios are generally viewed (by relevant-alternative theorists) as being too remote or far-fetched to be relevant.

It's easier to see the point if you describe an example from a third person point of view. Let's say there's a birdwatcher somewhere, and he sees a Gadwall Duck in a pond, which he correctly identifies. Now, there is a logical possibility that the bird might have instead been a species called a Siberian Grebe which looks just like a Gadwall when it's in the water. The birdwatcher himself has never heard of the Grebe, so he hasn't had any thoughts about this possibility. If it *had* been a Grebe, the birdwatcher would not have been able to tell the difference. (But, to repeat, we're stipulating that the bird was *in fact* a Gadwall.) Now, the question: Would we say that the birdwatcher "knows" that he saw a Gadwall Duck?

Well, it depends. If Siberian Grebes actually exist, and there are some nearby, then the birdwatcher does not know what he saw. He doesn't *know* because, even though it is in fact a Gadwall, it *could* easily have been a Grebe, and his evidence couldn't rule that out. On the other hand, if Siberian Grebes exist, but they are only ever found in Siberia, and the birdwatcher is very far from there, then it's plausible that the birdwatcher counts as knowing that he saw a Gadwall Duck. Even more clearly, if Siberian Grebes don't exist at all but are a purely made up possible species, then the birdwatcher is fine. He doesn't have to have evidence sufficient to rule out made up alternatives. This example motivates the idea that distant possibilities not tied to any real-world circumstances are not relevant when it comes to assessing knowledge claims.

Note that the relevant-alternatives (RA) theorist is not giving *an argument against* the skeptical scenarios. The RA theorist is saying that one *does not need* any argument against the skeptical scenarios in order to count as knowing stuff about the external world. (Compare: When I say the warehouse is empty, I'm not denying that there is dust on the floor. I'm saying that it was not necessary to remove the dust in order for the warehouse to count as "empty".)

* * *

Now I'm going to give you my take on the RA theory. I think it might be a fair response to the certainty skeptic. We'd have to get into some annoying details to see whether RA gives a correct analysis of "know". However, I'm not going to get into that, because the dispute between the RA theorist and the certainty skeptic is simply not very *interesting*. It's not very interesting because it's semantic: They're just disagreeing about the use of the word "know".

The interesting debate would be with the justification skeptic. Against that skeptic, I think RA theory fails. That's because the RA theory is all about how the word "know" works, but the justification skeptic's argument doesn't depend on that (see §6.2.4). The justification skeptic claims that our external-world beliefs are *unjustified*, which is bad enough.

Okay, what if we tried adopting a relevant-alternatives theory of *justification*? We could claim that a belief is justified as long as one has evidence against each of the relevant alternatives, and one need not have any evidence at all against the *irrelevant* alternatives, such as skeptical scenarios.

One *could* say that. I just don't think that is very plausible or helpful. I think that if a belief is to be justified, the belief must at least be very likely to be true. Now, if a belief is highly probable, then *any* alternative to it (whether "relevant" or not) must be highly *improbable* – that's just a theorem of probability. (If $P(A) > x$, and A entails $\sim B$, then $P(B) < (1 - x)$.) So, if our beliefs about the external world are justified, and they entail (as they do) that we're not brains in vats, then the probability of our being brains in vats must be very low.

But then, it seems that there must be some explanation of *why* the BIV scenario is very improbable given our current evidence. Whatever that explanation is, *that* would be the important response to give to the skeptic. We could then add that alternatives that are highly improbable are "irrelevant", but all the intellectual work would lie in showing that the skeptical scenarios are in fact highly improbable – and the RA theory wouldn't help us with that.

6.3.2. Contextualism

This is another semantic response to skepticism, closely related to the RA theory. **Contextualists** think, roughly speaking, that the meaning of "know" shifts depending on the context in which the word is used.

There are many words that work like this. For example, the word "here". Let's say that I'm at a philosophy conference. I heard that Daniel Dennett was going to be attending, and I want to talk to him (mostly to find out whether he is a zombie^[27]), but I don't know what he looks like. At the conference dinner, I go around to different tables, asking, "Excuse me. Is Daniel Dennett here?" In this context, by "here", I mean "at this table".

Now take another context. Say I'm visiting Fort Hays State University in Kansas. While I'm talking to one of the professors there, the professor

says, “Hey, did you hear that we’ve hired Daniel Dennett?” I respond (with a bit of surprise), “Wait, Daniel Dennett is *here* now?” In this context, “here” means “working at this university”. It does not mean “sitting at this table”. So the meaning of “here” shifts: In the first context, it includes a much smaller physical area than in the second context.

So maybe there is something like that with “know” – in different conversational contexts, we get more or less demanding standards for something to count as “knowledge”. What skeptics do is that they *raise the standards* for “knowing”. They do this mainly by talking about far-fetched skeptical scenarios and treating them seriously. In a conversation about skepticism, the standards are so high that almost nothing counts as “knowledge”, because “knowledge” in these contexts requires absolute certainty or something like that. Now, if we don’t realize that the standards shift with context, we might then be misled into thinking that we don’t know things in the *ordinary* sense of “know”, the sense that applies in normal contexts (outside of discussions of skepticism). When we stop talking about skepticism and have a more mundane conversation, the standards for “knowing” go back down, and so we are perfectly correct to say that we know all kinds of things about the external world.

You might think this is a conciliatory position: When skeptics say that we “don’t know” anything about the external world, they are correct; and yet, when in ordinary life you say that you “know” what the capital of Alaska is, you “know” how many people live in China, and so on, you are *also* perfectly correct. These seemingly incompatible claims can all be correct, as long as the meaning of “know” shifts.

But skeptics do not like this. They do not think they’re raising the standards for “knowledge”; what they think is that the standards for knowledge are always very high, and we don’t satisfy them, and thus our knowledge claims in ordinary life are false. Contextualism is thus a skeptic-unfriendly diagnosis of what’s going on.

* * *

My take on contextualism: There is some plausibility to it. If someone comes up to you on the street and asks, “Do you know what time it is?”, you can just look at your watch and answer “Yes.” (Then you should probably tell them the time.) There are low standards for knowing the time on the street. On the other hand, if you’re supposed to be checking the rocket that’s going to Mars, and if anything goes wrong the rocket is likely

to explode, then the requirements for “knowing” that the rocket is safe are higher – you need to be a lot more careful, you need to check and double-check everything, etc.

Again, though, I don’t think the contextualist response to skepticism is super-interesting, because it is too semantic – it just raises a dispute about the use of a particular word, “know” – and the response only applies to certainty skeptics who grant that the BIV scenario is highly improbable. Contextualism doesn’t tell us how to respond to a justification skeptic, and it doesn’t explain why the brain-in-a-vat hypothesis is unreasonable.

6.3.3. *Semantic Externalism*

Here’s another idea. Maybe the BIV hypothesis has to be rejected because it’s self-refuting.^[28]

Why would it be self-refuting? Maybe because, in order for us to have the concepts required to entertain the BIV hypothesis, we would have to have had some contact with the real world, which the BIV hypothesis says we have never had. So if we were BIV’s, we couldn’t be thinking about whether we were BIV’s. But the person who advances the BIV hypothesis presumably cannot think that we are not entertaining the very hypothesis that he is putting forward. So the hypothesis is self-undermining.

Of course, all the work is going to be in showing that we couldn’t have the concept of a BIV if we were BIV’s. Why is that?

First, we need to think about a property that philosophers call “**intentionality**”. This is the property of representing something, of being “of” or “about” something. (Note: This is a technical use of the word “intentionality”. It does *not* refer to the property of being intended by someone! Please don’t confuse intentions in the ordinary English sense with intentionality.) Examples: Words, pictures, and ideas in the mind all refer to things. When you have a picture, it is a picture *of* something; when you have an idea, it is an idea *of* something or *about* something.

When you think about this phenomenon of intentionality, a good philosophical question is: What makes one thing be *about* another? I.e., under what conditions does *x* refer to *y*? Of particular interest to us here: What makes an idea in your mind refer to a particular thing or kind of thing in the external world?

Here is a partial answer (partial because it only gives a necessary condition, not a sufficient condition): In order for an idea, *x*, to refer to an external phenomenon, *y*, there has to be the right kind of causal connection

between x and y . Example: You have certain visual sensations, such as the sensation of red. What this sensation represents is a certain range of wavelengths of light (or a disposition to reflect such wavelengths, or something like that). What makes your sensation count as referring to that physical phenomenon? There is no intrinsic similarity between the sensation and the underlying physical phenomenon. The answer is: The sensation refers to that range of wavelengths of light because those are the wavelengths of light that normally *cause* you to have that sensation when you look at things.

Here is a famous thought experiment (that is, famous among philosophers, which is really not very famous overall):

Twin Earth: There is another planet somewhere that is an almost exact duplicate of Earth. It has a molecule-for-molecule duplicate of every person on Earth, doing the same things that people here do, etc. There is just one difference between Twin Earth and our Earth: On our Earth, the rivers, lakes, clouds, and so on, are filled with the chemical H_2O . By contrast, on Twin Earth, the rivers, lakes, clouds, and so on are filled with a different chemical, which I will call XYZ. XYZ looks, tastes, feels, etc., just like H_2O . It's completely indistinguishable from H_2O to normal observation, though it has a different chemical formula. Now, assume that we're at a time before the chemical composition of water was discovered, say, the time of Isaac Newton. (Remember that people used to think that water was an element! The people on Twin Earth at that time also thought that their "water" was an element. The composition of water was only discovered in the 1800s.) Let's say Isaac Newton on Earth thinks to himself, "I want a cup of water." At the same time, Twin Isaac Newton on Twin Earth thinks to himself a corresponding thought, which he would also express by, "I want a cup of water." Question: What is our Isaac Newton referring to by his "water" thought? And what is Twin Isaac Newton referring to?

You're supposed to think that Newton is referring to H_2O (even though he does not *know* that this is what water is), because H_2O is in fact what fills the rivers, lakes, and so on around him. If someone gives him a glass filled with chemical XYZ, they would be tricking him (though he wouldn't know it): They'd be giving him something that looks like water but isn't real water.

At the same time, Twin Newton is referring to XYZ, not H₂O. If someone gives *Twin* Newton a glass filled with H₂O, they'll be tricking him.

Why does Newton's word "water" refer to H₂O, while Twin Newton's word "water" refers to XYZ? Answer: Because Newton's idea of water was formed by perceiving and interacting with rivers, lakes, etc., that were in fact made of H₂O. In brief, Newton's idea was *caused by* H₂O. Twin Newton's idea, on the other hand, was caused by interactions with XYZ. This is meant to show that the referents of your ideas are determined by what things you have actually perceived and interacted with, that caused you to form your ideas. This is known as the **causal theory of reference**.

By the way, you may think the Twin Earth scenario is pretty silly (like many philosophical thought experiments). A perfect duplicate of Earth is ridiculously improbable, and obviously there is no compound like XYZ. But philosophers generally don't care how improbable our scenarios are. We don't even really care if they violate the laws of physics or chemistry.

Still, why posit such an outlandish scenario? Well, the purpose is basically to rule out alternative theories about intentionality. Assuming that you agree that Newton is referring to H₂O, while Twin Newton is not, we have to find some *difference* between Newton and Twin Newton that could explain that. Since the two people are perfect duplicates of each other, with qualitatively indistinguishable experiences, the relevant difference cannot be anything in their minds, or even in their bodies.^[29] If we didn't make Twin Newton a perfect duplicate of Newton, then someone could have said that they are referring to different things because maybe their thoughts are intrinsically different, or something else in them is different.^[30]

Anyway, back to the BIV argument. If you buy the causal theory of reference, what would happen if there was a brain that only ever lived in a vat, and only had experiences fed to it by the scientists? This brain has never perceived, nor interacted in any normal way with, *any* object in the real world. So none of the BIV's concepts can refer to real-world things. All of the BIV's concepts are going to refer to virtual objects, or perhaps states of the computer that stimulates the brain, since that is what causes the BIV's experiences. E.g., when the BIV thinks, "I want a glass of water", it is referring to a *virtual* glass of water. It can't be referring to a real glass of real water, since it has no experience with such.

If that's so, what would the brain mean if it thought to itself, "I wonder whether I am a brain in a vat?" What would "brain in a vat" refer to? It would have to refer to a *virtual* brain in a vat, not an actual brain in a vat. The BIV cannot think about actual brains in vats.

Now there are two ways of formulating the argument against the BIV scenario. First version:

1. I'm thinking about BIV's.
2. A BIV cannot think about BIV's.
3. Therefore, I'm not a BIV.

The skeptic who advanced the BIV scenario can't very well deny (1), since the central point of the skeptic's argument is to make you think about BIV's. This makes the skeptic's argument self-undermining.

Here's the second version:

1. If a BIV thinks to itself, "I'm a BIV", that thought is false.
Explanation: By "BIV", it means a virtual BIV. But the BIV is not a virtual BIV; it's a real BIV. So the thought would be false.
2. If a non-BIV thinks to itself, "I'm a BIV", that thought is false.
3. So "I'm a BIV" is always false. (From 1, 2)
4. So I'm not a BIV. (From 3)

Notice that this response to skepticism does what the earlier responses avoided: It directly tries to show that you're not a BIV.

* * *

My take: One obvious problem is that the above response only applies to *some* skeptical scenarios. It can't be the case that *all* of your experiences to date have been BIV experiences, since that would prevent you from having a concept that refers to actual BIV's. However, this does nothing to refute the hypothesis that you were kidnapped just *last night* and envatted – you could have formed the concept of a brain in a vat *before* being turned into one.

Possible response: True, but if my life before last night was normal, then I can use my knowledge of the world gained up to that point to argue that humans do not actually possess the technology for making BIV's.

Counter-reply on behalf of the skeptic: Maybe after they kidnapped and envatted you, the scientists also erased all your memories of all the news items you read reporting on how we actually *do* have the technology for creating BIV's. They would have done this to trick you into thinking that you couldn't be a BIV.

Another problem with the response to skepticism is that the BIV would still have many important false beliefs. When it sees a glass of water on a table, perhaps, the BIV is not deceived, because it thinks there is a virtual glass of water, and that is what there is. But when the BIV talks to other people (that is, it has virtual conversations with virtual people), the BIV will be thinking that these “other people” are conscious beings just like itself – that they have thoughts, feelings, and so on, like the thoughts, feelings, and so on that the BIV itself experiences. But that will be false; they’re just computer simulations. And of course, a huge amount of what we care about has to do with other people’s minds. The skeptic will claim that all of that is doubtful. And so semantic externalism doesn’t do such a great job of saving our common sense beliefs.

6.3.4. *BIVH Is a Bad Theory*

When we started talking about responses to skepticism, you might have been hoping for an explanation of why the BIV Hypothesis is not a good theory for explaining our experiences, and why our ordinary, common sense beliefs are a better theory. At least, that’s what *I* was hoping for when *I* first heard the BIV argument. Yet almost all responses by philosophers try to avoid that issue. Given the confidence with which we reject the BIV Hypothesis (many people find it laughable), we ought to be able to cite some extremely powerful considerations against it. The BIV Hypothesis should be a *terrible* theory, and our ordinary, common sense world view (the “**Real World Hypothesis**”, as I will call it) should be *vastly* superior (or, as Keith DeRose once put it in conversation, the Real World theory should be kicking the BIV theory’s ass). So it’s pretty odd that philosophers seem to have trouble citing anything that’s wrong with the BIV Hypothesis as an explanation of our experience.

Sometimes, when you tell people (mostly students) about the BIV Hypothesis, they try claiming that the Real World Hypothesis should be preferred because it is “simpler”. I can’t figure out why people say that, though. It’s obviously false. In the Real World Hypothesis, there are *lots* of separate entities involved in the explanation of our experiences – motorcycles, raccoons, trees, comets, clouds, buckets of paint, etc., etc., etc. In the BIV Hypothesis, you only need the brain in its vat, the scientists, and the apparatus for stimulating the brain, to explain all experiences. Vastly simpler. Of course, there *may* be lots of other things in the world, but the BIV hypothesis makes no commitment regarding other things; it does not

need to cite them to explain our experiences. If we just care about simplicity, then the BIV theory is *vastly better* than the Real World theory!

Here is something else you might have heard about judging theories: A good theory must be **falsifiable**. That means there has to be a way to test it, such that if the theory were false, it could be proved to be false. The BIV theory is unfalsifiable: Even if you're not a BIV, there is never any way to prove that.

For the skeptic, though, this is a feature, not a bug. The skeptic would say, "Yeah, I know it's unfalsifiable. That was my *point* in bringing it up! How is that problem for *me*?" So now we have to explain what is wrong with unfalsifiable theories.

The idea of falsifiability was most famously discussed by the philosopher of science Karl Popper. Popper's idea was that falsifiability is essential to scientific theories. So if you advance an unfalsifiable theory, then you cannot claim to be doing science. But so what? The skeptic will just say, "Yeah, I never claimed to be doing science. Again, how is this a problem for me?" We need a better answer than "You're not doing science."

A better answer can be found in probability theory. The way a theory gets to be probabilistically supported is, roughly, that the theory predicts some evidence that we should see in some circumstance, we create that circumstance, and the prediction comes true. More precisely, evidence supports a theory provided that the evidence would be *more likely* to occur if the theory were true than otherwise. The theories that we consider "falsifiable" are those that make relatively sharp predictions: That is, they give high probability to some observation that is much less likely on the alternative theories. If those observations occur, then the theory is supported; if they don't, then the theory is disconfirmed (rendered less probable). "Unfalsifiable" theories are ones that make weak predictions or no predictions – that is, they don't significantly alter the probabilities we would assign to different possible observations. They allow pretty much any observation to occur, and they don't predict any particular course of observations to be much more likely than any other. (On this account, "falsifiability" is a matter of degree. A theory is *more* falsifiable to the extent that it makes more predictions and stronger predictions.)

Now there is a straightforward, probabilistic explanation of why falsifiability is important. (Popper, by the way, would *hate* this explanation. But it is nevertheless correct.) A highly falsifiable theory, by definition, is

open to strong disconfirmation (lowering of its probability), in the event that its predictions turn out false – but, by the same token, the theory is open to strong *support* in the event that its predictions turn out *true*. By contrast, an unfalsifiable theory cannot be disconfirmed by evidence, but for the same reason, it cannot be *supported* by evidence either. (This is a pretty straightforward point in probability theory.)

Suppose that you have two theories to explain some phenomenon, with one being much more falsifiable than the other. Suppose also that the evidence turns out to be consistent with both theories (neither of them make any false predictions). Then the falsifiable theory is *supported* by that evidence, while the unfalsifiable theory remains unsupported. At the end of the day, then, the highly falsifiable theory is more worthy of belief. This will be true in proportion as it is more falsifiable than the other theory.

All of that can be translated into some probability equations, but I'm going to spare you that, since I think most readers don't like the equations so much.

Now, back to the BIV theory versus the Real World theory. The Real World theory, which holds that you are a normal human being interacting with the real world, does not fit equally well with every possible sequence of experiences. The Real World theory predicts (perhaps not with certainty, but with reasonably high probability) that you should be having a coherent sequence of experiences which admit of being interpreted as representing physical objects obeying consistent laws of nature. Roughly speaking, if you're living in the real world, stuff should fit together and make sense. The BIV theory, on the other hand, makes essentially no predictions about your experiences. On the BIV theory, you *might* have a coherent sequence of experiences, *if* the scientists decide to give you that. But you could equally well have *any logically possible* sequence of experiences, depending on what the scientists decide to give you. You could observe sudden, unexpected deviations from (what hitherto seemed to be) the laws of nature, you could observe random sequences of colors appearing in your visual field, you could observe things disappearing or turning into completely different kinds of things for no apparent reason, and of course you might observe random program glitches. In fact, the overwhelming majority of *possible* sequences of experience (like, more than 99.999999999999%) would be completely uninterpretable – they would just be random sequences of sensations, with no regularities.

Our actual evidence is consistent with both theories, since we actually have coherent sequences of experience. Since the Real World theory is falsifiable and the BIV theory is not, the Real World theory is supported by this evidence, while the BIV theory remains unsupported.

* * *

My take: That's all correct. [\[31\]](#) The BIV theory is a very bad theory.

6.3.5. *Direct Realism*

“**Realism**” in philosophy generally refers to a view that says that we know certain objective facts. There's “realism” about different things – e.g., “moral realism” says that there are objective *moral* facts, which we can know; realism in the philosophy of perception (or, realism about the external world) says that there are objective facts in the external world, which we can know. In this chapter, we're interested in realism in the philosophy of perception, so that will be what I mean by “realism”.

Traditionally, there are two forms of realism: direct realism and indirect realism. It's easiest to explain **indirect realism** first. According to indirect realists, when we know stuff about the external world, that knowledge is always dependent upon our knowledge or awareness of something *in our own minds*. (Compare Descartes' idea that all knowledge must be built up from the knowledge of one's own existence and of the contents of one's own consciousness. See the discussion of him in §6.2.1 above.) For example, say you see an octopus. You aren't directly aware of the octopus; rather, you have an image of the octopus *in your mind*, which is *caused* by the octopus. The real octopus reflects light to your eyes, then a bunch of electrical activity goes on in your brain, which does a lot of information processing that you're not aware of, and then all this causes you to consciously experience the octopus image. You might mistake the image for the real object (that's what Hume said we were doing, anyway). But the image is the thing you're first aware of. Then you form the belief that there is an octopus in the outside world. Now, this is the crucial bit for our present purposes (remember, we're ultimately interested in addressing skepticism): On the indirect realist view, your belief about the physical octopus is *justified* on the basis of facts about *the image in your mind*. In that sense, your knowledge about the physical octopus is “indirect”. The only thing you know *directly* is the mental image. (Some indirect realists would say you are aware of appearances, or sensory experiences, or “sense

data”, or something else like that. The key point is that what you’re directly aware of is supposed to be something mind-dependent.)

(Terminological note: “Indirect realism” sometimes refers to the idea that we are directly aware of our mental states and only indirectly aware of the external world; sometimes, it refers to the idea that we have *non-inferential justification* for beliefs about our mental states and only inferential justification for beliefs about the external world^[32]; and sometimes, it refers to the conjunction of both theses. Here, let’s assume indirect realism includes both theses. Similarly, direct realism will include both the claim that we’re directly aware of external things and the claim that we have non-inferential justification for beliefs about the external world.)

Indirect realism, by the way, is by far the majority opinion in the history of philosophy, among philosophers who have addressed the issue at all. The alternative position (if you’re a realist) is **direct realism**: Direct realists think that, during normal perception, we have direct awareness of the external world. That is, we are aware of something in the external world, and that awareness is *not* dependent on the awareness of anything in our own minds. Also, direct realists think that we have immediate justification for at least some external-world beliefs. That is, we are justified in believing some things about the world around us, and that justification does not depend upon our knowing or having justification for believing anything about our own minds.

By the way, please don’t confuse direct realism with any of the following completely dumb views: (i) that there are no causal processes going on between when an external event occurs and when we know of it, (ii) that our perceptions are always 100% accurate, (iii) that we’re aware of external things in a way that somehow doesn’t involve our having any experiences. People sometimes raise objections to “direct realism” that are only objections to one of the above views (actually, that’s almost 100% of the objections). No one, however, holds those dumb views, so we’re not going to talk about them.

There are long arguments to be had about direct versus indirect realism. We’re not going to have those arguments now, though. I’m just going to say that there are interesting, plausible, and well-worth-reading defenses of direct realism (especially my own book, *Skepticism and the Veil of Perception!*). The point I want to talk about here is this: *If* you’re a direct

realist, then you have an escape from BIV-driven skepticism that is not available to the indirect realist. The skeptic's argument is really only an argument against indirect realists, not against direct realists. So now let me try to make out that point.

Indirect realists regard our beliefs about physical objects as something like theoretical posits – we start with knowledge of our own subjective experiences, and we have to justify belief in physical objects as the best explanation for those experiences. If you're doing that, then you have to be able to say why the belief in real, physical objects of the sort we take ourselves to be perceiving provides a *better explanation* than the theory that one is a brain in a vat, or that God is directly inducing experiences in our minds, or that we're having one really long dream, etc. The latter three theories (the skeptical scenarios) would seem to explain the evidence equally well, if "the evidence" is just facts about our own subjective experiences.

On the other hand, direct realists regard our perceptual beliefs as foundational: They do not need to be justified by any other beliefs. When you see an octopus, you're allowed to just start from the belief, "There's an octopus." You do not have to start from "Here is a mental image of an octopus." There is thus no need to prove that the Real World scenario is a better explanation for our experiences than the BIV scenario. Another way to phrase the point: According to indirect realists, our evidence is our experiences; but according to *direct* realists, our evidence consists of certain observable *physical facts*. For instance, the fact that there is purple, octopus-shaped object in front of you. Now, the BIV scenario is a competing explanation of *our experiences*, but it is not an explanation of the *physical facts* about the external world that we're observing. The BIV theory would explain, for example, the fact that you're having an octopus-representing mental image, but it does not even attempt to explain the fact that there is an octopus there. So if you regard our initial evidence as consisting of physical facts, then the BIV theory is a complete non-starter, as are all skeptical scenarios.

* * *

My take: Yep, this is also a good response to skepticism. I'm not going to defend direct realism right now, but I'll just mention that it is supported by a general theory about justified beliefs (the theory of Phenomenal Conservatism) that we're going to discuss in chapter 7.

You might think direct realism offers a cheap response to skepticism. It's almost cheating. *Of course* you can avoid skepticism if you get to just posit that external-world beliefs are immediately justified. I don't think this is cheating, though, because I think direct realism is the view that most of us start with before we run into the skeptic, and I think that what skeptics are trying to do is to refute our common sense views. No one starts out just being a skeptic for no reason. What's interesting is when the skeptic has an *argument* that seems to force us to give up our ordinary beliefs. But what we've just seen is that the skeptic only really has an argument against indirect realists. So maybe if you're an indirect realist, you should give up your beliefs (but see again §6.3.4). If you start out as a *direct* realist, as (I assume) most normal people do, then the skeptic hasn't given you any reason to change your belief system.

6.4. Conclusion

Most responses to external-world skepticism are unsatisfying. But that's okay; there are at least two responses that are pretty good. One response points out that the BIV argument only works against indirect realists, so if you're a direct realist, you're okay.

The other response works even for indirect realists. It argues that, due to its unfalsifiability, the BIV Hypothesis cannot be evidentially supported by our experiences. The Real World Hypothesis, however, can be and is supported, because it predicts a coherent course of experience, which is what we have.

7. Global Skepticism vs. Foundationalism

Philosophy may be the only field of study in which a major part of the discourse is arguing about whether the things we're studying even exist. Epistemologists are supposed to study knowledge, but we spend a good deal of our time talking about whether there is knowledge to begin with. Scientists don't do that – e.g., biologists don't spend their time arguing about whether there is any life.

Anyway, in this chapter, we consider arguments for **global skepticism**, the view that no one knows anything (not even facts about one's own mind, not necessary truths either, and not even the truth of skepticism itself!).

7.1. The Infinite Regress Argument

In order to know something to be true, it seems that you have to be justified in believing it, and that requires that you have some *reason* to believe it. For example, it could not be that you *know* that there are unicorns living on Mars, unless you have at least some reason to believe that. If you believe it for no reason at all, then even if it turns out to be true, the belief will merely have been a lucky guess, not knowledge.

Furthermore, it seems that the reason for your belief must itself be something that you know, or at least are justified in believing. But this leads to an infinite regress: Say you know A. Then there must be a reason for A, call it B. There must also be a reason for B, which we can call C. There must then be a reason for C. And so on.

But we can't complete an infinite reasoning process, nor do we have an infinite number of distinct beliefs that we could use to supply reasons. Furthermore, circular reasoning is fallacious, so, as the series of reasons goes on, we may not repeat any reason cited earlier in the chain. So we're screwed. (That's a technical term. In this context, it means that we have no way of acquiring knowledge.) Here's a concise statement of the argument:

1. You can know P only if you have a reason for P.
2. A chain of reasons must have one of three structures:
 - a. It's infinite.
 - b. It's circular.
 - c. It ends in one or more propositions for which there are no reasons.

3. You can't have an infinite series of reasons.
4. Circular reasoning can't generate knowledge.
5. If you don't know a premise, then you can't know anything on the basis of that premise.
6. Therefore, you can't know anything.

Something close to this argument goes back to the ancient skeptics (especially Agrippa and Sextus Empiricus), and it is frequently rediscovered independently by undergraduate philosophy students.

7.2. The Reliability Argument

When we acquire our putative knowledge, we do so using one or more *cognitive faculties*. These are faculties that are supposed to give us knowledge, such as vision, hearing, taste, touch, smell, memory, introspection, reasoning, and intuition. Maybe you have other things you'd like to add to that list, or maybe you'd like to remove some of the items on it. Let's not worry too much about the proper list of cognitive faculties; that's not going to matter for our purposes.

Now, here is something that seems like a plausible requirement for knowledge: Since all your putative knowledge comes from one or more of these faculties, it seems that you should first verify that your faculties are reliable, before you rely on them. If you don't know whether your faculties are reliable, then you can't really know whether any of the beliefs that you form using them are correct.

Here is an analogy. I have a toy called a Magic 8-Ball. It looks like a large, plastic 8-ball, and it's meant to be used like this: You hold the 8-Ball in your hand, ask it a yes-no question, and then turn it over. An answer to your question floats up to a window in the bottom. Possible answers include, "Yes, definitely", "My sources say no", "Outlook good", and so on. Now suppose I were forming lots of beliefs using the Magic 8-Ball. It seems that this would not be a way of acquiring knowledge. I wouldn't actually *know* any of the 8-Ball beliefs to be true, because I have no reason to believe that the 8-Ball is a reliable source of information. I have to first verify the 8-Ball's reliability, *then* I can use it to find out other things.

Furthermore, if I wanted to verify the 8-Ball's reliability, I obviously cannot simply ask the 8-Ball. That would be what philosophers call "**epistemic circularity**" (using a method or source to test itself). I need an independent source of information to check on the 8-Ball. And of course I must already know that other source to be reliable.

You can see where this is going. Just as with the Magic 8-Ball, in order for me to know anything using any of my natural cognitive faculties, I must first verify that my faculties are reliable. And I can't use a faculty to verify its own reliability; I must have an independent source. But I don't have an infinite series of faculties, and I can't rely on epistemic circularity (e.g., using two faculties to "verify" each other's reliability). So once again, I'm screwed.

Take, for example, the five senses. How do we know we can trust them? You could try taking an eye exam to see if your vision is good. But to collect the results of your exam, you would have to either use vision (to read the results) or use your sense of hearing (to hear the doctor telling you the results). These things only work if you already know you can trust vision or hearing.

Even more troublingly, suppose we ask how we know that *reason* is reliable. We could try constructing an argument to show that reason is reliable. But if we did that, we would be using reason to verify reason's reliability. There is not really any way around this. So once again, it seems that there is no way for us to know anything.

7.3. Self-Refutation

The first thing that comes into most people's minds after hearing the thesis of global skepticism is that global skepticism is self-refuting. There are better and worse versions of this objection. A bad version: "Global skeptics claim to know that we know nothing. But that's contradictory!" Reply: No, they don't. Global skeptics may be crazy, but they are not stupid. They say that we know nothing; they don't say that anyone *knows* that.

Slightly better version: "Okay, they don't explicitly say that we know global skepticism to be true. But they *imply* this. Because whenever you make an assertion, you are implying that you know the thing that you're asserting. That is why, e.g., if you say, 'Joe Schmoe is going to win the next election,' it is totally appropriate for someone to ask, 'How do you know?' It's also why it sounds nonsensical to say, 'I don't know who is going to win the election, but it's going to be Joe Schmoe.'"

Reply on behalf of the skeptic: So maybe there is this rule of our language which says that you're not supposed to assert P unless you *know* P. Then the skeptic can be justly charged with violating the social conventions and misusing language. (He's only doing that, though, because

our language provides no way of expressing his view without violating the rules. Language was invented by non-skeptics for use by non-skeptics.) Big deal, though. That doesn't show that the skeptic is substantively wrong about any philosophical point.

Counter-reply: "No, it's not just a linguistic convention that the skeptic is violating. There is an inherent norm of *rational thought*. That's why it seems nonsensical or irrational to think – even silently to oneself – such things as, 'Joe Schmoe is going to win the next election, but I don't know who is going to win.' It is likewise irrational to think, 'Global skepticism is true, but I don't know whether global skepticism is true.'"

That counter-reply seems pretty reasonable to me. Anyway, here is another version of the self-refutation charge: What exactly was supposed to be going on in sections 7.1 and 7.2? The skeptic gave *arguments* for global skepticism. An argument is an attempt to justify a conclusion. That's the main thing about arguments. (And by the way, if the skeptic *didn't* give any arguments, then we wouldn't be paying any attention to skepticism in the first place.) So, if the skeptic's arguments are any *good*, they are counter-examples to their own conclusions. The arguments are supposed to show that we lack knowledge *because* we can never justify our beliefs. But if the arguments show that, then we *can* justify at least some beliefs, because those very arguments justify the skeptic's belief in skepticism.

If, on the other hand, the skeptic's arguments are not good and don't show anything, then presumably we should disregard those arguments.

Finally, there is a general norm of rationality that one should not hold unjustified beliefs (the skeptic is relying on this norm to get us to give up our common sense beliefs). But since, again, the skeptical arguments claim that *no* belief is justified, this would mean that we should not believe either the premises or the conclusions of those arguments. So the arguments are self-defeating.

This objection to skepticism is so obvious that the (very few) skeptics in the world cannot have failed to notice it. Usually, their response is something along these lines: "Yep, that's right: Global skepticism itself is unjustified. I never said it was justified. I only said it was true."

It's hard to see how this is supposed to address the objection at all, though. It's really just granting the objection and then moving on as if granting the objection is the same as refuting it. The best I can figure is that the skeptics who say things like this are assuming that there is only one

possible objection that someone might be making, and that would be to claim that skepticism is literally an explicit contradiction, i.e., a statement of the form “A & ~A”.

But that’s not the objection. The objection is that skepticism is irrational for the reasons stated above; none of those reasons are rebutted by merely agreeing that skepticism is irrational.

7.4. The Moorean Response

The **Moorean response** to skepticism, also known as “the G.E. Moore shift”, was pioneered by the twentieth-century British philosopher G.E. Moore.^[33] I’m going to illustrate it for the brain-in-a-vat argument, but it works for any skeptical argument. Consider the following three propositions:

A. I know that I have hands.

B. To know that I have hands, I must know that I’m not a brain in a vat.

C. I don’t know that I’m not a brain in a vat.

Each of those propositions has some **initial plausibility**. That is, before hearing arguments for or against any of them, each of them (at least sort of) sounds correct. But they are jointly incompatible (they can’t all be true). Therefore, we have to reject at least one of them.

The skeptic thinks we should reject (A) because it conflicts with (B) and (C). That is the point of the BIV argument (§6.2.2). However, one could instead reject (B) on the grounds that it conflicts with (A) and (C), or reject (C) on the grounds that it conflicts with (A) and (B). We have to think about which of the three logically consistent options is most reasonable. Just because you heard the skeptic’s view first, doesn’t mean that it is the most reasonable.

Plausibility comes in degrees: Among propositions that are initially plausible, some are *more* plausible (they are more obvious, or more strongly seem correct) than others. So, if you have an inconsistent set of propositions that each seem plausible, you should reject whichever proposition has the *lowest initial plausibility*. Surely you shouldn’t reject something that’s *more* plausible, in order to maintain a belief that is *less* plausible; that would be unreasonable.

Now, it is just *extremely* initially plausible (it seems totally obvious to almost everyone) that a person in normal conditions can know that they have hands. It is not *as* obvious that a person can’t know they’re not a BIV, or that knowing one has hands requires knowing one isn’t a BIV. So the

latter two assumptions ((B) and (C) above) would each be more reasonable to reject. The skeptic's approach is actually the *least* reasonable option: The skeptic is rejecting the *most* initially plausible proposition out of the inconsistent set, rather than the least initially plausible one.

As I say, the Moorean response can be applied to pretty much any skeptical argument. When you look at skeptical arguments, you see the same thing with all or nearly all of them: The skeptic keeps asking us to reject the *most* initially plausible proposition out of the inconsistent set. Prior to considering skeptical arguments, such propositions as "I know I have hands", "I know how much two plus two is", and "I know I exist" are pretty much *maximally* initially plausible – i.e., you can't find anything that more strongly seems correct. Yet those are the sort of propositions that skeptics want us to reject. Meanwhile, the skeptics' premises generally include abstract, theoretical assumptions that are much less obvious. In the case of the Regress Argument, these include the assumptions that knowledge always requires reasons, that circular reasoning is always illegitimate, and that infinite series of reasons are impossible. In the case of the Reliability Argument, it includes the assumptions that all knowledge is produced by cognitive faculties, that all such knowledge requires prior verification of the faculty's reliability, that epistemic circularity is always unacceptable, and that there are only finitely many faculties. Plausible as each of these may be, they are not *as* plausible as the proposition that I know I exist. Certainly "I know I exist" would not be the least plausible proposition in any of these sets.

* * *

Now, despite that global skepticism is self-defeating and implausible, it is still an interesting topic of discussion, and more needs to be said about it. The reason we (that is, most epistemologists) are interested in skepticism is not that we think we have to figure out whether skepticism is true. It's obviously false. (*Pace* the handful of skeptics out there.) The reason we're interested in it is that there are some initially plausible premises, each of which is accepted by many, which lead to this unacceptable conclusion. So the problem is to figure out exactly what went wrong. Which premise is wrong, and why? That should shed some light on the nature of knowledge and rationality.

7.5. Foundationalism

7.5.1. The Foundationalist View

There are a few different responses to the regress argument. Some people think that circular justification is sometimes okay, sort of. A few people think that an infinite series of reasons is possible, sort of. But I'm not going to discuss those right now. There is one reaction that is by far the dominant reaction, both in the history of philosophy and among philosophy students and others who hear about the regress argument. The reaction is **foundationalism**.

Foundationalism rejects premise 1 in the regress argument:

1. You can know P only if you have a reason for P.

Foundationalists think that there are certain items of knowledge, or justified beliefs, that are “foundational” (also called “immediately justified”, “directly known”, “self-evident”, “non-inferentially justified”). Here, I'll talk in terms of *justification* rather than knowledge, for ease of exposition. (Foundational knowledge can be simply defined as knowledge whose justification is foundational.) **Foundational justification** is, by definition, justification that does not rest on reasons. In other words, sometimes you can rationally believe something in a way that doesn't require it to be supported by any other beliefs. Foundationalists think that some justification is foundational, and all *other* justification depends on support from foundational beliefs.

What would be examples of foundational propositions? “I exist” is typically viewed as foundational; also, propositions describing one's own present, conscious mental states (e.g., “I am in pain now”); also, simple necessary truths (e.g., “2 is greater than 1”). The great majority of foundationalists would accept all those examples. There is controversy about other examples. *Some* foundationalists (the direct realists) would add that propositions about the physical world around one are also foundational, when one directly observes them to be the case (e.g., “there is a round, red thing in front of me now”).

7.5.2. *Arguments for Foundationalism*

Why would one believe in foundationalism? One reason is the regress argument: We have knowledge, we can't acquire it by either infinite or circular chains of reasoning, so it has to be that we have some starting knowledge that doesn't require reasons. That was Aristotle's argument. (Notice that the foundationalist's regress argument is closely related to the skeptic's regress argument; the foundationalist and the skeptic have simply

chosen different propositions to reject out of the same set of jointly inconsistent propositions.)

Here's the other reason: Just think of some examples. When you think about the paradigm examples of putatively foundational propositions, it just seems that you can know those things directly; you don't have to infer them from something else.

Example: Say I go to the doctor. "Doctor," I say, "I think I have arthritis." The doctor asks, "Why do you believe you have arthritis?" Now, that is a completely reasonable question to ask; I need a reason to think I have arthritis. So I give my reason: "Because I'm feeling a pain in my wrist." And now suppose the doctor responds, "And why do you believe that you're feeling pain?"

Though his first question was reasonable, this second one is simply bizarre. If someone actually asked me that, I'm not sure how I should respond. I'd probably assume that either I'd misunderstood the question or the person was making a strange joke. If you asked an ordinary person this sort of question, they would probably either confusedly ask, "What do you mean?" or else just indignantly insist, "I know I'm in pain!" They wouldn't start citing evidence for their being in pain.

On its face, then, premise 1 of the skeptic's argument –

1. One knows that P only if one has a reason to believe P.

is unmotivated. It may sound plausible when stated purely in the abstract (which would probably be because *most* knowledge requires reasons). But when you think of cases like the belief that one is in pain (when one is in fact consciously experiencing that pain), (1) just doesn't seem plausible at all, at least not if a "reason" for P has to be distinct from the fact that P. It is unclear why anyone would believe (1).

This is worth remarking on, by the way, because this sort of thing is very common in philosophy. Some generalization sounds plausible when stated purely in the abstract, before you start thinking about all the cases that it subsumes. But then when you start thinking about specific cases, the generalization seems obviously false as applied to certain cases. When that happens, some philosophers stick to their initial intuition formed when thinking in the abstract. They may contort themselves trying to avoid the implications for particular cases, or they may just embrace all the counter-intuitive consequences.

Other philosophers, the rational ones, quickly reject the generalization and move on.

7.5.3. *The Argument from Arbitrariness*

Why should we believe the skeptic's premise (1)? The skeptic can't claim that it's self-evident, since the premise itself tells you that there is no such thing as a self-evident claim. (Anyway, (1) just doesn't look self-evident. It's highly controversial, as indicated above.) If you think you can just see (1) to be true (as people sometimes think), you're having a self-refuting thought. So the skeptic needs an *argument* for (1).

Skeptics have one main argument for (1). It goes something like this:

- 1a.If one lacks reasons for believing P, then P is arbitrary.
- 1b.Arbitrary propositions are not justified. Therefore,
- 1c.If one lacks reasons for P, then P is unjustified.

This is an extremely common argument among both professional and amateur skeptics. (Usually, it's not stated that explicitly, but the use of the specific word "arbitrary" is very common.)

What does the skeptic mean by "arbitrary"? They seldom say; in fact, I don't think I've ever heard a skeptic explain that. But I can think of three interpretations:

First interpretation: "Arbitrary" means "unjustified". In this case, the skeptic is begging the question in a very obvious way: Premise (1a) is just a paraphrase of (1c), so it can't be used to argue for (1c).

Second interpretation: "Arbitrary" means "not supported by reasons". In this case, premise (1b) is just a paraphrase of (1c), so again we have a circular argument.

Third interpretation: "Arbitrary" describes propositions that do not have any descriptive feature that distinguishes them from unjustified propositions. Notice that this is not quite the same as the first interpretation. In this interpretation, the argument is telling us that, in order for one proposition to be justified and another unjustified, there has to be some factual difference that *explains* why one is justified and the other not. The skeptic thinks that there can't be any such difference, if the putatively justified proposition isn't supported by reasons. In other words, the skeptic thinks the only feature of a belief that might explain its justification is the feature of *being supported by reasons*.

The third interpretation is the only one I can think of whereby the skeptic is not blatantly begging the question. Instead, the skeptic is merely making a false assumption – the assumption that all beliefs that aren't supported by reasons are relevantly alike, so that there is nothing to distinguish a “foundational” belief from a belief in any randomly chosen proposition. The skeptic would say, for instance, that if we may believe things that we lack reasons for, then I can just decide to believe that purple unicorns live on Mars. Why can't I just declare that to be “foundational”?

There are a variety of different forms of foundationalism, which give different accounts of which propositions are foundational. To illustrate, consider Descartes' view: The foundational propositions are the propositions that correctly describe one's present, conscious mental states. E.g., if you're presently in pain (which I assume is a conscious state), then you can have foundational knowledge that you're in pain. Perhaps also, simple necessary truths that one fully understands count as foundational, like “ $1+1=2$ ”. That is a very traditional foundationalist view, and it obviously does not allow [Purple unicorns live on Mars] to count as foundational.

By the way, I'm not saying that is *the correct* view. There are other (and better) foundationalist views. That is just to illustrate the general point that the foundationalist theories people have actually held are not subject to the charge of arbitrariness. None of them endorse just any randomly chosen belief.

7.5.4. Two Kinds of Reasons

I have an idea about why skeptics make the mistake that I just criticized. It's easy to conflate two kinds of “reasons”:

(a) S's reason for believing P.

(b) The reason why S's belief that P is justified.

How do (a) and (b) differ? *S's reason for believing P* has to be another proposition that S believes, it must seem to S to support P, and S's belief in that other proposition must (at least partly) cause S's belief that P. None of that has to be true of *the reason why S's belief that P is justified*. The reason why S's belief is justified is simply a fact that explains why the belief is justified. S does not have to know that fact, it needn't seem to S to support P, and, even if S happens to believe in that fact, S's belief in that fact needn't cause S's belief that P.

Example: Sue sees an empty cat food bowl, which she remembers filling earlier in the day. Sue infers from this that the cat has eaten. Let C = [The cat has eaten]. Sue's *reason* for believing C is:

[I filled the cat food bowl earlier today, and now it is empty].

On the other hand, *the reason why* Sue's belief that C is justified (assuming that it is) is something more like this:

[Sue inferred C from another belief that seemed to support it, that other belief was justified, and Sue does not have any reasons for doubting the reliability of the inference].

That whole thing in brackets has to be *true* for Sue to be justified in believing C. But Sue herself does not have to *believe* it, or have justification to believe it, or have inferred C from it. Sue just needs evidence for C itself; she doesn't need evidence that she herself is justified in believing C.

Now take an example of a putatively foundational belief. Let's say Sue has a headache, which is very noticeable to her. She believes proposition P: [I am in pain] (where the "I" of course refers to herself, not to me). What is *her reason* for this belief? It's plausible that she doesn't have any *reason* for believing she's in pain; she just is immediately aware of the pain. But it doesn't follow that there isn't any *reason why* her belief is justified. Here is why her belief is justified: because she is having a conscious pain. Plausibly, the fact that someone is in pain explains why it's reasonable for them to think that they are in pain. (Notice that this couldn't be described as *her reason* for the belief, because that would ascribe circular reasoning to Sue – she'd have to infer "I am in pain" from "I am in pain".)

So, I suspect that the global skeptic (at least the one who makes the "arbitrariness" argument) might be conflating these two kinds of reasons, and thus in effect assuming that if a person doesn't have a reason for a belief, then there can't be any reason why the belief is justified, and hence that the belief isn't justified.

Okay, that's enough about the regress argument for skepticism. Now let's turn to the reliability argument from §7.2.

7.5.5. A Foundationalist Reply to the Reliability Argument

The skeptic thinks that, to know P using some belief-forming method (or faculty) M, one must first verify that M is reliable. This leads to either an infinite regress or an epistemic circularity problem. How would foundationalists avoid this?

Most foundationalists would distinguish between different kinds of belief forming methods. The foundationalist is going to have some account of the belief-forming methods that generate foundational knowledge. Call these the “**foundational methods**”. Examples might include introspection, intuitions about simple necessary truths, and memory. If you’re a direct realist, you would include observation by the five senses as another foundational method of forming beliefs. Now, if you’re using a foundational method, then you do *not* need to first verify that it is reliable. You get to just rely on it, unless and until you acquire specific reasons for *doubting* that it is reliable. For instance, if you introspectively observe that you’re in pain, then you get to believe that you’re in pain (i.e., this belief would be rational or justified for you). You don’t have to first construct an argument that introspection is reliable.

Notice, by the way, that this just follows from the core tenet of foundationalism. If some method M generates foundational beliefs, then by definition, the person using M does not need to gather any other evidence in order for the M-generated beliefs to be justified. That’s just what “foundational” means.

On the other hand, if you’re using a *non*-foundational method, then the skeptic’s premise would apply – then you must first verify that the method is reliable. So if you’re going to be using the Magic 8-Ball to form beliefs, you have to first gather evidence of the 8-Ball’s reliability before you can know anything using the 8-Ball. The skeptic’s mistake is that of overgeneralizing: *Most* possible belief-forming methods are non-foundational, so if you look at some randomly chosen belief-forming method (like Magic 8-Ball reasoning), it’s going to seem plausible that the reliability of the method has to be independently verified. Skeptics mistakenly generalize from there to *all* belief-forming methods, applying what is true of non-foundational methods to the foundational methods as well.

7.6. Phenomenal Conservatism

7.6.1. *The Thesis of Phenomenal Conservatism*

Phenomenal Conservatism (PC) is a version of foundationalism that I have defended elsewhere (see my book, *Skepticism and the Veil of Perception*). PC holds that *appearances* (mental states wherein something seems to you to be the case) are the source of foundational justification. (The word “phenomenal” derives from the Greek word *phainomenon*,

which means “appearance”. That’s why I named the view “phenomenal conservatism”). Note that an appearance is not to be confused with a *belief*, since it is possible for a person to either believe or not believe that things are the way they appear.

There are several species of appearances, including at least the following:

- a. **Sensory experiences**, the experiences you have when you see, hear, taste, touch, or smell things. (Hallucinations and illusions also count as “sensory experiences”).
- b. **Memory experiences**, the experiences you have when you seem to remember something.
- c. **Introspective appearances**, the experiences whereby you are aware of your own present, conscious mental states. Note: These appearances, unlike all other appearances, need not be separate from the things they represent. The appearance of being in pain, for example, may just be the actual pain.
- d. **Intuitions**, the experiences you have when you reflect on certain propositions intellectually, and they seem correct to you.

Examples: It seems to me that there is a table in front of me (sensory experience), that I ate a tomato today (memory experience), that I am happy (introspective appearance), and that nothing can be completely yellow and also completely blue (intuition).

PC does not hold that all appearances are in fact true. There are such things as illusions, hallucinations, false memories, and so on. Nevertheless, according to PC, the rational *presumption* is that things are the way they appear, unless and until you have specific grounds for doubting that. These grounds would themselves have to come from *other* appearances, since appearances are the only foundational source of justification. For instance, if you submerge a stick halfway in water, the stick will *look* bent (visual appearance). But if you *feel* the stick, you can feel that it is straight (tactile appearance). Since you consider the tactile appearance more trustworthy, you reject the initial visual appearance. You would not, however, reject an appearance for no reason; there must be some other appearance that shows the first appearance to be defective.

7.6.2. *The Self-Defeat Argument*

I’m going to tell you my favorite argument for PC. Other philosophers don’t like it as much as I do, but I continue to think it’s a great argument. I

claim that alternative theories to my own are self-defeating.

Think about how you actually form beliefs when you're pursuing the truth. You do it based on what *seems true* to you. Now, there are some cases where beliefs are based on something else. For instance, there are cases of wishful thinking, where someone's belief is based on a *desire*; you believe P because you *want* it to be true. But those are not the cases where you're seeking the truth, and cases like that are generally agreed to be unjustified beliefs. So we can ignore things like wishful thinking, taking a leap of faith, or other ways of forming unjustified beliefs. With that understood, your beliefs are based on what seems right to you.

You might think: "No, sometimes my beliefs are based on *reasoning*, and reasoning can often lead to conclusions that initially seem wrong." But that's not really an exception to my claim. Because when you go through an argument, you're still relying on appearances. Take the basic, starting premises of the argument – by stipulation, we're talking about premises that you did not reach by way of argument. (There must be some such, else you would have an infinite regress.) To the extent that you find an argument persuasive, those premises *seem* correct to you. Each of the steps in the argument must also *seem* to you to be supported by the preceding steps. If you don't experience these appearances, then the argument won't do anything for you. So when you rely on arguments, you are still, in fact, relying on appearances.

Notice that all this is true of *epistemological* beliefs just as much as any other. For instance, beliefs about the source of justification, including beliefs about PC itself, are based on appearances. The people who accept PC are those to whom it seems right. The people who *reject* PC do so because it doesn't seem right to them, or because it seems to them to conflict with something else that seems right to them.

Now, in general, a belief is justified only if the thing it is based on is a source of justification. So if you think that appearances are not a source of justification, then you have a problem: Since that belief itself is based on what seems right to you, you should conclude that your own belief is unjustified. That's the self-defeat problem.

If you want to avoid self-defeat, you should agree that some appearances (including the ones you're relying on right now) confer justification. If you agree with that, it is very plausible that the appearances

that confer justification are the ones that you don't have any reasons to doubt – which is what PC says.

You might try adding other restrictions. Suppose, e.g., that you said that only abstract, intellectual intuitions confer justification, and sensory experiences do not. (External world skeptics might say that.) You could claim that this view itself is an intuition, not something based on sensory experience, so it avoids self-defeat. It is, however, pretty arbitrary. If you accept one species of appearances, why not accept all? There is no obvious principled rationale for discriminating.

Some philosophers hold that appearances provide justification for belief, but *only* when one first has grounds for believing that one's appearances in a particular area are reliable. E.g., color appearances provide justification for beliefs about the colors of things, *provided* that you know your color vision is reliable.

I disagree; I don't think one first needs grounds for thinking one's appearances are reliable. I think we may rely on appearances as long as we *don't* have grounds for thinking they *aren't* reliable. Why do I think that? See §§7.2, 7.5.5 above. If you require positive evidence of reliability, then you're never going to get that evidence, for the reasons given by the skeptic (the threat of regress or epistemic circularity).

7.6.3. *PC Is a Good Theory*

Anyway, PC is a good epistemological theory because it provides a simple, unified explanation for all or nearly all of the things we initially (before encountering skeptical arguments and such) thought were justified. It accounts for our knowledge of the external world, our knowledge of mathematics and other abstract truths, our knowledge of moral truths, our knowledge of the past, and so on. These are all things that philosophers have had a hard time accounting for, and it is *very* hard to find a theory that gives us *all* of them. At the same time, it is not overly permissive or dogmatic, because it allows appearances to be defeated when they conflict with other appearances. The theory seems to accord well with how we form beliefs when we are seeking the truth, and also with how we evaluate other people's beliefs.

Like all forms of foundationalism, PC avoids the skeptical regress argument by rejecting the skeptic's first premise. If it seems to you that P, you do not need a reason to believe P; you can presume its truth until you get a reason to doubt it. Note that we do not consider the seeming that P

itself to constitute a “reason”, because “reasons”, as we understand them here, have to involve other beliefs, and an appearance is not a *belief*. The appearance, in turn, is not the sort of thing that one could have reasons for, or that could be either justified or unjustified, since it is just an experience that one undergoes. (Compare: If you see a flash of light, that visual experience cannot be “justified” or “unjustified”, nor can it be based on a reason.)

7.7. Conclusion

Global skeptics think that we know nothing because we cannot complete an infinite chain of reasoning, and because we cannot verify the reliability of all our cognitive faculties before using them. This, however, is unreasonable. Besides being self-refuting, the skeptic’s arguments ask us to give up the *most* initially plausible of an inconsistent set of propositions, rather than giving up the least plausible as a rational person would do.

The leading alternative view is foundationalism, which holds that some propositions can be known or justified directly, and hence need no reasons. Probably the best version of foundationalism is Phenomenal Conservatism, which says that we are entitled to presume that whatever *seems* to us to be the case is in fact the case, unless and until we have reasons to think otherwise. We normally form beliefs in accordance with this principle all the time, including when we’re evaluating this very theory. The theory of Phenomenal Conservatism also accounts for all the beliefs we normally consider justified, which is otherwise really hard to do.

8. Defining “Knowledge”

8.1. The Project of Analyzing “Knowledge”

Since we’re talking about the theory of knowledge, maybe we should define “knowledge”. Unfortunately, a lot of people have tried to do this, and it’s a lot harder than it sounds. I’d like you to put up with some complexities in the next three sections, so that we can get to some interesting lessons at the end.

The goal here is to correctly analyze our current concept of knowledge, or the current use of “know” in English. To correctly analyze a term, you have to give a set of conditions that correctly classifies objects in all possible circumstances. Thus, if someone gives an analysis of “know”, it is considered legitimate to raise any conceivable scenario in which someone could be said to “know” or “fail to know” something, and the analysis has to correctly tell us whether the person in the scenario counts as knowing. In assessing this, we appeal to linguistic intuitions – that is, if normal English speakers would (when fully informed) call something “knowledge”, then your analysis should classify it as knowledge; if not, then not.

Aside: Analysis & Analytic Philosophy

Since the early 20th century, there’s been a style of philosophy known as **analytic philosophy**. Analytic philosophers emphasize clarity and logical argumentation (like this book!). At its inception, analytic philosophy was also largely devoted to analyzing the meanings of words. (They had a bad theory according to which this was the central job of philosophers.) Since then, “analytic philosophers” have drifted away from that emphasis, but there’s still a good deal of attention paid to word meanings.

This might seem unimportant to you – who cares about semantics? Why not just stipulate how *you* intend to use a word, and forget about the standard English usage? There are three reasons for not doing that. First, this causes confusion for other people who are familiar with the ordinary English use of the word.

Second, ordinary usage usually serves important functions. Human beings, over the millennia, have found certain ways of grouping and distinguishing objects (that is, certain **conceptual schemes**) to be useful and interesting. These useful conceptual schemes are embodied in our language.

Current usage reflects, in a way, the accumulated wisdom of many past generations.

Third, it is actually almost impossible to escape from the conceptual scheme that you've learned from your linguistic community. If you use a common word, such as "know", it is almost impossible to not be influenced in your thoughts by the actual usage of that word in your speech community. People who try to come up with new concepts usually just confuse themselves; they sometimes use the word in the new way they invented, but then slip back into using it in the normal way that others in their speech community use it.

So, all of that is to defend analytic philosophers' interest in the current, actual usage of "know" in our language. It also further explains, by the way, why I hate the technical uses of "valid" and "sound" in philosophy (see §2.8).

The main contrast to analytic philosophy is a style of philosophy known as **continental philosophy**, mainly practiced in France and Germany, which puts less emphasis on clear expression and logical argumentation. But we're not going to talk about continental philosophy here.

8.2. The Traditional Analysis

Here is a traditional definition, which they say that a lot of people used to accept: Knowledge is justified, true belief. ^[34] That is:

S knows that P if and only if:

i. S at least believes that P,

ii. P is true, and

iii. S is justified in believing that P.

A word about each of these conditions.

About condition (i): Naturally, you can't know something to be the case if you don't even believe that it is. Almost all epistemologists regard knowledge as a species of belief. Some people think that "belief" is too weak and that knowledge is something better and stronger than belief. The "at least believes" formulation accommodates this: You have to believe that P, or possibly do something stronger and better than believing it. (Note: The negation of "S at least believes that P" is "S does not even believe that P.")

About condition (ii): You can't know something to be the case if it isn't the case. In case that's not obvious enough, here is an argument: Knowing *that* P entails knowing *whether* P; knowing whether P entails being right about whether P; therefore, knowing that P entails being right about

whether P. Similar points can be made in various cases using the notion of knowing *when* or *where* something is, knowing *why* something is the case, knowing *what* something is, and so on. Example: If John knows that cows have 4 stomachs, then John knows how many stomachs cows have. If John knows how many stomachs cows have, then John is right about the number of stomachs they have. Therefore, if John knows that cows have 4 stomachs, then he has to be correct in believing that.

Sometimes, people say things that seemingly conflict with (ii), such as: “Back in the middle ages, everyone knew that the Sun orbited the Earth.” This sort of statement can be explained as something called *imaginative projection*. This is where you describe a situation *from the standpoint* of another person, pretending that you hold their views. When you say, “People in the middle ages knew that the Sun orbited the Earth”, what this really means is something like: “People in the middle ages would have described themselves as ‘knowing’ that the Sun went around the Earth.” They didn’t genuinely know it, though.

By the way, those first two conditions are uncontroversial in epistemology. Some people reject condition (iii), and some people add other conditions, but almost no one rejects (i) or (ii).

(i) and (ii) are only necessary conditions for knowledge. You can see that they are not *sufficient* for knowledge because of cases like the following:

Lucky Gambler: Lucky has gone down to the racetrack to bet on horses. He knows nothing about the horses or their riders, but when he sees the name “Seabiscuit”, he has a good feeling about that name, which causes him to confidently believe that Seabiscuit will win. He bets lots of money on it. As chance would have it, Seabiscuit does in fact win the race. “I knew it!” the gambler declares.

Did Lucky really know that Seabiscuit would win? I hope you agree that he did not. He just made a lucky guess, and lucky guesses are not knowledge. So we need another condition on knowledge besides belief and truth.

That’s where condition (iii) comes in. Lucky’s problem is that he had no justification for thinking Seabiscuit would win. He just liked the name, but that’s evidentially irrelevant. Other ways to put the point: Lucky’s confidence was not *reasonable* or *rational*; it was *groundless* (when it

needed grounds); it didn't *make sense* to be so confident. That's why he didn't count as "knowing".

8.3. Gettier Examples

In 1963, Edmund Gettier published a short article that became famous among epistemologists. The article refuted the "justified true belief" (JTB) analysis of knowledge, showing that conditions (i)-(iii) are not *sufficient* for knowledge. (Gettier doesn't dispute that they are *necessary* for knowledge, though.) Here's an example from the article:

Jones and Brown: You have extremely good reason to believe that Jones owns a Ford car. You decide to start inferring other things from this. You have no idea where Brown is, but you randomly pick a city, say, Barcelona, and you think to yourself: "Jones owns a Ford, *or* Brown is in Barcelona." That's justified, because the first disjunct is justified, and you only need one disjunct for the sentence to be true. Later, it turns out that Jones actually didn't own a Ford (he sold it just that morning), but coincidentally, Brown *was* in Barcelona. Q: Did you know [Jones owns a Ford or Brown is in Barcelona]?

Intuitively, no. But you satisfied the JTB definition: You had a belief, it was true, and it was justified. Therefore, justified, true belief doesn't suffice for knowledge.

(Note: You *do not think that Brown is in Barcelona* in this example. So please don't talk about this example and complain that the person is unjustified in thinking that Brown is in Barcelona. Don't confuse [Jones owns a Ford, *or* Brown is in Barcelona] with [Brown is in Barcelona]!)

Gettier's argument uses three assumptions: He assumes that if you're justified in believing something, and you correctly deduce a logical consequence of that, then you're justified in believing that consequence too. [35] He also assumes that you can be justified in believing a false proposition, and that you can validly infer a true conclusion from a false premise (you can be right about something for the wrong reason). Given these principles, you can construct Gettier examples.

Many students hate examples like Jones and Brown, because the proposition it has you believing is so strange. People don't make up random disjunctions like that. So here is a less annoying example:[36]

Stopped Clock: You look at a clock to see what time it is. The clock reads 3:00, so you believe that it's 3:00. This seems justified. Unbeknownst to you, that clock is stopped. However, coincidentally, it happens to be

3:00 at the time that you look at it (as they say, even a stopped clock is right twice a day).

Here, you have a justified, true belief, but we would not say that you *knew* that it was 3:00.

8.4. Other Analyses

Philosophers have tried to improve the definition of knowledge. Some add other conditions onto JTB; others try replacing the justification condition with something else. Other philosophers then come up with new counterexamples to the improved definitions. Then people try to repair the definitions by adding more complications; then more counter-examples appear; and so on.

8.4.1. No False Lemmas

Here's the first thing you should think of (but you probably didn't): Just add a fourth condition that stipulates away the sort of cases Gettier raised. Gettier raised examples in which you infer a correct conclusion from a false (but justified) belief. E.g., you infer "Jones owns a Ford or Brown is in Barcelona" from the false but justified belief "Jones owns a Ford." So just add a condition onto the definition of knowledge that says something like:

iv. No false beliefs are used in S's reasoning leading to P.

By the way, this does not require P to be based on reasoning; if P is foundational, then it's okay. What is required is just that it *not* be the case that S reasoned to P from one or more false beliefs. This condition is also known as being "**fully grounded**", or having "**no false lemmas**". (Note 1: Please don't confuse "fully grounded" with "justified". To be fully grounded is to fail to be based on any false propositions; this is neither necessary nor sufficient for justification. Note 2: A "lemma" is an intermediary proposition that is used in deriving a theorem.)

That takes care of the Jones and Brown example: The belief in that case violates condition (iv), so it doesn't count as knowledge. So the improved definition gets the right answer. Likewise, in the Stopped Clock case, we could say that the belief "It is 3:00" is partly based on the false belief that the clock is working.

But now, here's a new counterexample:

Phony Barn Country: Henry is driving through Phony Barn Country, a region where (unbeknownst to Henry), there are many barn facades facing the road, which look exactly like real barns when viewed from the road, but they have nothing behind the facade. There is exactly one

real barn in this region, which looks just like all the facades from the road. Henry drives through, thinking, as he looks at each of the barnlike objects around him, “There’s a barn.” Each time, Henry is wrong, except for the one time he happens to be looking at the real barn.

Obviously, when he looks at the fake barns, he lacks knowledge – he doesn’t know they are real barns (since they aren’t), and he doesn’t know they are fake barns (since he doesn’t believe they are). But what about the one real barn in the area – does he *know* that that one is real?

You’re supposed to have the intuition that Henry does *not* know. He’s correct that time, but just by chance. The no-false-lemmas principle doesn’t help with this case, since the belief that that object is a barn does not seem to be inferred from anything that’s false. If it is inferred from anything, it is inferred from the visible features of the object (its shape, size, distribution of colors), perhaps together with some background beliefs about what barns normally look like – but those beliefs are all *true*. So Henry satisfies all four proposed conditions for knowledge. There must be some other condition on knowledge that he is missing.

Here’s another counter-example that some people find more convincing. *Holographic Vase*: Henry comes into a room and sees a perfect holographic projection of a vase. The projection is so good that Henry believes there is a real vase there. Oddly enough, someone has put a real vase, looking exactly like the holographic projection, in the same location. So Henry is correct that there is a real vase there, though it is the holographic projection that is causing his experience, and if the real vase were removed, things would look exactly the same.

Does Henry *know* that there is real vase there? Intuitively, he does not, though he has a justified, true belief, which he did not infer from any false beliefs.

8.4.2. Reliabilism

According to one influential view (“**reliabilism**”), knowledge is true belief that is produced by a reliable belief-forming process, where a reliable process is one that would generally produce a high ratio of true to false beliefs. Some people regard the reliability condition as simply *explaining* the notion of justification; others would view it as a *replacement* for the justification condition on knowledge. Let’s not worry about that, though.

The biggest problem for reliabilism: How do we specify “the process” by which a belief was formed? There are more and less general ways to do

this, and you can get the same belief being “reliably” or “unreliably” formed depending on how you describe the process. (This is known as “the generality problem”.)

Take the Jones and Brown case: If the belief-forming method is (as you might expect) something like “deduction starting from a justified belief”, then that’s a reliable process.^[37] So in that case, we have a counter-example to reliabilism – the same counterexample that refuted the JTB definition. If the belief-forming method is instead described as “deduction starting from a false belief”, then it’s unreliable, so we get to say that the Jones & Brown case *isn’t* a case of knowledge.

For another example, take Phony Barn Country. If Henry’s belief-forming process is described as “visual perception”, that’s highly reliable. But if it is described as “looking at barnlike objects in Phony Barn Country”, that’s unreliable.

It’s not clear what general, principled way we have to describe “the process” by which a belief was formed. Without such a principled way, you can get pretty much any belief to count as either reliable or unreliable.

Another objection is that reliabilism allows people to count as “knowing” things that, from their own internal point of view, are not even reasonable to believe. For example:

Reliable Wishing: Don believes that he is going to become King of the Earth. He has no evidence or argument for this whatsoever. His reason for believing it is pure wishful thinking: He likes the idea of being King of Earth, so he tricks himself into believing it. Unbeknownst to Don, there is a powerful demon who likes Don’s irrationality, and this demon has decided that whenever Don forms an irrational, wishful belief, the demon will make it come true. The demon thus orchestrates a sequence of bizarre events over the next two decades that wind up making Don the King of Earth. Q: Did Don *know* that he was going to become King of Earth?

In this case, Don’s belief was true and formed by a reliable (for him) method. (Of course, wishful thinking is not reliable in general. But it is reliable for Don, due to the demon.) But it does not seem right to say that he *knew* that he was going to become King.

8.4.3. Proper Function

Another analysis:

S knows that P if and only if:

- i. S believes that P,
- ii. P is true,
- iii. The belief that P was formed by one or more properly functioning cognitive faculties,
- iv. Those faculties were designed to produce true beliefs,
- v. S is in the environment for which those faculties were designed, and
- vi. These faculties are reliable in that environment.

Note: The notions of “design” and “proper function” could be explained in terms of a divine creator, or they could be explained in terms of evolution (evolutionary psychologists often speak of how evolution designed various aspects of us to serve certain functions – of course, this is a sort of metaphorical use of “design”).

Notice that this analysis is similar to Reliabilism, except that the **Proper Function** analysis avoids the problem with cases like Reliable Wishing. Don doesn’t have knowledge, because he didn’t form his belief by a properly functioning faculty that was designed for producing true beliefs. It’s not clear, in fact, that Don was using a *faculty* at all; in any case, certainly there is no human faculty that was designed to produce truth via wishful thinking. So the Proper Function theory works for this case.

Problem: The analysis falls prey to the original Gettier example. When you form the belief [Jones owns a Ford or Brown is in Barcelona], you do so by inferring it from [Jones owns a Ford]. There is no reason to think any of your cognitive faculties are malfunctioning here (you made a valid deduction, after all), or that they weren’t designed for getting true beliefs, or that they were designed for some other environment, or that they’re not reliable. So the analysis incorrectly rules this a case of knowledge.

8.4.4. Tracking

Intuitively, knowledge should “track the truth” – i.e., when you know, you should be forming beliefs in a way that would get you to P if P were true, and get you to something else if something else were true. This leads to the **tracking analysis** of knowledge:

S knows that P if and only if:

- i. S believes that P,
- ii. P is true,
- iii. If P were false, then S would not believe that P.
- iv. If P were true, then S would believe that P. [\[38\]](#)

Clause (iii) is commonly understood among philosophers to mean something like the following: “Take a possible situation as similar to the way things actually are as possible, except that P is false in that situation. In that situation, S does not believe that P.” (There’s a similar interpretation of (iv), but I’m not going to discuss condition (iv) because it won’t matter to the problems I want to raise below.^[39])

This accounts for the Gettier example (Jones and Brown) from §8.3. According to condition (iii), you know [Jones owns a Ford or Brown is in Barcelona] only if:

If [Jones owns a Ford or Brown is in Barcelona] were false, then you would not believe [Jones owns a Ford or Brown is in Barcelona]

That condition is *not* satisfied. Rather, if [Jones owns a Ford or Brown is in Barcelona] were false, it would be false because Brown was somewhere else while everything else in the example was the same. Since your belief had nothing to do with Brown and was solely based on your belief about Jones, you *would* still believe [Jones owns a Ford or Brown is in Barcelona]. That’s why you don’t count as knowing [Jones owns a Ford or Brown is in Barcelona].

The tracking account has other problems, though. The theory implies, for instance, that you can never know that a belief of yours is not mistaken. Let’s say you believe P, and you also think that you’re not mistaken in believing P. To *know* that you’re not mistaken, you must satisfy the following condition:

If [you are not mistaken in believing P] were false, then you would not believe [you are not mistaken in believing P].

This is equivalent to:

If you *were* mistaken in believing P, then you would not believe you weren’t mistaken.

But you never satisfy that condition: If you were mistaken in believing P (whatever P is) you would still think you weren’t mistaken, since, by definition, you would believe that P was the case. (Note: Being “mistaken in believing P” means believing P when P is in fact false.)

Another problem: The theory implies that you can know a conjunction without knowing one of the conjuncts. For instance, I know [I’m not a brain in a vat, and I’m in Denver]. I satisfy condition (iii) because, if [I’m not a brain in a vat, and I’m in Denver] were false, it would be false because I was in another city, most likely Boulder. In that situation, I would know I

was in that other city, so I would not believe [I'm not a brain in a vat, and I'm in Denver]. However, according to the tracking account, I *can't* know [I'm not a brain in a vat], because if I were a brain in a vat, I'd still think I wasn't one. (See §6.3.3. above.) Hence, I can know (P & Q) but not know P. That seems wrong.

8.4.5. Defeasibility

Let's conclude with the most sophisticated and most nearly adequate analysis of knowledge, the defeasibility analysis. The **defeasibility theory** states:

S knows that P if and only if:

i. S believes that P,

ii. P is true,

iii. S is justified in believing that P, and

iv. There are no genuine defeaters for S's justification for P.

To explain: In this context, a “**defeater**” for S's justification for P is defined to be a true proposition that, when added to S's beliefs, would make S no longer justified in believing P.^[40]

This theory easily explains all the examples that we've discussed so far. In the Jones and Brown case, there is the defeater, [Jones does *not* own a Ford]. It's *true* in the example that Jones doesn't own a Ford, and if you believed that Jones doesn't own a Ford, then you would no longer be justified in believing [Jones owns a Ford or Brown is in Barcelona], since your only reason for believing [Jones owns a Ford or Brown is in Barcelona] was that you thought Jones owned a Ford. Thus, [Jones does not own a Ford] is a defeater for [Jones owns a Ford or Brown is in Barcelona]. That explains why the person in the example does not know [Jones owns a Ford or Brown is in Barcelona].

Now I'll just list the defeaters in other cases:

<u>Example</u>	<u>Defeater</u>
Stopped Clock	The clock is stopped.
Phony Barn Country	Most of the barnlike objects around here are not real barns.
Holographic Vase	There is a holographic projection of a vase there.
Reliable	(Doesn't need a defeater since justification

Wishing condition is violated.)

I'll leave it to you to think those through.

Interlude: A Bad Objection to Defeasibility and Some Other Analyses

Here's something students sometimes ask: "But how could I *know* whether there are any defeaters for my belief?" I think this is intended as a sort of objection to the defeasibility theory, but the question/objection rests on two misunderstandings. (People often have the same misunderstandings when they hear about reliabilism and some other analyses.)

The first misunderstanding is to think that the analysis is intended to *help you decide*, in an individual case, whether *you know* some particular proposition. That's not what we're doing. We're just trying to explain *what it means* to say that someone knows something. Whether and how you can know that you know something is a separate question.

Second, students often assume that it would be impossible to know whether there are defeaters, and thus that, even if you knew P, you could never know that you knew it. But that's not true. You can know that there are no defeaters for P, as long as (i) you believe that there are no defeaters for P, (ii) there are no defeaters for P, (iii) you're justified in believing that there are no defeaters for P, and (iv) there are no defeaters for the proposition [there are no defeaters for P]. There's no reason why all those conditions couldn't hold.

So do we finally have a correct analysis? Well, before you get too excited, I have two more examples:

Tom Grabit: You're sitting in the library one day, when you see Tom Grabit, whom you know well, grab a book and hide it in his coat before sneaking out of the library. You conclude that Tom Grabit stole that book. Unbeknownst to you, Tom has an identical twin brother, John Grabit, who is a kleptomaniac. If you'd seen John, you would have mistaken him for Tom. However, it was really Tom that you saw.

Comment: In this case, it seems that you *don't* know that Tom stole the book, because it *could* have been John, and your evidence can't distinguish Tom from John, so you're only right by luck. The defeasibility theory handles this case: The defeater is [Tom has a kleptomaniac identical twin]. But now here's a different variation:

Deluded Mrs. Grabit: Things are as in the above case, except that Tom *doesn't* have a twin; however, his crazy mother has been going around

saying that Tom has an identical twin who is a kleptomaniac. You are unaware that Mrs. Grabit has been saying this, and also unaware that she's crazy.

Comment: In this case, it seems that you *do* still know that Tom stole the book. You don't need evidence capable of distinguishing Tom from a hypothetical twin, given that in reality there is no twin.

Deluded Mrs. Grabit poses a problem for the defeasibility theory. [Mrs. Grabit says that Tom has a kleptomaniac identical twin] is a defeater. It's true (although it's false that Tom *has* the twin, it's *true* that his mother *says* that he has one), and if you believed that Mrs. Grabit says Tom has a kleptomaniac identical twin, then you wouldn't be justified in believing that Tom stole the book. That is of course because you would justifiably suspect that it was the twin who stole the book. So it looks like the defeasibility theory would (incorrectly) rule this *not* a case of knowledge.

So now we have to refine the definition of knowledge. Defeasibility theorists like to distinguish "misleading" defeaters from "genuine" defeaters (where "genuine" just means "non-misleading"). They then say that knowledge requires that there be no **genuine defeaters**, but it doesn't require that there be no **misleading defeaters**. [Mrs. Grabit says that Tom has a kleptomaniac identical twin] is said to be a misleading defeater, when Tom doesn't actually have a twin.

How should we define misleading defeaters? There are several possible approaches; I'll just mention three. Bear in mind that we want it to turn out that there is a *genuine* defeater in the Tom Grabit case with the actual twin, but only a *misleading* defeater in the Deluded Mrs. Grabit case. Now here are three suggestions:

a. A misleading defeater is one for which there exists a restorer, where a restorer is another true proposition such that, if you added it to your beliefs after adding the defeater, you would get back your justification for believing P.

Example: In Deluded Mrs. Grabit, the restorer is [Mrs. Grabit is deluded in thinking she has two sons].

Problem: In the Tom Grabit case with the actual twin, there is also a restorer: [It was not the twin that you saw].

b. A misleading defeater is one that would defeat your justification by supporting something false, as, for example, [Mrs. Grabit says Tom

has an identical twin] supports the false claim [Tom has an identical twin].

Problem: Then the original Tom Grabit case would also have a misleading defeater, because [Tom has a kleptomaniac identical twin] defeats by supporting the false proposition that it was the twin whom you saw.^[41]

c. A defeater for P is genuine if the subject's justification for P depends on the subject's being justified in denying that defeater. A defeater for P is misleading if the subject's justification for P does *not* depend on the subject's being justified in denying that defeater.

Example: In the Jones and Brown case, there is the defeater [Jones does not own a Ford]. This is genuine because your justification for believing [Jones owns a Ford or Brown is in Barcelona] depends upon your being justified in denying [Jones does not own a Ford] (i.e., in believing Jones owns a Ford).

But in Deluded Mrs. Grabit, your justification for believing [Tom stole the book] does not depend upon your being justified in denying [Mrs. Grabit says that Tom has an identical twin]. So this is a case of a misleading defeater.

Problem: We could also say, with about equal plausibility, that in Tom Grabit (with the actual twin), your justification for believing [Tom stole the book] does not depend upon your being justified in denying [Tom has an identical twin]. So this would be (incorrectly) ruled a misleading defeater.

8.5. Lessons from the Failure of Analysis

8.5.1. *The Failure of Analysis*

There are many more attempts to analyze knowledge, and also some more complications that people have added to the ones I listed above. I've just given you five relatively clear examples so that you can see what the discussion of the meaning of "know" is like.

That discussion has been going on since Gettier published his paper in 1963. One book that appeared twenty years later surveyed the discussion up to that point.^[42] It listed 98 different hypothetical examples that had been used to test the dozens of different analyses offered up to that point. To this day, no consensus has been reached on the proper analysis of "knowledge". Indeed, for every analysis, there are counter-examples that most epistemologists would agree refute the analysis.

Isn't that weird? Philosophers are not dumb people, and they understand the word "know" as well as anyone. A fair number of them have been working on this project, and they've been going for decades, almost 60 years as of this writing. And all they're trying to do is correctly describe how we use the word "know". You'd think they would have done it by now. Why can't we just introspect, notice what we mean by "know", and write it down?

Epistemology is not unique in this respect. Philosophers have tried to analyze many other concepts, such as those expressed by "good", "justice", "cause", "time", and "if". Philosophers have tried to define things for as long as there have been philosophers, but they tried especially hard in the twentieth century, because there was a popular school of thought at the time according to which analyzing concepts was the central job of philosophers. But none of our attempts has succeeded. Ever.

Caveat: Some philosophers would dispute that; some claim to have correctly analyzed one or more important concepts. So here is a more cautious statement: No philosophical analysis that has ever been given has been generally accepted as correct among philosophers. "Knowledge" isn't special; it is merely the most striking example of the failure of analysis, because philosophers tried extra hard with this particular word.

8.5.2. A Lockean Theory of Concepts

Why did we think that defining "knowledge" would be tractable? Here is one reason you might think of:

1. We understand the word "know".
2. Understanding a word is knowing its definition.
3. Therefore, we know the definition of "know".

So it should be a simple matter to state that definition. Here is another reason you might have:

1. Concepts are introspectively observable mental items.
2. Most concepts are constructed out of other concepts.
3. Introspection is reasonably reliable.
4. Therefore, for most concepts, we should be able to see how the concept is constructed from other concepts.

A definition would then simply describe how the given concept is constructed from other concepts.

The above two arguments embody an understanding of words and concepts that is pretty natural at first glance.^[43] It's what you would

probably think if no one told you otherwise. To illustrate, take the concept triangle (it's common to use all caps to denote a concept). This concept is constructed from the concepts three, straight line, closed, etc. If one possesses the concept triangle, one can directly reflect on the concept and see that it contains these other concepts, arranged in a certain way. One can then simply describe one's concept by saying: "A triangle is a closed, plane figure bounded by three straight lines." Furthermore, if you *don't* know one of these things about triangles – for instance, you don't know that they have to be *closed* (so you think they might allow gaps in the perimeter), then you don't understand the concept triangle.

Similarly, provided that one understands knowledge, one ought to be able to introspectively examine this concept, and (unless it is among the few *simple* concepts that aren't formed from any other concepts) one should be able to see what other concepts it is composed of and how those concepts are combined to form the concept knowledge.

Granted, our minds are not always transparent to us. Sometimes we confuse one thought or feeling with another, sometimes we deceive ourselves about our own motivations, sometimes we fail to notice subtle mental states, and sometimes we have unconscious mental states that we are entirely unaware of. All that has been known for a long time. But none of those phenomena should make it extremely difficult to analyze typical concepts. It's not as if our understanding of the concept knowledge is completely unconscious (is it?), or we're trying to deceive ourselves about it, or we're not noticing it. Epistemologists have attended very deliberately to this concept, and epistemologists are not generally very confused people. So it shouldn't be *that* hard to describe the concept.

8.5.3. A Wittgensteinian View of Concepts

The Lockean theory of concepts implies that linguistic analysis should be tractable: We should be able to state many correct definitions, without too much trouble. The experience of philosophy, however, is the opposite: After the most strenuous efforts over a period of decades (if not centuries), we've managed to produce approximately zero correct definitions.

Qualification: Not *all* words are indefinable. Most *mathematical* expressions can be correctly defined ("triangle", "prime number", "derivative", etc.). Also, of course it is possible to create a technical term and give it a stipulative definition, as scientists sometimes do. There may even be a very small number of terms from ordinary life that have

definitions – for instance, a “grandfather” is a father of a father. But there are, at most, *very few* definable concepts, and none that are philosophically interesting (in the way that “justice”, “knowledge”, “causation”, and so on are philosophically interesting).

Anyway, the inference seems unavoidable: The Lockean theory of concepts is wrong. In its place, I would advance a broadly Wittgensteinian view of concepts.^[44]

Interlude: Ludwig Wittgenstein

Ludwig Wittgenstein was a famous philosopher from the early 20th-century who wrote a lot of very confusing stuff about language, logic, and the mind. I think what I have to say about concepts is like some stuff that Wittgenstein said, but I don’t actually care how well it matches Wittgenstein’s views. I also don’t care, by the way, whether the “Lockean theory” matches Locke’s views. You have to add in caveats like this whenever you mention a major philosophical figure, because there are always people who have devoted their lives to studying that figure and who, if you let them, will give you all sorts of arguments that the famous philosopher has been completely misunderstood and never really said the things they’re famous for saying.)

There are three things wrong with the Lockean view. First, understanding a word (except in rare cases) is not a matter of knowing a definition. Understanding a word is a matter of having the appropriate *dispositions* to use the word – being disposed to apply the word to the things that it applies to, and not apply it to the things it doesn’t apply to. Accordingly, the way that we learn words is hardly ever by being given a verbal description of the word’s meaning. We learn almost all of our words by observing others using the words in context. We then attempt to *imitate* others’ usage – to apply the word in circumstances that seem to us similar to those in which we have previously observed others using the word. Each time we hear the word used, that slightly influences our dispositions. You understand a given word to the extent that you can successfully imitate the accepted usage.

Second, because concepts are *dispositional* mental states, most features of a given concept are not directly introspectively observable. Our main access to the implications of a concept comes, not through directly reflecting on the concept, but through *activating* the dispositions that constitute our understanding. When we confront a particular situation

(whether in reality or in our imagination), we find ourselves inclined to describe that situation in a particular way, using particular words. That reveals the contours of the concepts expressed by those words. This, by the way, explains the common method of testing definitions by consulting our linguistic intuitions regarding specific scenarios. If the way we applied concepts were by applying a pre-given definition, then this methodology would be backwards; we would reject intuitions about cases that conflict with our pre-given definitions, rather than the other way around.

Third, concepts are typically *not* constructed out of other concepts. The way we acquire the great majority of our concepts is by learning new words, i.e., being exposed to a word sufficiently that we ourselves gradually become disposed to apply that word in a similar range of conditions to those in which it is generally applied by other people. There is in this nothing that resembles combining and rearranging simple concepts to form compound ones.

What does all this imply for the project of defining words like “knowledge”? In order to successfully define a word (by the lights of contemporary philosophers), one would have to find a way of combining *other* concepts in such a way as to pick out exactly the same range of cases. There is no reason for thinking that this should generally be possible.

You could think about it like this: Think of an abstract space, the “quality space”, in which the dimensions are different characteristics that things could have. The exact characteristics of an object determine its “location” in that space. And you could think of concepts as drawing lines around regions in the quality space, where the things inside the lines fall under the concept. Each concept corresponds to a particular region with a particular shape. The boundaries can also be fuzzy, which is an extra complication (that is, there can be cases in which a concept only *sort of* applies, also known as “borderline cases”).

Note that you can generalize this idea: You can imagine a quality space for actions, states of affairs, properties, or whatever we may have concepts of. You can also imagine adding dimensions for *relationships* that things bear to other things. So the spatial metaphor can be applied to concepts generally.

Now, what factors affect the contours of a concept? Again, most concepts correspond to specific words in a natural language; my concept of knowledge is, more or less, just my concept of that which the word

“knowledge” applies to. Thus, the concept’s contours are shaped by the pattern of usage of the word in our language. That usage pattern is affected by all sorts of messy factors. It’s affected by the predilections of the people using the language, what they care about, what they find interesting. It’s also affected by the distribution of things in the outside world. For instance, we have the concept cat because we keep seeing objects in which certain characteristics tend to cluster together. If there were a lot fewer cats, the word might fall into disuse or change its meaning. If there were many objects that were intermediate between cats and dogs on various dimensions, then we’d probably have different concepts – perhaps we’d have one concept covering those objects together with cats and dogs. We decide how to divide up the world based partly on what classification schemes enable efficient transmission of information, given background experience that most people in our linguistic community have.

Meanings also drift over time. You can tell this because most words in our language originated from words in other languages with very different meanings. For instance, the word “happiness” derives from the Norse *hap*, meaning chance. “Phenomenon” derives from the Greek *phainomenon*, meaning appearance. “Virtue” comes from the Latin root *vir*, meaning man. The uses of these words drifted quite a ways over time, and our current words are probably still in flux, on their way to some other set of meanings that people in 10,000 years will have. There is no one directing this process; each word’s meaning just drifts in unsystematic ways according to the needs and interests of all the various individuals using the language.

All of this is to explain why the particular “shapes” of our concepts (the contours of the regions in the quality space that they include) are normally complex, fuzzy, and unique. There is no reason at all to expect a nice, neat set of conceptual relations, whereby you could take a concept and restate it using some different set of concepts. There is nothing in our language or in the world that requires there to be, for any given concept, some other collection of concepts that happen to match the exact boundaries of the given concept. There is no need for the boundary drawn by one concept to exactly coincide, for any distance, with the boundary of another concept.

That is why most concepts are undefinable. In order to define a concept, you have to find some way of combining other concepts (using something like set union and intersection operations) so as to exactly reproduce the unique contours of the original concept.

Now you can see why I like to teach a bit of the literature on the analysis of “knowledge” to epistemology students. They don’t end up learning the correct definition of “know”, but they might end up learning that there isn’t one and that we don’t need one. If I just asserted those things at the outset, you probably wouldn’t have believed me, nor should you have. You have to go through some of the attempts to define knowledge and see what’s wrong with them, in order to understand what I’m talking about when I say that philosophers can’t define anything.

This is no counsel of despair. The same theory that explains why we haven’t produced any good definitions also explains why it was a mistake to want them. We thought, for instance, that we needed a definition of “know” in order to understand the concept and to apply it in particular circumstances. Luckily, that isn’t how one learns a concept, nor how one applies it once learned. We learn a concept through exposure to its usage in our speech community, and we apply it by imitating that usage. Indeed, even if philosophers were to one day finally articulate a correct definition of “knowledge”, free of counter-examples, this definition would doubtless be so complex and abstract that telling it to someone who didn’t already understand the word would do more to confuse them than to clarify the meaning of “knowledge”.

To be clear, all of this concerns the ordinary senses of words that appear in ordinary language – thus, it is a mistake to seek a definition for the ordinary English sense of “know”, or “dog”, or “happy”. Technical terms are another matter. Naturally, if someone introduces a new term, or proposes to use a term in a new sense, they still need to explain what it means.

Part III: Metaphysics

9. Arguments for Theism

We turn to the most popular metaphysical issue: the existence of God. (I know it's popular because when I post about God on my blog, I get lots of page views.) In this chapter, we review the three main traditional arguments for the existence of God. But first, some definitions.

9.1. Views About God

God is traditionally defined in the Western philosophical tradition as a being with the following characteristics:

- i. **Omniscient** (all knowing): The being knows everything.
- ii. **Omnipotent** (all powerful): The being can do anything (or: The being can do anything that it is metaphysically possible for a being to do).
- iii. **Omnibenevolent** (all good): The being is as good as it is possible for a being to be; the being is morally perfect.
- iv. **Creator of the universe**: The being created the universe (or: It created all the contingent things other than itself).

(Contingent things are things that exist but could have failed to exist. See section 2.4.) For short, you can describe God as “the O3 world-creator”.

Note: It is far from obvious that conditions (i)-(iv) give the only interesting or relevant conception of God. Nevertheless, that's how a lot of people think about God, so we're starting with that. Some people give other definitions that are in the general vicinity, e.g., God is the greatest conceivable being, or a supremely perfect being. Those are reasonable alternative definitions.

I would like to insist, however, that we should *not* define God in some *completely different* way that doesn't even refer to a conscious being. For instance, please don't say that God is love, or God is nature, or God is goodness in general. Those statements are category errors or abuses of language. If I define “God” to refer to my couch, then we can easily prove that “God exists”, but this definition would be unhelpful, because it bears no relation to how the word “God” is normally understood, it doesn't help to illuminate anything, and it only serves to sow confusion. Similar problems beset the attempts to define God as love, nature, goodness, etc.

There are three positions one might take about the existence of God:

Theism: God exists.

Atheism: God does not exist.

Agnosticism: I don't know whether God exists.

We're going to discuss arguments for theism in this chapter, and for atheism in the next chapter. The main motivation for agnosticism would be the failure of the arguments for theism and atheism, or their being roughly evenly balanced. The three main arguments for theism that people discuss in the Western philosophical tradition are known as the Ontological Argument, the Cosmological Argument, and the Argument from Design.

9.2. The Ontological Argument

9.2.1. *Anselm's Argument*

"Ontology" refers to the branch of metaphysics that studies what exists. So the "Ontological Argument" is an argument about existence – which is really not very informative. It's a lame name, actually, since *all* the arguments for the existence of God are concerned with existence, yet only one of them is called "**the Ontological Argument**".

Anyway, the argument comes from St. Anselm in the 11th century. It goes something like this (I say "something like" because this is my paraphrase/interpretation):

1. God is defined as a being so great that nothing greater than it can be conceived. (Premise/definition)
2. So nothing greater than God can be conceived. (From 1)
3. One can conceive of a god that exists. (Premise)
4. A god that exists is greater than one that doesn't exist. (Premise)
5. Therefore, if God doesn't exist, then something greater than God can be conceived. (From 3, 4)
6. Therefore, God exists. (From 2, 5)

Granted, the concept of "greatness" is pretty vague, so you might object to that. But it's still sort of plausible to define God as the greatest conceivable thing, and also sort of plausible to say that he'd be "greater" if he existed than otherwise.

Opinions about the merits of this argument vary drastically. Some smart people have hailed the argument as a powerful proof, while others (myself included) consider it the epitome of sophistry.

9.2.2. *Descartes' Version*

Even René Descartes (who was no dummy) was taken in. His version of the ontological argument went something like this:

- 1'. God is defined as a supremely perfect thing. (Premise/definition)
- 2'. So God is supremely perfect. (From 1')
- 3'. Any supremely perfect thing possesses all perfections. (Premise/definition)
- 4'. Existence is a perfection. (Premise)
- 5'. So God possesses existence. (From 1'–4')

Again, the notion of “perfection” is pretty vague; nevertheless, there’s some plausibility to claiming that an existent God would be “more perfect” than a non-existent one (which, I suppose, would really not be very perfect at all).

9.2.3. *The Perfect Pizza Objection*

The most popular type of objection people come up with when they hear about the Ontological Argument is the *reductio ad absurdum*: You can use the same form of argument to prove ridiculous things. The original *reductio* had to do with a perfect island,^[45] but I prefer to imagine a perfect pizza, which I will name “Spizza”.

To parody Anselm’s argument: Let Spizza be defined as a pizza so great that no greater pizza can be conceived. Spizza can be conceived to exist, which would be greater than not existing. So if Spizza doesn’t exist, then one can conceive a pizza greater than Spizza. But this is by definition impossible; hence, Spizza must exist.

To parallel Descartes’ argument: Let Spizza be defined as a supremely perfect pizza. A supremely perfect pizza must possess all perfections, including existence. Therefore, Spizza possesses existence.

In response, Descartes claims that his argument is importantly different from the pizza argument (though he discussed a different example in place of a pizza): Descartes’ definition of God, he says, relies only on a single, simple concept, that of *perfection*; all other properties of God logically follow from the idea of perfection. By contrast, the definition of Spizza conjoins two separate concepts, that of perfection and that of pizza-hood. Since these concepts aren’t logically connected, one can’t guarantee that anything satisfies both concepts. Perfection entails existence, so something perfect must exist; but perfection doesn’t entail pizza-hood, so no perfect things need be pizzas.

Pace Descartes, I think this reply fails. Descartes may have identified a *difference* between his argument and the pizza argument, but the difference he cites doesn't have anything to do with *the logic of the argument*. It doesn't explain how any specific step of the Spizza argument would be defective in a way that wouldn't apply to the God argument. Descartes gives a reason why Spizza need not exist, which may well be a good reason. But that doesn't help his case, because the critic *is not saying that Spizza exists*. The critic is saying (a) Spizza obviously *doesn't* exist, (b) this fact shows that the *form of argument* we used to prove that Spizza exists must be no good, and therefore (c) the parallel argument for God must also be no good. One can't respond to that by simply giving an independent argument that Spizza need not exist; that just supports (a). Rather, the theist would need to show that the Spizza argument *isn't parallel* to the God argument.

Let's compare the God argument to the pizza argument:

<u>God Argument</u>	<u>Pizza Argument</u>
1'.God is defined as a supremely perfect thing.	1".Spizza is defined as a supremely perfect pizza.
2'.So God is supremely perfect.	2".So Spizza is supremely perfect.
3'.Any supremely perfect thing possesses all perfections.	3".Any supremely perfect thing possesses all perfections.
4'.Existence is a perfection.	4".Existence is a perfection.
5'.So God possesses existence.	5".So Spizza possesses existence.

Even if Spizza is a compound concept, how does that stop the argument on the right from going through? Exactly what step doesn't work if you use a compound concept? Step 1"? That's just a stipulative definition. You can define compound concepts just as well as simple ones. Step 2"? That's just a deduction from 1". If the move from 1' to 2' is valid, then so is the move from 1" to 2"; there's no claim about simplicity or complexity of concepts used in that inference. Step 3"? No, that's exactly the same as 3', which Descartes accepts. Step 4"? No, that's identical to Descartes' step 4'. Step 5"? That's just a deduction from 2"–4". The form of that inference is the same as the inference from 2'–4' to 5'. The fact that Spizza is a complex concept doesn't change that.

Conclusion: The *reductio* works. And if it works against Descartes, it works against Anselm too.

9.2.4. *Existence Isn't a Property*

The *reductio* shows that the Ontological Argument is unsound, but it doesn't explain *what* exactly went wrong in the argument. Let's keep looking for that.

Another traditional response to the Ontological Argument, due to Immanuel Kant in the 1700s, claims that the argument mistakenly treats *existence* as just another property that a thing might have. Existence isn't one of the properties that a thing has, because you have to first exist in order to have any properties at all. There's something confused about the idea that God is greater or more perfect if he exists than if he doesn't. It requires imagining two things, an existing God and a non-existent God, and finding them to have different qualities. But if God doesn't exist, that doesn't mean there is a thing, a "non-existent God", which has some different properties from an existing God; it means, rather, that there isn't anything relevant there at all. So there's no room to talk about the properties that God has.

I think there is an important insight in this objection. Existence, indeed, is not one of the qualities a thing can have or fail to have. At the same time, I don't think this gets to the heart of the error in the Ontological Argument. The objection directly challenges step 4 in Anselm's argument and 4' in Descartes' version:

4. A god that exists is greater than one that doesn't exist.

4'. Existence is a perfection.

Let's say that those are both false because existence isn't a property. It's still possible to construct a version of the Ontological Argument. Kant's objection tries to stop the argument too late: It lets steps 1–3 go by without resistance. But if steps 1–3 are all okay, and only 4 is the problem, then the following argument should be fine:

1*. Let God be defined as a perfect being *who exists*.
(Premise/definition)

2*. God is a perfect being who exists. (From 1)

Kant objected to the idea that "perfection" or "greatness" entails existence, but we can circumvent this objection by explicitly building existence into the definition of "God", so that there is no debate about whether existence is implied in the definition.

Of course, no one would advance the above argument, because it's too obvious how to *reductio* it. But if you buy the beginning of Anselm's and

Descartes' arguments (steps 1–2 and 1'–2'), then 1*–2* looks to be equally valid. So the argument must have gone wrong right at the beginning.

9.2.5. Definitional Truths

I think the Ontological Argument's basic error is a confusion about how definitions work. The first inference is of this form:

1**.x is defined to be F.

2**.So x is F.

That is invalid. A definition just explains what a given word means. No fact about the meanings of words can guarantee facts about the external, non-linguistic *world*, which is what 2** is about. A definition tells you what must be true of an object *in order for a given word to apply to that object*; it cannot tell you whether there is in fact anything that the word applies to. So the valid inference would be:

1**.x is defined to be F.

2**.So *if there is something that "x" applies to*, then that thing is F.

Or: *If x exists*, then it is F.

(Notice that "There is something that 'x' applies to" is equivalent to "x exists".^[46]) Examples:

- a) Let Santa Claus be defined as a jolly, fat man who distributes gifts to the world's children on Christmas. It follows from this definition that if "Santa Claus" applies to anyone – that is, if Santa Claus exists – then that person is jolly, distributes gifts, etc. And that's *all* that follows. In particular, it does *not* follow that Santa Claus in fact delivers gifts on Christmas.
- b) Let Spizza be defined as a perfect pizza. It follows (only) that if Spizza exists, then it is a perfect pizza.
- c) Let God be defined as a perfect being. It follows (only) that if God exists, then He is perfect.

Accordingly, let's fix the Ontological Argument: In Anselm's version, step 2 should read "*If* God exists, then nothing greater than Him can be conceived." If you plug that in, then the conclusion we get at the end is, "If God exists, then God exists." Similarly, step 2' in Descartes' version should read, "If God exists, then He is supremely perfect", and the conclusion would then be "If God exists, then He possesses existence." Those conclusions are trivially true. They don't help the theist.

9.3. The Cosmological Argument

9.3.1. The Kalam Cosmological Argument

The **First Cause Argument**, also called the **Cosmological Argument**, claims that we need to invoke God to explain why the universe exists, or why we and all the contingent things around us exist. Here is a natural statement of the argument:

1. Whatever begins to exist has a cause. (Premise)
2. The universe began to exist. (Premise)
3. Therefore, the universe has a cause. (From 1, 2)
4. If the universe has a cause, that cause is God. (Premise)
5. Therefore, God exists. (From 3, 4)^[47]

Premise 1 seems plausible intuitively – it doesn't seem that things can simply pop into existence for no reason, especially universes. It also seems to be supported by experience – all observations so far show that everything that comes into existence is caused by something. E.g., when a house comes into existence, this is caused by a construction crew putting it together; when a baby giraffe comes into existence, this is caused by a mama and papa giraffe; etc.

Objection: In quantum mechanics, there is a phenomenon where virtual particles randomly appear, seemingly out of nothing, and then disappear. Also, there is a theory that our universe might have arisen spontaneously from a quantum vacuum fluctuation. (I'm not going to try to explain that, though.^[48]) *Reply:* These are not examples of stuff arising out of nothing. They are examples of stuff arising from *quantum fields*. The "vacuum state" in quantum mechanics is not *nothing*; it is just a particular configuration of the fields. We still need an explanation for why the fields exist.^[49]

What about premise 4 – why think the cause of the universe would be God? Well, let's think about what the cause of the universe would have to be like. To create the universe, it would have to be outside the universe. Also, we should assume that this cause did not itself have a beginning in time, since otherwise we would just need to explain what caused *it*, leading to an infinite regress. So this would have to be a non-physical entity outside the universe that either exists eternally or is outside time. It must also be immensely powerful, since it created this vast universe. No one has thought of any plausible candidates for what this cause could be, other than a god. (But note that this does not get you the traditional "omni" properties from §9.1.)

Some people would question those arguments. But they're not the most interesting things to talk about. The most interesting step to discuss in the

argument (and the main step that philosophers in fact discuss) is step 2: Why think that the universe had a beginning? Here is a sub-argument for 2:

- a. It is impossible to complete an infinite series.
- b. It is also impossible for there to be an actual infinity.
- c. If the universe had no beginning, then the history of the universe is a completed, actually infinite series.
- d. Therefore, the universe had a beginning.

Explanation for (a): An infinite series is endless, and you can't get to the end of a series that is endless. E.g., imagine trying to count all the natural numbers. You could not finish this task, because there are infinitely many of them. Or, imagine meeting a strange person in the wilderness somewhere. As you approach, he is saying, "... 4, 3, 2, 1. Finished!" You ask what he's doing. He says that he has been counting all the natural numbers, backwards – he counted down from infinity – and he's just finished. This seems impossible.

Explanation for (b): Aristotle distinguished actual infinities from potential infinities, arguing that only *potential* infinities are possible. You have a potential infinity when some thing or process has no definite limit to what it could be or do. For instance, there might be no limit to how high you can count, so your counting is potentially infinite. An actual infinity would be a situation in which something actually has a quantity larger than any finite number – for instance, if you had actually counted to infinity, then your counting would be actually infinite.

One type of argument against actual infinities is that they lead to all sorts of paradoxes. Here are two examples:

Hilbert's Hotel: Imagine that you have a hotel with infinitely many rooms, all of which are full. When a new guest arrives, you can still fit them in: You send an announcement out to all the guests, asking them to move to room number $(n+1)$, where n is their current room number. You then put the new guest in room 1. The next day, an *infinite* number of new guests arrive, yet you still fit them in: You tell each guest to move to room number $2n$, where n is their current room number. This frees up all the odd-numbered rooms, which you use to accommodate the infinity of new guests.

This is counter-intuitive – it should not be possible to fit new guests into a hotel when the hotel is already full.

Thomson's Lamp: Imagine that you have a lamp, which starts out on. After $\frac{1}{2}$ minute, it gets switched off. After another $\frac{1}{4}$ minute, it gets switched back on. After another $\frac{1}{8}$ minute, back off; then on, then off, etc. At the end of 1 minute, it has been switched infinitely many times. Does the lamp end up on or off?

There is no good answer to this. The lamp started out on, and it was never switched off without subsequently being switched back on; so it should be on. On the other hand, it was switched off after $\frac{1}{2}$ minute, and after that it was never switched *on* without subsequently being switched back *off*; so it should be off.

The point: If we assume the existence of actual infinities (an actually infinite hotel, a lamp that is actually switched infinitely many times, etc.), we get absurd consequences; therefore, actual infinities are impossible. That supports premise (b) above. The Thomson Lamp scenario also illustrates the idea that one can't complete an infinite series, per premise (a). There are many more paradoxes that arise from imagining various infinite things, but we don't have time to go through all of them, so we'll just rest with those two examples.^[50]

Now, if the universe had no beginning, that would mean that its history contains an infinite series of past events, stretching back forever. This would be an actual infinity, and it would also be a *completed* infinite series: Infinitely many events or states of affairs must have transpired, in sequence, in order to reach the present state of the universe. Since this is impossible, we must conclude that the universe had a beginning.

9.3.2. Reply: In Defense of Some Infinities

My take: The Kalam cosmological argument is mistaken. There are many examples of actual infinities, including: the infinite extent of space; the infinite number of sub-regions contained in any given region of space; the infinite number of sub-intervals contained in any interval of time; the infinite number of natural numbers, sets, propositions, and other abstract objects.

There is also at least one type of infinite series that we know is regularly completed. It is called "the Zeno series" (after the Greek philosopher who first thought of this): Suppose you want to move from point A to point B. To do so, you must first travel half the distance, then half the remaining distance, then half the remaining distance, and so on. This infinite series of halfway motions must actually be completed for you to arrive at point B (it

would not be sufficient that you merely be *capable* of completing each one). So that looks like an actual infinity, and a completed infinite series. If you deny that actual, completed infinities are possible, then you'll have to deny that anything ever moves.

By the way, here's a variation on Zeno's series: To get to point B, you have to go half the distance, as we said. But *before* going half the distance, you must first go one *quarter* of the distance. And before that, one eighth of the distance, etc. That's another infinite series. If actual infinities are impossible, then you cannot even start your motion.

Now, as to the paradoxes of the infinite: What is the difference between the Zeno series (which is completable) and the Thomson Lamp series (which is presumably uncompletable)? It's not that one is potential and the other actual; both are actual. The rough answer is that in the Thomson Lamp series, an infinite energy density is required: Each time the lamp is switched, the switch must be moved a certain distance, in half the time as in the previous switching. This means that the force applied to the switch must increase without bound. In addition, the total distance through which the switch is moved is infinite. Therefore, an infinite amount of energy is expended in flipping the switch. This energy expenditure must occur in the vicinity of the switch during the minute that it is being switched – hence, you have infinite energy density. Another problem is that the switch would need to have infinite material strength; otherwise, it would break apart at some point in the series, when it was subjected to too great a force.

By contrast, the Zeno series does not require any physical magnitudes to be infinite. There are *infinitely many stages* of the series, yet all the physical magnitudes – total distance, velocity, energy density, etc. – are finite. That's why the Zeno series is possible, yet the Thomson Lamp scenario is impossible.

To be a little more precise, we can contrast three kinds of quantities:

- i. *Cardinal numbers*: These are possible answers to a “how many” question. E.g., 1, 2, 3, and so on. “Infinitely many” is also a possible answer.
- ii. *Extensive magnitudes*: These are continuous quantities that add across the parts of a thing. E.g., length, duration, mass. If the left half of an object is 1 foot long, and the right half is 1 foot long, then the whole thing is 2 feet long.

iii.*Intensive magnitudes*: These are continuous quantities that do *not* add across the parts of a thing. E.g., temperature, energy density, material strength. If the left half of an object is at a temperature of 100 degrees, and the right half of the object is also 100 degrees, it is *not* the case that the whole object is 200 degrees.

Now, the correct theory of the infinite is this: Infinite cardinal numbers are possible, and infinite extensive magnitudes are possible, but infinite intensive magnitudes are impossible. Examples: There are infinitely many numbers (cardinal number); there are infinitely many stages in the Zeno series (cardinal number); space is infinite (extensive magnitude); time is infinite (extensive magnitude); but there *cannot* be infinite energy density in any spacetime region (intensive magnitude), nor can any object have infinite material strength (intensive magnitude).

We don't have time to go into further detail about this theory of the infinite. If you want to know why infinite intensive magnitudes are bad, while extensive magnitudes and cardinal numbers are okay, get my book, *Approaching Infinity*.

What about the example of Hilbert's Hotel, which was supposed to be paradoxical? On my view, the infinite hotel *is possible*. The hotel has infinitely many rooms (cardinal number) and therefore must be infinitely large (extensive magnitude), but there is no need for any intensive magnitudes to be infinite in order for this hotel to exist. But isn't it weird that you could fit more people into the hotel, even if it was initially full? Sure, that's kind of weird, but I think it's just simply true – there's a compelling argument for it, and there isn't a compelling argument against it.

9.3.3. *The Principle of Sufficient Reason*

Another version of the Cosmological Argument claims that all contingent things must have an explanation. This is known as the **Principle of Sufficient Reason**, because it says that there must always be a reason sufficient to explain why things are as they are, rather than one of the other ways they could have been.^[51] Now, you might think that all contingent truths could be explained by other contingent truths. For instance, one can explain why you exist by saying that your parents created you. You can explain why *they* existed by saying that *their* parents created them. And so on. Of course, the human species hasn't existed forever, but maybe *some* causes have always existed, and each cause is explained by another cause before it, going back forever.

But even if this is possible (if you accept actual, completed infinities), this still doesn't satisfy the Principle of Sufficient Reason, because we can just take the *whole infinite series* of causes, and ask why *it* exists. The series is itself contingent (it could have been otherwise), so, if the Principle of Sufficient Reason is correct, then there has to be an explanation for it.

What could explain the entire series of causes and effects in the history of the universe? It can't be any ordinary, physical phenomenon. It would have to be something very powerful that exists outside the universe, and maybe outside time (if that makes sense). Again, the only plausible candidate anyone can think of for such an explainer is God.

Now, you might well wonder: Why don't we then need an explanation for why God Himself exists? And then an explanation for whatever explains God, and so on, *ad infinitum*?

The answer is that the Principle of Sufficient Reason only applies to contingent things, and God is said (by people who endorse this argument) to be a *necessary* being.

That is, God, unlike you and me and all the contingent things around us, could not have failed to exist. God is also sometimes said to be "self-existent", i.e., the explanation for His existence is contained within His essence (roughly, it's just inherent in what God is that He has to exist).

How could this be true? Think about the Ontological Argument again: If it were sound, it would show that God *had to* exist, and the explanation for why He must exist would be contained in the very definition of God. The Cosmological Argument is supposed to give further support for this idea of a self-existent, necessary being: There has to be such a being, because otherwise there would be no way to satisfy the Principle of Sufficient Reason.

Notice that this version of the Cosmological Argument, unlike the one from §9.3.1, does *not* claim that the universe had a beginning in time, nor does it reject actual infinities.

9.3.4. Reply: Against the PSR

The problem with the preceding argument: It assumes the Principle of Sufficient Reason, which we have no reason to assume. Granted, *most* things have an explanation. But why assume that *everything* (including, e.g., the conjunction of all contingent facts) must have an explanation? Why can't there be things that *just are* ("**brute facts**", as philosophers sometimes say)?

There is in fact at least one compelling reason to reject the Principle of Sufficient Reason (PSR): The PSR seems to entail that there are no contingent facts. For suppose that C is a contingent fact. According to the PSR, there must be a sufficient reason for C, call it “D”. D must be another fact, since facts are what provide “reasons” for other facts. Furthermore, to be a *sufficient* reason for C, D must presumably *entail* C.

D, in turn, must be either contingent or necessary. If D is necessary, then C itself must be necessary (whatever is entailed by a necessary truth is itself necessary). If D is contingent, then there must be some sufficient reason for D, call it E. And so on.

Now, even if there is an infinite series of contingent reasons standing behind C (D, E, F, etc.), we still won’t have satisfied the Principle of Sufficient Reason, because there will have to be a sufficient reason for that whole infinite series. And even if you introduce another infinite series, we still won’t be done. Indeed, if you introduce an infinite series of infinite series, we won’t be done, because we can still ask why the whole infinite series of infinite series of contingent reasons exists.

So the only way out is to posit a *necessary* fact at some point in the series of sufficient reasons. But, again, this will mean that the entire series must be necessary. So there are no contingent facts. This conclusion is really not cool, because this is undermining the whole distinction between necessary and contingent truths.

Why can’t we just say that God is the sufficient reason for all contingent truths, as the theist wants us to say? Well, to begin with, that’s a category error. A *reason* is a kind of fact, not a person. For instance, the reason why there are eight planets in the solar system could be “God willed that there be 8 planets”, but not simply “God”.

Okay, why not say that “God wills C” is the sufficient reason for C, where C is any contingent fact? (Or perhaps a contingent fact not produced by the free will of other beings.) Well, then we’d have to ask whether it was necessary that God will C, or contingent. If it was necessary, then C itself is necessary. If it was contingent, then there would (according to the PSR) have to be a sufficient reason why God willed C. Here we embark on a regress again. And again, the only way to fully satisfy the Principle of Sufficient Reason is going to be to cite some non-contingent fact at some point. And that’s going to just render the whole series necessary.

It's pretty unreasonable to hold that everything is necessary. The whole concept of metaphysical possibility, necessity, etc., is standardly motivated in the first place by giving intuitive examples of things that "could have" and "could not have" been otherwise. Also, the idea that you and I and the various objects we see around us are contingent was kind of important in motivating this version of cosmological argument to begin with.

Conclusion: The Cosmological Argument doesn't work. Sorry. There could be an infinite series of contingent causes, with no explanation of the whole series.

9.4. The Argument from Design

9.4.1. *Design and Life*

The **Argument from Design**, most famously defended by William Paley in 1802, claims that you can just see indications of intelligent design if you look around the universe. Paley compares our situation to someone finding a watch on the ground. Back in the 19th century, I guess a lot of people had mechanical pocket watches, and if you looked inside, you could see a complicated, very precise mechanism working. Even if you'd never seen a watch before, you would immediately know that this thing had to have been designed by someone. It's too intricately ordered to have *just happened*.

Where do we find similar complex order, apart from human artefacts? For a while, the best example for theists was life: When you look at living organisms, you see much more intricate and impressive order than that found in a mere watch. Therefore, they were probably designed too, but by a being much more intelligent and powerful than humans. The obvious candidate is God.

I'm not going to say much about this instance of the Argument from Design, though, because it has not aged well. Paley wrote before the Theory of Evolution was developed (Darwin published *On the Origin of Species* in 1859), so in Paley's time, there really was no good explanation for life that anyone had thought of, other than "God did it." Paley was correct to think it implausible that life just appeared by chance.

Today, though, only the most hardcore religious fundamentalists claim that "God made it" is a superior explanation for life to the Theory of Evolution. (I *think* most scholars who are theists would say that God *used* the process of evolution to bring about the variety of life that we see. Maybe God started off the process, *maybe* he guided the process at some

points, but he didn't just make each species as it currently is.) According to Evolution, the intricate mechanisms found in living things developed over long stretches of time by a process in which organisms that were better designed would leave behind more copies of their genes. Occasional, random mutations would introduce variability; beneficial mutations would survive, while harmful ones would be weeded out. Over millions of years, complicated collections of traits can develop that work well together for a given organism. Very briefly, this is a good explanation because:

- i. Living things show the kind of patterns of similarities that you would expect if some species had common ancestors and if different species had split off from common ancestors at different times.
- ii. The fossil records show living things progressing (changing certain traits in a consistent direction) over geologic time.
- iii. Many organisms today possess features that serve no current function but that look like remnants of traits possessed by species that existed earlier. For instance, some snakes have degenerate hind limbs, too small to be of any use, which could be explained if they had evolved from lizards.
- iv. When we look around, humans and other organisms seem to have traits that, in their evolutionary past, would have caused them to leave behind more copies of their genes. They do *not* generally exhibit traits that maximize justice, or their own happiness, or the survival of their species, or the good of the planet as a whole, or *any other value* that you can think of that an intelligent being might want them to maximize. They come closest to maximizing reproductive success. This is predicted by the Theory of Evolution.

These are the sorts of reasons why virtually all biologists endorse Evolution. Since this isn't a biology book, I won't go on about that. If you want to know more, I would recommend looking up some of Richard Dawkins' books. [\[52\]](#)

9.4.2. *Fine Tuning*

The Argument from Design is not dead yet, though. It has gotten a new lease on life from a phenomenon sometimes called the **fine tuning** of the universe.

There are certain quantifiable general features of the universe which we can call the universe's "parameters". For example: the gravitational constant, which determines how strong the force of gravity is; the mass of a

proton; the charge of an electron; the ratio of the strength of the electromagnetic force to that of gravity; the overall mass-energy density of the universe; the initial entropy of the universe. Those are all very important features of the universe or the physical laws governing it. They make a huge difference to how things go in our world, and we don't know why any of those parameters have the specific values they do.

Now, it turns out that many of the universe's parameters happen to fall within a narrow range of possible values that would make it possible for life to exist. For instance, if the energy density of the early universe were slightly larger, then the universe after the Big Bang would have re-collapsed long before life had time to evolve. If it were slightly smaller, then the universe would have expanded too quickly for stars to form, so that again there would be no life.

This sort of phenomenon is what people call "fine tuning". (Note that "fine tuning" does *not* presuppose that a conscious being adjusted the parameters – the people giving the fine tuning argument aren't idiots, so they're not just presupposing the conclusion. "Fine tuning" just refers to the empirical phenomenon that the universe's parameters are within the narrow ranges that allow life to exist.)

Perhaps the most impressive example is the initial entropy of the universe, which was ridiculously low. According to the traditional Big Bang theory, the universe originated in a giant explosion about 14 billion years ago. At its beginning, the universe had an incredibly low entropy. According to one estimate, if you randomly picked a possible initial state for a universe, the probability of picking one with such a low entropy is about 1 in $10^{10^{124}}$ (ten to the power of ten to the power of 124).^[53] The low initial entropy, in turn, is crucial to explaining life and everything else in the universe that we care about.

If you don't know what all that means, it's too complex to adequately explain now (you need another book for that^[54]). Very briefly, though, entropy is commonly described as a measure of the amount of disorder in a physical system (there is a precise scientific way of quantifying it). The entropy of the universe is constantly increasing, and all life processes, as well as everything else interesting, require doing stuff that increases entropy. The universe is heading toward a state of maximal entropy; when it reaches that state, known as "thermal equilibrium", all life processes will be

impossible. If you imagine all the possible ways of arranging stuff in the universe, an extremely high proportion of them would put the universe in thermal equilibrium immediately. Life is only possible (for a while – it will eventually run down) because our universe luckily started out very far from thermal equilibrium, i.e., with extremely low entropy, 14 billion years ago, for some unknown reason.

So, given that almost all ways of assigning values to the universe's parameters would be unfriendly to life, why does the universe in fact have life-friendly parameters? The theist says: Because an intelligent, benevolent, and immensely powerful being *set* the parameters of the universe that way, *in order* to make life possible.

If there is such a being, it sounds pretty close to what we would call “God”. Note, however, that this argument can't establish the traditional “omni” properties of God – it can't show that the creator is omniscient, omnipotent, and omni-benevolent, since those properties are not necessary to create this universe, though the creator is presumably at least *very* powerful. [\[55\]](#)

9.4.3. *Bad Objections*

I'm going to quickly go through some lame objections, so we can spend more time on the one that matters.

Weak Objection #1: “How do you know there wouldn't be life if the parameters of the universe were different? Maybe there would just be a different kind of life from the ones we're familiar with, with different requirements.”

Reply: That's utterly implausible. The fine tuning arguments do not turn on specific features of life on Earth – e.g., that it's carbon-based, or that it uses DNA to replicate. They turn on *extremely broad* features of the universe. And we don't need to have a precise listing of all the kinds of life possible in order to know that *certain* situations are inhospitable to life. For instance, no scientist would take seriously the suggestion that maybe there is life on the sun, or in a black hole, or in the middle of interstellar space. Similarly, if the universe had recollapsed before planets ever formed, or if there were no elements heavier than hydrogen, then obviously there wouldn't be life. Don't be foolish.

Weak Objection #2: “If there's a universe-creator, it's a totally different kind of being from us, so how can we know what it would want? Why assume it would want there to be life?”

Reply: The Fine Tuning Argument does not claim that you can *logically deduce* that there would be life from the assumption of an intelligent creator. All the argument requires is that life is more likely to occur if there is an intelligent designer than if there isn't. Without the designer, there is (allegedly) no explanation for the amazing coincidence of all these life-friendly parameters of the universe. With the designer, there is at least a plausible explanation. That's enough for us to say that the life-permitting parameters constitute *evidence for* intelligent design. The strength of this evidence is proportional to the *ratio* of (the probability of life-permitting parameters given intelligent design) to (the probability of life-permitting parameters given no design). And it is in fact very easy to see why an intelligent designer *might plausibly* prefer a life-containing universe to a universe with nothing but hydrogen, or a universe that quickly collapses into a giant black hole, etc.

Weak Objection #3: "If there's an intelligent designer, and it values life, then the universe should be filled with life. But there is only a tiny amount of life, compared to all the non-living matter and empty space."

Reply: If there is an intelligent designer, we don't know what its plans are, nor what its capabilities are. The Fine Tuning Argument supports the existence of a powerful creator, but it can't tell us whether the creator is *all-powerful* or merely *very powerful*. Perhaps it is much more difficult to create a universe with a lot of life than to create one with a small amount of life, and our creator did the best it could. Or perhaps there *is* a lot of life in the universe, but it is simply spread out, and perhaps there are good reasons to want different life-bearing planets to be far apart. (Maybe problems occur when species from different planets interact, etc.)

Since we don't know the details about the designer, we cannot confidently predict that there would be more life than we see. It remains a plausible possibility that the universe would be about the way we observe it to be. That's enough for the argument to work (compare the reply to Weak Objection #2).

Weak Objection #4: "There is no good theory about how to assign probabilities to the possible values of the universe's parameters, prior to measuring them. Therefore, we can't say *anything* about such probabilities, so we can't say it's improbable that the parameters would fall in these incredibly tiny ranges that permit life to exist. And thus we can't make any probabilistic inferences from this evidence."

Reply: The premise of the objection is true: There is no generally accepted account of how to assign initial probabilities to possible values of physical parameters. That poses a challenge for probabilistic reasoning about those parameters. Notice, however, that there is nothing about this problem that is specific to reasoning about laws of nature or intelligent designers of universes; there is, in fact, no accepted theory of how to assign probabilities to *any* kind of evidence prior to gathering it. So this objection really poses a *general* skeptical problem about probabilistic inference from evidence to a theory. But you probably don't actually want to reject all such inference.

Let me give you a hypothetical example to illustrate how extreme Weak Objection #4 is:

Made by God: While exploring the surface of Mars, astronauts discover a new kind of crystal. When they look at it under a microscope, they find that the molecules of this crystal spontaneously arrange themselves into patterns that look exactly like the English words, "Made by God" in Times New Roman font. Everyone who looks in the microscope sees it. Scientists are able to figure out that this is actually a complicated, hitherto-unnoticed consequence of some very specific features of the laws of nature, features that no one has any explanation for. Over the next few decades, many more crystals are discovered, scattered across all the planets of the solar system, which, when looked at under microscopes, look like the phrase "Made by God" spelled out in each of the languages of Earth. Again, the laws of nature just happen to be arranged to ensure that this happens.

I know, this is a pretty fanciful hypothetical. But just for the sake of argument, imagine that that happened. According to Weak Objection #4, this would provide *no evidence whatsoever* for any kind of intelligent design. We should just shrug our shoulders and chalk it up to coincidence. After all, the laws of nature had to be *some* way; why *shouldn't* they just happen to be designed so that crystals on all the planets spell out "Made by God" in every language? We can't say that's improbable, because we have no good theory of how to assign probabilities to different possible laws of nature.

I hope you agree that that's ridiculous. If your opposition to theism is so extreme that your position wouldn't even admit that there was evidence for

theism in the “Made by God” story, then I think you need to step back and take a break.

Weak Objection #5: There is said to be an important principle called the **Anthropic Principle**, which says something like, “Observers only exist in conditions that allow observers to exist.” If the parameters of the universe *didn’t* allow for life to exist, then we wouldn’t be here to talk about it. When reacting to the Fine Tuning Argument, people frequently invoke this tautology and vaguely suggest that this provides some sort of objection to the argument. I’m not sure what the objection is supposed to be, though. Perhaps the idea is either: (i) “We could not exist in a universe that didn’t allow for life. So *that explains* why this universe allows for life” or (ii) “We could not exist in a universe that didn’t allow for life. So we *don’t need* any explanation for why the universe allows for life.” (Alternately, the objection might depend on the Multiverse Theory; see §9.4.4.)

Reply: The Anthropic Principle is trivially true. It does not, however, provide any explanation at all for the fine tuning, nor does it show that we don’t need an explanation. Here is an example to help make the point:

Firing Squad: You’ve been convicted of treason (a result of one too many intemperate tweets about the President) and are scheduled to be executed by firing squad. When the time of your execution arrives, you stand there blindfolded, listening to the fifty sharpshooters lift their rifles, fire, and then . . . Somehow you find yourself unscathed. All fifty shooters have apparently somehow missed. Wondering how this could have happened, you start entertaining hypotheses such as: Maybe someone paid all the soldiers to deliberately miss, maybe someone broke into the armory last night and loaded all the guns with blanks, etc. [\[56\]](#)

It is completely reasonable for you to entertain such hypotheses; surely the shooters did not all miss purely by chance. Yet according to Weak Objection #5, you have no reason to think that there was any sort of plan by anyone to save you. Instead, you should say to yourself: “If the shooters hadn’t all missed, then I wouldn’t be here to think about it. So that explains why they missed.” Or: “If the shooters hadn’t all missed, then I wouldn’t be here to think about it. Therefore, I don’t need any explanation for why they missed.” That’s parallel to the objection to the Fine Tuning Argument.

[9.4.4. The Multiverse Theory](#)

The leading alternative to the intelligent design theory is the **multiverse theory**: There exist a large number of parallel universes, maybe infinitely many, with different values for their parameters. Many cosmologists actually believe that this is true or likely to be true. There are theories in physics that explain how new universes might periodically come into existence.^[57] Obviously, “universe” here does not mean “everything that exists”. Universes in the present sense are, roughly speaking, bunches of stuff (physical material, fields, space and time) that are causally isolated from each other. I.e., the stuff in one “universe” can affect other stuff in the same universe but can’t affect stuff in any other universe. This could be because the “other universes” exist in their own, separate spaces, or they are separated from our space along a dimension that we can’t move through, or they are just too far away for any signal from us to ever reach them.

Now, if there are such parallel universes, especially if there are infinitely many, that would make it much more likely that there would be at least one universe in which all the parameters fell within the range required for life to evolve. *Most* universes, perhaps, have parameters that exclude life, but we of course would not expect to ever observe one of those universes, due to the anthropic principle. We can only find ourselves in one of the life-permitting universes. So, on the multiverse theory, it is unsurprising that we see life-permitting universe parameters. So we don’t have to posit intelligent design.

Now for the leading objection to the multiverse theory. To explain the objection, we need to introduce a hypothetical example:

Coin Flip: You flip a coin, and it comes up heads ten times in a row. In general, if a coin is flipped ten times, the odds of its coming up heads every time is $1/1024$. This is sufficiently improbable that you start looking for an explanation. One hypothesis is that there are thousands of other coins somewhere, each getting flipped ten times.

Question: Does the hypothesis of many other coins *explain* the fact that you observed ten heads in a row? And do you therefore *have evidence* for lots of other coins being flipped?

I hope you can see that the answer is “no” to both. (If the answer were “yes”, then every time you observed *anything* unlikely, that would be evidence for other universes.^[58]) Now, why is the answer “no”? Well, it is true that the existence of many other coin flippers would make it more likely that *at least one* coin would come up heads ten times in a row.

However, that's completely irrelevant, because your evidence in the story is not [At least one coin came up heads ten times]. Your evidence in the story – what you observed to be true – was [*This* coin (the one you just observed) came up heads ten times]. And the probability that *this specific* coin would come up heads ten times in a row is completely unaffected by the existence or non-existence of other coins flipped by other people. Thus, the observed outcome with this coin also provides no evidence whatsoever for the existence of other coin-flippers.

Now draw the parallel to the multiverse theory. We observe some seemingly improbable parameter values in the universe. The hypothesis of many parallel universes would indeed make it much more likely that *at least one* universe would have life-friendly parameters. However, the existence of all these other universes would have no effect whatsoever on the probability that *this* universe would have life-friendly parameters. And our evidence is that *this* universe has life-friendly parameters. So our evidence provides no support at all to the multiverse theory since the multiverse theory does nothing to explain the evidence. That is what I call the “this universe” objection to the Multiverse theory.

I used to think that objection was decisive and therefore that the Multiverse theory was a complete loser. I was wrong, though; there is a way to defend the Multiverse explanation, but it is one that most atheists are not going to like. Assume that you could have been born (i.e., you had some chance of being born) into *any universe* that was hospitable to your kind of life. In that case, the probability that you would find yourself alive now would be greater if there were many universes than if there were only one universe. So the multiverse theory can explain the fact that you are alive now, and so your current existence is evidence for the multiverse.

On the other hand, suppose you *could not* have been born in any other universe besides this one, even if there were other universes. In that case, your probability of finding yourself alive now would be completely unaffected by the existence of other universes, and so you have no evidence for other universes.

So the way to defend the multiverse theory as an explanation for our evidence is to claim that persons are not world-bound: A given individual could have existed in any universe that was sufficiently hospitable. Any time a sufficiently hospitable universe exists, you have a chance of being born into it. And maybe that's true, so maybe the multiverse theory is okay.

Why did I say most atheists wouldn't like this? Because *most* atheists are also *physicalists*, i.e., they think that the only things that exist are physical things. In particular, they believe that *persons* are purely physical things. If you're a purely physical thing, it is hard to understand how *you* could possibly have existed in another universe. It doesn't seem as if any *particular* physical object could have been in another universe instead of the one it is actually in. For instance, the chair I'm sitting on couldn't have been in another universe. Of course, there could be *a chair* with the same shape, color, and other qualities in another universe, but it could not literally be *this* chair; it could only be *another* chair just like this one. Similarly, if persons are just particular physical objects, then it seems that a specific person could not have existed in another universe.

On the other hand, suppose you believe that persons have an immaterial component, a "soul", which determines their identity such that *you* are wherever your soul is. Then it is possible that a particular soul could have inhabited a body in a different universe.

So, if you think people are just physical objects, you should probably reject the multiverse theory as an explanation for fine tuning, and therefore you have a good reason to *accept* intelligent design instead. But if you think persons have immaterial souls, then the multiverse theory is a viable explanation for fine tuning, and thus you have *less* reason to accept intelligent design. This is the opposite of what you might have thought, since most intelligent-design advocates also believe in souls, and most intelligent-design skeptics also think we are mere physical objects.

9.5. Pascal's Wager

9.5.1. Pascal's Argument

The 17th-Century French philosopher/mathematician/scientist Blaise Pascal gave a famous argument to show that it is rational to believe in God. The argument ("**Pascal's Wager**") doesn't claim to provide *evidence* for theism, though. Instead, it tries to show that it is *prudent* to be a theist. According to a common version of Christianity, all atheists end up in hell, which is infinitely bad (you suffer for eternity), whereas Christians who properly solicit forgiveness can go to heaven, which is infinitely good (you have good experiences for eternity). Therefore, if there is even a *chance* that Christianity is true, it would be smart to become a Christian.

Aside

I phrase this as an argument about Christianity specifically, rather than theism generally, because not all theists believe in heaven and hell. Actually, not even all Christians believe in hell, but nearly all did in Pascal's time. By the way, Muslims also have a similar view, wherein non-Muslims are tortured eternally in hell.

Granted, becoming a Christian has some costs. You'll probably spend time in church that could have been spent doing something more fun. You might have to give up some sins that are really fun. If there really is no God, you'd be taking on these costs for nothing. But these would only be finite costs, which are nothing compared to the *infinite* cost of going to hell. In standard rational choice theory, when you are uncertain what outcome a given action will produce, you multiply the *probability* of each possible outcome by the benefit or cost that you would get if that outcome happens, then add these products together. (You can treat costs as negative benefits for these purposes.) So, e.g., a 10% chance of getting 10 units of good is exactly as good as a 100% chance of getting 1 unit of good. This gives you a number called the "expected value" of the action. You then choose the action with the highest expected value.

In the case of deciding whether to become a Christian, you can do these calculations:

Expected value of being a Christian =

(Probability that Christianity is true)(Benefit you get if you're a Christian and Christianity is true) + (Probability that Christianity is false)(Benefit that you get if you're a Christian and Christianity is false)

Expected value of being an atheist =

(Probability that Christianity is true)(Benefit you get if you're an atheist and Christianity is true) + (Probability that Christianity is false)(Benefit that you get if you're an atheist and Christianity is false)

To illustrate, suppose the probability of Christianity is 10%, the costs of being a Christian in this life are 100 units of lost pleasure, and the benefits of being an atheist are 0 (we treat atheism as the baseline). Then the calculations go like this:

Expected value of being a Christian = $(0.1)(\infty) + (0.9)(-100) = \infty$

Expected value of being an atheist = $(0.1)(-\infty) + (0.9)(0) = -\infty$

Since the expected value of being a Christian is infinitely greater than the expected value of being an atheist, it's rational to be a Christian. Notice

that you get the same results if you replace the “0.1” with *any* number other than zero. You can also replace the “-100” and “0” with *any* finite numbers – you’ll still get the same result.

9.5.1. Objections

Here are some objections that you might think of.

Objection 1: A good God wouldn’t send people to hell just for not believing in him. That would be cruel and unjust!

Reply: Are you absolutely certain of that? If there’s any chance at all that God will send you to hell for not believing in him (whether or not doing so would be good, just, etc.), then Pascal’s argument works.

Objection 2: Pascal is assuming that all you have to do to get into heaven is believe in God. But don’t you have to do some other stuff, like following God’s commandments, or being baptized, or confessing your sins...?

Reply: No, Pascal is *not* assuming that all you have to do is believe in God. He is not assuming that belief is *sufficient* to get into heaven. He is assuming that belief is *necessary* to get into heaven – hence, you should believe. If other things are also required, then obviously you should do those other things *too*. That in no way conflicts with the advice to embrace theism.

Objection 3: I can’t just *decide* to be a Christian; I don’t have that kind of control over my beliefs. I can’t believe something that I have (almost) no evidence for.

Reply: If you’re having trouble believing, there are steps you can take. Try attending church regularly. Surround yourself with Christians, listen to all their arguments as sympathetically as possible, and avoid listening to any atheists. Act as if you believe. Try repeatedly telling yourself Christianity is true. Etc. After a while, there’s a good chance that you’ll start to believe.

Objection 4: Christianity is super-implausible for such-and-such reasons. [Fill in your favorite arguments against Christianity here.]

Reply: Read the original argument again. It works as long as you give a nonzero probability to Christianity. It could be one in a googol, and you still get an infinite expected value for becoming a Christian. [\[59\]](#)

Objection 5: Maybe there's a god who sends people to hell for *believing* in Christianity, rather than for not believing!

Reply: Okay, that's possible. However, it's *more likely* that people go to hell for not believing rather than for believing. There is at least *some* evidence for Christianity, e.g., the testimony of the Gospels, people who have religious experiences, people who claim to have seen miracles that would fit with Christianity. And Christianity (at least one version of it) holds that people go to hell for not believing in God. By contrast, there is *no* evidence for the opposing theory that we go to hell for believing. Indeed, no one even seriously advances *that* theory. From the standpoint of self-interest, you should bet on the theory that is more likely to be true.

Counter-reply: Here's a reason for thinking that we might go to hell for being Christian: *If* God exists, then He gave us our rational, cognitive faculties. Plausibly, He intended us to use them properly, which would include basing our beliefs on the best available evidence. Also, if He wanted us to believe in Him, He could easily have showed Himself in an unambiguous way such that we'd all believe. Since God hasn't showed Himself, and hasn't given us adequate evidence of His existence, we can infer that either there is no God *or* there is but He doesn't want us to believe in Him. Therefore, we are more likely to be punished for frustrating His wishes by believing in Him than for rationally withholding belief.

Objection 6: Maybe Christianity has probability zero. Some arguments suggest that the traditional Christian God may be metaphysically impossible or inconsistent with observations (see ch. 10).

Reply: Again, there's at least *some* evidence for Christianity (see reply to objection 4). Whatever arguments you have against Christianity, you shouldn't be *100%* confident that you're correct. After all, people are sometimes wrong even when they *think* they have a really great argument (even a conclusive one). There's no argument against Christianity that all experts will agree on; there are always some smart, informed people who disagree. So there's always some chance that the argument is wrong and Christianity is really true.

Counter-reply: Actually, for a theory to have probability zero, there doesn't have to be a conclusive argument against it. If a theory is one of infinitely many alternatives, each of which is initially

equally plausible, then all of these alternatives should start out with probability zero. (It can't be anything greater than zero, since then the total probability for all the alternatives would add up to more than 1.) Furthermore, it's plausible that Christianity is in fact one of an infinite number of possible religions. In particular, the theory about heaven and hell that Pascal invokes is one of infinitely many alternatives: Maybe atheists go to hell for only 1 day, or for 2 days, or for 3 days, etc. Since "infinitely long" is one of infinitely many alternatives and no more plausible than any of the others, it has to start out with probability zero.

Objection 7: The argument only considers Christianity and atheism. But there are many other religions one could adopt.

Reply: True. But note that the only other religions you need to consider are ones that promise an infinite reward (like heaven) or infinite cost (like hell). Among those religions, you should pick the one that is most likely to be true. (More precisely, pick the course of action that is most likely to get you an infinite reward. If there are multiple compatible courses of action that *might* get you an infinite reward, do as many of them as possible.) Pascal would need a separate argument to show that Christianity was the most likely – e.g., that it's more likely than Islam. But that's beyond the scope of this text.

My comment: It looks to me like most of the replies above work, and so the objections fail. Indeed, some are embarrassingly confused. However, I find the counter-replies under objections 5 and 6 pretty reasonable.

9.6. Conclusion

No argument proves that there's a god. The Ontological Argument and the Cosmological Argument both fail. The complex order found in living things provides little evidence for a god, since the theory of evolution gives a better explanation for this order. However, the fine tuning of the parameters of the universe provides better evidence for a creator, since that hypothesis could explain why the parameters are fine tuned. Another viable theory, however, is that there is a multitude of parallel universes, and a given person could find himself in any sufficiently hospitable universe. It is unclear which of these theories, if either, is correct.

10. Arguments for Atheism

10.1. Cute Puzzles

Recall that God is traditionally supposed to be all-knowing, all-powerful, and perfectly good. These properties lead to some puzzles, suggesting that maybe the traditional definition of God is self-contradictory or contradicts some very widely accepted views about God.

10.1.1. Omnipotence and Immovable Stones

The most famous puzzle: Can God create a stone so heavy that he couldn't lift it? If no, then God is not all-powerful, since we just found something he cannot do. If yes, then God is not all-powerful, since he would be unable to lift that stone. More generally, we can ask: Could God create a state of affairs that he himself could not alter? This seems to show that the concept of omnipotence is contradictory.

The theist's best response is to refine the definition of "omnipotence". Rather than saying that an omnipotent being can do "anything", we could define an omnipotent being as one who *can bring about any outcome that it is metaphysically possible for a being to bring about*.^[60] Why is this a reasonable revision to make? We want to address a reasonable theist, not a crazy one (who cares if you can refute the craziest version of a view?). And the reasonable theist wants to ascribe to God the maximum possible amount of power, not an *impossible* amount of "power". So if we find contradictions in the definition of "omnipotence", that shows that we chose a poor definition. The new proposed definition seems to capture the most power that a being could have.

Given this understanding, God could not create a stone so heavy that he could not lift it. There is no weight that a stone could have that would render it metaphysically impossible for it to be lifted, and therefore no weight that would prevent an omnipotent being from lifting the stone. So it is metaphysically impossible for any being to create a stone that God could not lift, and hence even God cannot create such a stone.

10.1.2. Omnipotence and Error

Since God is a perfect being, he presumably cannot *make a mistake*. But then, he is not omnipotent, since there is something He can't do that it is possible for a being to do. Similarly, God cannot do evil, nor can He

commit a sin (an offence against God). So there are things we humans can do that God cannot.

Reply: Oops, we mis-defined “omnipotence” again. We meant to ascribe to God the maximum possible amount of power. Intuitively, the ability to make mistakes is not the sort of ability that renders one more powerful; quite the opposite, in fact. So we should actually interpret “omnipotence” to exclude that particular “ability”.

Regarding the ability to do evil, or commit moral wrongs, one could argue that God *has the power* to do such things; he simply does not choose to exercise that power. What makes a being morally good, after all, is not that it is *unable* to do evil, but rather that it freely chooses not to do the evil things that it could do.

Regarding the ability to sin, one might think this is ruled out definitionally, since no being can commit an offence against itself. But then it would not really be a limitation of God’s power to be unable to do this contradictory thing. Alternately, if you think a being *can* commit an offence against itself, then there is no reason why God wouldn’t have the power to do so – though, again, he wisely chooses not to exercise that power.

How could we modify the definition of “omnipotence” to accommodate these points? We might try saying that an omnipotent being can bring about any possible arrangement of the universe (God’s committing a sin, or an evil act, or making a mistake, are not possible arrangements of the universe). Or perhaps we should just define an omnipotent being as a being with the greatest possible amount of power, or a being such that no being could be more powerful than it.

10.1.3. Omniscience and Free Will

Does God know the future? If he does not, then he is not all-knowing. But if he *does* know the future, then this seems to entail that no one has free will, since we have to do that which God already knows we are going to do. By the way, if that’s right, then even God himself wouldn’t have free will, which might contradict his omnipotence.

There are two ways a theist might respond, other than giving up free will. One is to say that God’s foreknowledge is actually compatible with free will, because neither God nor any other external factor is *making* us choose in the way we are going to. Suppose you’re going to choose between A and B, and God knows that you’re going to select A. You *could* freely choose B instead; if you did, you wouldn’t thereby make God wrong;

rather, if you chose B, then God would have always known that you were going to do *that*. This may sound better if you think God is somehow outside of time (if you can make sense of that), and from his perch outside time, he simply *sees* you choosing A, without himself interfering.

The second response would be to say that omniscience does not require knowing the truth-value of every proposition; rather, it only requires knowing the truth value of every proposition *that has a truth-value* (i.e., that is determinately true or false). Propositions about your future free choices do not have determinate truth-values at the present time. In other words: Suppose you're going to freely choose between A and B. Then there is, as of now, *no fact of the matter* about which you will choose (that, maybe, is what makes the choice "free"). Since there is no such fact, God does not know it. This isn't a genuine limitation on His knowledge, since he still knows every fact that exists.

* * *

My take: These sorts of puzzles are entertaining but of no deep import. The most they should be able to accomplish is to make the theist refine his definitions, which means that they do not get to the heart of the dispute between theists and atheists. It is reasonable for the theist to respond to these sorts of puzzles by refining definitions so as to eliminate the puzzles, because in general, a good definition should (usually) not be contradictory. Thus, if someone succeeds in showing that the traditional definition of "God" is contradictory, that just shows that it's a poor definition, and we should replace it with a definition that isn't contradictory.

There are limits to this. If the traditional conception of God is so confused that there is nothing anywhere close to it that is logically coherent, then we should just declare that there is no God. But if there is something in the neighborhood that is coherent, then we should interpret "God" as having one of the coherent meanings that is close to what is traditionally said about God. That's because our purpose is to learn, not merely to score points against people. The way to learn is to address the most interesting defensible views, not to spend our time discussing trivially false ideas.

10.2. The Burden of Proof

So far, it looks like there is no definitive proof either that there is or that there isn't a God. In such a situation, what should one believe? One obvious answer is that one should suspend judgment, i.e., be an agnostic.

That's not what most atheists say, though. Most atheists (or at least many) invoke a **burden of proof principle**, that is, a principle that is supposed to tell us which side in a disagreement has the obligation to provide evidence, and which side is presumed to be correct until proven wrong.

You're probably familiar with the *legal* burden of proof principle – that defendants in a criminal trial are presumed innocent, and prosecutors have the obligation to provide evidence of guilt. In a similar way, many hold that those who make *positive* assertions in a discussion (claims that something exists or that something has some property) have an obligation to provide evidence for their view, whereas those who make only *negative* claims (denying that something exists or that something has some property) are presumed correct until proven otherwise. Theism is a positive claim, so, according to this idea, theists have the burden of proof, while atheism is the default position if no proof can be provided either way.

You periodically hear people say things like that in popular discourse, though not so much in professional philosophy. I should note, though, that it's not always clear exactly what the "burden of proof" people are saying – they're usually not as explicit as I was above, probably because they themselves are not sure what they believe. (That happens a lot in philosophy.) So people will often make remarks that are ambiguous between (a) and (b) below:

(a) In the absence of evidence, you should believe *x does not* exist.

(b) In the absence of evidence, you should *refrain from believing* that *x* exists.

Notice that (b) is obvious, while (a) is a lot more puzzling and less obvious. I'm going to address (a). (Notice that if you *merely* accept (b), there is no asymmetry between positive and negative claims, so there would be no point to speaking of a "presumption" or "burden of proof" that applies to only one side in a disagreement.)

Why would there be a burden of proof for positive claims? The most common answer is "Because it is impossible to prove a negative." This answer has two main problems. First: It's obviously false. It is often possible to prove negative claims. For instance, if I want to prove that there is no beer in the refrigerator, I can open the refrigerator and do an exhaustive search of it. If there really is no beer in there, that will prove it. Likewise, there are arguments that purport to prove, or at least provide

strong evidence, that there is no God (see §§10.1, 10.3). We certainly can't rule out such arguments merely from the logical form of the conclusion, "There is no God."

Here's the second problem: How on earth would the impossibility of having evidence for a proposition make it rational to assume that the proposition was *true*? That's what the "burden of proof" argument is saying – that because it's impossible to have evidence for "there is no God", we get to start out by *assuming* that there is no God.

Maybe the idea is that, since we couldn't have evidence of atheism even if it were true, the absence of evidence for atheism isn't a mark *against* it; however, since we *might* have evidence for theism if theism were true, the absence of evidence for theism *is* a mark against theism.

There is a legitimate point to be made in this vicinity. But let's be more precise. Suppose there is some possible evidence, E, which would be more likely to occur if God exists than if God doesn't exist. (For instance, E might be an event in which Jesus appears and performs several miracles in front of you.) If E occurs, that'll be (at least some) evidence for theism. If E *doesn't* occur, then its failure to occur will be (at least some) evidence for atheism. How *strong* the evidence is will depend on the probabilities (i.e., the probability of E given theism compared to the probability of E given atheism).

Note, however, that there is no *asymmetry* between positive and negative claims here, so there's no reason to speak of a "burden of proof" or "presumption" applying to one side. You can say exactly parallel things, interchanging theism and atheism, thus: Suppose there is some possible evidence, D, which would be more likely to occur if atheism is true than if theism is true. (For instance, D might be some terrible thing like the Holocaust.) Then if D occurs, that will be some evidence for atheism. If D doesn't occur, that will be some evidence for theism. So there's no asymmetry between theism and atheism: Both views can have evidence for or against them. Indeed, the evidence for either view would just be the negation of whatever would be evidence for the other.

Here is another confused argument that people sometimes state: The more positive claims a theory contains, the more likely it is that the theory will go wrong somewhere. It's a trivial theorem of probability that $P(A \ \& \ B) \leq P(A)$ (read: The probability of A and B both being true is less than or equal to the probability of A being true – almost always, it is strictly *less*),

so if you can explain all your evidence using only A, don't introduce B in addition.

This argument deploys a trivially true premise, that $P(A \& B) \leq P(A)$. This premise, however, provides no support at all to the burden of proof principle. You can see this because the premise, again, *doesn't identify any asymmetry* between positive and negative claims. $(A \& B)$ is generally less probable than A, *regardless of whether B is a positive or a negative claim*. So the argument provides exactly as strong a reason for eschewing *negative* assertions as it does for eschewing positive ones. Here it is of crucial importance to avoid confusing a *negative claim* with the *absence* of a claim. To make a negative claim is to say, for example, that *x does not exist*; this is not merely a failure to assert that *x exists*.

To apply this to the issue of God: Suppose we can explain all our experiences using some collection of scientific theories, T. According to the above argument, we should believe only T; we should not add on G (the proposition that God exists), since $P(T \& G) < P(T)$. But by exactly the same reasoning, we also should not add $\sim G$, since $P(T \& \sim G) < P(T)$. Atheism is $\sim G$. So the probabilistic argument doesn't support atheism at all; it supports *agnosticism*.

Conclusion: I don't know why there would be a burden of proof on this issue, so I'll assume there isn't one.

10.3. The Problem of Evil

There is one leading argument against the existence of God, which spontaneously occurs to almost everyone who thinks about God: If there is a God, then why is there so much awful stuff going on in the world? The problem for theists of explaining the existence of evil is known as **the problem of evil**, and the argument for atheism based on the existence of evil is known as the **argument from evil**.

In the context of the "problem of evil", the word "evil" is used to refer to *anything bad*. Thus, not only malicious people are "evil", but cancer, earthquakes, global warming, and so on are all called "evils". This usage often confuses students; even after I explicitly explain this, they keep saying that cancer isn't "evil" because it's not a conscious being. So I will now just start referring to the Argument from Badness (which is what the argument from evil should really be called). Here are two formulations.

Argument from Badness, version 1

- 1.If God existed, then He would be all-knowing, all-powerful, and perfectly good.
- 2.An all-knowing being would be aware of everything bad in the world.
- 3.An all-powerful being would be able to prevent or remove everything bad.
- 4.A perfectly good being would be willing to prevent or remove everything bad.
- 5.If God were aware of all bads and both willing and able to prevent or remove them, then there would be no bads.
- 6.Therefore, if God existed, then nothing bad would exist. (From 1–5)
- 7.Bad things exist.
- 8.Therefore, there is no God. (From 6, 7)

Argument from Badness, version 2

- 1.A morally perfect creator would create only the best world that it could create.
- 2.An omnipotent being could create the best possible world.
- 3.Therefore, an omnipotent, morally perfect creator would create only the best possible world. (From 1, 2)
- 4.This world is not the best possible world.
- 5.Therefore, this world was not created by an omnipotent, morally perfect being. (From 3, 4)

Note: Both of those are deductive arguments, so they purport to show that our evidence is *inconsistent* with the existence of God. However, many atheists would take a weaker position: Many would just say that all the bad things we've seen constitute *evidence against* the existence of God (they call this **the evidential argument from evil**).

Now let me give an example to illustrate the argument from evil. Suppose there is a serial killer on the loose. (The FBI estimates, by the way, that there are between 20 and 50 serial killers loose in the U.S. at any given time.) He lives in a house up in the mountains, and that is where he takes his victims to torture and kill them. The killer has one neighbor: me. Would I do anything to stop the killer? There are a number of reasons why I might not do anything: Perhaps I don't *know* that my neighbor is a killer. Or perhaps I know but I have no way of stopping him. Or perhaps I'm afraid that if I try to stop him he will kill me. But suppose nothing like that is the case: I'm fully aware of my neighbor's evil deeds, as I can see him

dragging unconscious people into his house, later I hear their screams, and then finally I see him dragging the bodies out of his house. I also know that I can easily stop him, at no risk to myself (for instance, I could just call the police). Then, obviously, I would stop my neighbor. If I didn't, I would be a horrible person.

If there is a God, He is regularly in a position just like me in that example. Terrible things are happening to people all the time – including being tortured and killed by psychopaths, among many other intolerable ordeals. God, if He exists, sees all these things happening with perfect clarity. He could easily, instantly put a stop to them, at no cost to Himself. He is morally perfect and loves all of us deeply. Given all this, it makes no sense that He stands idly by. There is one obvious explanation for why God isn't stopping horrific suffering: He doesn't exist.

Theists want to argue that the existence of badness is *not* incompatible with God's existence, and is not very strong evidence against His existence either. We turn to their views presently.

10.4. Theodicies and Defenses

A **theodicy** is an attempt to explain why a perfect God would allow bad things to exist. There are many theodicies, some much better than others. We start with some terrible ones.

10.4.1. *How Do We Know What God Values?*

If you talk about the problem of evil in a class, you're pretty much guaranteed to hear at least one student say something like this: "How do we know what God considers good or bad? Maybe what is 'bad' to us is good to God."

In order for this suggestion to succeed as a theodicy, it would have to be that *everything* that happens that we consider bad God considers good. It would have to be, e.g., that God found the Holocaust delightful, that He likes childhood cancer, torture, war, and so on. (It can't be that he just values suffering, though, because then we'd have no answer to why he sometimes allows us to *not* be suffering.) This is just a thoroughly implausible and unmotivated suggestion.

Now, it may seem that philosophers pretend to doubt the obvious all the time (see, e.g., chapters 6–7) and therefore that this absurd suggestion should be right up our alley. I guess that's what students think. Here's what they're missing: We don't just pick any absurd thing and say it for no reason. When we're going to say some ridiculous thing, we give an

argument for it, starting from premises that seem at least somewhat plausible. The “how do we know what God values?” theodicy doesn’t do that – it just arbitrarily suggests something absurd, with no explanation or justification. Don’t do that. There’s no point to that.

By the way, another problem with this theodicy is that it tries to defend theism with a suggestion that no theist would accept. Theists in general have some strong views about what is good and bad, and, in my experience, none of them think that literally everything that happens is good. No theist thinks, for example, that sin is good, that all murders are good, etc. So they can’t respond to the problem of evil by saying that what humans consider bad is really good.

10.4.2. *How Would We Know What Goodness Is?*

Here’s another thing you’re guaranteed to hear in a class discussion on the problem of evil: “Without evil, how would we know what good was?” I guess the implied argument is that (a) if badness didn’t exist, then we would not know what goodness was, and (b) *knowing* what goodness is is itself such an enormous good that it outweighs all the bad things in the world.

This theodicy has three big problems. First, if He is omnipotent, God could simply will us to know what goodness is, and we would immediately know it. Then He wouldn’t have to allow all this badness.

Second, if for some reason God can’t do that (?), He could merely expose us to a *small* amount of badness, the minimum necessary to understand the difference between good and bad. We have no explanation for why He would allow the stupendous quantities of awfulness that we actually see around us.

Third, there is plenty of suffering that doesn’t help anyone to understand or appreciate goodness. The famous atheist biologist Richard Dawkins writes:

The total amount of suffering per year in the natural world is beyond all decent contemplation. During the minute it takes me to compose this sentence, thousands of animals are being eaten alive; others are running for their lives, whimpering with fear; others are being slowly devoured from within by rasping parasites; thousands of all kinds are dying of starvation, thirst and disease. It must be so. If there is ever a time of plenty, this very fact will automatically lead to an increase in population until the natural state of starvation and misery is restored. [\[61\]](#)

Dawkins presents a particularly gruesome example: There is a species of wasp that stings its prey, injecting it with a neurotoxin that paralyzes but does not kill it. When it is time to reproduce, the wasp paralyzes a caterpillar, then lays its eggs inside the caterpillar's body. When they hatch, the babies eat their way out. It is unknown whether the caterpillar can feel itself being devoured from within when this happens. But that's the sort of thing that goes on in nature. Nobody learns any valuable moral lessons from that. It's not as though the caterpillar is going to have an epiphany in which it understands the nature of good and evil and suddenly appreciates all the good times in its life, during its last moments of being devoured from within.

10.4.3. The Lord Works in Mysterious Ways

Here's another thing that people say: "God's ways are mysterious to us, because he is so much greater and more knowledgeable than us. We should just trust that he has a good reason for all the evils he allows, even if we can't see it."

Granted, this is *possible*. We are limited – we don't know everything, and we can't work out all the logical implications of what we know. So we cannot rule out with 100% confidence that there is a good reason for all the evils that we see. This calls into question premise 4 in both versions of the Argument from Badness ("A perfectly good being would be willing to prevent or remove everything bad"; "This world is not the best possible world").

That being said, this really is an extremely weak response to the argument from evil. If we allowed this sort of response to deter us from drawing conclusions, then essentially no theory would ever be rejected. Whenever there was evidence against a theory, the theory's advocates would just say, "Maybe there is some unknown explanation for the evidence that's compatible with our theory." They could say this as long as the evidence wasn't *deductively conclusive* (which it never is). We would still have people advocating Ptolemaic astronomy, the theory of the four elements, the four-humors theory of disease, etc. ^[62]

Granted, if you are absolutely convinced of God's existence to begin with, then it would indeed make sense for you to simply infer that God must have a good reason for allowing everything that he allows. But in a response to the argument from evil, you can't assume such a position – you

can't assume that we already know that God exists, because that's exactly what the atheist is denying, so you would be begging the question.

If, then, we start from an open-minded position (not assuming that God exists), and we see many evils for which we can find no explanation of why God would allow them, then we have no reason to assume that there is any good reason for allowing those evils. The simplest explanation is that there is no God.

At best, the present response to the problem of evil reminds us that the existence of evil is not *deductively conclusive proof* that there is no God. It does not block the conclusion that the existence of evil is *strong evidence* that there is no God.

10.4.4. Satan Did It

Perhaps all the bad things in the world are the responsibility of the Devil rather than God. This theory, however, really doesn't solve anything. If God is all-knowing, then he presumably knows about anything bad that Satan is about to do before Satan does it. Or *at least* he must know *while* Satan is doing it. If God is all-powerful, then God could always prevent Satan from committing any evil acts, or at least could immediately remedy them as soon as they happened. So we are basically left with the same problem we had at the start.

10.4.5. God Will Fix It

Sometimes theists say: Okay, there are some bad things going on in the world. But don't worry – God will make everything right in the afterlife. All misdeeds will be punished, all good deeds rewarded, and all unjust suffering compensated. Just be patient – God will get to it.

This response again fails to address the problem. It doesn't explain why any evil should ever exist in the first place. Remember, the philosophical problem is not to show that God is not horrible. The problem is to defend the idea that there is a *perfect* being – a being that is both *maximally* powerful and *maximally* good. Such a being would not have to allow bad things to happen and then fix them later; He could prevent anything bad ever happening to begin with. And if He somehow allowed something bad to happen, He wouldn't have to wait before fixing the problem; He could remedy the problem *instantly*. Presumably, it is better to avoid a problem than to allow it to happen then try to remedy it; and it is better to remedy a problem immediately than to wait. So a perfect being would prevent or immediately alleviate all problems. If, for example, existence in Heaven is

the best possible state, then God could simply put everyone in Heaven immediately, without taking a detour through corporeal existence.

10.4.6. Evil Is a Mere Absence

The above were all things that you sometimes hear from lay people but usually not from philosophers. Now let's get to some defenses that philosophers have devised.

Some theistic philosophers have tried arguing that evil isn't a thing. That is, they think that what we call "bad" is really only an absence of goodness, rather than something that exists in its own right. This goes back to Saint Augustine around 400 A.D. The theory was designed to avoid the conclusion that God created evil: God (according to Christianity) created everything that exists, so to avoid the conclusion that God created evil, we have to deny that evil is a thing. If evil is a mere *absence* of something, then there is no need for it to have been created by God.

There are two problems with this. First, the claim, if applied to all the bads in the world, is just about maximally implausible. Maybe *some* of the things we call "bad" are merely absences of good, but not all of them. Pain, for example, is not just an absence of pleasure, nor is it just an absence of goodness. Similarly, hatred is not a mere absence of love – there is a big difference between merely failing to love someone and hating them. Etc.

Second, the theory doesn't address the problem of evil as stated in §10.3. If the problem was to acquit God of committing a moral wrong by creating evil, then maybe the "absence" theory works. But a perfectly good being does not merely fail to create intrinsically bad things. It does not merely fail to perform wrongful acts. A *perfectly good* being *does the best that it can*. Regardless of whether evil is a presence or an absence, it remains the case that this world is not the *best possible* world. Hence, it is not the world that an omnipotent, perfectly good being would have created.

10.4.7. Evil Is a Product of Free Will

Perhaps the most popular response to the problem of evil is the **Free Will Defense**: Evil is a result of human free will, and free will is so valuable that it was worth it for God to give us free will, even knowing that it could result in enormous amounts of evil.

It is plausible that free will is intrinsically valuable. It is also important because having free will seems to be a necessary condition for having moral virtue or performing morally praiseworthy actions. Nevertheless, there are at least three important objections to the free will defense.

First, many atheist philosophers claim that God could have given us free will *and* at the same time ensured that we would never act wrongly. Allegedly, He could have done this by giving us only good desires to begin with. Since we would never *want* to do anything bad, we would freely choose to act rightly every time. You would be surprised by how many philosophers think you can “have free will” and also have your actions be completely predetermined. (This view is known as “compatibilism”. Compatibilism was also invoked in section 10.2.3 as part of a possible defense of theism.) However, it’s not worth spending time discussing this view now, because it’s very counter-intuitive, and it would take us far afield into discussions of the problem of free will. (We’ll return to compatibilism in ch. 11 on free will.)

A stronger objection would be that God could have given us free will and arranged things so that wrongful actions would be very *improbable*, though not impossible. The theist can’t plausibly deny that this is possible, since even we humans are often able to predict each other’s behavior with high *probability*, and we can often do things to make it more likely (though not guaranteed) that other people will behave rightly, despite that we all have free will.

Here is the second problem for the free will defense: There are many bad things that are not produced by anyone’s free will. Natural disasters are not anyone’s fault, nor are most diseases. (Granted, *some* diseases are caused by poor choices. But most are not.) Furthermore, the vast majority of suffering is endured by animals in the natural world (as discussed in §10.4.2) and is unrelated to human activities. So the free will defense could explain at most a fraction of the world’s suffering.

Finally, even for the evils that are produced by human free will, it seems that a just god would do something about many of them. For instance, once it became clear that Josef Stalin was a murdering, evil dictator, it seems that a just and benevolent god would stop him, perhaps after his first few murders, rather than leaving him in power until he died of old age, allowing him to kill tens of millions of innocent people and oppress and impoverish many more. It does not seem that the value of Josef Stalin’s freedom was so great as to outweigh all the suffering of the millions of people he harmed.

10.4.8. Evil Is Necessary for Virtue

Maybe God permits suffering in the world because it causes people to respond in virtuous ways and to develop better moral character. Many

virtues can only be exercised in response to something bad. For instance, one can only exhibit the virtue of compassion in response to the suffering of others. One can only exhibit the virtues of fortitude and perseverance in response to some form of adversity. One can only exhibit courage in response to danger. And so on.^[63] Furthermore, one might think (and many theists *would* think) that moral virtue is the most important good, so perhaps this outweighs the badness of all the suffering in the world.

This is the most plausible and insightful response to the problem of evil so far. Nevertheless, it too has problems.

First, we face problems with the notion of omnipotence again: If He is all-powerful, then God should be able to simply make us virtuous immediately, by divine fiat. Then He wouldn't need to allow suffering. In response, the theist would likely argue that virtue only has value – and perhaps only truly counts as virtue – if it is built up by the person through difficult choices. If God just made you a certain way, then you deserve no moral credit for it.

Second, although suffering can call forth virtuous responses, it also often calls forth vicious responses. Suffering can cause people to become bitter, callous, and resentful. When we look around at how people react to suffering (either their own suffering or that of others), it is not obvious that the virtuous responses outweigh the vicious responses.

Third, if this theodicy were correct, one would expect that the world's suffering would be distributed in such a way as to maximize moral development – e.g., the people who were most likely to respond virtuously would be the ones most likely to encounter suffering. But this does not seem to be the case at all; the world's suffering seems to be *randomly* distributed, as far as we can tell. For instance, there is a good deal of suffering by babies, small children, and mentally disabled people, who are unlikely to develop virtuous responses, as well as by people who react to suffering by becoming more vicious. The majority of the world's suffering is, again, endured by non-human animals, who can hardly be thought to develop moral virtue in response. (Nor do humans exhibit much virtue in response to animal suffering; quite the opposite, in fact.)

Finally, the amount of human suffering in the world has changed drastically over the course of history – people in previous centuries suffered a lot more than people today. How could this be explained? Call the typical level of suffering of past generations the “high level” of suffering, and call

the typical modern level of suffering the “low level”. If people need high levels of suffering to develop virtue, and virtue is the most important thing, then God should be making us have the high level of suffering today. On the other hand, if people only need low levels of suffering to develop virtue, then God should have ensured only low levels of suffering for the *previous* generations. Either way, the pattern of suffering that we see makes no sense.

10.4.9. *God Creates All Good Worlds*

In the Argument from Badness (version 2), we argued that God would create only the best possible world. Maybe this isn’t true. Maybe what a perfect being should do is to create *every* possible world that’s good. Furthermore, maybe God actually did that (sort of) – maybe he created an infinity of parallel universes, including all the possible universes that would be better than nothing. We find ourselves, perhaps, in a typical better-than-nothing universe.

Problem: God could not, in fact, realize every desirable possibility. That’s because some good possibilities are incompatible with each other. For instance, there is a desirable possible life I could have in which I am a computer game programmer (full-time, for my entire working life). There is another desirable life I could have in which I am a philosopher (full-time, for my entire working life). These two lives are incompatible with each other; I could not have both at the same time. (Just accept that these are incompatible jobs.) Likewise, a *world* in which I am, at the present time, a full-time programmer is incompatible with one in which I am, at this time, a full-time philosopher.

Granted, God could create *another person*, otherwise just like me, who was a game programmer at the present time, in another universe. But that person could not be *me*, since I’m here. I can’t literally be living two lives in two different places at one time. God could also give me a desirable life as a game programmer *at another time* (perhaps after I die, I’ll be reincarnated and live to be a game programmer!). But that could not be happening *now*, given that I’m now a philosopher.

So, given that God cannot give a person more than one desirable life at a given time, it seems that what He should do is give to each person, at each time, the best life that that person could experience. If you reflect on your life, and that of others around you, I think you’re going to quickly conclude that God has not in fact done that.

Complication: Maybe there are interactions between different lives, such that the best life for one person isn't always compatible with the best life for everyone else. In that case, what God should do is create the best possible overall collection of lives, across all of time. Now, granted, we can't see the rest of the distribution of sentient lives, apart from a small portion of the history of this planet. Nevertheless, what we see definitely does not look like part of the best possible distribution of lives. It seems that if the rest of the total (infinite) distribution of lives remained the same, but the lives people experienced *here and now* contained a lot less suffering, then the total would be better.

10.4.10. *There Is No Best World*

Maybe the reason God hasn't created the best possible world is that there *is no* best possible world. No matter how good a world was, it would always be possible to make it better. E.g., however many desirable lives there are, the world could always be made better by there being one more valuable life. Even if there are infinitely many desirable lives, you could imagine the same set of infinitely many desirable lives, but add one more desirable life not contained in the original set; that would seem to be an improvement.

Since there is no best possible world, God could not be expected to create the best possible world. Therefore, we should not object to theism merely on the ground that this world is not the best possible.

This reply is meant to preserve theism. However, it might instead show that traditional theism is impossible. Traditionally, God is supposed to be the greatest possible being. But just as (according to the argument) there is no maximum goodness that a world could have, maybe there is no maximum degree of greatness that a being could have. That would be plausible because, if a god creates a good world, it seems that you could always imagine another god who would have created a *better* world, and it seems that the second god would be greater than the first (perhaps the second god would be more powerful or more benevolent than the first). Thus, if there is no best possible world, it's plausible to conclude that there also is no greatest possible being.

But suppose you think that there could still be a greatest possible being, even though there is no best possible world. The greatest possible being wouldn't create the best possible world, since there isn't one. Still, it seems plausible that the greatest possible being would at least create an *extremely*

good world – one much better than *this*. A morally perfect being, surely, would not stand idly by while innocent people were tortured to death, for example. So the theist still has a good deal of explaining to do.

10.4.11. *The World Has Infinite Value*

Here’s another sort of reply that only a philosopher would think of: Maybe this *is* the best possible world; more precisely, it is tied for best with many other possible worlds. Suppose, plausibly, that the universe is infinite: Maybe there are infinitely many planets out there containing sentient life, though most are too far away for us to see. Or maybe there is infinite time, so an infinite number of valuable lives will exist during the universe’s infinite future. Or maybe there are infinitely many parallel universes that contain sentient life. In any of these cases, it is plausible that the total value of all of creation would be infinite. Since it is infinite, the value of this world that God (allegedly) created is as high as it could be; it is equal to the value of any other infinitely good world. God thus had no reason to prefer any other world over this one.

In response: It’s plausible that the universe has infinite total value. (Indeed, I am confident that the universe, including all of space and time, is *either* infinitely good *or* infinitely bad. I’m not sure which, though.) However, it remains very implausible to conclude that the universe could not have been improved upon in any way. Suppose that there is an infinite number of people, whom we can call “p1”, “p2”, “p3”, etc. Suppose that the goodness of a life can be measured on some numerical scale, where “1” represents a life that is just barely worth living, “100” represents a wonderful life, and so on. (Negative numbers represent bad lives, such that one would be better off not living at all than having such lives.) Now imagine two possible worlds, in which the infinitely many people would have the following welfare levels:

		Persons				
		p1	p2	p3	p4	...
Welfare levels	<u>World 1</u>	1	1	1	1	...
<u>World 2</u>	100	100	100	100	...	

In World 1, everyone has 1 unit of good in their life, so their lives are just barely worth living. In World 2, the same people exist, but they all have 100 units of good, making their lives wonderful. Which world is better?

Both of them have infinite total value. But come on, obviously World 2 is better. If you had a choice between those two options, and you went for World 1, you'd be an asshole. We wouldn't buy the excuse that the worlds were "equally good" since they both had infinite total value. Similarly, even if the value of this world is infinite, God could have made it better by giving everyone wonderful lives.

10.4.12. Weakening the Conception of God

We've now been through eleven solutions to the problem of evil. I think they all failed. But there is one approach that might succeed: Weaken your conception of "God". Instead of insisting that God must be all-knowing, all-powerful, and maximally good, you could say that God is merely *very* knowledgeable, powerful, and good. This opens up the possibility that God doesn't know about all the evils in the world, or he can't eliminate them, or for some reason he accepts a certain amount of badness. Maybe creating universes is hard, and God did the best he could. (*You try creating a better universe some time!*)

Most theists dislike this approach because they are attached to the traditional "O³ world-creator" conception. However, we should not limit ourselves to traditional conceptions – nobody ever promised us that the truth about the universe had to be some simple idea found in our historical traditions. If you buy the Fine Tuning Argument, for example, it suggests that there is an intelligent designer, but it doesn't tell us that the designer has to be a "triple omni" god.

Nevertheless, here's a problem for this approach: Even if God can't eliminate *all* the evils in the world, we should expect him to at least show some signs of trying to reduce the world's evils. Many human beings are working to alleviate suffering, and they seem to be doing a better job of it than God. Bill Gates and Warren Buffett, for example, have saved literally millions of lives through their poverty relief efforts. God, as far as we can tell, hasn't saved *anyone*. (Of course, you could claim that maybe God is invisibly preventing lots of disasters that we don't know about, but there's no evidence of this.) If there's a God, surely He ought to be able to do a better job of helping people than a mere human can do. And if he did, surely we would expect to see at least some evidence of this.

Counter-reply: God has bigger fish to fry. The above argument is very anthropocentric – we think God isn't doing anything because we don't see Him doing anything *for us*, here on Earth. But the universe is a big place. If

God created it, then he has *100 billion galaxies* to attend to, and in *each* of these galaxies, 100 billion stars. Some of those (we don't know how many) contain entirely different species of sentient life, each with their own tribulations. We don't even know if that is all of creation; there might be more out there that we can't see. If I had 10^{22} stars to attend to – even if I was *extremely* powerful, and knowledgeable, and good – you probably wouldn't see me around much either.

10.4.13. The Case of the Serial Killer

Return to the example from §10.3: I have a serial killer neighbor, whom I know about, yet I do nothing to stop him. This is a good example for illustrating not only the initial problem of evil, but also the implausibility of most replies to the problem. Let's say that after the police catch the killer, they come over to my house to question me. Perhaps the following dialogue would ensue:

Cop: Professor Huemer, did you know that your neighbor was a serial killer?

Me: Oh yes, I knew everything.

Cop: Well, didn't you know that you could stop him by calling the police?

Me: Yes, of course, I could easily have stopped the killings. I just chose not to.

Cop: Why the hell not? Were you afraid of your neighbor or something?

Me: Oh no, I knew he couldn't harm me.

Cop: Then are you some kind of sicko? Did you actually want those people to be horribly tortured and killed?

Me: Oh no. In fact, I love all people very much.

Cop: What, then? Why didn't you do something??

How might this dialogue continue, such that you'd agree that I was still a decent person? Imagine filling it in with one of the above proposed responses to the problem of evil. I might say one of the following...

(1) How do you know that torture and murder are really bad? Maybe they're good. (See §10.4.1.)

(2) I wanted to help people truly understand the difference between good and evil. Also, I thought the murders would help people appreciate all the good things in the world. (§10.4.2.)

(3) I had good reasons that I refuse to explain to you. (§10.4.3.)

- (4)The murders were the fault of the serial killer, not me! (§10.4.4.)
- (5)It's okay because all those victims are in heaven now. (§10.4.5.)
- (6)Allowing the murderer to continue wasn't an evil act by me; it was merely the *absence* of a *good* act. So get off my case! (§10.4.6.)
- (7)I didn't want to interfere with the murderer's free will. (§10.4.7.)
- (8)I wanted people to have virtuous responses to the murders. People reading about them in the newspapers experienced compassion. You cops got to exhibit perseverance and industriousness in looking for the killer. The victims had a chance to exhibit courage and hope during their ordeal. (§10.4.8.)
- (9)The total value of the world is infinite, and I knew that it would still be infinite regardless of whether I stopped the serial killer. So there was no reason to stop the killer. (§10.4.11.)

I submit that all of the above are terrible reasons. (A couple of the theodicies don't lend themselves well to parodies in this example, though.) If I said any of those things, the cops should conclude that I'm a terrible person. On the face of it, then, they are not good explanations for why God allows horrible suffering either.

The Free Will defense and the Virtuous Response defense (#7–8 above) might be exceptions, however. You might think that if *I* stop a serial killer, that would have different effects from the effects if *God* stops him. If God does it, maybe that would cause humans to become complacent and stop caring about problems in the world, since they'd assume that God would take care of everything. That won't happen if I do it, because everyone will still know that I can't be expected to fix all the world's problems.

10.5. Conclusion

The strongest argument for atheism is the evidential argument from evil: If there were a God, we'd expect the world to be a lot better than it is. As far as we can tell, suffering is randomly distributed in the world, just as it would be if there were no one in charge. There is no pattern suggesting *any* larger purpose – e.g., we don't see suffering distributed according to who would learn the most from it, or who justly deserves to suffer, or who has or hasn't endorsed a particular religion, *or any other pattern at all*. This doesn't deductively entail that there's no God, but it is, evidentially speaking, as bad for theism as it could be.

Most theistic responses to the problem are awful. Most of them give reasons for God to allow evil that, if sound, would also be reasons why *you*,

for example, shouldn't stop evil when you have the chance. That's unacceptable. However, not *all* responses are open to this objection. Some evils might be explained by human free will, and God might allow some others to enable us to develop moral virtue. The most adequate response, which applies to all the evils we see on Earth, would be to deny that God is really all-powerful.

A form of theism that rejects divine omnipotence is thus defensible. However, the evidence still weighs against it to some degree, simply because we do not see any sort of pattern in the distribution of suffering in the world, nor any clear indications of divine intervention in the world, in pursuit of any goal. I here invoke the following theorem of probability:

$P(h|e) > P(h)$ if and only if $P(h|\sim e) < P(h)$.

In words: The *truth* of *e* would *raise* the probability of *h*, if and only if the *falsity* of *e* would *lower* the probability of *h*. (I'm not going to give the proof of that right now, but that's a very straightforward theorem.) We *could have* seen some meaningful pattern in the distribution of suffering, or some indications of divine intervention, and if we did, theists would certainly cite that (correctly) as evidence *for* the existence of God. Therefore, our failure to see those things is evidence *against* the existence of God.

My own attitude is agnostic: There is some reason to think the universe has a creator (§9.4), but also some reason to doubt this, as just explained. None of the arguments on either side is decisive. If there is a creator, though, it's probably *not* a triple-omni being.

11. Free Will

11.1. The Concept of Free Will

Free will is a capacity that humans are often thought to have, an ability to choose freely among competing alternatives. There are two conditions that are required for free will:

- i. *Alternate possibilities*: You must sometimes have more than one thing that you can do.
- ii. *Self-control*: You must have control of your own actions.

These are widely accepted requirements for free will, though people disagree about how to further analyze those conditions (e.g., how to analyze “control” or “can”). To see the plausibility of both conditions, contrast two examples of things that lack free will. First, a robot: The robot, in an obvious sense, controls its own actions (its computer “brain” selects a course of action and makes the robot body carry it out). However, the robot does not truly have *alternative possibilities*: At any juncture, it must always do whatever its programming dictates. That is why, intuitively, it doesn’t have free will.

Second, consider a radioactive atom that may either decay or fail to decay in the next hour. According to some interpretations of quantum mechanics, which thing happens is completely undetermined – there is no sufficient cause that explains why it does one thing rather than another, and there is no way in principle, even with complete information about the initial state of the atom and its environment, to predict when the atom will decay. So the atom has alternative possibilities. It does not, however, *control* its own decay; it just decays randomly. And that is why, intuitively, the atom doesn’t have free will either.

Human beings are different – or at least, so it seems. It appears that we often have many options open to us and that it is up to us which option is realized. You could either keep reading this chapter, or get up and take a walk, or throw the book out the window, or ... lots of other things. And which of those things happen is under your control. (But please keep reading.)

Note that, obviously, free will doesn’t require being able to do *anything you want*. I’d like to fly to Alpha Centauri this weekend, but I can’t. That

doesn't mean I don't have free will. Having free will also doesn't require that *all* of your choices be completely free; maybe sometimes you lose control of yourself (you know, like when someone misuses the expression "beg the question", and you fly into a fit of rage – you couldn't be blamed for this). Rather, to "have free will", it is only required that *sometimes* you have *more than one* option that you can freely choose among.

Why is free will important? Partly, it's important simply because it's a key aspect of how we experience life – we seem to observe ourselves making free choices all the time, and that's central to our self-understanding. Also, it's important because free will is presupposed by various other judgments. If you say a person *should* perform one action rather than another, that seems to presuppose that the person can freely choose between those options. Likewise, it seems that it would be inappropriate to either *praise* or *blame* a person for what they have done unless their actions are up to them. If you have a retributive theory of punishment (i.e., you think people who do bad things should be punished because they *deserve* it), that also seems to presuppose free will, since people don't deserve punishment for what they couldn't help doing.

11.2. Opposition to Free Will

There are two main threats to free will. Traditionally, people have worried that maybe we lack free will because everything that happens is *determined*. Today, some people also worry that even if our actions are *undetermined*, we would *still* lack free will.

11.2.1. The Theory of Determinism

Let's start with the threat of **determinism**. There are a couple of possible ways of defining "determinism":

Determinism(1): The view that every event has a *sufficient cause*, i.e., a cause such that the event had to happen, given that the cause happened.

Determinism(2): The view that the state of the universe at any given time, together with the laws of nature, determine a unique future.

These are both reasonable ways of defining "determinism", and they both lead to a threat to free will if one thinks that determinism is true. Determinism(1) is a more traditional version of the view; determinism(2) is a more modern and more precise formulation.

To explain determinism(2) some more: Suppose you were given the exact state of every particle in the universe (and, if there are things other

than particles, the complete state of everything else too) – including the exact value of each particle’s mass, electric charge, velocity, location in space, and so on. Suppose also that you knew all the laws of nature (such things as the conservation of momentum, the law of gravity, the conservation of energy, etc.) I mean the *complete and correct* account of the laws of nature, not merely the laws thus far discovered by human beings. Determinism(2) says that, from those things, it would be possible to deduce precisely everything that would ever happen after that time.

Clarification

The point here is really not about what *a human being could figure out*. So don’t complain that it would take too long to do the deduction, or that you’d need a computer the size of the universe, etc. We can restate the point without referring to a person doing predictions: The essential point is that there is one and only one complete future evolution of the world that would be consistent with both the current state of the world and all the laws of nature. So, in order for the future to be any different from what it is actually going to be, either the past would have to be different or the laws of nature would have to be different from what they are.

That’s the doctrine of determinism. You’d be surprised at how many people believe something like that. (At least, I’ve been surprised by it.) Back in ancient times, some people thought that everything was predetermined because the gods, or the Fates, had already planned everything out. During the middle ages, some thinkers worried that everything was predetermined because events in our lives could be accurately forecast by astrology. (No joke!) Also in the middle ages, Christian theologians worried that God’s omniscience might mean that everything had to be predetermined (see §10.2.3). After Isaac Newton founded classical physics in the 17th century, people worried that everything was predetermined because everything could in theory be predicted from the laws of physics and the physical state of the world. People are still worrying about that one today. Also, in the twentieth century, some overconfident psychologists started claiming that everything a person does could be predicted from the environmental influences on the person, or from this together with the person’s genes, or something like that. (No one claims to actually be able to *do* these predictions, mind you; people just speculated that in theory it would be possible.) Some people today worry that all our choices are determined by microscopic events in our brains.

I'm telling you all this so that you can appreciate the variety of reasons why people have doubted free will throughout the ages, most of which have been very lame reasons. I think this makes it reasonable to speculate that many human beings have some sort of natural bias toward determinism, which they then rationalize by coming up with theories about what forces, other than themselves, are responsible for everything that happens.

11.2.2. Evidence for Determinism?

Be that as it may, what evidence do we today have for determinism? Some people appear to view it as a self-evident truth, in need of no argument. People used to talk about “the law of causality” (which is basically determinism(1)) as a necessary truth.

Also, the physics of the seventeenth through nineteenth centuries was deterministic – the laws enabled you to predict exactly how a physical system would evolve, given its initial conditions. Since human bodies are made of atoms just like all the other physical objects in the world, you might think that the movements of human bodies (and thus, all overt human actions) are determined by the initial conditions of the physical particles in and around them, given the laws of nature.

You might wonder: “What about our thoughts, feelings, desires, and so on? Surely they have something to do with our actions; we're not just inanimate objects like billiard balls.” Well, the determinists would say that those thoughts, desires, etc., are either identical with, or completely caused by, physical events in our brains. E.g., you have some pattern of neurons firing in a certain part of your brain, and that's what makes you move your arm. The thought, “I'm going to pick up this cup” either *just is* the pattern of neural activity or is *caused by* the pattern of neural activity, which also causes the arm motion.

The biggest problem with this view is that most physicists abandoned determinism in the twentieth century. The most prominent interpretations of quantum mechanics are indeterministic – they allow for events to occur that lack sufficient causes. E.g., the decay of a radioactive atom is said to be completely random. This does not seem to be metaphysically impossible, nor is it ruled out by modern science, so the traditional arguments for determinism (left over from before the 20th century) are weak.

To be clear, indeterminism hasn't been proven either. Right now, there are multiple different *interpretations of quantum mechanics*, as we say – basically, these are different theories in physics that use the same equations

and explain the same experimental results, but they give different explanations for what the equations mean and why the experimental results happen.^[64] Of these different theories, some are deterministic and some are indeterministic. So there's no consensus among physicists on whether determinism is true. This makes it pretty weird that there are lots of amateur philosopher people who seem to be totally convinced of determinism.

11.2.2. No Free Will Either Way

Anyway, here's a stronger argument against free will: Some people say that we couldn't have free will even if determinism is false. They say that if our actions are undetermined, that merely means that they are *random*, like the events described in quantum mechanics. By the way, some smart people think that quantum events in the brain actually set off chain reactions that lead to macroscopic human actions.^[65] So it's not totally crazy to think that quantum events actually render human actions indeterministic. However, the anti-free-will people would say this just means that some human actions are random. But a random event isn't under our control, so we still wouldn't have free will.

Notice that there is a tension between the two conditions for free will: (i) alternate possibilities and (ii) self-control. If determinism is true, then we lack alternate possibilities. But if indeterminism is true, then we lack self-control since our behavior is merely random. So we lack free will either way. Or so one might argue.

11.3. Deterministic Free Will

11.3.1. Compatibilism

The most popular view of free will among philosophers is one that virtually no one who isn't a philosopher likes. The theory is known as **compatibilism**. It holds that free will and determinism are compatible – and of course, the point of claiming that these things are *compatible* is generally so that you can say that we in fact have free will even though determinism is true. I.e., even though every action is completely determined by antecedent causes, even though everything you do was necessitated by events in the remote past together with the laws of nature, still, many of your actions are perfectly free, up to you, under your control, etc. (This view is sometimes called “**soft determinism**”, to distinguish it from “**hard determinism**”, which holds that we lack free will due to the truth of determinism.)

Almost every student who hears about compatibilism thinks it's absurd, even contradictory. In fact, if you don't think it's nonsensical at this point, you probably misunderstood it, so read the preceding paragraph again. Nevertheless, you'd be surprised by how many philosophers have historically held this view. It's probably the plurality view among philosophers today, too. How could someone hold this seemingly contradictory view?

11.3.2. *Analyses of Free Will*

Here's the first thing to say about compatibilism, which really is some evidence for it: Every time a philosopher tries to clearly and precisely analyze the notion of free will, they end up with a compatibilist analysis. You might think this is false because my account of free will in section 11.1 includes the "alternate possibilities" condition, which looks incompatible with determinism. What I mean, though, is that when we try to further analyze what it means to *be able* to do something, it generally turns out that it's compatible with determinism that a person should be able to do more than one thing.

The key point, for compatibilists, is that there are different senses of "possible". So an action may be impossible in one sense but possible in another. Let's say that you're thinking about throwing this book (or e-reader, or computer) out the window. You're not actually going to do it, but, as we would normally say, you *could* do it. What does it mean that you "could" throw it out the window?

Maybe it means something like this: If you tried to throw it out the window, you would succeed. Notice that if *that's* what "could" means, then it's perfectly compatible with determinism that you could either throw or not throw the book. Determinism(1) says that your non-throwing of the book will have a sufficient cause, just like all other actual events (or non-events?). This cause will include something like the fact that you want to keep reading the book. Because you want to keep reading it, you won't try to throw it out the window, and so you won't in fact throw it. All that is perfectly compatible with saying that *if you tried* to throw it out the window, you would easily succeed. That proposition doesn't require there to be any uncaused events.

Likewise, consider determinism(2), which says that the future is fully predictable (in principle) from the complete state of the world at any time plus the laws of nature. Of course it's compatible with that to say that if you

imagine a *different* state of the world, then you'd predict a different future. The hypothetical, "If you tried to throw the book, you would succeed" asks us, in effect, to imagine a different state of the world from the one that actually obtains at present. It may well be that, given the current actual state of the world, the laws of nature determine that you don't throw the book out the window, but, at the same time, if the state of the world were different in such a way that you currently had the intention to throw the book out the window (with everything else being normal), the laws of nature would determine that you *do* throw the book out the window in the near future. Thus, despite the truth of determinism, it may be true that "you could throw the book out the window" and also that "you could refrain from throwing the book out the window".

That was a simple illustration of a compatibilist analysis. No one thinks that analysis is *exactly* right – i.e., that "x can do A" means precisely "if x tried to do A, then x would succeed". Like all philosophical analyses, there are going to be counter-examples to that. (See §8.5 on this sort of problem.) Even so, I think this illustrates how it's not completely ridiculous to think that determinism might be compatible with free will. If freedom means *anything like* what we've suggested above, then it's going to be compatible with determinism.

Here's another idea. Maybe free will is a matter, not of *whether* one's actions are caused, but of *how* they are caused. E.g., are they caused *internally*, by your own desires and values, or are they caused by *external* forces? To illustrate this idea, compare two imaginary bits of dialogue (these are taken from a compatibilist philosopher).^[66] First exchange:

Jones: I once went without food for a week.

Smith: Did you do that of your own free will?

Jones: No. I did it because I was lost in a desert and could find no food.

Second exchange:

Gandhi: I once fasted for a week.

Smith: Did you do that of your own free will?

Gandhi: Yes. I did it because I wanted to compel the British Government to give India its independence.

(This alludes to the hunger strikes that Mahatma Gandhi famously engaged in when he was fighting for Indian independence from Britain.) Notice that both of these exchanges sound perfectly natural. That shows

that we're using the ordinary notion of "free will", not some weird philosophers' notion. Notice also that in the first exchange, the "action" (going without food – if you want to call that an action) is unfree because it was imposed on the person by external circumstances. In the second exchange, the action is free, not because it had *no* cause, but because it was caused by the person's own beliefs, values, and desires.

And notice, again, that if something like this is the right way of understanding free will, then free will is perfectly compatible with determinism. It doesn't require any event to be uncaused. It just requires that your actions have *the right sort* of causes.

There's more to say about what the right kind of internal causes are. E.g., we probably want to exclude psychological compulsions (even though they are, intuitively, "internal"). We might further want to require that the person's actions be caused by motives that they endorse on reflection, or something like that (thus excluding, e.g., the behavior of most addicts). There's a lot of literature on working out those sorts of details. But we won't go into that now. The important point is that, according to many philosophers, a free action is an action with the right kind of causes, not an action with no cause.

11.3.3. Freedom Requires Determinism

Here's another reason for thinking that freedom is compatible with determinism: Arguably, freedom actually *requires* determinism. Here, we hark back to section 11.2.2. If you perform an action that has no cause, that means that the action must not be caused by your own desires, beliefs, or other motives. It's just some random event. That seems to entail right away that the action wasn't under your control and wasn't free.

Now, the opponents of free will would like to declare at this point that since free will is *also* incompatible with determinism, it's impossible to have free will. But that's not the most plausible conclusion. Once we realize that indeterminism conflicts with free will, it is most plausible to think that – since free will doesn't seem to be a blatantly inconsistent or meaningless idea – it must be that free will is compatible with determinism instead ... maybe because one of the analyses of free will suggested above (§11.3.2) is correct, or because of some other compatibilist analysis. If free will requires determinism, then it is probably compatible with determinism.

11.4. Libertarian Free Will

In the free will literature, “libertarianism” denotes the view that we have free will *and* this free will is incompatible with determinism. This is the most commonsensical view, though it is not the most popular among philosophers. (*Note*: Please do not confuse this with libertarian political philosophy.)

11.4.1. Incompatibilism

Due to the popularity of compatibilism among philosophers, we need an argument to show that freedom and determinism are incompatible with each other. Here’s one sort of argument, which people call “the Consequence Argument”:

If determinism is true, then our acts are the consequence of laws of nature and events in the remote past. But it’s not up to us what went on before we were born, and neither is it up to us what the laws of nature are. Therefore, the consequences of these things (including our present acts) are not up to us.^[67]

Here’s a more elaborate statement. Let “P” stand for a complete description of the state the universe was in at some time in the remote past, e.g., at the time of the Big Bang. Let “L” stand for a complete statement of all the laws of nature. And let “A” stand for a (true) statement describing some action that I perform in the present. It seems that no matter what I do, P will remain the case; that is: For any action that I can perform, if I perform it, the past will remain as it in fact was. Similarly, it seems that no matter what I do, L will remain the case (I can’t do anything such that the laws of nature might be different, or the laws might be violated, if I did it). Therefore, no matter what I do, (P & L) will be the case. But, if determinism is true, then (P & L) *entails* A. That’s just how we defined determinism(2). So, if determinism is true, then no matter what I do, A; that is, for any action that I can perform, if I perform it, A will remain the case. That means that (given determinism) I have no alternative possibilities, so I lack free will. In brief:

- 1.No matter what I do, P.
- 2.No matter what I do, L.
- 3.Therefore, no matter what I do, (P & L).
- 4.If determinism, then (P & L) entails A.
- 5.Therefore, if determinism, then no matter what I do, A.

That’s the best known argument for incompatibilism. Now here is a related but slightly different argument.^[68] In general, it seems that if, in

order for me to do x , something would have to have happened in the past that did not in fact happen, then I cannot now do x . Here are a couple of examples to illustrate this principle:

Class Grade: A student asks me, near the end of the semester, “Hey Prof, how can I get an A in your class?” I reply: “In order to get an A in my class, you would have to have scored at least 87.5% on the first four tests. But in fact, you scored only 73%.” The student responds, “So you’re saying I can’t get an A.” This would obviously be a correct inference.

Global Warming: A scientist is testifying before Congress about global warming. A senator asks him, “What can we do to ensure that there is no further global warming after today?” The scientist replies: “Well, there is a 40-year delay between when greenhouse gases are released into the atmosphere and when they have their full effect on global temperatures. Therefore, in order to avert any further global warming after today, we would have to have stopped greenhouse gas emissions 40 years ago – which, as you know, we did not do.” The senator could obviously infer the conclusion: We cannot avert further warming.

You can think of any number of examples like that. Now, if determinism is true, then in order for me to do *anything* different from what I’m actually going to do, the past would have to have been different from what it in fact was. Therefore, if determinism is true, then I *cannot* do anything different from what I’m actually going to do. Which means I don’t have any alternative possibilities, so I don’t have free will.

11.4.2. For Free Will: The Appeal to Introspection

We just argued that determinism conflicts with free will. Some people would conclude that we in fact have no free will (as in §11.2). Not the libertarians, though. The libertarians instead infer that determinism must be false, since we obviously *do* have free will.

Why believe in free will? One reason is introspection: It just seems, introspectively, that we can often observe ourselves freely choosing among multiple options. And nearly everyone, even the determinists, agrees that our choices generally *feel* free.

In response, hard determinists often say this is an “illusion”, adding that we are determined to experience this illusion. This reply, however, is very lame by itself – that is, if not supplemented with actual evidence that *shows* free will to be illusory. In general, if we seem to observe something, it is

very lame to simply say, “Oh, that’s an illusion” and move on. Rational people assume that what we seem to observe is real, unless there is evidence to the contrary; they don’t assume that whatever we seem to observe is illusory until proven real (see §7.6).

Example: Silly Sam claims that there are no blue foods in the world. Smart Sally shows Sam a blueberry, then eats it in front of him, thus refuting Sam’s claim. Sam declares: “Oh, that’s just an illusion.” He gives no further argument, evidence, or explanation. When asked why he thinks it was an illusion, Sam replies, “Because there are no blue foods.” In this case, Sam’s defense is dialectically pathetic. That’s just like the determinist who declares “free will is an illusion”, with no further justification.

Here is a better reply on behalf of the (hard) determinist: Free will requires that you have *unrealized possibilities* – things that you could have done but *did not actually do*. But the only things a person can ever observe, even in principle – whether it be introspective observation or observation by the five senses – are things that actually happen or actually exist. For instance, if there is a cat on the table, you can observe that, but if there *isn’t* actually a cat but there merely *could have* been one, you can’t observe *that*. There’s nothing that a merely possible cat looks like. Similarly, you can’t observe merely possible actions, so you can’t observe that you have alternative possibilities, so you can’t observe that you have free will.

Nevertheless, it strongly seems as if we have free will. Perhaps this sense of our freedom isn’t exactly an *observation* (since, as claimed above, one can’t observe unrealized possibilities); perhaps we should call it an “intuition”, or something like that. Regardless of whether it counts as “observation” or not, one should generally assume that things are as they seem to be, until proven otherwise – and there is no dispute that we *seem* to have free will.

Another thought is that perhaps the availability of various actions is a purely negative condition: When you are choosing among some set of (apparent) options, each option is available as long as there *isn’t* anything to *stop* you from choosing it. And perhaps you can introspectively detect the absence of any impediments that would stop you from choosing any of multiple different options. By the way, maybe there are also some *positive* requirements for free will, like self-control, but they aren’t what is at issue here – the argument from the previous paragraph was only questioning the possibility of observing unrealized possibilities. So the idea here would be

that we can observe that there are certain unrealized possibilities by observing the absence of relevant impediments.

11.4.3. Free Will and Other Common Sense Judgments

Here's another reason to accept free will: As noted earlier, it's implied in lots of other judgments that seem obviously correct. Example: During his reign over the Soviet Union, Joseph Stalin sent about 14 million people, including millions of political prisoners, to the gulags – forced labor camps in which the conditions were so harsh that about 1.6 million of them died. No one made Stalin do this, nor did he have any justification for it; he did it partly because he wanted free labor, and partly because he wanted to terrorize all who might question his rule. Now, it seems that Joseph Stalin was obviously blameworthy for doing that. Yet those who deny free will would seemingly have to say that Stalin was completely blameless, since he couldn't help doing it. (Implicit assumptions: If you lack free will, then you're not responsible for your actions; if you're not responsible, then you're not blameworthy.)

Here's another example. During World War II, the Nazis rounded up Jews in Germany and shipped them off to concentration camps, where, near the end of the war, millions were ultimately executed. Oskar Schindler was a German industrialist at the time, who convinced the Nazis to allow some Jews to work in his factories instead of going to the concentration camps. As the war went on, Schindler repeatedly interceded on behalf of his workers, at great risk to himself. He ultimately expended his entire fortune on bribes to Nazi officials to keep them from deporting and killing his workers. He is credited, in the end, with saving a total of about 1200 Jews from execution.^[69] Now, it seems that Oskar Schindler was worthy of praise for this behavior. Yet those who deny free will would seemingly have to claim that Schindler deserves no credit, since he wasn't responsible for any of this.

By the way, not all of the absurd judgments implied by determinism are *moral* judgments. Here's a third example. I have a fork on the table in front of me. Now, it seems to me that I should not stick that fork in my eye right now. (Not because doing so would be immoral, but because it would be rather imprudent. It would hurt, it wouldn't satisfy any of my goals, etc.) But to say that someone either "should" or "should not" do some action seems to imply that they have a choice about whether they do it. There's no sense in saying that I *should* do x if I cannot do it, or I cannot avoid doing

it, or I have no control over whether I do it. So those who deny free will would seemingly have to disagree with the judgment “I should not stick this fork in my eye.”

That should help to explain why I personally find it crazy that so many intellectuals deny free will.

11.4.4. *For Free Will: The Self-Defeat Argument*

Many have found something self-defeating about being a determinist. This general thought goes back at least to Epicurus around 300 B.C.:

The man who says that all things come to pass by necessity cannot criticize one who denies that all things come to pass by necessity: for he admits that this too happens of necessity.^[70]

The British philosopher J.R. Lucas argued similarly:

Determinism ... cannot be true, because if it was, we should not take the determinists' arguments as being really arguments, but as being only conditioned reflexes. Their statements should not be regarded as really claiming to be true, but only as seeking to cause us to respond in some way desired by them.^[71]

In general, problems arise when you try to apply the doctrine of determinism to our thoughts about determinism itself. There are various ways the determinist might then be in a self-defeating position:

- a) As Epicurus says, if the determinist *criticizes* the belief in free will, or argues that we should accept determinism instead, this seems incoherent. If determinism is true, those who believe in free will are determined to do so; hence, they couldn't be criticized for this, nor could it be true that they *should* believe something else.
- b) Even presenting an argument for determinism may be self-undermining. An argument is an attempt to give one's audience a *reason* for believing a certain conclusion. But if determinism is true, any given person either cannot believe determinism or cannot avoid believing it. It doesn't make sense to try to give someone a *reason for* doing something that's impossible, or unavoidable, or that the person one is addressing has no control over.
- c) When we reason about free will and determinism, we are deliberating about what to believe. But deliberation presupposes the existence of alternatives that one has control over. It makes no sense to deliberate about whether to do A if you don't think you have any choice about it.

- d) It seems that any rational belief or thought process has to be governed by norms, such as that one should prefer to believe true things over false things – or at least things that are likely to be true over things that are likely to be false, or something like that. But any norm about what should be done presupposes that there are alternatives that we have some control over.

Now, because most determinists are confused people (like people in general, I guess), I have to clarify some things about that sort of argument. The argument does not claim that it's logically impossible for determinism to be true. The argument claims that it's impossible for you to *rationaly endorse* determinism. I.e., if you believe determinism, then you're being irrational.

Determinists sometimes respond to the above sort of argument by saying merely that it doesn't *prove* that determinism *can't be true*. That's right. But that's also not a defense. It's not as if the only possible criticism of a view is that the view couldn't be true. One can also criticize a view by arguing that it can't be justified or rational. E.g., suppose someone announces that the number of stars in the universe is even. Call this view "even-numberism". I object that even-numberism is arbitrary and it's equally likely that there are an odd number of stars. The even-numberist retorts: "That doesn't prove that my view isn't true!" I respond: "Yeah, I know. But it does show that it's irrational for you to believe it. Try addressing the actual criticism."

To head off another misunderstanding: I'm not saying that determinism is exactly like even-numberism. The point of that example is that *to criticize a view, you don't have to prove that it's false*. It's also a good criticism if you can show the view to be unjustified. Now, the reason why even-numberism is unjustified is different from the reason why determinism is unjustified, but they're both unjustified views. The reason even-numberism is unjustified is that there's no evidence for it and the base probability is only 50%. The reason why determinism is unjustified is that it contradicts other beliefs that all rational thinkers must have.

Here is a way of making the self-refutation explicit, which I thought of when I was an undergraduate student a few decades ago:

1. About the free will issue, we should believe only what is most likely to be true. (Premise – presupposed by reasoning)

2. In general, if S cannot do A, then it is not the case that S should do A. (Premise)
3. Therefore, with respect to the free will issue, we *can* believe only what is most likely to be true. (From 1, 2)
4. If determinism is true, then there is only ever one thing that a person can do; we never have alternative possibilities. (Def. of determinism)
5. Therefore, if determinism is true, then we believe only what is most likely to be true. (From 3, 4)
6. I believe in free will. (Premise based on introspection)
7. Therefore, if determinism is true, then the belief in free will is what is most likely to be true. (From 5, 6)

Step 7 shows how belief in determinism is self-refuting: If we assume determinism, then we can infer that the doctrine of free will is probably true instead.

In my experience, almost all philosophers who hear this argument *hate* it. They immediately feel as if it's just some sort of cute trick, and therefore they assume there's something wrong with it even if they can't say what it is, and they aren't very interested in examining it.

That's why I started with the discussion of other self-refutation arguments (e.g., Epicurus), because no one says *those* are mere tricks. Rational thought presupposes things like "we should only believe what is likely to be true", and the 7-step deduction above just draws out how any such norm conflicts with determinism.

11.5. Other Reflections

11.5.1. How Does Libertarian Free Will Work?

I've argued that we have free will (§11.4.2–11.4.4) and it's incompatible with determinism (§11.4.1). So determinism is false. That doesn't tell you much, though, about how free will works, and in fact I don't know how free will works.

Some libertarians say that the key to free will is **agent causation**. (Btw, an "agent" in this context is just a being who performs actions.) Ordinary causation involving the physical world is usually thought of as **event causation** – that is, you have one event causing another event to happen. But in the case of the choices of beings with free will, allegedly, an event is caused by *the agent* (the person), not by another event. Your choice to eat a banana is not caused, e.g., by certain thoughts, feelings, or desires going

through your mind; it is caused by *you*. Non-libertarians usually say they can't understand this view – they don't know what it means for an event to be caused by a person rather than by an event or state of affairs. And indeed, there is not much that can be said to further explain it. But anyway, this view might reconcile free will with determinism(1), but not with determinism(2): It may be that all events have causes, but what happens isn't completely predictable from the prior state of the world because some of these causes are *agents*, rather than states or events.

By the way, some people think that even in the purely physical world, it is objects rather than events that cause things. Thomas Reid put the point by saying that *events* do not *act*; only *things* can act.

Here's another vague view about how free will works: People have two kinds of characteristics and states: *physical* characteristics and states, and *mental* characteristics and states. Examples of the latter include feeling hungry, being honest, and believing in global warming. These mental characteristics and states are not reducible to the physical. They are generally *caused by* physical states and events, but they are not themselves physical. Also, these mental states and characteristics can cause physical events. They do this by directly causing some electrical activity in your brain, and that electrical activity then causes your body to move in certain ways. In this view, the effects of these mental states are something over and above the effects of the purely physical states that you have, so in principle this would be detectable by very close observation of the brain. Now, if you look at purely physical systems with no mental properties, there is no room for free will – their behavior is either determined or random. But maybe the presence of these non-physical, mental properties enables people to perform actions that wouldn't be predictable based purely on our physical states and the (physical) laws of nature. And maybe that's what free will is about.

Aside

This theory posits something called **downward causation**. Roughly, this is a phenomenon in which properties of a whole system influence the properties or behavior of the system's *parts*. It is controversial whether this is possible; most philosophers believe that causation can only go the other way: The properties and behavior of the parts must determine the properties and behavior of the whole.

I have some sympathy with both of the above views, i.e., the “agent-causation” and “downward causation” views. I suspect both have something

to do with how free will works. But they still leave free will seeming mysterious.

11.5.2. Degrees of Freedom

Most discussion of free will treats freedom as a binary property (a property having only two values): Either you have free will, or you do not have it. Philosophers give qualitative conditions for an action to be free, where satisfying the conditions presumably makes you fully free, and failing to satisfy them makes you not at all free.

However, it seems that one can have varying *degrees* of freedom. Among actions that are to some degree free, some are more free than others. For example, suppose a person's behavior is partly explained by their poor upbringing. Or a psychological disorder. Or the influence of drugs or alcohol. I say "partly explained" here because nobody's behavior is ever *completely* explained by those things. But the influence of extraneous, non-rational factors can be more or less difficult to resist, and a person can have a more or less active role in their own decision-making process. So these outside influences would diminish one's free will.

These different degrees of freedom lead to different degrees of blameworthiness, in the event that one acts badly. This is why, for example, if you kill someone in a fit of rage, you get a less harsh sentence (for second-degree murder) than you do if you plan everything out beforehand (as in first-degree murder). Of course, you also get different degrees of praise in the event that you do something good.

Freedom is a good thing, and it's generally better to have a higher degree of freedom. The way we can increase our degree of freedom is by being more self-aware – if you are aware of the factors influencing your emotions and desires, you are less likely to fall prey to influences that you would not endorse. This is why it is good to reflect periodically on why you make the choices you do.^[72]

12. Personal Identity

12.1. The Teletransporter

Imagine that you have to take a trip from New York to London (say, because you need to receive the prize for the World's Greatest Philosophy Student, which is awarded in London). You're really not looking forward to the boring, 7-hour plane ride in coach, though. It's cramped seating, they don't have any good vegan meals, and the airline only has crappy middle seats available at this time.

Fortunately, you've just learned that scientists have invented a new transportation method known as "the teletransporter". It's like the transporter in *Star Trek*: It scans your body, disassembles it, then sends a beam of energy to some distant location, where somehow your body (or at least, a body indistinguishable from yours) is reassembled. They're touting this as the future of transit. Instead of taking 7 hours, you can make the trip in 0.02 seconds using the teletransporter (which transmits energy at the speed of light).

There's only one downside: Philosophers can't agree on whether the person who steps out on the other side will really *be you* or instead be merely an incredibly accurate *replica* of you. In other words, they don't agree on whether the "teletransporter" is a device for transporting people or a device for destroying people and then making new people who look just like them. Of course, you don't want to use the teletransporter if doing so will end your life.

There's no observation or experiment that we could do to test this. When you observe the people who emerge from the teletransporter, they behave exactly like the original people. They have memories qualitatively just like those of the original people. They claim to be the original people. But none of that shows that they are really the same people who got into the device, rather than distinct people who merely have the same brain configurations. How could we settle that question?

This looks like a job for ... philosophy.

12.2. The Problem of Subject Identity

12.2.1. Basic Question

The philosophical *problem of personal identity* concerns the question: Under what conditions is A *the same person* as B? For instance, if you permanently lose all your memories, will *you* continue to exist, or will there be a different person where you once were? If you die, but then someone produces a clone of you with the exact same brain configuration as your brain at the time of your death, would the clone *be* you, or would it just be another person very similar to you?

12.2.2. *Persons and Subjects*

Now I have to make some notes to clarify the issues. Philosophers use the term “person” in perhaps a nonstandard sense. As we use the term, it means something like a conscious, intelligent being. (There’s a lot more that philosophers say about persons, but we don’t need to worry about that.) Note that *persons* do not have to be *human beings*. Here are some examples that science fiction fans can appreciate: the Klingons, the android Commander Data, Jabba the Hutt, and Rocket the intelligent raccoon. All of those are persons in the philosophers’ sense, though none of them are human. Also, human beings need not be persons. For instance, a brain-dead human would be considered no longer a person in the philosophers’ sense.

I don’t think that matters much for thinking about the problem of personal identity, though. Or rather, I think discussion of “personal identity” should be replaced with discussion of “subject identity”, or “conscious being identity”. (A “subject” is just a thing that can have experiences.) We should ask: Under what conditions is A *the same being* or *the same subject of experience* as B, regardless of whether either A or B is a person? That’s the real question of interest. It’s possible that a being might either become a person or cease to be a person, while continuing to exist as the same subject of experience. E.g., maybe you started out life as a non-person (like when you were a fetus) and grew into a person some time after birth. In that case, what we care about is under what conditions *you* continue to exist – e.g., what made that fetus an earlier stage *of you* – regardless of whether you are *a person* at all the times you exist.

12.2.3. *Two Kinds of Identity*

Now I have to make an important conceptual point to head off a confusion that happens virtually every time you talk about personal identity with a class. There are two senses of “identical” in English: **qualitative identity**, and **numerical identity**. A and B are *qualitatively* identical (or: indistinguishable) when they have the same qualitative characteristics. E.g.,

you could have two electrons in a box with equal mass, equal charge, equal energy, and even the same quantum state (if you know what that means). They're still *two* electrons, though, not one (e.g., the total mass inside the box would be twice the mass of an electron). Then the electrons are qualitatively but not numerically identical.

On the other hand, we say that A and B are *numerically* identical when A just is B, i.e., there is only one object that we're talking about. Example: George Washington is identical with the first President of the United States: Those are just two ways of referring to the same person. Of course, numerical identity *entails* qualitative identity, but not vice versa. George Washington is numerically (and therefore also qualitatively) identical with the first U.S. President.

Crucial point for understanding the problem of personal identity (or subject identity): The issue is about *numerical* identity. We are not asking when A is qualitatively identical to B (which would be trivial). We're asking when A is numerically identical with B, that is, when A and B are one and the same being.

12.2.4. Identity over Time

Most of the philosophical discussion has to do with persistence over time. You start with a person, something weird happens to them, and then we ask whether the original person survived, i.e., is the being who exists at the end the same being who existed at the beginning? Example: You get your memories erased; is the post-erasure person numerically identical with the original person, or are these two people?

You might be tempted to think: "Oh, that's trivial: The thing at the end is never numerically identical with the thing at the beginning. Numerical identity, as you just told me, entails qualitative identity. The being who exists at the end doesn't have exactly the same properties as the one at the beginning. E.g., in your example, the post-erasure person doesn't have the same memories as the pre-erasure person. So they're not qualitatively identical, so they're not numerically identical."

This reasoning of course is not correct; it's not true that persons always cease to exist whenever they undergo any changes (that would be crazy). We can avoid the crazy conclusion by being clearer about the properties that one has. Call the person who exists initially "Early You", and call the person who exists after the memory erasure "Late You". Early You might be numerically identical with Late You (that is, this isn't logically ruled

out), for we could say that Early You has certain memories *at time t1*, but Early You does *not* have those memories *at time t2*. Meanwhile, Late You lacks those memories at time t2, but (perhaps) *had* those memories at time t1. That's all logically consistent. And it's consistent with holding that Early You is qualitatively identical to Late You, since they have the same time-indexed properties: Both have the property of *remembering such-and-such at t1* and also *not remembering such-and-such at t2*. At any rate, that's what we would say if we think that Early You = Late You. And that explains why we can't just trivially conclude that you never have numerical identity over time.

12.3. Theories of Personal Identity

There are several natural views of personal identity that can all be straightforwardly refuted with counter-examples. This is why personal identity poses a philosophical problem. In fact, when you think about it, it almost seems that every possible view is unacceptable.

12.3.1. The Body Theory

Here's a theory: A and B are the same person (or the same being) provided that they have the same body. And here's a counterexample to that:

Brain Transplant: John undergoes a brain transplant operation, in which his brain is removed from his body and put into Jane's body. The operation is completely successful, so that the brain is able to control the new body, receive signals from the body's nerves, and so on, just like normal.

Intuitively, John continues to exist, occupying what used to be Jane's body. We could describe the operation by saying that John has gotten a new body. Yet the Body Theory implies that this is impossible. Instead, the theory implies that Jane got a new brain (i.e., that the person occupying Jane's body at the end is in fact Jane).

12.3.2. The Brain Theory

Because of examples like Brain Transplant, nobody believes the body theory. Here's a more plausible view: A and B are the same person provided that they have the same *brain*. Now here's an apparent counterexample to that:

Mind Transplant: In this case, John and Jane keep the same physical brains. However, a very sophisticated operation is done to *reconfigure* the neurons in the two brains. John's brain gets reconfigured to have exactly

the pattern of neural connections that was originally in *Jane's* brain. At the same time, Jane's brain gets reconfigured to have exactly the pattern of connections that was originally in John's brain. (The two brains luckily happen to have the same number of neurons.) When the two wake up, the person in John's body (the male body) has Jane's memories, Jane's personality traits, thinks their name is "Jane", etc. Similarly, the person who wakes up in Jane's body thinks their name is "John", etc.

In this case, the Brain theory implies that the person in John's body at the end of the story is the same person (John) as the person who was in that body at the beginning of the story, since the same physical brain is present. But this seems false. Intuitively, a person does not survive having their memories, personality, and all other mental qualities erased and replaced with someone else's mental qualities.

Whether the person who wakes up in John's body is really Jane or just a person with an illusory sense of being Jane is another question. We don't have to answer that for present purposes. It's enough to say that that person *isn't John*. Thus, the Brain Theory is false.

Another potential problem: Both the Body Theory and the Brain Theory imply that reincarnation is metaphysically impossible. If you are just your brain (/body), then it would be metaphysically impossible for you to exist with a different brain (/body). But we can clearly imagine being reincarnated, and there just doesn't seem to be anything impossible about it. Note that the argument here isn't "There is reincarnation; therefore, you are not your brain/body".^[23] The argument is "It's *possible* that there should be reincarnation; therefore, you are not your brain/body." That's valid.

12.3.3. The Naïve Memory Theory

Maybe continuity of memory is the key to personal identity: Maybe A counts as a later stage of B as long as A remembers B's experiences. That would imply that in Mind Transplant, John and Jane switch bodies, which perhaps is the intuitive result.

This theory, however, faces a decisive objection: It conflicts with the **Transitivity of Identity**, the principle that if $x=y$, and $y=z$, then $x=z$ (where "=" denotes identity). Consider the case of a person who, at the age of 30, remembers some stuff that happened when he was 10. Also, at the age of 80, he remembers some stuff that happened when he was 30. But at the age

of 80, he no longer remembers any of the stuff that happened when he was 10. All of this could very well happen. The memory theory implies:

The 10-year-old = the 30-year-old.

The 30-year-old = the 80-year-old.

But the 10-year-old $\neq$ the 80-year-old.

(where “=” denotes the “same person as” relation), thus violating transitivity. (This, btw, was how Thomas Reid refuted John Locke’s theory of personal identity.)

12.3.4. *Psychological Continuity*

The naïve memory theory could be modified to be less naïve. We could say: For A to be a later stage of B, it is only required that there be an unbroken chain of person-stages in between A and B, such that at each stage, the person remembers the previous one. Thus, for example, 80-year-old Jabba counts as the same being as 10-year-old Jabba because: 80-year-old Jabba remembers being 79-year-old Jabba, and 79-year-old Jabba remembers being 78-year-old Jabba, and ... and 11-year-old Jabba remembers being 10-year-old Jabba. That’s basically what I mean by continuity of memory. (But note that the real condition requires *continuity*, so it must be that at *every* time, the person remembers his immediately preceding stages.)

There is a broader notion of *psychological continuity* that we might want to invoke. This would involve continuity in psychological traits in general – that is, no large, discontinuous changes in a being’s personality, mental capacities, beliefs, or other traits. (Those are just some possible examples of psychological traits that you might think are important. I lump together all theories that identify some set of psychological traits that must be not-discontinuous, including the memory continuity theory of the preceding paragraph, under the category “psychological continuity theories”.)

Back to memory continuity. The memory continuity theory is refuted by the following example:

Temporary Amnesia: As a result of a traumatic event, Revan suffers retrograde amnesia: He is knocked unconscious and wakes up with no episodic memories of anything that happened before he woke up, and he has no idea who he is. Over time, however, the memories come back.

By the way, this sort of thing actually happens – people sometimes wake up with amnesia after an accident, and the amnesia does in fact

usually clear away and they regain their memories. Now, according to the memory continuity theory, we would have to say that the person who exists after the accident is *not* the same person as the one who existed before the accident. *Even after the memories return*, and the person returns to “his normal life”, the theory implies that it’s still a different person. The amnesia event constitutes a break in continuity, which means the original person is gone. And this just seems false.

You might try modifying the theory: We could say that as long as B has memories of being A, B is the same person as A. So once Revan’s memories return, he becomes the original person again. But notice that this theory violates transitivity again. Let “Revan1” refer to the person who exists before the traumatic event, let “Revan2” denote the person who wakes up with no memories, and let “Revan3” be the person who exists after the memories are recovered. Our latest theory implies that $\text{Revan1} = \text{Revan3}$, and $\text{Revan3} = \text{Revan2}$ (since Revan3 has memories of both Revan1 and Revan2), and yet $\text{Revan2} \neq \text{Revan1}$ (since Revan2 lacks memories of Revan1). Surely this cannot be.

Another problem:

Perfect Clone: Someone makes a perfect clone of you while you’re asleep and configures its brain exactly like yours. You and the clone both wake up at the same time, and you both have exactly the same mental states at that time. You both have the same tastes, interests, memories, etc., and both of you think they are “the real you”.

In this case, you and the clone both satisfy the memory continuity condition. (Aside: You *could* claim that the clone merely has *false memories* since those things didn’t really happen *to him*, whereas you have genuine, true memories since the events happened *to you*. But that begs the question: That reply requires *assuming* that “the copy” is not really you. What we need is an independent criterion that tells us why the copy counts or does not count as the same person as the earlier you.)

So the memory continuity theory implies that there can be two of you existing at the same time, in two different places. Both copies satisfy the criterion for being you, so they are *both* you. But this cannot be: It violates the basic definition of numerical identity (which, again, is what we are interested in). Numerical identity is the relation that a thing bears to itself and to nothing else. If $x=y$, then, by definition, x and y are *one*, not two (and, I assume, one thing cannot be in two places at once). E.g., if you put x

on a scale and you put *y* on that scale at the same time, then there is one thing on that scale, not two.

These problems for the memory continuity theory can be generalized to all psychological continuity theories. You can modify the Temporary Amnesia case to include a temporary change in personality, beliefs, mental capacities, etc. You can also apply the Perfect Clone case straightforwardly to any psychological continuity theory. The amnesia case shows that psychological continuity is not *necessary* for identity across time, and the clone case shows that it isn't *sufficient* either.

12.3.5. Spatiotemporal Continuity

The Perfect Clone case could be avoided by adding a *spatiotemporal* continuity condition for personal identity. This would be the condition that A and B are stages of the same person only if there is a spatiotemporally continuous series of person-stages connecting A to B. That is: There should be a continuous path through spacetime from A to B, or vice versa, where every point on that path is occupied by a person-stage. (A “person-stage” is a temporal stage of a person, also called a “person time-slice”. It's the part of a person's life that exists at a particular point in time.)

In the Perfect Clone case, we could claim that the real you is the person in the body that is spatiotemporally connected to the original. The clone is not really you because it came into existence in a different place, separated from the original body that was you. Hence, the clone is just a person with false memories of living your life. (The memories would be generally correct as to the qualitative events that happened in your life, but the clone would be falsely remembering that those events happened *to it*, when in reality they happened to a different person.)

But we can devise a new case that makes trouble for spatiotemporal continuity too.

Fission: Imagine a species of intelligent beings who reproduce by *dividing*, rather than by getting pregnant or laying eggs. This is how amoebas and other single-celled organisms actually reproduce – the amoeba divides into two smaller amoebas, which subsequently try to grow to normal size. So imagine that there are intelligent amoebas who reproduce this way.

In this case, after the amoeba divides, the two “daughter amoebae” both satisfy the criterion for being identical to the original amoeba – they're both connected to the original amoeba by spatiotemporally continuous

sequences. We could also imagine that there is psychological continuity and continuity of memory. (Of course, after the division, the two will swiftly start to diverge both physically and psychologically.) But the two daughter amoebae cannot both be identical to the original, since, again, numerical identity requires that there be only one object. This shows that spatiotemporal continuity, even when combined with psychological continuity, does not suffice for identity across time.

12.3.6. *The No-Branching Condition*

There's a cheap way of avoiding fission and cloning cases. We could just stipulate that identity across time requires there to *not* be any copying or dividing going on, or anything like that. That is, we could say:

A is the same being as B provided that: There is psychological continuity between A and B *and* there is no other being existing at the same time as A that also has psychological continuity with B, nor is there any other being existing at the same time as B that also has psychological continuity with A.

That's the "no-branching" condition. It basically just rules out counterexamples like Perfect Clone and Fission by fiat.

The problem with this is that it makes personal identity **extrinsic** rather than **intrinsic**. That is, according to this theory, whether a particular future entity is *you* can depend on things going on *elsewhere*, outside of that entity. As I have stated the theory, it implies that there is a peculiar way of murdering a person without having any physical effect on their body: Just produce a perfect copy of them. Once you do that, the no-branching condition for personal identity is violated, which means that *neither* of the two people who exist at the later time count as the same person who existed originally. So the original person has been annihilated. The perfect crime! If you can keep the clone secret, no one will even know that your target was murdered! This just seems plainly wrong.

12.3.7. *The Closest-Continuer Theory*

The same problem afflicts the Closest-Continuer Theory. This is a theory that says that in order for person A to be a later stage of person B, A has to (i) pass a certain threshold degree of similarity to B (e.g., A might have to have sufficiently similar memories, character traits, or brain configuration to B), *and* (ii) also be *more* similar to B than anyone else who exists at the time A exists. (Note: You can add other conditions to the "similarity", e.g., spatiotemporal or psychological continuity conditions.

The important point is that identity with the original person depends upon being the *best candidate*, or the person with the best claim to being the original.)

This theory, again, makes identity extrinsic. In other words, whether A is the same person as B depends not just on the facts about A and B, but also on whether there are any *other* beings existing at the same time as A that have certain characteristics.

Take the Fission case. Let's name the original intelligent amoeba "Meba1". Call the two amoebas that exist after division "Meba2" and "Meba3". Since Meba2 and Meba3 are *equally* good candidates for being the continuer of Meba1, the Closest Continuer theory implies that neither of them is the same being as Meba1. So Meba1 is gone. If, however, there is a slight error during the division process, so that, say, Meba3 has one tiny difference from Meba1, then the theory tells us that *Meba2* is the same being as Meba1. (Knowing this, Meba1 would have a self-interested reason to try to ensure that one of its offspring is defective!) This seems wrong: Whether Meba2 is the same conscious being as Meba1 is, so to speak, entirely between Meba1 and Meba2. It shouldn't be changed by changing someone *else* without doing anything to either Meba1 or Meba2.

12.3.8. *The Soul Theory*

Some believe that persons have a special, immaterial component called "the **soul**", which determines one's identity. (This view is sometimes called "**mind-body dualism**".) The soul is also said to be the subject of mental states (it is your soul, rather than your body, that experiences thoughts, feelings, and so on). In all the thought experiments about personal identity, you go wherever your soul goes. The unobservability of souls makes it possible to account for any intuitions about personal identity you want; we can just suppose that a person's soul goes wherever we intuitively think the person is.

For instance, in the Temporary Amnesia case, you continue to exist as yourself after the amnesia, provided that *your soul* is still there, in the same body.

In the Mind Transplant case, John goes wherever his soul goes. Perhaps reconfiguring his brain causes his soul to be expelled from that body, in which case the person who wakes up in John's body is not John.

In the Perfect Clone case, the real you is whichever person has the original soul. E.g., maybe the original body continues to be animated by the

same soul (as seems plausible), and the clone has its own, new soul.

In the Fission case, the original soul can go with either of the descendent amoebas, or it could be destroyed (or just cease to be in a body); there's no way of knowing since we can't see other people's souls. So Meba1 *could* be the same being as either Meba2 or Meba3, or neither.

The problem with the soul theory: It's controversial whether souls exist. We can't observe them by the five senses, we can't detect them with scientific instruments. We have no good theory of where they come from or how (if at all) they can affect the physical world. They seem to be *ad hoc* posits designed to let us maintain whatever views we want about identity of persons, while making no definite predictions about persons or anything else. A large majority of contemporary philosophers are *physicalists*, that is, they think everything in the world is physical, so they can't accept souls. For these reasons, the soul theory is highly unpopular among contemporary philosophers.

12.4. In Defense of the Soul

Now, despite what I just said, I think the Soul Theory is the only plausible theory of personal identity that anyone has devised. Every other theory has obviously false implications.

12.4.1. Objections to the Soul

Before discussing why the soul theory is good, let's just briefly talk about some objections people raise to the notion of a soul, and why I'm not persuaded by them.

Objection/question 1: Since we can't observe souls with the senses, how could we know about them, even if they existed?

Reply: We are directly aware of our souls by introspection. More precisely, you are introspectively aware of yourself as a subject of experience. The best account of what a subject of experience is is that it is a soul (see below for discussion). Hence, there are souls.

Objection 2: We have no good account of where souls come from or how they can affect the physical world.

Reply: One can know that something exists without knowing where, if anywhere, it came from. If one is seemingly directly aware of something, it would be unreasonable to deny that it exists merely because one does not know how it came to be. Similarly for not knowing how the thing affects other things.

Objection 3: It's not just that we *don't know* how souls affect the physical world; it's that it *seems impossible* that they should do so, given that they are themselves entirely non-physical.

Reply: This may be the most popular argument against the soul in contemporary philosophy. It's common to appeal to "the causal closure of physics", which is the principle that non-physical things can't affect physical things. But I have no idea why that's supposed to be true. (Is it an a priori axiom? Did we test a lot of non-physical things to see whether they affected physical things?)

How do we know what causes what? We don't generally know it a priori; we know it by experience: If A's are generally followed by B's, we hypothesize that A's cause B's. (One can also do tests to rule out confounding variables, but let's leave aside such complications, which are irrelevant here.) Certain mental states are followed by certain physical states – e.g., an intention to move your hand is usually followed by your hand moving – so it's reasonable to conclude that mental states affect physical events. That's true even if mental states aren't physical. I don't know why we're supposed to insist that this couldn't be the case.

Likewise, I don't know why there couldn't be laws of nature that say that when certain mental events happens, then certain physical events follow. We don't know why such laws should exist, but we also don't know why *any* of the fundamental laws of nature exist.

Objection 4: The soul theory is ad hoc.

Reply: The soul theory is the most natural interpretation of the phenomenon of consciousness. This is shown by the fact that it is the view taken by nearly all humans throughout history, across cultures. So I don't think it's ad hoc.

Objection 5: Souls are weird, mysterious, spooky, etc.

Reply: "Spooky" is a term of abuse, not an objection.

Perhaps the soul is weird in the sense of being very different from other things. But the same is true of every other basic category of thing that exists (time is weird, space is weird, fields are weird, numbers are weird, etc.). There is no reason to assume that weird things don't exist.

Perhaps the soul is also mysterious in the sense of being poorly understood. But there is no reason to assume that poorly

understood things don't exist.

By the way, this "objection" is extremely widely felt among contemporary philosophers, though they rarely put it in print. I, however, find it empty. I think it's basically just a negative emotional reaction masquerading as an argument.

12.4.2. Principles of Identity

Here are four interesting principles about identity of persons (and other subjects of experience):

- i. Identity is a **one-to-one relation**:* Every being is identical with exactly one being; no one is ever identical with *two* beings.
- ii. Identity is **transitive**:* If *x* is identical with *y* and *y* is identical with *z*, then *x* is identical with *z*.
- iii. Identity is **intrinsic, not extrinsic**:* Who a given being is depends solely on facts about that being; it does not directly depend on facts about *other* beings. You cannot, for example, end a person's existence solely by creating another person with certain characteristics who never interacts with the original person.
- iv. Identity is **objective**:* If *A* is a person and *B* is a person, there is an objective fact as to whether $A=B$. It is not subjective, indeterminate, or a matter of convention whether, for example, *I* exist in any given possible scenario.

I think all of those are very obvious truths. Notice that none of those turn on dubious intuitions about controversial cases, e.g., about whether a person can survive erasure of their memories. We discussed the first three principles in section 12.3 above, where we saw how they rule out several alternative theories of personal identity.

Let me say more about (iv), the objectivity condition. If you look at people purely from the outside, it might seem as though it is sometimes indeterminate whether *A* is the same person as *B* (say, because the person has changed a lot). But if you think about it from the first person perspective – from the inside, so to speak – it is hard to understand what it would mean for identity to be indeterminate. If some experience is going to occur in the future, I will either feel that experience or not. If it can be indeterminate whether a future being is me, then it would be indeterminate whether I will have the experiences that being has – i.e., I will neither have nor fail to have those experiences. What would it be like for me to neither have nor fail to have an experience? I just can't make sense of this.

Now consider the view that facts about identity are a matter of convention. This is plausible when you're talking about identity of inanimate objects. E.g., as Theseus progressively replaces planks from his ship, at what point does it cease to be the same ship? (See §1.1.) It's plausible to answer: "That is a semantic question. It just depends on what we decide to *call* 'the same ship'". So suppose someone holds that sort of view about persons as well. Suppose someone says that in general, it's purely a matter of convention whether A is "the same person" as B.

Think about the implications. Imagine that you've just been diagnosed with a deadly disease that will kill your body within a year. If the above theory of identity is correct, you have a possible plan for surviving: Just change the conventions for talking about you. Find some baby and try to convince everyone to start *calling* that baby you. Obviously, it would be hard to get everyone to do this, but let's suppose you succeed. Then, since facts about identity are supposedly purely conventional, that baby will *become* you. So you won't have to die. Woohoo!

That, I take it, is ridiculous. Intuitively, you have no reason at all to try to get people to start talking like that. (At least, no reason from the standpoint of your desire for longevity.) There's no reason to do that because it won't change *the facts* about your own mortality. People can treat the baby as you – e.g., start calling it by your name, treat it as the owner of your property, hold it to your promises, etc. – but it just won't in fact be you. What this example shows is that it's not a matter of convention what is or is not *you*. It's a matter of objective fact.

Now, you might say the theory of identity that we just refuted is a pretty silly one. Of course no one would suggest that facts of personal identity are *completely* conventional in all cases. At most, some might say that facts about personal identity are convention-dependent in certain borderline cases. A randomly chosen baby is, objectively, not you. But if we find a person more similar to you, with a tighter causal connection to you, *then* perhaps it could be a matter of convention whether that person is you. E.g., if we created a clone with your memories, then maybe it's a semantic question whether the clone is you.

But the same sort of reasoning shows the implausibility of this. Suppose again that you have a terminal illness. You also know that a clone of you is going to be made and implanted with your memories. Now, you may well have reason to want that clone to be made – that's plausible, since the clone

might really be you (we don't know for sure). But assume you know the clone will be made no matter what you do, so that's not in question. The question is: Do you have a really strong reason to try to make other people *call* that clone you? That is, if you could alter the *linguistic conventions* of your speech community, should you try to alter them so as to ensure that the clone is you, and thus to avert your own demise?

I take it, again, that the idea is obviously confused. No, you don't have any longevity-related reason to do that. You're obviously not made any better off by a mere change in linguistic conventions. This example shows that, even in cases where the facts of personal identity are *unclear* (we have no firm intuitions), they still are not *conventional*.

12.4.3. *Only the Soul Theory Satisfies the Principles of Subject Identity*

Now I'm finally going to tell you why I claimed that the soul theory is the only plausible view. Note that it satisfies principles (i)-(iv): If personal identity is determined by the soul, then personal identity is a one-to-one, transitive relation, and facts about identity are intrinsic and objective.

No other reasonable theory satisfies these principles. The body theory and the brain theory (§§12.3.1–12.3.2) violate condition (iv), the objectivity of identity. This is because the identity of material objects across time is often conventional. This is shown by the Ship of Theseus case: There is no objective fact about whether it is “really” the same ship at the end as at the beginning of the story; it's a semantic question. And note that ordinary material objects undergo similar changes all the time: Individual atoms in your body (including your brain) are periodically replaced. The configuration of your brain and body can also change. Exactly how much change is compatible with still having “the same” material object is a semantic question, to be settled (if at all) by linguistic conventions. But since personal identity is objective and non-conventional, personal identity cannot turn on the identity of the brain or the body. The argument generalizes to any material object made up of many parts.

This argument could be avoided by identifying a person with an elementary particle, since (I assume) the identity of elementary particles is objective and non-conventional. However, this is a super-implausible view that no one holds. There is no particular elementary particle in the human body that is so special as to plausibly determine the identity of the person.

The memory, psychological continuity, and spatiotemporal continuity theories all violate condition (i), that identity is one-to-one. This is because it is metaphysically possible for more than one entity to have continuity-of-memory, psychological continuity, and/or spatiotemporal continuity relations to some initially existing person. Note that the argument generalizes to *any* theory that uses any repeatable properties or relations as criteria of identity. E.g., a theory could say that B is a later stage of A provided B has the same personality traits as A, or the same feelings and desires, or the same DNA sequence, or the same arrangement of atoms in the brain, etc. All of those are *repeatable* – it’s possible for multiple entities to have those personality traits, those feelings and desires, that DNA sequence, that arrangement of atoms in the brain. Therefore, *all* of those theories violate condition (i): They imply that multiple distinct things could all be identical with the same person. (They also violate transitivity; any violation of (i) is also a violation of transitivity.)

One can avoid the above problem by stipulatively constructing an unrepeatable property. Take some condition that you thought was a plausible condition for identity before hearing the objection in the preceding paragraph. Call the condition “C”. (E.g., C could be “the condition of having the same configuration of neurons in the brain as A”.) Now, we can construct a new theory of personal identity that says B is the same person as A whenever B is *the only thing* that satisfies condition C. That’s a one-to-one relation: By definition, only one thing can ever be “the only thing” that satisfies condition C! So that avoids the objection of the preceding paragraph. However, any theory of this kind violates condition (iii), that identity is intrinsic, not extrinsic. On such a theory, you could kill a person without touching them, simply by creating another entity that satisfies condition C.

That rules out almost everything. Say we’re trying to articulate the conditions for some future person to count as *you* (that is, a later stage of you, or “the continuer” of you). We can’t rely on any purely qualitative conditions, because that violates condition (i), that identity is a one-to-one relation. We can’t rely on a qualitative condition together with a no-branching clause to exclude duplicate “you”s, because that violates condition (iii), that identity is intrinsic. We can’t rely on the identity of any composite physical object, because that violates condition (iv), that identity is objective. We can’t rely on the identity of a simple (having no parts)

physical object, because, well, that's just super-implausible. What's left? It seems that the only remaining possibility is to appeal to the identity of some *non-physical*, simple entity. If there is a simple, non-physical entity that determines the identity of persons – well, that's pretty much the definition of a “soul”.

12.4.4. *Unanswered Questions*

One thing that people don't like about the soul theory: It leaves a lot of personal identity questions unanswered. These questions *have* objective answers, according to the theory; it's just that we have no apparent way of *knowing* the answers, because we can't directly observe anyone else's souls.

In the teletransporter case (§12.1), we asked: Is the person who emerges in the distant location the same person who got into the teletransporter? The soul theory answers: It is the same person if and only if it has the same soul. There's no way of verifying whether the person on the other end has the same soul, though, because you can't see their soul. In fact, even if you yourself get into the teletransporter, you still won't be able to figure out whether it is a transportation device or an annihilation-followed-by-copying device. Immediately before you step in, you won't know if you're about to simply be destroyed. Immediately after stepping out (if indeed you do), you still won't be able to know whether you actually have the same soul as the person who got into the device, or whether instead you merely have false memories of being that person. You can introspectively observe your own soul at the moment, but you can't directly observe whether it's the same as the soul that existed at a given earlier time.

There's a variant of the teletransporter case in which the transporter malfunctions and produces two copies of “you” (that is, two persons with bodies indistinguishable from the body that went into the device). In that case, we can deduce that not both of those persons is you. But the soul theory doesn't tell us *which*, if either, is the real you – because, again, you can't inspect souls.

Similarly, in the Perfect Clone and Fission cases, the soul theory doesn't enable us to know which person, if any, is identical to the original. Come to think of it, even in ordinary life, with no weird stuff going on, you can't be sure that you're the same person who occupied your body five years ago. Maybe there was a different soul in your body then and it somehow got replaced, but the brain still retains the same memories, so that you have the

memories of what that other soul experienced. There's no reason to think that's the case, but there's no way to refute it either.

Some people think this is a bad feature of the soul theory. And it is indeed annoying. But I don't think this is a flaw in the theory; I think it's just an annoying feature of reality. There's no reason to think that there shouldn't be such unknowable facts. I would note that you can raise parallel worries about material objects too. For instance, the table in front of me *might* have just disappeared and been immediately replaced by another table that looks exactly like the original. There's no reason to think that happened, but there's also no way of proving that it didn't. (And note that this isn't just a semantic question – I'm imagining completely different matter, distinct individual particles, coincidentally appearing at the same location just as the original particles in the original table disappear. I take it that, if that happened, it would in fact not be the same table.) Since it's perfectly possible to have unknowable facts about identity of physical objects, we shouldn't in principle object to precisely the same sort of unknowable facts about identity of persons.

Part IV: Ethics

13. Metaethics

13.1. About Ethics and Metaethics

13.1.1. Ethics

We're now going to talk about the nature of ethics. Ethics studies what is good, bad, right, and wrong. Statements about good, bad, right, and wrong are known as “**evaluative statements**”.^[74] Statements that are *not* evaluative are called “**descriptive**”.

Here's an example of a famous ethical problem, known as “the **Trolley Problem**”:^[75]

Trolley: A runaway trolley is headed for a fork in the track. On the left track, there are five people in the way who will be killed if the trolley goes that way. On the right track, there is one person who will be killed if the trolley goes that way. There's no way to move any of the people out of the way in time. The trolley is presently set to go down the left track, toward the five. You can flip the switch to send the trolley down the *right* track instead, toward the one person. Or you can do nothing and let the trolley continue toward the five. What should you do?

Now, before you propose one of the annoying distractions that students like to use to avoid the main issue, let me clarify that all six of the people on the tracks are normal, average people. None of them is Adolf Hitler, nor is any of them a scientist about to discover the cure for cancer. None of them is responsible for the trolley being out of control, nor for their being trapped on the tracks. Nor are there any other weird circumstances that could be used to change the subject and that I would obviously have mentioned if they were intended to be part of the story. The dilemma is just the simple, obvious one suggested by the above description.

With that understood, most people think that it is *morally permissible* to switch the trolley, possibly even obligatory. A minority, however, think that it's *morally wrong* because it amounts to murder, and you can't murder an innocent person even to save five others.

13.1.2. Metaethics

We're not going to try to answer here what is the right thing to do in the Trolley case. I'm just using that as an example of an ethical question. I want to raise questions like this: Is there a *fact* about what is the right thing to do

in that case? If there is, how could we know it? What does it mean that turning the trolley is “morally right”? These are known as **meta-ethical** questions.

Basically, meta-ethics studies the nature of ethics and ethical questions. (Note: I use “moral” and “ethical” interchangeably, as is common in ethics.) For instance:

- a) Do ethical questions have objectively correct answers?
- b) How, if at all, do we know what is right or wrong?
- c) What does “good” mean?
- d) What reason do we have for being moral?

Notice that those aren’t themselves ethical questions, since they don’t call for evaluative statements as answers; they’re just questions *about* ethics, about what is going on when we make moral judgements and such. Question (c) might have some direct implications for what is good, but in itself, it’s a descriptive question. It could be resolved by looking at how people in fact use the word “good” (regardless of whether anything is in fact good).

13.1.3. Objectivity

The most discussed question of metaethics is the first one I listed above: Are there objective values, or objective ethical truths?

But wait, what does that question mean? Basically, it means: Are there ethical statements that are true independent of the attitudes of observers toward the things that are being evaluated? For instance, you might think “murder is wrong” is an objective ethical truth: This would be to say that murder is wrong regardless of anyone’s attitudes toward it. It’s wrong independent of whether we approve or disapprove of it, like or dislike it, etc. If our society has a sudden change of conventions and people start approving of murder, murder won’t become morally okay; rather, our society will just be wrong. That’s what it means to say that murder is “objectively” wrong.

Most people will basically accept that definition of “objectivity”, but then they will proceed to use the word in ways that don’t cohere with the definition we just gave. Here are four questions that people often confuse with each other:

- a. Is the truth of an evaluative statement independent of observers’ attitudes?

b. Can evaluative statements be conclusively proved, or known with certainty?

c. Are there some categories of action (say, “murder”) that are always wrong, regardless of the circumstances?

d. Are there evaluative statements that everyone would agree with?

Notice that those are all *completely different* questions. Only (a) is the question of objectivity. I’m telling you this because almost everyone, on first being introduced to metaethics, immediately confuses all those questions with each other, and then starts using “objectivity” to express some vague jumble of answers to all four questions. (As I noted in the preface, almost everyone is insanely confused whenever they talk about philosophy. If you’re a beginning student, you’re undoubtedly insanely confused right now. Don’t worry, though; it gets better if you keep studying philosophy.)

Anyway, we’re going to spend most of the chapter on the question of objectivity.

13.1.4. Five Metaethical Theories

Now I’m going to tell you the five main views in meta-ethics, and why there are exactly five of them.

Suppose you think that there are *no* objective ethical truths. This view is sometimes called ethical (or moral) **anti-realism**. As an anti-realist, there are the following things you could believe:

i. Ethical statements are neither true nor false.

ii. Ethical statements are always false.

iii. Ethical statements are sometimes true, but their truth depends on the attitudes of observers.

Notice that those are *all the logical possibilities*. If none of (i), (ii), and (iii) is the case, then, logically, there must be objective values. (If ethical statements are true or false, and they’re not always false, then they’re sometimes true. If they’re sometimes true, and their truth *doesn’t* depend on the attitudes of observers, then they’re sometimes “objectively true”, on the definition of objectivity given above.)

The alternative to anti-realism is **moral realism**, which holds that there are objective ethical truths. These ethical truths are either **reducible** to descriptive truths, or they are irreducible. What does it mean to be “reducible”? Basically, to say that goodness is reducible is to say that you could explain *what goodness is* using only descriptive (non-evaluative)

terms. E.g., suppose you think that “*x* is good” just means “*x* promotes pleasure”. That some action promotes pleasure is a matter of descriptive fact (it’s also objective). So that would be a reductionist view – truths about what is good would be reducible, as we say, to truths about pleasure. By contrast, the non-reductionists say that there is no way of explaining what goodness is except using other evaluative expressions (e.g., “desirable”, “positive”, “worthy of pursuit”).

So there are five views relating to the objectivity of morality – three forms of anti-realism and two forms of realism. Each of these logical possibilities has been defended by some philosophers. Here are the five possibilities again, with the names of the theories that embrace each possibility:

(Anti-realism)

- i. Ethical statements are neither true nor false: **Non-cognitivism** / **Expressivism**.
- ii. Ethical statements are always false: **Nihilism** / moral **Error Theory**.
- iii. Ethical statements depend for their truth on the attitudes of observers: **Subjectivism** / **Relativism**.

(Realism)

- iv. There are objective ethical truths, which are reducible to descriptive facts: **Ethical Naturalism**.
- v. There are objective ethical truths, which are irreducible: **Ethical Intuitionism**.

Below, we’ll add some complications to those views.^[76] We’ll also see the problems with each of them.

13.2. What’s Wrong with Non-Cognitivism

13.2.1. The Non-Cognitivist View

How could ethical statements be neither true nor false? You may recall the Law of Excluded Middle, ($A \vee \sim A$) (see §5.2.2). It seems like the following would be a valid instance of the law: Either it is right to divert the trolley in the Trolley case, or it is not right to divert the trolley in the Trolley case. If it is right, then the sentence “It is right to divert the trolley in the Trolley case” is true. If it isn’t right, then “It is right to divert the trolley in the Trolley case” is false. So, the sentence is either true or false.

How could one avoid this conclusion? Well, the Law of Excluded Middle only applies if you have a sentence that is used to *assert a proposition*. Not every sentence does that. For instance, “Hurray for kale!”,

“What year is this?”, and “Please pass the tequila” are neither true nor false, since they don’t assert propositions. Instead, they: express an emotional attitude, ask for information, and direct someone’s behavior, respectively.

Maybe moral statements are like that: They don’t really assert propositions, i.e., they do not purport to describe the way the world is. Maybe, instead, they are used to express certain emotions, so, e.g., “Abortion is wrong” is sort of like “Boo on abortion!” Or maybe they are just sophisticated ways of telling other people what to do, so that “Abortion is wrong” is like “Don’t have an abortion!” The first of those views is called “emotivism”; the latter is called “prescriptivism”. Emotivism and prescriptivism are two (once-popular) forms of non-cognitivism.^[ZZ] Notice how they both explain why moral statements would fail to be either true or false. You could also have a hybrid view (maybe “Abortion is wrong” expresses a feeling of disapproval toward abortion *and* directs other people not to have abortions). You could also introduce other non-cognitive attitudes that moral statements express. The key point is that, according to non-cognitivists, moral terms don’t express genuine beliefs, in the sense that there is no particular way that you’re thinking the world (or anything else) is when you make a moral claim. By the way, this is probably the most popular view in the meta-ethics literature.

13.2.2. The Linguistic Evidence

The most obvious problem with non-cognitivism is that moral statements act exactly like proposition-asserting statements in all known respects. They do not act like interjections (like “Ouch!”), commands (like “Pass the tequila”), or any other non-assertive sentences. Here are some examples of what I mean by acting like a proposition-asserting statement:

- i. Proposition-expressing sentences can appear in the antecedents of conditionals, i.e., you can say “If P, then ...”, where “P” expresses a proposition. You cannot insert a non-propositional expression for P. E.g., “If it is raining, then we’re going to get wet” is fine, but “If hurray for kale, then ...” and “If please pass the Cholula, then ...” make no sense, regardless of how you fill in the ellipses.
- ii. You can call propositions “true”, “false”, “possible”, “probable”, etc. You can’t do that with non-propositional phrases. E.g., “Probably it’s going to rain” makes sense, but “Probably please pass the Cholula” is nonsense.

- iii. You can add “It’s not the case that” in front of a proposition-expressing sentence. You can’t do that with non-propositional phrases. E.g., “It’s not the case that Obama is a Muslim” is okay, but “It’s not the case that hurray for democracy” is nonsense.
- iv. You can convert propositional sentences into yes/no questions – e.g., from the statement, “It is raining”, you can form the question, “Is it raining?” But from the interjection, “Boo on Marx”, you can’t form a sensible question, “Boo on Marx?”
- v. Words that are used to ascribe properties can be converted into abstract nouns, such as “redness” and “squareness”. There are no similar abstract nouns for emotive terms (e.g., “booness” or “hurayness”).
- vi. Certain attitude verbs take proposition-expressing complements. E.g., “I believe that ...”, “I hope that ...”, and “I wonder whether ...” all have to be followed by proposition-expressing sentences. E.g., you can say, “I wonder whether it’s going to rain”, but not “I wonder whether boo for the broncos”.

Items (i)-(vi) above identify standard marks for whether a sentence asserts a proposition. Leaving aside moral sentences, I know of no exception to any of them. You can see the obvious reasons why all of the above observations are true. E.g., “I believe that ...” has to be followed by a proposition-expressing sentence because you can’t believe something that’s not a proposition; that just makes no sense.

Now, here’s the point: Moral sentences clearly pass every test for being proposition-expressing sentences. To illustrate:

- i. “If abortion is wrong, then God is going to be upset with us.”
- ii. “It’s true that murder is wrong.” “Watching pornography is probably okay.” “It’s possible that abortion is wrong.”
- iii. “It’s not the case that contraception is murder.”
- iv. “Is abortion wrong?” “Is pleasure the sole good?”
- v. “I am questioning the act’s rightness, not its goodness.”
- vi. “I believe that stealing is wrong.” “I wonder whether abortion is wrong.” “I hope that pornography is okay.”

All of those make sense. They wouldn’t make sense if moral statements were emotive expressions, or imperatives, or any other non-proposition-asserting thing. So moral statements express propositions.

Interlude: The Frege-Geach Problem

The mid-twentieth-century philosopher Peter Geach raised an objection like the one above. He said he got the idea for it from the late 19th-/early 20th-century philosopher Gottlob Frege. So the problem for non-cognitivism has come to be called the **Frege-Geach Problem**. By the way, Frege's name is pronounced "fray-guh" not "freej" – take note of that so that philosophers don't laugh at you.

Most of the literature treats the problem as being: How can a non-cognitivist explain what sentences like "If abortion is wrong, then God is going to be upset" *mean*? (Similarly for all the others, like "It's true that abortion is wrong", "I believe abortion is wrong", etc.) Suppose you say that "Abortion is wrong" means something like "Boo on abortion!" (expressing a non-cognitive, emotional attitude about abortion). Then what could be meant by "I wonder whether abortion is wrong"? Does it mean "I wonder whether boo on abortion"? But that makes no sense.

Non-cognitivists have tried to answer questions like that. They give increasingly complicated and confusing theories that we don't have time to talk about in detail. But basically, the non-cognitivist has to say that the expressions used in tests (i)-(vi) above – "if...then", "true", "probable", "possible", "not the case", "believe", "hope", "wonder" – each have two meanings, i.e., that they have different meanings when applied to ordinary, descriptive statements from what they mean when applied to evaluative statements. They then need to come up with a new interpretation, ad hoc, for each of these expressions (there's no systematic way of explaining how all of them change when you apply them to an evaluative sentence). This has to be done for *every* expression in the language that gets applied to a proposition-expressing clause. All to explain why every test that can normally be used to distinguish proposition-expressing sentences from non-propositional sentences gives the wrong answer when applied to ethical sentences.

13.2.3. *The Introspective Evidence*

I'm only going to touch on this briefly. Introspectively, moral judgments seem exactly like beliefs, and not like emotions or desires. That's why we call them "moral judgments", not "moral feelings". We can hold them with more or less confidence, we often consider arguments for and against them, we find ourselves wondering what the right conclusion is, and so on.

We often have strong feelings about moral issues, which might tempt you to think that moral "judgments" are really just feelings. But note that

people often have strong feelings about descriptive questions too. E.g., many people have strong feelings about whether they themselves are better looking than their peers, or whether white people are smarter than black people, or whether President Trump colluded with the Russians in the 2016 election – even though those are all perfectly descriptive, non-ethical questions. So it's not particularly weird that people have strong feelings about morality too, and it hardly shows that moral statements don't express propositions.

Note also that our emotions and desires don't always track our moral judgments. You can be upset about some action that you know wasn't wrong, and you can be unperturbed by an action that you know *was* wrong. When Emperor Nero killed Agrippina, that was *very* wrong, and I *know* that, yet I have no particular desires or feelings about it. (Nero was an ancient Roman emperor widely known for being a total psycho. Agrippina was his mother.) I just have a cool, cognitive attitude about it: I intellectually recognize the wrongness of Nero's action.

As far as I can tell, then, all the evidence is against non-cognitivism.

13.3. What's Wrong with Subjectivism

13.3.1. The Subjectivist View

According to **ethical subjectivists**, ethical truths depend on observers' attitudes. Whether a thing is good or bad, right or wrong, depends on how people feel about that thing, or what people think of it, or in some other way how people react to it.^[78] One prominent form of subjectivism is **cultural relativism**, the view that right and wrong depend on what practices are accepted in one's culture.

Why would ethical truths depend on someone's attitudes? Basically, subjectivists think that "x is right" just means that the speaker (whoever is saying that x is right) has a sentiment of approval toward x, or that society approves of x, or something like that. By the way, notice that if you give this kind of theory, you cannot then analyze "approving of x" in terms of believing that x is right (that leads to circularity). Rather, you would claim that "approval" is just a particular positive emotional attitude that people can have. Notice that on a subjectivist view, what is right and wrong will vary from one society to another and possibly even from one individual to another.

Notice also that subjectivism *is not non-cognitivism*. On the non-cognitivist view, as discussed above, moral statements *do not assert*

propositions and therefore are *neither true nor false*. On the subjectivist view, moral statements *do* assert propositions and *are* either true or false. For instance, “murder is wrong” is *true* (according to most subjectivists) since society, as a matter of fact, disapproves of murder, and so do I, and so does God (according to people who believe in God).

13.3.2. *Motives for Subjectivism, 1: Tolerance*

Subjectivism (including cultural relativism) is very unpopular among philosophers, but fairly popular among non-philosophers who have philosophical views. (But note that non-philosophers usually hold some incoherent jumble of views in which they confuse subjectivism, non-cognitivism, and nihilism with each other.) There are two prominent motivations for subjectivism, both of which are pretty awful, intellectually speaking.

Here’s the first one: Tolerance is a virtue. Being a subjectivist makes you tolerant, because you won’t want to impose your views on other people if you don’t think your views are objectively right. Therefore, we should be subjectivists.

What’s awful about that? First, it provides no evidence for subjectivism. It just tries to argue that being a subjectivist might make you nicer, but that doesn’t mean that subjectivism is actually true. (Compare this argument: “Mormons are nice people. Therefore, you should be a Mormon.” That doesn’t provide any evidence that the tenets of the Mormon faith are actually true.)

Second, subjectivism is not a reliable way of promoting tolerance. Suppose that I personally approve of imposing my views on others, and suppose my society also approves of imposing its culture on other societies. Those attitudes, by the way, have historically been very prominent, which is what gave rise to this whole concern about tolerance in the first place. Now, in this situation, subjectivism tells me that I *should* impose my views on others, or that my society *should* impose its culture on other societies. That’s morally right, because that is what I (/my society) approve(s) of. So in this case, subjectivism leads to the opposite of its intended result. It only works out if you happen to be in a culture that already supports tolerance (in which case, you probably don’t need the subjectivist doctrine).

Third, there are obviously better and more honest ways of promoting tolerance. Like, you could argue that tolerance is objectively good. That view would support tolerance even if you are in an intolerant society.

13.3.3. *Motives for Subjectivism, 2: Cultural Variation*

The other main motivation for subjectivism (especially cultural relativism) is the wide variation in moral codes among human beings, especially between different cultures. E.g., polygamy is considered bad in our society, but has been considered perfectly fine in most primitive societies. Therefore, it is said, right and wrong vary from one society to another.

What's awful about this argument? Well, it just seems to confuse beliefs with truths. The inference seems to be something like this:

P1.Moral beliefs vary from one society to another.

C.Therefore, moral truth varies from one society to another.

That's invalid on its face. To make it valid, you'd have to add another premise, something like this:

P2.All beliefs are true.

P1 and P2 together get you C. But P2 is ridiculous.

To be more charitable, perhaps the intended argument is an inference to the best explanation: The *best explanation* for why there is so much variation in moral beliefs is that there aren't any facts there for us to discover. The way that people manage to converge in their beliefs – for example, in science – is that reality constrains us. Everyone tries to figure out the facts, and to some degree we succeed in doing so, and so we wind up thinking roughly the same things. But in morality, perhaps, there aren't any facts independent of our attitudes for us to figure out.

There are two problems with this argument (or maybe a single, two-part problem): First, there has actually been very wide disagreement about many non-moral matters of fact. Different cultures, in addition to having different practices and norms, also have drastically different views about things like medicine, the origin of the Earth, the structure of the cosmos, how many gods there are (if any), what the gods want, and so on. (Examples: They might think that diseases are caused by evil spirits, rather than by germs; that Earth was created by some gods, rather than by gravitational accretion; that the Sun orbits the Earth rather than the other way around.) No one concludes that therefore all those things are entirely dependent on our attitudes and that there are no objective facts about them.

Granted, there has been more convergence on scientific beliefs in modern times – that is, societies that are exposed to modern science and under its influence tend to agree on the origin of the Earth, the structure of

the cosmos, and so on. But this brings us to the second problem for the relativist's argument: The same is true of morality. The same societies that have converged in their scientific beliefs (technologically and economically advanced societies) are *also* converging in their *values*: They are moving toward liberal, democratic values. E.g., they have been moving and continue to move more toward belief in equality, respect for the dignity of the individual, opposition to needless violence, and so on. (Most primitive societies are extremely illiberal.)^[79]

The societies that continue to have very different values from ours tend to be primitive societies – those are the ones that anthropologists are always raising to show widely different cultures – and they have very different *descriptive* beliefs from us as well. So if you don't think that descriptive facts are subjective or relative, you shouldn't conclude that moral facts are subjective or relative either.

The above doesn't prove that subjectivism is false. But it shows that the leading *arguments* for it *don't work*; they don't give us any good reason to believe it.

13.3.4. The Nazi Objection

We turn now to the case against subjectivism. Imagine that you live in Nazi Germany. Your society approves of rounding up Jews and sending them to concentration camps. According to cultural relativism, what is morally right is what society approves of. Therefore, according to cultural relativism, it is morally right for you to help round up Jews to send them to concentration camps. Meanwhile, people like Oskar Schindler, who tried to save Jews from the concentration camps (see §11.4.3), would have to be judged as villains. This, to put it mildly, does not seem correct. It is in fact hard to think of how a theory about morality could go more wrong than that.

A similar point applies to a more individualistic subjectivism which holds that what is right for an individual is whatever that individual approves of. Just add to the above example the stipulation that you yourself happen to be a Nazi at the time. Then the subjectivist view implies that it is morally right for you to round up Jews, and it would be wrong for you to instead try to help them.

The point can be generalized to any subjectivist view. Suppose the subjectivist view says:

x is right if and only if: G takes A toward x.

where G is some person or group and A is some attitude. Then imagine a case where G takes A toward something horrible – say, torturing babies for fun. The theory implies that in that situation (when G takes that attitude), it is morally right to torture babies for fun. But obviously that isn't right. So the theory is false.

Notice, btw, how the Nazi argument above is just an instance of this: Cultural relativism says that x is right if and only if *society* takes the attitude of *approval* toward x . Then for x , you could plug in the act of sending Jews to concentration camps, and the theory says: If society approves of sending Jews to concentration camps, then doing so is morally right.

13.4. What's Wrong with Nihilism

13.4.1. The Nihilist View

Subjectivism and non-cognitivism can both be refuted *empirically*, that is, by evidence that comes from observation. Both theories contain false views about what moral language means, so we can refute them by looking at how people talk about morality. We just do not talk in the way that we would if moral sentences were used to express emotions, describe social conventions, etc.

By contrast, nihilism can't be refuted in that way. The nihilist agrees with the realist that when we talk about "right", "wrong", "good", and "bad", we are *trying* to describe objective facts. The sentence "Murder is wrong" is really used to ascribe an objective property of wrongness to the act of murder. But the nihilist disagrees with the realist about whether these objective moral properties *exist*: In the nihilist's view, we are just *mistaken* when we attribute objective moral properties to things. Thus, nothing is either right or wrong, good or bad. Statements like "Murder is wrong", "Hitler was evil", and so on are simply false.

Why would someone hold this view? Basically, because they're convinced by the above arguments against non-cognitivism and subjectivism, but they don't believe in objective values. But why not believe in objective values? One argument was discussed above (§13.3.3), so I won't repeat that here. But there are a couple of additional arguments ...

13.4.2. Against Objective Values: The Humean Argument

The following is the most popular argument in meta-ethics, and possibly in all of contemporary philosophy; I've seen it in more places than any

other philosophical argument. I call this the “Humean Argument”, and the people who endorse it “Humeans”:

P1.If there were objective values, then beliefs about them would motivate us to act, independent of our desires.

P2.Beliefs alone can never motivate action; only desires can motivate.

C.Therefore, there are no objective values.

Other philosophers must think this argument is super-persuasive, since so many of them have rested so much on it. Or at least, anti-realists find it persuasive.^[80]

Interlude: The Humean Theory of Reasons

David Hume was a very famous Scottish philosopher from the 1700s. He spent most of his time being skeptical about stuff (you know, like religion, morality, the external world, and everything that isn’t directly observed). Also, though, he is widely thought to have held a certain view about reasons for action, which is hence called the “Humean Theory of Reasons”. This is the view that reasons for action always depend on desires: If you want x , and you believe action A (which you can do) will increase your chances of getting x , then you have a reason to do A . If you *don’t* want anything that you think A will make more likely, then you don’t have any reason to do A . (You can also say the same thing with “motive” substituted for “reason”.)

Btw, “desire” here is supposed to refer to a particular sort of mental state that you can be introspectively aware of. “Desire” can’t really be defined (like most things; see §8.5), but it is important that it is separate from belief (or other cognitive states) – i.e., no belief can be a desire. Also, “desire” does not merely mean “motive”, which would turn the Humean theory into the trivial thesis that motives for action always depend on motives and would (incorrectly) imply that if a belief motivated action then the belief would be a desire.

The Humean theory of reasons is controversial, since some philosophers think that you can have a reason to do A merely because you believe A is morally right, even if it won’t get you anything you want.

Now, why accept premise P1? Generally, little is said to defend that; it’s just supposed to be intuitive. Note that P1 is generally understood to mean that moral beliefs would provide us with *at least some* motivation to act in the ways we consider moral; that is compatible with our also having

conflicting motives (like, say, self-interest), which can sometimes be stronger than our moral motives.

You can give examples to illustrate the idea. The comedian Louis CK has a comedy bit in which he talks about why he eats meat:

People talk about like, “Don’t eat dolphins when they’re in the tuna.” ... Why not kill and eat a dolphin? No, I don’t fuckin’ get it. If you’re a tuna, fuck you, we’re eating you. So I don’t really see the difference. And I think it’s *wrong* to eat tuna, and dolphin, and cows, and everything. But I eat them. I eat them all. Because I *don’t care* that it’s wrong. I totally think it’s *terrible*, but that’s not important to me. So what if it’s wrong?^[81]

There is something incongruous about his insistence that eating meat is terrible but that this isn’t important to him. That’s part of what’s funny. (You have to hear his delivery to get the full comedic effect.) You might even think it’s inconsistent to say that you think something is wrong, terrible, etc., and yet that you don’t care and it isn’t important to you. (How could this be inconsistent? It must be because recognizing an objective moral fact *necessarily motivates you* to act in accordance with it!) Note: That part in parentheses there is *the anti-realists’ interpretation of the realist’s position*. That’s what anti-realists think that moral realists are thinking. If that’s right, then the idea of an “objective moral value” is something like this: It’s an objective fact with the peculiar feature that, as soon as you become aware of it, it automatically motivates you to behave in a particular way. Hence, we arrive at

P1.If there were objective values, then beliefs about them would motivate us to act, independent of our desires.

My assessment of the case for P1: This is extremely dubious. Moral realists do not in fact all agree that objective values necessarily motivate you to act independent of your desires. Some moral realists say that you are motivated to act morally only to the extent that you *desire* to be moral. For instance, they’d say that the belief that stealing is wrong does not *by itself* motivate you to avoid stealing; it only motivates you to avoid stealing if you already want to avoid doing wrong. About the Louis CK comedy bit: Part of the joke lies in the fact that many people *are* like Louis CK – they know that meat-eating is wrong, but they just don’t have the motivation to do anything about it. It’s just that people usually aren’t that honest about it. The joke isn’t that he’s contradicting himself; the joke is that he’s an

immoral person. All of which is to say that the anti-realist's interpretation of "objective values" is questionable – it's far from obvious that objective values would have to be inherently motivating.

Anyway, what about premise P2? Why think that only desires can motivate people? When you think about some mundane examples, you can see that a belief by itself doesn't tell you what to do. Say you are standing on the platform at Union Station in Denver. Someone comes up and tells you, "Hey, the train for the airport is leaving in 1 minute." Would that information motivate you to do something – say, to get on the train? Not by itself, no. Sure, if you already want to get to the airport, *then* the fact that the train for the airport is about to leave would motivate you to get on that train. But if you want to avoid the airport, then the same fact would motivate you to stay away from that train. The information about the train *by itself* doesn't give you a goal, since it doesn't tell you whether getting to the airport is desirable or not.

You can give lots of examples like that. Humeans generalize to say that you only have a reason to act if you start out with some goal, and just believing something about how the world *is* can't give you a goal. Only your desires give you goals. Hence, we have

P2. Beliefs alone can never motivate action; only desires can motivate.

Now here's a possible objection. Sometimes, we say things that seem to imply that we are motivated by something other than desire. E.g., you spend the night studying even though you "don't want to" – you'd much prefer to be out drinking, tipping cows, or whatever it is that college students do when they're not studying. Yet you study anyway. So you must be motivated by something other than your desires, right?

Well, the Humeans would say: No, that just shows that you *wanted* to get a good grade in your class *more* than you wanted to go out drinking. In general, they would say, all cases in which people claim to do something that they "don't want to do" are just cases in which they had conflicting desires, not cases where they really did something that they didn't expect to satisfy *any* of their desires.

My assessment of the case for P2: All of that is also extremely dubious. What the Humean says about beliefs is indeed true of *descriptive* beliefs, but not true about *evaluative* beliefs. If you have a purely descriptive piece of information – say, that the airport train is about to leave – that doesn't tell you what to do, because, e.g., it doesn't tell you whether getting to the

airport is good or not. That's right. But the proposition [Getting to the airport is good] obviously *does* tell you whether getting to the airport is good. Evaluative propositions just are propositions that tell us what we should do, or what is good or bad. The Humean theory only seems intuitive when you are thinking about purely descriptive facts.

Now, you might of course question whether there ever is a fact such as [Getting to the airport is good]. But remember where we are: The Humeans were supposed to be giving us an *argument* to *show* that there aren't any such facts. So they can't *assume* that there aren't any. If we don't start by assuming that there are no evaluative facts, then it just is not at all obvious why we would say that beliefs can never motivate action, regardless of whether they are descriptive or evaluative.

My assessment of the argument overall: It just isn't clear why we should accept either premise. They're not *obviously false*, but nor are they obviously true. Either premise could reasonably be rejected. So I don't find the argument persuasive.

13.4.3. Against Objective Values: The Argument from Weirdness

Okay, I think this might be what is really motivating nihilists and other anti-realists: Objective values are *weird*. In fact, one famous argument against moral realism is officially named "the argument from queerness". [82] If there are objective values, they are very different from all the things that science studies. It's weird that they're not part of our best scientific theories about the world. It's weird that we can't detect them by the five senses, nor by any scientific instruments. People will say stuff like this in conversation, though usually not in print (actually, they usually give even less explanation than I just did).

Aside

Variations of the argument from weirdness appear all over philosophy. People say that moral value is weird, the soul is weird, libertarian free will is weird, abstract objects (numbers, sets, etc.) are weird, synthetic a priori knowledge is weird – and therefore, that these things don't exist. (Compare §12.4.1.) This sort of "argument" seems to have an *enormous* impact on the prevailing philosophical views. I personally think it's an embarrassment that philosophers rest so much weight on such a vague, inarticulate "argument".

Let's think about what the charge of weirdness really means. First interpretation: Maybe it means "counter-intuitive". In that case, the premise

of the argument from weirdness is just false: Objective values are not counter-intuitive at all. You can tell this from the fact that almost all societies throughout history seem to have regarded values as objective, most thinkers in the history of ethics have done likewise, and *even the nihilists themselves admit* that moral realism is built into ordinary language. (The standard nihilist view is that words like “good”, “bad”, “right”, and “wrong” are intended to refer to objective moral properties. That’s why the nihilists think that all moral claims are false.) So it’s hard to see how you could claim that moral realism is counter-intuitive. By the way, you’ll find lots of *much weirder* things if you start studying modern physics.

Second interpretation: Maybe weirdness just amounts to being very different from other things. But then, lots of things are weird in that sense. Matter, space, time, numbers, fields, and consciousness are all weird (different from other things). Why should we believe that weird things don’t exist? This is just a very lame argument.

13.4.4. Nihilism Is Maximally Implausible

The main reason for rejecting nihilism is its extreme initial implausibility. In other words, the Moorean Response to skepticism (§7.4) applies to this case. Take an uncontroversial moral statement, the most obvious you can think of – say, “You shouldn’t torture babies for fun.” That is extremely plausible on its face; it is indeed difficult to think of any statement that is more plausible. The nihilist wants us to reject that statement – he says it’s *false* that you shouldn’t torture babies for fun – on the basis of the sort of arguments discussed above. So consider the following four propositions:

1. You shouldn’t torture babies for fun.
2. “You shouldn’t torture babies for fun” entails that there are objective values.
3. If there were objective values, then beliefs about them would motivate us to act, independent of our desires.
4. Beliefs alone can never motivate action; only desires can motivate.

Those four propositions are jointly inconsistent (1+2 entail that there are objective values; 3+4 entail that there aren’t). So we have to reject at least one of them. Which should we reject? Whichever one is the least initially plausible (the least obvious on its face). Which one is that? I’m not sure, but it’s definitely not (1). Proposition (1) is in fact the *most* obvious of the four. The nihilist wants us to reject (1) on the basis of (2)-(4). That’s irrational –

it's irrational to reject the most plausible proposition rather than one of the less plausible ones. I would reject (3) or (4), maybe both. (For defense of (2), see sections 13.2–13.3.)

The same point will almost certainly apply to any argument for nihilism. The nihilist would have to find some premises that are more plausible than “You shouldn’t torture babies for fun”. But “You shouldn’t torture babies for fun” is just about the most plausible statement I can think of.^[83] No philosopher has come up with any premises that are more obvious than that that could be used to argue against it. For example, suppose the nihilist uses the premises “Moral values are weird” and “Weird things don’t exist”. Well, those are *much* less obvious than “You shouldn’t torture babies for fun.” So they couldn’t be used to refute the proposition that you shouldn’t torture babies for fun.

13.5. What’s Wrong with Ethical Naturalism

13.5.1. The Naturalist View

Ethical naturalists think that moral facts are real and objective, but that they are **reducible** to descriptive facts. Basically, this means that you can explain what it is for something to be good (or right, etc.) using some more fundamental terms, without using any moral terms.

Here’s an example of something that’s reducible: water. You can explain what water is in more fundamental terms, without using “water” or any other word that is close in meaning to “water”. Here is how you would explain it: Water is *a substance composed of molecules each of which contain one oxygen atom bonded to two hydrogen atoms*. The italicized phrase is the “reduction” of the concept water – it explains the underlying nature of water without using water or any similar concept.

Maybe moral rightness is like that; maybe you could explain what it is without using moral concepts. For instance, maybe rightness is *the property of an action where the act increases the total amount of enjoyment and/or reduces the total amount of suffering in the world*. Notice that the phrase in italics doesn’t use “right” or any other moral term.

How would we know moral truths, on this view? One of the central motivations of ethical naturalism is to avoid having to appeal to things like “intuition” in explaining ethics. That’s part of why they appeal to analogies like the scientific theory of water. The naturalists want to say that ethical theories could be justified in the same way as scientific theories are – which is to say, because they help to explain our observations. Here’s an actual

example from the literature: How do we know that Adolf Hitler was evil? Well, it's an empirical (based on observation) fact that he ordered the Holocaust, and the theory that he was evil helps to explain why he did this. Therefore, we have empirical justification for thinking he was evil.

So ethical naturalists hold two views: (i) that ethical truths are reducible to objective, descriptive facts, and (ii) that ethical knowledge is justified on the basis of observations.^[84]

13.5.2. A Point About Meaning

Ethical naturalists (and maybe everyone else) used to be confused about meaning. They would say or imply that you could explain the *meaning* of the words “good”, “right”, etc., using non-moral terms. E.g., they might say that “right” *just means* “increases the total enjoyment in the world”.

You can see that that's false since it makes sense to ask, “Is it right to increase the total enjoyment in the world?”, and that obviously does not mean the same as “Does increasing the total enjoyment in the world increase the total enjoyment in the world?”

Interlude: G.E. Moore's Open Question Argument

G.E. Moore was a 20th-century British philosopher who wrote a book called *Principia Ethica*. (He also gave a famous reply to skepticism in epistemology; see §7.4.) In the book, he argued that “good” could not be defined. His argument for this came to be called “the **open question argument**”, mostly because there was a passage in which he briefly used the phrase “open question”. The context was something like this: Suppose someone claims that “x is good” means “x promotes pleasure”. G.E. Moore says it's an open question whether promoting pleasure is good. (Roughly, what this means is that you could make sense of someone wondering whether promoting pleasure is good – that's not inherently confused.) But it's not an open question whether promoting pleasure promotes pleasure. Therefore, “good” can't simply mean “promotes pleasure”. Moore thinks you can give similar arguments against *any* reductionist definition of “good”.

These days, naturalists generally do not claim to explain the *meaning* of moral terms anymore. They generally recognize that you can't do that (at least, not without using other evaluative terms). But they still claim that we can explain the underlying nature of moral properties or moral facts. That's why the “water/H₂O” example above is useful. Notice that the H₂O theory explains the underlying nature of water, even though it does *not* explain the

meaning of the word “water”. People understood the word “water” long before modern chemistry arose. Before the late 18th century, people thought that water was an element. It’s not the case that they didn’t understand what “water” meant, though. So it can’t be that “water” *means* the same as “H₂O”.

Similarly, the ethical naturalists would say, most people today understand the meaning of “good”, “right”, etc., but they don’t know the correct theory that explains the underlying nature of moral properties.

13.5.3. Bad Theories

The theory mentioned in §13.5.1, about the right being whatever increases the total amount of enjoyment, is not a very good theory, since there are many counterexamples to it. Example: Suppose you could torture an innocent person, and a very large number of sadistic people could then watch this torture on video, each getting a little sadistic pleasure from it. There would be so many sadistic viewers that the total amount of pleasure would be greater than the amount of suffering felt by the victim. Would it be morally right to do this? It seems not.

That was just one possible reductionist theory. Here’s another one. Suppose someone says, “The good is that which promotes one’s survival”. (This implies that what is good is relative to an individual, since an event can promote one person’s survival but not another’s.) Counter-examples: Imagine an event that causes you to live slightly longer but also causes you to be in agony the entire time. Or imagine an action that increases your life expectancy by 5 minutes but also kills 500 other people. Are these things good?

There are similar points to be made about every reductionist theory that anyone has devised; all have apparent counter-examples. In general, it seems that there are many different things that are good, and many different kinds of action that are right.

Ethical naturalists today often talk in the abstract about how a reductive theory of goodness might exist, even though they don’t actually have the theory. This is odd. There’s no reason to think goodness is reducible if you can’t plausibly say what it might reduce to. Think about the water example: Why do we think water is reducible? Because we have a specific reductive theory that we have empirical evidence for. People didn’t just claim that water was reducible in general before there was any good theory about it.

People thought it was an element *until* they had specific evidence that it could be decomposed into hydrogen and oxygen. Of this, more below.

13.5.4. A Bad Analogy

Contemporary ethical naturalists motivate their theory using examples like the water/H₂O example above. (Also: heat and molecular kinetic energy; sound and compression waves in the air; color and spectral reflectance distributions; etc. There are many examples of reductionist theories in science. But we'll stick with the water example for simplicity.) But the water/H₂O analogy is a bad one – naturalist theories about ethics are not like that at all. When you see why, you'll see that the basic idea of ethical naturalism is confused.

How do we know that water is H₂O? There are two main pieces of evidence. In one experiment (the electrolysis experiment), you apply an electric voltage to some water. Bubbles of hydrogen gas form at the cathode and oxygen gas at the anode, and the mass of water present declines by the same amount as the mass of hydrogen and oxygen that is produced. In another experiment, you burn some hydrogen in the presence of oxygen, and you get some water that condenses on the sides of the container. The mass of water produced equals the mass of hydrogen and oxygen consumed.^[85] The best explanation for these experimental results is that water is a compound of hydrogen and oxygen. A person who doesn't already believe the water=H₂O theory can see how the theory would explain these results. Furthermore, they can independently (i.e., without using the water=H₂O theory) detect water, hydrogen, and oxygen using the observable properties of these substances, they can do the experiments, and they can verify that the results are as stated.

Now, exactly what comparable evidence could be given for a naturalistic ethical theory? For simplicity, let's say someone has the theory that goodness=pleasure. What experiments could we do to test that, in the way that the electrolysis experiment tests whether water=H₂O? First produce some goodness, then see whether it's pleasurable? How would we do that? Make some pleasure, then see whether it's good? Well, we can certainly cause some pleasure, but we don't have any way of detecting the moral goodness or non-goodness of it, without relying on our pre-existing ethical views. Notice how this is different from the water/H₂O example.

You might say we can just *intuit* that pleasure is good; it will just *seem good* to us. This, however, is relying on another meta-ethical view, intuitionism (discussed in §13.6), which naturalists reject. Ethical naturalists are specifically trying to *avoid* appealing to intuition, which is why they use analogies to scientific theories in the first place.

What of the idea that we could explain, e.g., why the Holocaust happened by appealing to the theory that Adolf Hitler was evil (§13.5.1)? Sure, that explanation works, *provided* that you already accept that genocide is wrong, bad, evil, or something like that. But if you thought that genocide was *good*, then you could equally well “explain” the Holocaust by citing Hitler’s *goodness*. In general, everything that you can explain by citing a moral claim could equally well be explained by someone who held an *opposite* value system; they would just use an opposing moral claim. Because of this, you can’t justify a particular value system by appealing to this sort of explanation (i.e., explanations like “Hitler ordered the Holocaust because he was evil”).

Notice that this is *not* like real scientific theories. The water= H_2O theory explains why you can convert a sample of water into hydrogen and oxygen gas, and you don’t have to already believe the theory (or any other theory of water) to see that it would explain that. If you hold an opposing theory – say, that water is a compound of uranium and radon – you *cannot* explain the experimental results equally well.

That’s how the “good=pleasure” theory is different from the “water= H_2O ” theory. That explains why “water= H_2O ” is justified, whereas “good=pleasure” is completely groundless. Once you see this, you can probably see that a similar point could be applied to any reductionist theory about goodness, or any other moral property. They’re all going to be like that – if you assume some completely opposite moral theory, it makes no difference to your empirical predictions. Basically, once you fix the *descriptive* facts, different moral views make no difference to what you should expect to observe in the physical world.

There is one plausible exception to that: Different moral views make a difference to what we should expect *people to believe*, *provided* that people have a faculty of ethical intuition, or something like that, which can detect moral facts as such. Naturalists don’t accept such a faculty, though. Fortunately, intuitionists accept it, which we will discuss presently.

13.6. Ethical Intuitionism

13.6.1. The Intuitionist View

Ethical intuitionists hold three main views: (i) There are objective moral truths, (ii) (some) moral truths are irreducible, and (iii) we know (some of) these truths through ethical intuition.

Why believe (i)? Because the only ways to not have objective moral truths would be for subjectivism, non-cognitivism, or nihilism to be true, and we saw the problems with all of those above (§§13.2–13.4). Why believe (ii)? Well, we saw the problems with reductionism above as well (§13.5). It remains to talk about (iii).

Intuitionists like to compare ethics to mathematics. Note: This does not mean that ethics is exactly like mathematics in all ways (if that were true, this wouldn't be a *comparison*; ethics would just *be* mathematics). Rather, we draw the comparison to highlight certain specific points. People sometimes ask, for example, *where* goodness is, or where it “comes from”. Goodness is not located anywhere, nor does it come from anywhere, any more than the number 2 is located or comes from somewhere.

More importantly, people sometimes find ethical knowledge weird because it is not based on observation. But mathematics is not based on observation either. Mathematics starts from certain self-evident axioms, from which you can then infer further conclusions. What is a “self-evident” proposition? Basically, it's one that is obvious when you think about it, in a way that doesn't require an argument; you can directly see that it's true. For instance, that 3 is greater than 1, that the shortest path between two points is a straight line, or that if $a=b$ and $b=c$, then $a=c$.

Similarly, perhaps the field of ethics rests on self-evident ethical axioms. For instance, maybe it's self-evident that enjoyment is good in itself; that one should not cause harm for no reason; or that if a is better than b and b is better than c , then a is better than c .

Now, what is an “**ethical intuition**”? Essentially, an *intuition* is a mental state that you have in which something just seems true to you, upon reflecting on it intellectually, in a way that does not depend upon your going through an argument for it. An *ethical* intuition is just an intuition that's about ethics. All the above are examples of intuitions. E.g., when you think about $[3 > 1]$, you should have a (mathematical) intuition that it's true; when you think about [It's wrong to cause harm for no reason], you should have an (ethical) intuition that that's true.

Why should we believe our intuitions? In an earlier chapter, we discussed the principle that it is rational to assume that things are the way they appear, unless and until one has specific reasons to doubt this (see §7.6). This, I argue, is the foundation of all reasonable beliefs. That includes the beliefs that we get from perception, memory, introspection, and reasoning, as well as intuition – in all of these cases, we believe what we believe because it seems correct to us and we lack sufficient reasons to doubt it. So, that’s also why it makes sense to believe, for example, that it’s wrong to cause harm for no reason: That seems true, and we have no good reason to doubt it. (Of course, some would claim to have good reasons to doubt it; see §13.4.)

13.6.2. *Objection: Intuition Cannot Be Checked*

Some object that intuition is not an acceptable way of forming beliefs because there is no way of *checking* a particular intuition to see whether it’s really true, and thus no way of knowing whether intuition in general is reliable. (Some would say this about intuition in general; others would only say it about *ethical* intuition.)

This is false in one sense but true in another. If you’re allowed to consult other intuitions – both your own and other people’s – then you *can* check on a particular intuition. For instance, if I intuit that murder is wrong, I can “check” that by asking whether other people also intuit that. I can also see whether my intuition that murder is wrong is consistent with my other ethical intuitions (say, my intuition that it’s wrong to cause harm for no reason, my intuition that life is valuable, and so on). So it’s just *false* that you can never check on an intuition. Many intuitions can be tested in these ways and will in fact pass the tests.

Of course, some would object to the idea of using intuitions to check other intuitions. If you’re *not* allowed to consult other intuitions, then indeed you generally cannot check on a particular intuition. That’s the sense in which it’s *true* that you can’t check intuitions. However, in that sense, you cannot check on any of the other basic types of cognition that we rely on either (see §§7.2, 7.5.5). For instance, there is no way of checking on observations made by the five senses, without relying on other observations. If you want to check on the reliability of your senses, you could, say, ask other people whether they perceive the same things you do. But that would depend upon your *perceiving* those other people, perceiving

the answers they give, and trusting those perceptions. A similar point applies to basically any test you might try to do.

Similarly, if some skeptic comes along and doubts whether memory is reliable, you have no way of settling that doubt without relying on memory. Let's say I want to test my memory. I seem to remember where I live. So I go to the address that I remember my house is at, and, lo!, I find a house there that looks just like the one I remember. I go inside, and there is a bunch of stuff there that looks just like the stuff I remember. Etc. This suggests that my memory is reliable. However, this test requires me to use my memory in the process of testing it – when I get to the house, I must *remember* that I previously remembered a house just like that. Furthermore, to construct a suitable inductive argument, I will need to *remember* many cases like this – i.e., many cases in which my memory proved reliable in the past.

Finally, you face the same problem with reason itself, if you want to verify that reason is reliable. You could try constructing an argument that reason is reliable, but that would require you to reason.

My point: Intuition is just like reason, observation, and memory in this respect: You can't check its reliability without using it. You probably don't think (and very few moral anti-realists think) that we should ignore reason, observation, and memory; therefore, you also shouldn't ignore intuition merely because it can't be checked without using intuition itself. According to Phenomenal Conservatism (§7.6.1), this is all okay, because we are permitted to *start* from the assumption that what seems to be the case is the case, unless and until we have specific reasons for doubting that. We don't have to first prove that appearances are reliable.

13.6.3. *Objection: Differing Intuitions*

Perhaps the most popular objection to intuitionism is that people sometimes have conflicting intuitions, and therefore (?) we should not believe intuitions. I put a question mark there, because it's not always clear what the argument is. I think there are a few different strands of thinking:

First, sometimes it seems that the concern is that, if we form ethical beliefs based on intuitions, then there may be intractable disagreements. We might be unable to resolve a disagreement with someone who has conflicting intuitions, and that is bad. In reply, I would note that this is not actually an objection to intuitionism, in the normal understanding of an objection – that is, the “objection” does not attempt to cite any evidence

that intuitionism is not actually true. It just cites something bad that could happen if intuitionism is true. And I agree that this bad thing could happen – you could find yourself in an intractable disagreement with someone. This doesn't mean intuitionism isn't correct, though; in fact, there *are* intractable disagreements, so this is just a *correct prediction* of intuitionism. There are also, by the way, equally intractable disagreements about many other things besides ethics. All areas of philosophy, as well as religion, as well as many (non-ethical) political questions occasion disagreements that appear about as difficult to resolve as the disagreements in ethics. No one should be surprised that there might be disagreements in ethics that we can't resolve.

Second, some would argue that disagreement undermines the claim to objectivity – an ethical claim can't be objectively true if we don't all accept it. This sort of argument, however, just rests on a conceptual mistake (see §13.3.3). Objective truths are defined as truths that don't depend on observers; they are not defined as truths that everyone agrees on. You can't stop a fact from existing objectively merely by refusing to accept it! On the intuitionist view, ethical facts are not dependent on our intuitions. Ethical facts exist independent of us; ethical intuition is merely our way of becoming aware of them. (Compare: Physical objects exist independent of your sensory experiences; your sensory experiences are merely your means of becoming aware of these independently-existing objects.)

Third, some believe that intuition just isn't very reliable, due to the variations in intuitions across individuals and across cultures. This is the most reasonable version of the objection. If there are objective ethical facts, then whenever people's intuitions disagree, someone has to be wrong. If there's a lot of disagreement, then intuitions go wrong a lot. But there's no reason to assume that *your* intuitions in particular are much better than other people's; therefore, if ethical intuitions in general are unreliable, yours are probably unreliable, and therefore you should stop trusting them.

In response, this objection is correct about *some* intuitions. For instance, suppose that abortion just seems wrong to you, intuitively. That's a highly controversial intuition – many people lack that intuition or even have an opposite intuition. Therefore, you should not rest much weight on that intuition. (Note, however, that there is probably a lot less *intuitive* disagreement than there is disagreement in *reasoning*. For example, people who think abortion is wrong almost certainly think that because they have an *argument* for that conclusion, not simply an intuition.^[86] But for the sake

of argument, let's just assume that someone has an intuition that abortion is wrong.)

However, there are other intuitions that are not controversial. For instance, that enjoyment is intrinsically good, that you shouldn't cause harm for no reason, or that if *a* is better than *b* and *b* is better than *c*, then *a* is better than *c*. That's why I used those examples above (and not the example of "abortion is wrong") in §13.6.1. You can't very well argue that *those* intuitions are unreliable due to widespread disagreement, when there *isn't* widespread disagreement about those things.

This is enough for the intuitionists. Intuitionists are not silly: They don't claim that we know *all* ethical truths. They only claim that we know *some* ethical truths.

By the way, something similar holds in many other areas, perhaps all areas of human intellectual endeavor: There are always things that are obvious and uncontroversial, and then other things that are widely disputed. That's true in other branches of philosophy, and in science, and in everyday life. It's reasonable to withhold judgment about the disputed questions (especially if you're a non-expert). But there is no reason to also withhold judgment about the obvious and uncontroversial points.

13.7. Conclusion

Ethical intuitionism was often ridiculed in twentieth century philosophy. It has, however, enjoyed something of a resurgence in recent decades, perhaps because it is in fact the most sensible view in metaethics.

The alternatives are non-cognitivism, subjectivism, nihilism, and naturalism. Non-cognitivism is no good because ethical sentences behave exactly like cognitive (proposition-asserting) sentences in all known respects. Subjectivism is no good because it implies that if you (or your society, or whatever) approve of torturing babies for fun, then torturing babies for fun is good. Nihilism is no good because it implies that recreational baby torture isn't wrong, which is too implausible on its face to be justified by speculative philosophical assumptions. And naturalism is no good because reductive theories of goodness do not make empirical predictions in the way that real scientific theories do, and thus naturalistic theories of goodness have no justification.

Turning to the common objections to intuitionism: The inability to check intuition without relying on intuition is not a major problem, since we similarly cannot check on memory without relying on memory, on

observation without relying on observation, or on reason without relying on reason. On the other hand, disagreements in ethics *do* provide a reason for doubting one's own ethical intuitions; however, this only applies to intuitions about controversial questions, not intuitions that virtually everyone agrees on.

I suspect that the main reason why many people are not comfortable embracing ethical intuitionism is that they vaguely sense that the view is “weird”. I think that is lame – I think that feeling of weirdness has no evidential value. So we should feel free to embrace the view that coheres with common sense ways of thinking about morality.

14. Ethical Theory, 1: Utilitarianism

14.1. An Ethical Puzzle

I opened the previous chapter with one of the most famous hypotheticals in ethics, the “**Trolley Problem**”:

Trolley: A runaway trolley is headed for a fork in the track. The trolley is currently set to go to the left, where it will collide with and kill five people. You can instead switch it toward the right track, where it will kill only one person. Should you switch the trolley?

Most people think that it’s okay, perhaps even obligatory, to switch the trolley, because doing so will result in only one death instead of five. Now here is a case where people react differently:

Footbridge: A runaway trolley is going to collide with and kill five people. There is a footbridge over the track, where you and a much larger man are standing. If you push the fat man off the bridge, he will land on the track in the path of the trolley. Due to his size, the man’s body will be enough to stop the trolley before it hits the five people on the track, thus saving the five while killing the fat man. (Your own body, by contrast, is too small to stop the trolley. ^[87]) Should you push the fat man?

In both of these cases, assume there are no other consequences of the action beyond the obvious (e.g., you won’t be prosecuted if you kill the one person, the one person wasn’t just about to cure cancer, the people you save are not serial murderers, etc., etc.). All that will happen is that one person will die and five will be saved. With that understood, most people think that killing the one person to save five is *right* in Trolley, but *wrong* in Footbridge.

This raises a puzzle: What’s the difference between these cases? Why might it be okay to kill one person to save five in one of these cases, but wrong in the other two?

Philosophers have no agreed upon answer to that. ^[88] Some claim that there actually *isn’t* any relevant difference, and that it’s okay to sacrifice the one person in both cases. Others think it’s *wrong* to sacrifice the one person in both cases. And still others offer reasons why Trolley is different from Footbridge, such that sacrificing the one person is okay only in the Trolley case.

I've raised this puzzle to illustrate the sort of thing that people in ethics talk about. In the remainder of this chapter, we'll discuss *ethical theory*. Ethical theories give general accounts of what actions are right and wrong, what outcomes are good and bad, and/or what traits are virtues and vices. So they could be used to answer questions like the one about the trolley. However, there is no consensus among philosophers on what ethical theory, if any, is correct. Instead, there are a number of competing theories and principles, which we're going to talk about presently.

14.2. The Utilitarian View

Let's start with the most popular comprehensive ethical theory among philosophers, though one that is still held only by a minority of philosophers: utilitarianism. Utilitarianism says that the right action in any circumstance is always the action that results in the greatest total quantity of well-being in the world, for all beings affected by the action, where well-being is understood here in terms of enjoyment or desire-satisfaction. (Some utilitarians would say "enjoyment"; others would say "desire-satisfaction". For present purposes, you can treat suffering as negative enjoyment, and desire-frustration as negative satisfaction.^[89])

Utilitarianism can be broken into three key controversial ideas:

1. **Consequentialism:** *The right thing to do is always whatever produces the best consequences overall in the long run. That is: You should always make the choice such that, if you make it, the greatest amount of good will exist, out of all the choices available to you.*
2. **Hedonism or preferentism:** *The only intrinsic good is pleasure (for the "hedonistic utilitarian") or desire-satisfaction (for the "preference utilitarian"). Hereafter, I'll use "utility" for pleasure and/or desire-satisfaction, for simplicity of expression.*
3. **Impartialism:** *The utility of all beings is equally important. One should not, e.g., privilege one's own interests over those of others, or one's own family over other families, or one's own species over other species. Just produce the greatest total benefits, regardless of who gets them.*

There are debates about all three of those, which we'll talk about below.

14.3. Consequentialism

14.3.1. Objections to Consequentialism

It makes sense, on its face, that the ethical thing to do would be to produce the most good. Why would anyone object to that?

Well, because of cases like Footbridge. And there are lots of other examples.^[90] Here are a few more:

Organ Harvesting: You are a doctor who has five patients who all need organ transplants, without which they will die. One needs a heart, one needs a lung, one needs a liver, and two need kidneys. At the same time, you have one healthy patient who just happens to somehow be compatible with all five of the sick patients. This healthy patient does not want to give up any of his organs, since that would kill him. You could nevertheless kill the healthy patient, harvest his organs, and thereby save five other patients. Should you harvest the organs?

Framing: A crime has been committed in a certain town that has caused great public outrage. The sheriff knows that if no one is punished, there will be riots. Unfortunately, the sheriff cannot find the actual criminal. But he can frame an innocent person, causing that person to be punished and thus forestalling the riots. The innocent person would be seriously harmed, but this would be a smaller total quantity of harm than the harm that would be caused by the riots. Should the sheriff frame the innocent person?

Promise: You and your best friend have gotten caught in a snowstorm in Antarctica. At some point, it becomes clear that your friend won't make it out alive. He asks you to promise that when you return to civilization, you will make sure that his entire fortune goes to his son. You make the promise. But when you get back to civilization, you realize that your friend's son is probably just going to waste the money, and it would do much more good if given to charity. Should you tell everyone that your friend's dying wish was to give his entire fortune to charity?

Electrical Accident: Jones has had an accident and gotten caught in some electrical equipment at a television station. He is currently suffering painful electrical shocks. In order to rescue him, you would have to turn off the TV station's transmitter for 15 minutes. This will relieve Jones' pain but will also interrupt the broadcast of the World Cup, which a large number of people are watching, thus causing a diminution in entertainment for many people. Alternately, you could wait until the broadcast is over, in which case Jones will suffer great pain for the next

hour, but no one will have their entertainment interrupted. There are so many people watching that the total decrease in welfare from interrupting the broadcast would be greater than the decrease in welfare suffered by Jones if he remains trapped. Should you rescue Jones now, or wait for the game to finish?

In all of those cases, we have an action that intuitively seems wrong, even though it would seemingly produce the best overall consequences (pushing the fat man off the bridge, harvesting the organs, framing the innocent person, breaking the promise, and leaving Jones trapped for an hour). I frequently discuss such cases in classes, especially the Trolley and Organ Harvesting cases. In Organ Harvesting, people's reactions are particularly strong – almost everyone is extremely confident that you may not kill the healthy patient. That's *more* intuitive than the claim that you *may* turn the trolley in the Trolley case.

Sometimes, utilitarians try to come up with other negative consequences that the action might have, in order to explain why you don't have to do the seemingly immoral thing. Maybe the organ-harvesting doctor would get sent to prison and be unable to treat any more patients ever again, and that might result in fewer lives being saved in the long run. Or maybe the public will find out about the organ harvesting, patients will become afraid to go to their doctors, and then fewer lives will be saved in the long run. Of course, the response to such speculations is to stipulate that these things are not the case. This is a hypothetical example, so we can stipulate what happens. Assume that the doctor *will not* get caught, other patients *will not* find out what he did, etc. (Similar things can be said about the other cases.)

Note about how to treat hypotheticals

Discussion of hypothetical examples is not like real life decision-making. In *real* situations, you should always look for ways out of a dilemma or ways of avoiding having to confront a hard issue. You should also try to consider all possible (realistic) consequences of an action. That's because in real life, you're just trying to do the right thing in that particular case.

But in discussing hypothetical examples in ethics, we're doing something very different. We're trying to illuminate a specific theoretical issue. Thus, in discussing hypotheticals, one should never try to avoid the dilemma or avoid addressing the hard issue that the example is trying to present. One also should not bring up possible consequences that are not

related to that issue. Doing so only makes the other person take up time tweaking the example to try to avoid the irrelevant issues, which is not a useful way to spend our time.

With that understood, utilitarians would generally “bite the bullet”, as we say: That is, they’d just accept the counter-intuitive consequences of their theory. (Indeed, I have one colleague whose diet seems to consist almost entirely of bullets. There’s no need to name him; everyone who knows him knows whom I’m talking about.) They would say it’s right to push the fat man in Footbridge, right to kill the healthy patient in Organ Harvesting, etc. There are a couple of ways they try to make this seem more palatable. One is by criticizing alternative theories that try to explain the difference between Trolley and Footbridge. If we try very hard to find a believable explanation of that, and we can’t, then at some point we might just conclude that our intuitions are mistaken and there is no relevant difference. We’ll talk about some alternative theories below. I myself am not very impressed with this approach, though, because (a) I’m not convinced that we’ve yet thought of all the theories about that difference, and (b) even if you say there *isn’t* any difference between Trolley and Footbridge, I think it’s at least as plausible to conclude that the action is *wrong* in both cases as to conclude (with the utilitarians) that it is *right* in both cases.

The other way that utilitarians try to make their bullet-biting palatable is by directly questioning the reliability of intuitions about hypothetical cases. Maybe our intuitions are attuned to real-world cases, and in the real world, removing people’s vital organs, pushing people off bridges, etc., normally has overall bad consequences. In the hypothetical scenarios, those actions have good consequences, but our unconscious evaluation mechanism still reacts negatively because those types of actions are *usually* bad. So we experience a negative emotional reaction. That emotional reaction then causes us to think “This action seems wrong”, when in reality the action is right.

I’m not so impressed with that argument either. It’s a *possible* explanation, but it seems like the simpler, more straightforward explanation for why the actions in the above cases seem wrong is that the actions in those cases are wrong.

A related argument is that intuitions about hypothetical cases must be unreliable since different people often have conflicting intuitions about a

given case. This point is fair, as far as it goes – the cases we’ve considered are all at least somewhat difficult cases, about which smart people may react differently. Therefore, we should not have very high confidence in our intuitions about those cases. Contrast the following alternative scenario:

Easy Trolley Problem: As in Trolley, except that there is no one on the right-hand track. If you switch the trolley to the right, it will run into a pile of sand which will safely stop it. What should you do?

The answer to the Easy Trolley Problem is completely obvious and uncontroversial. The other cases we’ve been talking about are not like that – you can see why people would give different answers in the other cases, in a way that you can’t see any different answers in Easy Trolley. So maybe we shouldn’t trust our intuitions about the hard cases.

That being said, it’s not clear that this point cuts in favor of utilitarianism. Utilitarians embrace the common intuition about Trolley (the original version, not just the Easy version), but *reject* the common intuitions about the other cases we’ve discussed. If intuitions about hypothetical cases are unreliable, you could just as well say that the common intuition about Trolley is unreliable. Why would we pick that one to rely on?

The utilitarian might say: “We shouldn’t rely on *any* of these intuitions, not even the intuition about Trolley. Rather, we should rely on general ethical *theories*, and deduce the correct action in particular cases from our theories.”

The problem: Those theories are just going to be based on more intuitions (see ch. 13). There’s no reason to think intuitions about abstract philosophical theories are more reliable than intuitions about hypothetical concrete scenarios. If anything, the level of disagreement about abstract theories is *greater* than the typical level of disagreement about hypothetical scenarios. That’s why the history of philosophy is filled with intractable disagreements among philosophers – most of these are abstract, theoretical-level disagreements, not disagreements about concrete cases.

14.3.2. For Consequentialism

Be that as it may, there are some important intuitions that support consequentialism. Put it like this: Let’s say you’ve got a choice between two worlds, world A and world B. You know that A is *better* than B. Now, which one are you going to choose? On its face, it just seems obvious that you should choose the better alternative over the worse alternative, right? Well, whenever you make *any* decision, that decision can be viewed as

choosing among a set of possible worlds – the way the world will be if you choose A, the way it will be if you choose B, etc. Shouldn't you obviously choose the best? That's what the consequentialist is saying.

Here's another line of thinking. Imagine one of these hypothetical scenarios that's about sacrificing one to save five. But this time, take yourself out of the decision-making role and imagine that you're just an observer. You see someone deciding whether to push the fat man off the footbridge, but you can't do anything yourself. You know that if the agent pushes the fat man, the fat man will die and the five others will live. If the agent does not push, the fat man will live but five others will die.

What would you *hope* happens? It seems that a rational, benevolent person would hope that the fat man dies, rather than that *five* other innocent people all die. Now, you might say that in the case where the fat man dies, there would have been a wrongful act by the agent that you're observing – in fact, a murder – whereas in the alternative where five people die, there will be no *murders*. And perhaps a murder is worse than an accidental death. But no way is it *more than five times* worse. So if you know that either one murder will occur or five accidental deaths will occur, you should still hope for the one murder.^[91]

All this is *logically consistent* with saying that it would be wrong for the agent in Footbridge to push the fat man. But there is a tension here. If you, the outside observer, should hope that the agent pushes the fat man, then it seems that *the agent* should also hope that he himself pushes the fat man. But if he should *hope* that he does that, then it also seems plausible that he should just *do* the thing that he hopes he's going to do. This illustrates the weirdness of non-consequentialist ethics.

There are other problems for non-consequentialist theories that we'll discuss later (§§15.2, 15.4). As is often the case for philosophical theories, the best argument for utilitarianism lies in the problems with the other views.

14.4. Hedonism & Preferentism

14.4.1. For Hedonism or Preferentism

Utilitarians say that the good (that is, the only intrinsic value) is enjoyment, or the satisfaction of one's desires, or something like that.

What is an "intrinsic value"? **Intrinsic** values, or intrinsic goods, are contrasted with **instrumental** values or goods. Something is instrumentally good when it is good *as a means* to something else that's good. For

instance, money is (only) instrumentally valuable: It's good because you can use it to buy *other* things that are good, such as chocolate, computer games, and Mike Huemer's books. On the other hand, something is intrinsically good when it is good *as an end in itself*. For instance, happiness is good for its own sake, not merely because it helps you gain something else, so we call happiness an intrinsic value.

Why would someone think that utility is the only intrinsic good? Basically, because if you think about other things that seem good, they all seem to be means to utility. For instance, life seems to be good. You might think it is *intrinsically* good. But one could claim, not implausibly, that life is good only as means to utility: If you're alive, then you can have some pleasure, whereas if you're dead, you can't (it's also a lot easier to have your desires satisfied if you're alive). So if you value pleasure, you should also value life, provided that life is more pleasant than painful. On the other hand, if you imagine a life that is devoid of pleasure, or in which none of your desires are satisfied, that doesn't really seem good. If you knew the rest of your life was to be more painful than pleasurable, and if you also couldn't do any good for anyone else, then you might rationally decide to commit suicide.

Here's another value: knowledge. One could argue that that, too, is valuable only as a means. Some topics are interesting to us (like philosophy!), and therefore knowing about them gives us intellectual pleasure. Also, when you have a lot of knowledge, you tend to be better at achieving your other goals. You have to have knowledge to get a good job, design a bridge that won't collapse, take over the world, or whatever it is that you want to do. This doesn't prove that knowledge isn't *also* intrinsically valuable (it could be good in itself *and* also good for obtaining other things), but it makes it more reasonable to deny that knowledge has intrinsic value, since we can explain why people value knowledge without ascribing it intrinsic value. Now, if you imagine some knowledge that lacks those benefits – let's say that you don't find the subject at all interesting, and the knowledge also will not help at all for attaining any other goal – then it's really not clear that the knowledge is still good. E.g., you could learn the numbers in the 1970 Cleveland area telephone book. That would be a bunch of knowledge, but it doesn't seem valuable.

Here's another value: friendship. Again, one could argue that this is only instrumentally valuable. People enjoy spending time with friends, so

friendship is a means to enjoyment. Friends also help each other when one of them is in need. But imagine you had a friendship that lacked these benefits: You don't at all enjoy knowing the other person, and you also never receive help from them (perhaps because you never need it). Would you still value that friendship? Probably not.

And so it goes. We can't address every thing that people value, but in general, for each thing that seems good, you can pretty well explain how it is generally a means to enjoyment or desire-satisfaction. You can also imagine cases in which the thing wouldn't be a means to enjoyment or desire-satisfaction, and in those cases it is usually much less clear that it's valuable.

14.4.2. Against Hedonism & Preferentism

What are the reasons for *not* regarding utility as the sole intrinsic good? There are many other things that, to many people, seem to matter. Here are some examples:

Cookie: Ted is an evil serial killer who has tortured and murdered many people. He is now spending his life in prison. By contrast, Theresa is a saintly woman who has spent her life helping others. Now, it happens that you have a cookie that you can give to either Ted or Theresa (those are your only choices). Ted likes cookies *slightly* more than Theresa does, so Ted would get slightly more pleasure out of it. Assume there are no other relevant consequences of the action. Which choice would be better?

Most people intuit that it would be better if Theresa gets the cookie, because she deserves it, even though this would produce less total pleasure and desire-satisfaction.

Equality: Alice and Bob are equally deserving people, who are presently equally well off. You have some benefits to distribute. You can either give 100 units of benefit to Alice and 0 to Bob, or give 45 units to both Alice and Bob. Which choice is better?

Most people intuit that it would be better if both receive 45 units of benefit, because this is more fair, even though there would then be a smaller total quantity of benefit (90 versus 100).^[92]

Experience Machine: Scientists have developed a device called "the Experience Machine", which is capable of producing any desired experiences by direct brain stimulation. Unfortunately, once you are hooked up, you can't be detached. You are given the option of being

attached to the experience machine and having pleasurable experiences for the rest of your life. You can get whatever type of enjoyable experiences you want, and you may also opt to have your memories of life before you plugged in erased. Your body will then lie inert in a bed with wires coming out of your head, but in your mind, it will seem like you're experiencing whatever you most want to experience (say, being a great movie star, or the ruler of the universe, or just having the pleasure center of your brain continuously stimulated). Should you plug into the machine?^[93]

Most people reject the experience machine, even though plugging in would obviously result in far more total pleasure over the rest of their lives – thus showing that we value things over and above pleasure, and indeed, over and above any experience. If hedonism were true, life in the experience machine would be the best possible life; yet intuitively, it does not seem to be a very good life at all, let alone the best.

A hedonist might try biting the bullet, insisting that you should give the cookie to Ted, give the 100 units of benefit to Alice, and plug into the experience machine. But it's not clear what the justification for this would be. It's certainly intuitive that pleasure is good. But it's also intuitive that you should give the cookie to Theresa, give the 45 units of benefit to Alice and Bob, and reject the experience machine. If we give any credit to ethical intuitions, there's a strong case that some things other than pleasure matter. If, on the other hand, we *don't* give any credit to intuitions, then there's no reason for thinking pleasure is good in the first place (nor for holding any other ethical views).

Preference utilitarians are in a slightly better position: They could justify rejecting the experience machine by saying that the machine would not actually satisfy our desires. This is because we have a desire to live in contact with reality, or something like that. Granted, the person in the experience machine might *feel* as if they were living in contact with reality, and they might (if they have their memories from before they plugged in erased) *believe* that they were living in contact with reality, but they would not *in fact* be living in contact with reality. Hence, their desires would not in fact be satisfied, though they'd be tricked into thinking that their desires were satisfied.

Still, the preferentist would have to bite the bullet on Cookie and Equality, just like the hedonist.

14.5. Impartialism

14.5.1. *Partial vs. Impartial Ethical Theories*

Compare the following two views:

Utilitarianism: The right action is the one that produces the greatest quantity of welfare for all beings affected.

Ethical Egoism: The right action is the one that produces the greatest quantity of welfare for oneself (i.e., for the person who is acting).

(The latter view has been held by a few thinkers in the history of ethics, including Epicurus, Thomas Hobbes, and Ayn Rand. It is very much out of favor, though, especially among nice people.) These two views have in common that they are both *consequentialist*: They say that you should do whatever produces the best consequences (in their interpretation of what is best). They also agree that welfare is the only intrinsic good. But they're diametrically opposed on another dimension: The egoist has the maximally *partial* view (the egoist privileges himself as much as possible), while the utilitarian has the maximally *impartial* view (the utilitarian values all beings equally, privileging no one).

Most people are somewhere in between these two extremes – hardly anyone (except perhaps psychopaths) acts like a pure egoist, and probably no one acts like a pure utilitarian. (Many philosophers endorse utilitarianism intellectually, but even they do not actually *act* in accordance with it.) We care a lot more about ourselves than about others; also, we care more about our own families than other families, our own country than other countries, and our own species than other species.

So there's an interesting question here: To what degree are we morally required to be impartial? Or: To what degree may we favor ourselves and those close to us, over other beings whom we are not close to? For instance, say you have \$1000 in your bank account. If you spend it on yourself, you'll gain 5 units of utility from it. If you donate it to a poverty-relief charity, some stranger in the developing world will gain 500 units of utility from it. Are you obligated to donate the money, or can you spend it on yourself?

Or imagine a variant of the Trolley Problem in which the one person on the right-hand track is your own child. Are you still obliged to turn the trolley, or may you let the five strangers die to avoid killing your child?

14.5.2. *For Partiality*

The argument for partiality is basically a direct appeal to intuition: It just seems, to most people, that it's okay to privilege yourself to some degree over others, at least in many circumstances. If you only have enough food to feed one person, you can use it to feed yourself; you don't have to give it to someone else (not even if that other person will get greater utility). In some cases, it even seems *obligatory* to privilege those close to you over strangers – e.g., a parent should feed her own children before feeding strangers. A parent might even be obligated, say, to buy decent clothes for her own children before buying food (a more important good) for strangers.

Btw, this isn't an argument for egoism – you don't have to go to the most extreme partiality possible. Common sense morality allows *some* degree of partiality to oneself and one's family and friends, but not the most extreme degree. E.g., of course you can't steal food from the poor in order to sell it so that you can buy crack and hookers for yourself.

It's worth taking a moment to appreciate how extreme the demands of utilitarianism really are. If you have a reasonably comfortable life, the utilitarian would say that you're obligated to give away most of your money. Not so much that you would starve, of course (because if you literally starve, that'll prevent you from giving away any more!). But you should give up any non-necessary goods that you're buying, so you can donate the money to help people whose basic needs are not met. There are always plenty of such people. To a first approximation, you have to give until there is no one who needs your money more than you do. (By the way, if you can get away with it, you should also steal other people's money and give it to charity! But that's another issue.)

Furthermore, utilitarians do not recognize any morally significant difference between *harming* someone and *allowing* a harm to befall someone, nor between harming and failing to benefit. So they think that killing someone is morally equivalent to failing to save someone's life when you have the chance. Thus, on their view, those of us who fail to save as many lives as we could are morally comparable to murderers. By the way, that's all of us – no one, not even those who believe in utilitarianism, actually saves as many lives as they can. Utilitarians tend to donate *more* to charity than other people do, which is great, but still not nearly as much as they should according to their view.

By the way, if you talk to them about it, most utilitarians will admit to not giving as much as they should. They will generally explain it by

confessing that they are bad people (like everyone else, though not quite *as* bad as most).

After reflecting on this, many people think that this just seems like an unreasonably demanding morality. If you think that literally every human being in the world, even those who are generally held up as the best and most admirable among us, is *morally horrible* – something comparable to mass murderers – then it seems like maybe your standards of judgment are off.

14.5.3. For Impartiality

Why would someone hold the extreme impartialist view? Basically, because there doesn't seem to be any relevant difference between you and other people that would explain why *your* welfare is more important or more valuable than the welfare of other people. What's special about you?

You might say, "My welfare is more important *to me*." But this just seems to be saying that you personally *care* more about yourself than about others, not that your welfare is genuinely *more valuable* than other people's welfare. The fact that you happen to care more about *x* than about *y* does not show that *x* is actually better than *y*, or that you have any *reason* to care more about *x* than about *y*. So the remark "my welfare is more important to me" doesn't address the question.

Ethical egoists have an answer to this. They would say that value is **agent-relative**. That is, there is *no such thing* as something's being good *in general*, or from the viewpoint of the universe. Things can only be good or bad *for* particular people. And it doesn't make sense to weigh different people's goods against each other, or add together different people's goods. The reason you should exclusively pursue your own happiness is that your own happiness is the only thing that is good *relative to you*, and a rational agent maximizes the good relative to that agent. [\[94\]](#)

Is all that true? I don't think so. I think the egoist's argument would prove too much. It lets us escape from the extreme demands of utilitarianism, which some would consider good, but it has deeply implausible implications about other cases. Think about this more extreme version of the Trolley Problem:

Extreme Trolley: There is a runaway trolley heading for New York City. The trolley is carrying a nuclear bomb set to detonate inside the city, killing 20 million people. You can turn the trolley away from New York and toward a lonely patch of land containing a single cabin. The owner

of the cabin is not there at the moment; he only occasionally uses the cabin during vacations. If the trolley goes that way, no one will be hurt, but the bomb will destroy the cabin and render the land unusable, which will be a minor inconvenience for the owner. (Also, the owner of that land is an asshole who doesn't care about New York, so he won't be at all harmed by the destruction of New York.) What should you do?

In this case, again, assume there are no other morally relevant factors that aren't obvious from the statement of the scenario. Your own interests are not going to be affected one way or the other (you don't know anyone in New York, no one is going to get mad at you over your decision, etc.). Notice first that on the egoist view, you have *no reason* to divert the trolley, since your own interests are not at stake. You might just flip a coin, or perhaps decide not to divert the trolley since doing so would require slightly more effort than sitting and doing nothing.

More importantly, if you think that only agent-relative value exists (whether or not you're an egoist^[95]), then you must think there is no reason to divert the trolley in Extreme Trolley. Granted, diverting would serve the interests of the 20 million people who live in New York. But it would harm the one person who owns the lone cabin. On the agent-relative theory of value, there is no way of weighing one person's interests against others – there is no such thing as what is *overall* better or worse. There is only what is better *for* a particular person. That is precisely the claim that would get us out of having to donate most of our money to charity. If that works (if it really shows that we don't have to donate to charity), then it also works to show that there's no reason to divert in Extreme Trolley.

That conclusion is absurd. It's obviously *better* if you divert the trolley. So there must be such a thing as an outcome being overall better (not merely better for some particular agent) – in other words, there is **agent-neutral value**, as we say.

If there is agent-neutral value (and we're consequentialists), it's hard to see how we're going to avoid the argument for giving most of our money to charity, since doing so increases agent-neutral value.

Interlude: Extreme Libertarianism

In political philosophy, libertarians generally believe that it's wrong to violate one individual's rights (including property rights) in order to benefit others. Some libertarians take a particularly extreme view, that one may *never* do this, no matter how large the benefits. Some argue for this by

saying that value is agent-relative and/or that it doesn't make sense to compare different people's benefits, nor does it make sense to add different people's benefits together. My Extreme Trolley case is a counterexample to this extreme view.

Of course, there are more moderate forms of libertarianism that are still okay.^[96]

14.6. Rule Utilitarianism

Rule utilitarianism is a variant on utilitarianism that is supposed to avoid some of the implausible implications of consequentialism (see §14.3). Traditional utilitarianism (as described above) is known as "act utilitarianism". This is what the two views say:

Act Utilitarianism: You should always perform that act, out of all acts available to you, that produces the most utility.

Rule Utilitarianism: You should always act in accordance with the set of general rules that would produce the most utility if everyone followed those rules.

Rule utilitarianism is supposed to avoid things like Organ Harvesting, Framing, etc. The rule utilitarian says that it would be best if people follow a general rule of not killing healthy patients and not framing innocent people. Sure, those things might be good in a small number of circumstances, but it would be bad if they were allowed in general.

The biggest problem for rule utilitarianism is that of specifying what count as legitimate "rules" for purposes of applying rule utilitarianism, i.e., exactly what rules are we supposed to be comparing? If just *any* general imperative counts as a "rule", then rule utilitarianism will be equivalent to act utilitarianism. The easiest way to see this point: Does act utilitarianism itself ("Always perform the act that maximizes utility") count as a rule? If so, surely that rule produces the most utility if everyone follows it. But then, rule utilitarianism just collapses into act utilitarianism.

Suppose we stipulate that act utilitarianism doesn't count as a legitimate rule. Here's a more general issue: Are the "rules" that we consider allowed to contain exception clauses? For instance, could we consider a rule such as, "Don't kill healthy patients, *except* when you have five other patients who need organ transplants, and the healthy patient's organs are compatible with the other five, and you're sure you won't get caught, etc."? It seems that that sort of rule would produce greater utility than the simpler rule, "Never kill healthy patients." So the rule utilitarian should endorse the rule

with the complicated exception clause – but then we’re back to the same bad result about Organ Harvesting that we were trying to avoid. And you can see that we’re going to get similar results for the other cases (Framing, Promise, Electrical Accident, and any other counter-examples to act utilitarianism). Rule utilitarianism hasn’t gotten us anywhere.

We could avoid this by stipulating that rules may *not* contain exceptions. I.e., the rule utilitarian view could be: “You should always act in accordance with the set of *exceptionless* rules that would have the best consequences if everyone followed them.” This would exclude making an exception to the “don’t kill patients” rule for cases where killing patients maximizes utility. However, this view would have other highly counter-intuitive consequences, because common sense morality recognizes *some* exceptions to general rules. For instance, there is a general rule that you should not kill other people, but there is an exception for cases of self-defense. The version of rule utilitarianism that we’re presently considering would apparently *reject* that exception, so we’d have to say it’s wrong to kill people even in self-defense.

Maybe you think that’s not so bad; after all, some people (known as “pacifists”) think that killing is wrong even in self-defense. But we’d wind up with lots of other crazy absolute rules, such as “never lie”, “never steal”, “never break a promise”. So you couldn’t tell a lie even to save someone’s life, etc.

The rule utilitarian wouldn’t like that. They’d want to allow *some* exceptions but not others. But it’s just very unclear on what principled grounds we could allow a rule like “Don’t kill except in self-defense” but not allow a rule like “Don’t kill healthy patients except when doing so saves a larger number of other patients.” These sorts of problems probably explain why most utilitarians are act utilitarians.

14.7. Conclusion

There are a fair number of utilitarians, and that reflects a certain obvious appeal to the theory. When you just think about the theory by itself, without considering concrete cases, the theory makes sense. Nevertheless, most philosophers reject utilitarianism because of the sort of counterexamples discussed above – e.g., utilitarianism implies that you should harvest organs from a healthy patient, plug everybody into the experience machine, and save two strangers rather than your own child.

Most utilitarians either haven't attended sufficiently to the counterexamples (some people aren't even aware of all the examples), or they dismiss the counter-examples for not-very-good reasons, such as "I don't trust intuitions." In general, ethical beliefs rest on intuitions. You can choose to prefer *some* intuitions over *others*, and maybe you can disguise your intuitions or refuse to call them "intuitions", but one way or another, your ethics is going to be based on intuitions. Now, if you accept intuition as a source of justified ethical beliefs, then it seems that you should reject utilitarianism because it conflicts with too many strong, widely-shared intuitions. On the other hand, if you *don't* accept intuitions as a source of justification, then you still shouldn't accept utilitarianism, because you'd have no reason to believe that pleasure is better than pain, or that people should choose better outcomes rather than worse outcomes. If you reject intuitions, then you have no reason to prefer utilitarianism over anti-utilitarianism, the view that we should always maximize suffering.

That being said, utilitarianism is not a crazy view (*pace* some of its opponents). I grow more sympathetic to it as time passes. As one reflects more about ethics, one comes to appreciate the very serious intellectual problems that other ethical views face that utilitarianism completely avoids. We'll get into some of those problems in the next chapter.

15. Ethical Theory, 2: Deontology

15.1. Absolute Deontology

15.1.1. Terminology

Deontological ethics – “**deontology**” for short – is defined as the denial of consequentialism.^[97] Deontologists, in other words, think that the right course of action is *not* always to maximize the good. They generally think this because of the sort of examples discussed earlier (§14.3.1) – you shouldn’t kill a healthy patient to distribute his organs to five others, you shouldn’t frame an innocent person to prevent riots, etc.

There are stronger and weaker forms of deontology. **Absolute deontology**, or **absolutism**, holds that there are certain types of action that are always wrong, regardless of how much good they might produce or how much harm they might avert.^[98] (But note that there are other uses of the term “absolutism” in other contexts.) A popular example would be the view that it is always wrong to intentionally kill an innocent person, no matter what the benefits. Even if you could save the world from certain destruction by killing one innocent person, on this view, you shouldn’t do it.

One could also have absolutist views about other moral prohibitions, such as the prohibition on stealing, lying, or breaking promises; however, absolutism about those things is less common.

Moderate deontology (my term) is defined to be the alternative to both absolutism and consequentialism. Moderate deontologists think that you shouldn’t always maximize the good; however, they also deny that there is any type of action that is always wrong regardless of the consequences. To return to the example of intentionally killing the innocent: A moderate deontologist would generally say you should *not* kill an innocent person to save just two other people, but he *would* accept killing an innocent person to save the entire world. (Where one draws the line will vary from one thinker to another.)

15.1.2. The Categorical Imperative, 1: Universalizability

The most famous absolutist is Immanuel Kant, whose ethical views are studied in pretty much every ethics course (in addition to utilitarianism, which is a natural contrast). Kant advanced a principle that he called **the**

Categorical Imperative, which is supposed to be the fundamental principle of morality, from which all other principles of right and wrong follow:

Categorical Imperative, 1st version: Always act in such a way that you could will that the maxim of your action should be a universal law.

Interlude: Immanuel Kant

Immanuel Kant was a very interesting German philosopher of the 1700s (1724–1804). He was born in Königsberg, East Prussia, and never left the town once in his entire life. He was so anal that people could set their watches by the time that Kant took his daily walk. His greatest work was the *Critique of Pure Reason*, which is among the most abstruse works in the history of philosophy that is still meaningful. The difficulty of following Kant's words is not entirely due to the abstractness of the subject matter and the profundity of his thoughts; it is also due to his incredibly awful writing. Anyway, in that work, he tried to explain how it is possible to have substantive knowledge that is not based on experience.

In ethics, his most famous work is the *Foundations of the Metaphysics of Morals* (a.k.a. *Groundwork of the Metaphysics of Morals*), which is often used to torture students in philosophy courses. It, too, is so incredibly hard to follow that it often drives students to tears. It is much better to read some contemporary exposition of Kant's ideas, rather than to attempt to read Kant himself.^[99]

Here's the origin of the "categorical imperative" terminology: First, an "imperative" is just a sentence that tells someone what to do. Like "Tie your shoes!" or "Don't murder!" Second, in logic we distinguish *conditional* (or hypothetical) sentences from *categorical* ones. A conditional statement or imperative is one that has an if-then form, for example, "If you want a delicious smoothie, go to the Watercourse." A categorical statement or imperative is one that is *not* conditional. In Kant's view, morality gives us *categorical* imperatives. It's not that we must behave morally *if* we want something else to happen; we just have to behave morally, period.

What is "the maxim of your action"? Basically, it's a rule that explains what you're doing and why. An action is wrong if you couldn't universalize the maxim, in that there would be some sort of contradiction (or something like that) in willing that everyone should follow it. For example, say you've borrowed \$50 from your roommate with a promise to repay it. Now you're thinking of breaking that promise out of pure selfishness – you just want the money for yourself. Could your maxim be universalized? No – you couldn't

coherently will that everyone should break promises whenever it's in their interests to do so, because then the whole institution of giving and accepting promises would collapse, in which case no one would loan money anymore, and it would no longer be possible to profit in the manner you're hoping to do. So there's something close to a contradiction involved in willing that everyone act like you. So that's supposed to show that it's wrong to break the promise. Kant actually concludes that it is *always* wrong to break a promise, no matter what the consequences. This makes him a particularly extreme deontological absolutist.

By the way, many people, on hearing about this, remark that Kant's Categorical Imperative sounds similar to the Golden Rule. The Golden Rule says, "Do unto others as you would have done to you." (More colloquially: Treat people the way you would want to be treated.) However, Kant's view is distinct from this. He does not say that it's wrong to break the promise because you *would not want* other people to break promises *to you*. He says it's wrong to break the promise because there is something *inconsistent*, or self-defeating, or something like that, involved in the desire for *everyone* to break promises whenever it serves their interests. It's not that you wouldn't like the consequences of everyone following your maxim; it's that everyone *cannot* follow the maxim since it would be self-undermining.

You can make a similar argument about lying, so Kant also thinks no one should ever lie, regardless of the consequences. For example, suppose that Jack is in your attic, hiding from his homicidal wife Jill, who plans to murder him. Jill shows up at your door and asks you where Jack is. On Kant's view, it's permissible to *refuse to answer*, but you can't *lie* to Jill – you cannot, e.g., tell Jill that Jack has taken a trip to New Jersey, even if doing so would save Jack's life. Many people regard this as an insane ethical view. Of course you should lie to the murderer! Hardly anyone agrees with Kant about this. But there are more people who agree with Kant about other issues, such as the absolute prohibition on murder.

Here's another example, which Kant actually discusses: Say you're sailing a cargo ship. Your ship has cargo that belongs to someone else, which you promised to deliver to its destination. The ship runs into a storm, and it is in danger of sinking unless some weight is thrown overboard. According to Kant, it would be wrong to throw any of the cargo overboard, since that would involve breaking your promise and intentionally destroying someone else's property. So you just have to take your chances.

Maybe the ship will sink, destroying the cargo and killing everyone aboard, but at least you would not have intentionally destroyed it.

As you've probably noticed, that's also crazy. I think all this is much crazier than utilitarianism.

Other duties Kant believed in: He thought it was always wrong to commit suicide, that everyone is obligated to donate to charity, that we're also obligated to develop our talents and improve ourselves, and that masturbation ("self-abuse", as he called it) is always wrong.

Interlude: Perfect & Imperfect Duties

Some of the examples above are "perfect duties" – roughly, duties that you have to be fulfilling at all times. E.g., the prohibition on lying means that at all times, you must refrain from lying (you can't just refrain *sometimes*). Imperfect duties, by contrast, only have to be fulfilled sometimes. For example, Kant says that you have a duty to be charitable. But of course, you don't have to be giving to charity constantly; you just have to do it sometimes. So the duty of charity is "imperfect".

15.1.3. The Categorical Imperative, 2: The End-in-Itself

Immanuel Kant had a second moral principle, which he also called "the Categorical Imperative" and which is also supposed to be the fundamental principle of morality. He claimed that this second principle was merely another formulation of the same principle we stated above, but it's pretty hard to see how that's supposed to be true. Anyway, here it is:

Categorical Imperative, 2nd version: Act so that you treat humanity, whether in your own person or in that of another, always as an end, never merely as a means. [\[100\]](#)

For example, suppose (again) that you're considering breaking a promise to return some money that you borrowed, because you want more money for yourself. If you break the promise, you would be treating the lender as a mere means. Roughly, this means you'd be making the other person a part of your plan of action, without regard to *their* choices or goals. There is something very intuitive about this: Other people are not mere tools at your disposal.

Notice that the principle isn't that you can't treat a person as a means *at all*; it's that you can't treat a person *merely* as a means. So it's okay to make another person part of your plan of action, as long as you respect the other person's autonomy in the process (and thus treat them *also* an end). The way to do this is to obtain their consent. So if the lender consents to your

not returning the money, you're okay! But if he doesn't consent, then you'd be violating the categorical imperative.

Now, here's something great about the Categorical Imperative. It offers an explanation for the difference between the Trolley Problem and the Footbridge case (§14.1). In the Footbridge case, it is wrong to push the fat man off the bridge, since doing so would treat the fat man as a mere means to saving the other five. By contrast, in the original Trolley case, diverting the trolley does *not* treat the one person on the right-hand track as a means. The person on the right track isn't a *means* to saving the five at all. You can see that because if the one person were not present, you would still divert the trolley, thereby saving the five on the left track in exactly the same way. The fact that your plan works just as well if the one person is not present shows that he is not a means to achieving the goal.

By contrast, of course, if the fat man on the bridge were not present, you could not carry out your plan to save the five in Footbridge. That's because the fat man's body is the actual means of stopping the trolley.


This kind of reasoning also works pretty well in the Organ Harvesting, Framing, and Promise cases from §14.3.1. It doesn't work well on Electrical Accident, though. (The guy who's caught in the electrical equipment isn't a means to entertaining the others; if he weren't present, everyone would get the same entertainment in the same way.)

15.1.3. *The Doctrine of Double Effect*

This brings us to a popular principle in ethics, known as the “**Doctrine of Double Effect**” (“DDE” to the in crowd). The DDE basically says that it's worse (harder to justify) to *intentionally* harm someone than it is to harm someone as a foreseen *side effect* of one's action. In fact, people who subscribe to the DDE often think that it's *absolutely prohibited* to intentionally harm the innocent (at least, in some very serious way), whereas it can sometimes be okay to harm the innocent knowingly but not intentionally.

Wait, what's the difference between harming someone *intentionally* and doing it *knowingly*? Well, when you *intentionally* harm someone, the harm to the other person is either the *end* that you're aiming at or a *means* to that end. (Notice how this connects up with Kant's notion of the obligation to treat persons as ends.) That's the thing that's super-bad. By contrast, when you harm someone as a mere side effect, the harm isn't aimed at, neither as

a means nor as an end, even though you might know it's going to happen. Here's a diagram to illustrate the causal relations:



Example: The DDE is often used in military ethics to distinguish between acceptable collateral damage and war crimes. If you *deliberately target* civilians, that's a war crime. It's widely regarded as immoral, even if the war is otherwise just. (Though it has happened a lot.) On the other hand, if you are aiming at a military target, you aren't required to hold off just because there happen to be some civilians nearby who are going to be hurt in the process. Your attack still has to pass a consequentialist test – the benefits must outweigh the consequences – but it's not subject to the absolute deontological prohibition on intentionally harming the innocent.

15.1.4. Rights

The notion of *rights* is a particularly popular deontological concept, especially in the U.S. There are two kinds of (alleged) rights: positive and negative. A **positive right** is a right to receive some benefit (as in “the right to an education”). A **negative right** is a right *not* to be harmed or interfered with (as in “the right to free speech”, which is a right not to be prevented from speaking). The right to property, by the way, is generally understood as negative – it's the right *not* to have your property taken or damaged.

What is it to have a “**right**” to something? Basically, we say that you have a right to something when other people are obligated to give you that thing (in the case of a positive right) or to not interfere with your having that thing (in the case of a negative right). This obligation is understood deontologically: People have to respect your rights even if slightly better consequences would follow in a particular case from violating your right.

E.g., in the Organ Harvesting case, we would say that the healthy patient has a *right* to his organs (this is a negative right). Thus, we may not interfere with his possession of those organs even if more overall good would be produced by taking them from him.

Rights are generally understood as *agent-centered constraints*. This means that to properly respect rights, you do not simply try to minimize the total number of rights-violations in the world; rather, you are obligated to ensure that *you yourself* don't violate any rights. So if you could violate a right and thereby prevent someone else from violating two rights, you shouldn't do it.

Interlude: Claims & Permissions

What I've just described is what is sometimes called a "**claim right**". By contrast, people sometimes talk about *permission rights*, which *don't* impose obligations on others. You have a permission right to do A provided that you're not morally obligated to refrain. The permission right, unlike the claim right, doesn't rule out other people trying to stop you from doing it. Example: If you're in a boxing match, you have a permission right to punch the other person. But you don't have a *claim* right – i.e., the other boxer doesn't have to let you punch him. But again, I'm interested in talking about claim rights in this section.

Why should we believe in rights? One reason is that the idea of rights explains our intuitions in cases like Organ Harvesting, Framing, and Promise. Also, some people think that the notion of rights sort of follows from the Kantian idea that individual persons are ends in themselves, not tools for others to use.^[101] The way that you treat a person as an end in himself is by respecting certain constraints on how that person can be treated without his consent. Those constraints are just the principles of individual rights.

I note, however, that the notion of rights does not do so well with the Trolley Problem. Diverting the trolley toward the one person on the right-hand track would seem to be a rights-violation. Surely that person has a right to life, and surely you violate the right to life when you kill him. So it looks like rights theorists would probably *reject* turning the trolley.^[102] This is counter-intuitive, though not crazy – a fair number of people think turning the trolley seems wrong, since it seems like murder.

Some people take an absolutist view of rights, holding that it is never permissible to violate someone's rights, no matter the consequences. This is

known as believing in “**absolute rights**”. Other philosophers think rights may permissibly be violated in some extreme circumstances (but of course not merely because slightly better consequences are produced); this is known as endorsing “**prima facie rights**”.

15.2. Objections to Absolutism

15.2.1. *Extreme Consequences*

There are many possible forms of absolutism, depending on what one thinks the absolute prohibitions are. But let’s suppose there’s an absolute prohibition on intentionally killing innocent people. (That is a particularly popular proscription – if you accept any absolute prohibition, you probably accept that one.) Here’s an apparent counter-example to that:

Alien Threat: Powerful aliens have shown up and, for whatever reason, they are determined to destroy the Earth, killing all 8 billion people, unless you kill one innocent person. There is no other way to stop them from destroying Earth. Should you kill the innocent person?

In that case, surely you should kill the one person.

There are even more compelling cases if the absolutist posits more absolute prohibitions. E.g., some people think that all rights are absolute. In that case, consider:

Miracle Hair: Humanity is suffering from a deadly disease that will shortly wipe out everyone. Only one little girl is immune. If you pluck a single hair from her head, you can use it to synthesize a medicine that will cure everyone else. For whatever reason, the girl will not consent to give one of her hairs. There is no way to persuade her. Should you take a hair without consent?

Taking someone’s hair without her consent is a rights violation (albeit a small one!). So, if all rights are absolute, then we’d have to let humanity perish. That seems crazy.

15.2.2. *Portions of a Life*

Suppose the absolutist bites the bullet on the Alien Threat example – the absolutist says you have to let humanity perish rather than killing one innocent person. Now let’s turn up the pressure. I forgot to mention one detail about the case: The innocent person that you’d have to kill is already 94 years old and has terminal cancer. He has only one month to live in any case; nevertheless, he wants that month and does not consent to be killed. Now, are you still going to say that we can’t rob him of one month, in order to save the rest of humanity?

Notice that there's no *qualitative* difference between this and any other killing – killing a person is always nothing other than *shortening* their life. Cases merely differ in *how much* a life is shortened. So if you're an absolutist, you'd presumably have to say we can't kill the 94-year old who only has one month to live.

Okay, what if it was only one day? Or one hour? Or a *tenth of a second*? Could we really claim that it would be wrong to cause someone to die 0.1 seconds earlier, if doing so would save the rest of humanity?

15.2.3. Risks to Life

A similar problem arises when we consider probabilities. If you say “it's always wrong to intentionally kill an innocent person”, then what should you say about intentionally creating a *risk* of killing an innocent person?

Suppose the absolutist says that that is morally *permitted*, provided the consequences are sufficiently good. This would defeat the point of being an absolutist, since no action ever creates a *100% certainty* of death. In the real world, actions only ever create some *probability* of any given outcome (sometimes, of course, the probability is extremely high).

So suppose instead that the absolutist says that any nonzero risk of killing the innocent is morally prohibited. This sounds more like what an absolutist would say. In that case, *all* actions are prohibited, since every action carries *some* nonzero probability of killing an innocent person. So one would presumably have to remain completely motionless until one dies.

Finally, suppose the absolutist says that creating an $x\%$ risk (with sufficiently good consequences) is allowed for small values of x but absolutely prohibited when x is above some threshold. I.e., you can create small risks but not big risks. But then you could create a small risk on multiple occasions, resulting in a large risk overall. Each of the individual actions would be okay, but the collection would be absolutely prohibited. This seems paradoxical.

15.3. Moderate Deontology

Moderate deontologists (as I call them) hold a middle ground position between consequentialism and absolute deontology. They think that some kinds of actions normally should not be performed, even if they produce better consequences; however, in some extreme cases it is permissible to perform them, to prevent something vastly worse from happening. Thus, the moderate deontologist would reject the idea of intentionally killing one innocent person to save five others, as in the Organ Harvesting case.

However, the moderate deontologist would accept killing one innocent person to save *millions* of others (or more, as in the Alien Threat case).

You might wonder, “What is the threshold? How much larger must the benefits of an action be than the costs, to justify violating someone’s rights?” There is no generally accepted answer to that, and no accepted way of deciding other than making an intuitive judgment. Different people will have different intuitive judgments.

Why would someone believe this view? Because it accommodates common sense ethical intuitions: It seems wrong to sacrifice one person to save five others (in certain sorts of cases, like the Organ Harvesting and Footbridge cases), but it seems alright to do so to save a million others. So we could simply embrace both of those judgments; we don’t have to take the extreme view that violating an individual’s rights is *always* wrong no matter the consequences, nor the opposite extreme view that individuals have no rights at all.

Moderate deontology is also inevitably a **pluralistic** view. What I mean by this is that it posits more than one distinct moral principle. There must be some deontological principles that tell us what kinds of actions are generally wrong for non-consequentialist reasons, plus a consequentialist principle that tells us we have reason to produce better consequences, other things being equal. Consequentialist considerations will have to be weighed against deontological considerations in particular cases. (Contrast utilitarianism and Kantianism, each of which posits a single moral principle.)

What would the deontological principles be? Most moderate deontologists would list several principles, e.g., “Other things being equal, you should keep your promises, avoid lying, avoid hurting others, avoid stealing, etc.” The “other things being equal” clause applies to all those types of actions, because sometimes one of those imperatives conflicts with another. E.g., in some circumstances, you would have to lie to avoid hurting others. There would also be a general principle that, other things being equal, you should produce more utility rather than less.

When two or more of these values (producing utility, keeping promises, not stealing, etc.) conflict with each other, there is no algorithm for resolving the conflict. One must simply make an intuitive judgment as to which value is more important in the circumstances. Often, one will be uncertain of what is the right thing to do; that is how moral life is. [\[103\]](#)

15.4. Objections to Moderate Deontology

15.4.1. Arbitrary Cutoffs

One thing people don't like about moderate deontology is that it posits moral lines that have no principled rationales. Since (for example) it's wrong to kill an innocent person to save five lives, but okay to do it to save a million lives, there must be some particular number of lives saved at which it first becomes permissible. What is that number?

The problem is not just that the theorist can't tell us authoritatively what the number is. The deeper problem is that there doesn't seem to be any possible explanation of what could make one number the correct one. There is nothing special about any particular number between 5 and 1,000,000 that could explain why it was the cutoff. And if there is no factor that could *explain* the cutoff, then, plausibly, there can't *be* such a cutoff.

Notice how utilitarianism and absolutism avoid arbitrary cutoff points. The utilitarian says you can kill an innocent person to save 2 or more lives. There is nothing arbitrary about this; it's easy to say what's special about the number 2: It's the first number such that the benefits of the action would be greater than the costs (assuming there are no fractional lives). The absolutist, on the other hand, says you can't kill an innocent person to save any number of lives. That also, obviously, avoids any arbitrary cutoff point.

15.4.2. The Aggregation Problem

Here's another problem, which appears more serious. Moderate deontology prohibits some actions that increase utility overall, actions that harm one (or more) persons but produce a greater total benefit for others. It's possible to have a *collection* of such actions that, overall, benefits *everyone* involved. Intuitively, such a collection of actions should be acceptable.

Say there are two people, A and B, who each have money in a bank. You want to give money to A, but, for whatever reason, you are unable to do so *directly*. What you can do is to hack into the bank's computer and run a program which will steal \$1 from B's bank account, then give \$2 to A. (The extra dollar comes from your own account, and you're happy to give it, as A needs the money and you don't.) The moderate deontologist says this is impermissible: Stealing is a rights-violation, which cannot be justified merely by producing slightly greater benefits for others. So far, so good – that's not an implausible view on its face.

Now, it happens that you also would like to give some money to B, but again you can't do so directly. You can only run a computer program which will steal \$1 from A and then give \$2 to B. (This is a separate program from the one that would steal from B to give to A.) Just as before, the moderate deontologist will presumably declare this impermissible.

However, if you do *both* actions, A and B are both better off. Intuitively, there's nothing wrong with *that*. Money is fungible, so the net effect on A and B will be exactly as if you had simply transferred \$1 to each of them, which would surely have been permissible. The fact that you carry out the transfer in this odd way, taking \$1 away and giving \$2 back, seems morally irrelevant.

Action	Effect on A	Effect on B
First transfer	+\$2	-\$1
Second transfer	-\$1	+\$2
Both transfers	+\$1	+\$1

So moderate deontology implies that two wrongs can make a right: You can have two actions which are each morally wrong, yet it can be morally right to do both. This seems problematic.[\[104\]](#)

Deontologists could avoid this example by, say, denying that the individual money transfers in the example are morally prohibited. To be a deontologist at all, you must think there are *some* cases where the utility-maximizing action is morally wrong. But you don't have to think this in *all* cases. So maybe the money transfer just happens not to be one of the cases where a deontological prohibition applies. However, we would need some general, principled way of ruling out *any* case in which two utility-maximizing actions are prohibited yet their combination benefits everyone concerned. It's not clear that one can do this while being true to the usual intuitions behind moderate deontology.

15.5. Conclusion

All ethical theories are problematic in one way or another. Utilitarianism supports seemingly grossly immoral actions in certain hypothetical cases – killing a healthy patient to steal his organs, framing an innocent person for a crime to prevent riots, pushing a man off a bridge so his body stops a runaway trolley, breaking a deathbed promise in order to appropriate a child's inheritance, and letting an injured man continue to suffer so as to avoid interrupting other people's entertainment.

Absolute deontology supports other seemingly terrible decisions, such as letting the rest of the world perish to avoid plucking a single hair from one person's head, or to avoid shortening one person's life by 0.1 seconds, or to avoid imposing an extremely tiny risk of death on one person. If we reflect on how all actions entail risk, it looks like absolutism might entail the absurd conclusion that literally *all* actions are absolutely prohibited. This is much worse than utilitarianism. Don't be an absolutist.

Finally, moderate deontology requires drawing seemingly arbitrary lines, and it also seems to create the possibility of cases in which two or more actions are each wrong, and yet the combination of them is morally okay.

Overall, I judge the problems for moderate deontology to be the least bad.

16. Applied Ethics, 1: The Duty of Charity

16.1. The Shallow Pond Argument

Here's a hypothetical that comes from the contemporary philosopher Peter Singer^[105]:

Shallow Pond: You're walking to class one day, passing by a shallow, ornamental pond on campus, when you notice that a small child has fallen into the pond. He appears to be drowning. You could wade into the pond and save the child. However, doing so will get your clothes all wet, possibly ruining your nice new suit, and cause you to miss class. Are you obligated to save the child?

Few people have difficulty with this: Obviously, you have to pull the child out of the pond. The inconvenience to you is trivial in comparison with the life of another person.

That was just a hypothetical scenario. Now here is a non-hypothetical scenario. You live in a world where many people are suffering from malnutrition or dying of malaria, tuberculosis, and other preventable diseases, due to extreme poverty. You are much better off than those people and are frequently able to buy for yourself various goods that you don't need. You could give some of your money to charitable organizations that aid the global poor, thereby helping to save some lives. For example, you could donate to UNICEF, GiveWell, or the Against Malaria Foundation. However, this would require giving up some luxuries that you enjoy. All of this is *in fact true* for almost all readers of this book (unless you're a poor person in the developing world, in which case I don't know how you would even have gotten a hold of this book). Yet most people give nothing to such causes. What is the appropriate moral assessment of this?

This non-hypothetical situation is analogous to the shallow pond in obvious ways. In both cases, you are aware of someone who is in great need. You can do something to alleviate their need, at very low cost to yourself. If you're obligated to help the child in the shallow pond, then, it seems that you are equally obligated to help the global poor by donating to charity.

That is Peter Singer's conclusion. Notice that his conclusion is not merely "It's preferable to give to charity rather than not giving", or "Giving to charity is praiseworthy." Those are pretty obvious claims, which almost everyone already agrees with. His claim is that giving to charity is *morally obligatory*, not optional. If you're not giving anything, then you're acting like the asshole who walks past the drowning child because he doesn't want to get his clothes wet – which is truly horrible behavior. Also notice that Singer isn't saying that "society" should do something to help the poor (though that might also be true). He is saying that *you personally* (and me, and each other well-off individual) have an obligation to contribute.

Singer says the Shallow Pond example illustrates the first premise in the following argument:

- 1.If you can prevent something very bad from happening without sacrificing anything of comparable significance, then you are obligated to do so.
- 2.You can prevent some very bad things from happening without sacrificing anything of comparable significance, by donating to poverty relief efforts.
- 3.Therefore, you are obligated to donate to poverty relief.

How much are you obligated to give? Well, you have to keep giving until the point at which giving more would require sacrificing something of comparable significance (of course this is vague, but it's still of some use as a guideline). Obviously, you shouldn't give so much that you yourself would starve, or go without needed medical care, or something like that (that would be sacrificing something of comparable significance). You also don't want to do things that would prevent you from being able to give in the future. So if, say, your job requires you to dress reasonably presentably, then you don't want to give away so much that you're unable to do that and then you get fired. However, if you're spending money on lots of trivial goods, as most people are, you should stop that and give the money to charity instead – e.g., if you're eating out at restaurants, going to movie theatres, or buying extra shoes when you already have a perfectly good pair. No reasonable person would consider those luxuries to be of comparable significance to another person's life.

Note: Absolute vs. Relative Poverty

People in wealthy nations (including, say, college students) often think of themselves as "poor" despite having adequate food, clothing, shelter, and

medical care. This is using a *relative* notion of poverty: We call ourselves poor when we are *poorer than* the other people in our society. By contrast, many of those in the developing world (i.e., the world's poorest countries) are **absolutely poor**, meaning that they don't have enough money to meet their basic needs – so they have inadequate nutrition, or they don't have medical care, and they're in danger of dying because of this. Many of the relatively poor (in wealthy nations) are far richer than the world's absolute poor.

16.2. Objections in Defense of Non-Giving

I'm going to assume that we agree about saving the child from the shallow pond. But is that really analogous to donating to charity? Here are some arguments you might give for why we need not donate to poverty relief, even though we *would* have to save the drowning child.

Objection #1: "I'm not certain that my money will really help the poor. What if the charity organization just keeps the money for itself?"

Reply:

- (a) Let me add another detail to the Shallow Pond story. You were about to wade in and pull the child out, but then you noticed that he had stopped moving. You were thus uncertain of whether you would actually save him or whether he was already dead. Since you didn't want to take the risk of possibly getting your clothes wet for nothing, you decided to just keep on walking. Is this okay?
- (b) Obviously, don't donate money to some random web site that some guy just started; donate only to reputable organizations. For example, UNICEF is an extremely famous poverty relief organization started by the U.N. in 1946, with a presence in 192 countries. It's not going to turn out that UNICEF is some giant hoax that's been running for 70 years without anyone noticing. (If you think that, you should consider seeing a shrink, because that is schizophrenic-level paranoia.) There are also charity review organizations, such as GiveWell, which monitor the cost-effectiveness of charities. So if you are concerned about ensuring that your donation really does some good, you can go to <https://www.givewell.org> and see the charities that are most efficient.

Objection #2: “There are so many poor people in the developing world that it’s impossible to save all of them. My contribution would just be a drop in the bucket.”

Reply: Imagine that just as you were about to pull the child out of the Shallow Pond, someone came by and told you that as it turns out, there are actually thousands of ponds, pools, rivers, etc., around the world where children are drowning. Realizing that you can’t save all of them, you decide there’s no point in saving the one here, so you just continue walking to class and let that kid die. Is this okay?

Objection #3: “There are lots of other people who could help even more easily than me. The millionaires and billionaires should donate their money, instead of me!”

Reply: I have another detail to add about the Shallow Pond. You’re just about to wade in to save the child, when you notice that there are several other people standing around the pond doing nothing. Any of them could save the child. Some of them don’t even have any class to get to (judging from how they are just lounging about on the lawn), and a few of them are closer to the child than you are. Nevertheless, none of them is in fact doing anything. You try shouting out, “Hey, someone save that child!” but they just ignore you. You say to yourself, “Well, if none of *them* is saving the child, I refuse to do it either!” Then you just keep walking, leaving the child to die. Is this okay?

Objection #4: “Why bother? People in poor countries have such bad lives that it’s hardly worth preserving them.”

Reply: You’re just about to pull the child out of the pond, but then someone comes by and tells you, “Oh, I know that child. He has a pretty bad life. In fact, he’s just visiting from Bangladesh; he’s scheduled to go back there next week and return to his life in the slums.” Hearing this, you figure there’s no point getting your clothes all wet to save *him*. So you keep walking. Is this okay?

Objection #5: “Saving lives in the Third World is futile or counter-productive. Their problem is that they have too many people. Saving lives there will just cause the population to increase, which will cause more starvation and suffering in the future.”

Reply:

- (a) Again, just as you were thinking about saving the drowning child, someone comes by and tells you that the kid is going to be sent to Bangladesh next week, if he's still alive. You reflect that there are too many people in Bangladesh as it is, so you decide to just let the kid die instead. Does this sound cool?
- (b) Actually, no, there's no evidence that poverty relief efforts cause an increase in population. Quite the opposite (see §16.3).

Objection #6: "This is too demanding. If we accept Singer's argument, we won't just be giving a little bit occasionally. We'll be giving *almost everything* we have. After I've saved one starving child, there will be another one I could save. And another one. No matter how many I save, there will always be this argument that I could save one more, by giving a little more money – all the way to the point where I have only just enough to meet my own basic needs. There would be no particular point at which I sacrificed anything of comparable significance to another person's life. But the *cumulative* effect of *all* the giving would be to pretty much ruin my life. But it's just not reasonable to ask people to make that much of a sacrifice for others. The Shallow Pond story is different, because you only have one child to save."

Comments:

This objection is better than the previous ones. What's right about it: I think the Shallow Pond Argument really is that demanding. Also, it's plausible that this makes it too demanding.

We could try modifying the Shallow Pond story to make it more analogous to the world poverty situation. Let's say the pond, while still shallow, is extremely large. There are thousands of children drowning in that pond, and more fall in every minute. After you pull the first one out, there will be another one to pull out, and another, and so on (it's a different child each time, though). If you spent the rest of your life pulling children out, you'd never finish, because more keep falling in. Are you then obligated to spend every waking moment for the rest of your life – apart from the minimum amount of time needed to sustain your own life – pulling children out of that pond?

I'm going to say not. Now, do you have to pull *any* of the children out of the pond? Or may you, in this scenario, completely ignore the

pond and all its drowning children?

It would be very strange if it was morally okay to completely ignore the pond. If there is only one child in the pond, you have to pull that child out. Presumably, if there are two, then you have to pull both of them out. If there are three, you have to save all three. Presumably, it goes on that way for a while. It would be very strange if, at some point, your obligation suddenly drops to zero, merely because there were more people in need.

In other words, say that if there were exactly n children in the pond, you'd have to pull all of them out. Suppose, in fact, that n is the *largest* number such that you'd be obligated to try to save all of them. And suppose that you have just in fact come upon a shallow pond, where you see exactly n children drowning. You're about to wade in and start saving all those children. But then, just as you approach the water's edge, you notice one more child that you hadn't seen before (he was hidden behind another child's hat, you see). It would be very strange if you suddenly concluded that it's now perfectly fine to turn around and walk away, saving *none* of the children, because of that one added child.

Don't say that you now have to save all $n+1$ children. That can't be, because we *stipulated* that n is the largest number that you'd have to save. I think what this shows is: If there are any number of children less than or equal to n , then you have to save all of them (that's by stipulation); if you see *more* than n children in the pond, then you still have to save n of them. It's **supererogatory** (praiseworthy and beyond the call of duty) to save additional children after that.

I don't know the correct value of n . It's a matter for intuitive judgment.

It's presumably significantly more than 3, but less than a million.

How does all this apply to donating to charity? Well, Objection 6 does not show that we're not obligated to give to charity at all. Rather, we're obligated to give the largest amount such that, *if* that amount would completely solve the problem, we'd be obligated to give it.

In practice, that principle is hard to apply, since we don't have very clear intuitions about what that largest amount is. I assume, though, that it would be an amount that feels like a significant sacrifice

(hence the saying, “Give till it hurts”), but not an amount that ruins your life.

Note about Utilitarianism

Utilitarians wouldn’t agree with what I just said. They don’t buy these “too demanding” objections. The utilitarian view would be that you have to give until the point at which giving more would actually cause more harm than good (either because you yourself need that money more than any of the people you could give it to, or because giving more now would somehow prevent you from giving in the future). This would be an enormous amount. If you’re a normal person in a prosperous society, it would probably mean giving upwards of 90% of your income away. And it would ruin your life. But that, says the utilitarian, is a small price to pay for all the other lives you would save.

Oh, by the way, you’re almost certainly also going to be obligated to *change your career*, from whatever it currently is to some more lucrative one, so that you can donate more to charity. E.g., if you’re a philosophy professor, you probably could become a lawyer instead, in which case you’d be able to give a lot more to charity. It doesn’t matter if you would hate practicing law; that, again, is a small price to pay for saving more lives.

Peter Singer is well known as a utilitarian. However, when he talks about famine relief, he doesn’t advance the utilitarian view (which many people find to be so insanely demanding that it just puts them off from donating anything). He thinks that his argument (§16.1) should work on anyone with *any* reasonable ethical view, including reasonable deontologists. That is, it works to show that we have an obligation to give *some* significant amount to charity, though it doesn’t establish the extreme utilitarian view of our obligations.

16.3. Poverty and Population

Way back around the start of the 19th century, there was an economist named Thomas Malthus. He wrote about the problem of population. His theory was that any society’s population would naturally increase exponentially, until they used up their food supply and other resources. The death rate would then have to increase until the food supply was just enough to sustain the population.

This view is apparently so intuitive that some people today still believe it. Some think that world poverty is caused by “overpopulation”, and that

relief efforts lead to more population growth. This leads some to conclude that there is no point in contributing to poverty relief (objection 5, §16.2).

Fortunately, this theory is just factually false. It is true that there is a correlation between fertility (birth rates) and poverty – the countries with high fertility tend to also be poor. This is not, however, because fertility causes poverty. It's the reverse: Poverty causes people to have more children. When people's income goes up, they do not generally increase the number of children they have; they *decrease* it. In fact, the most prosperous nations in the world have birth rates *below* the death rate and may have a problem with a dwindling population. (But they can prevent population decline via immigration.)

How can this be? When people get more money, they can *afford* to have more children, so why doesn't their fertility increase? The answer is basically that children take up your time, and, in wealthy nations, people have other things they'd like to do with their time. For one thing, if you have lots of kids, that can interfere with your career; so the better your career prospects are, the greater the deterrent to having kids. You also just have more fun things to do in a wealthy country – sailing your boat, smoking marijuana, playing *Call of Duty XII*, etc. – which children can interfere with. *Most* people are going to still have kids, but some number of well-off people are going to decide it's not worth it.

One factor that is particularly anti-correlated with fertility is women's education. When education is more available, women postpone having children and have fewer children – partly because they're in school during some of their child-bearing years, and partly because after finishing school they have more desirable career options available. So, again, they have other things to do that compete with producing offspring.

In poor nations, on the other hand, people may have more children because they want the children to work during childhood, or they want the children to help support them when they are older. Another key factor is child mortality: In poorer nations, infants and children are more likely to die. Parents don't know how many of their children will die before reaching adulthood. Therefore, if they want to be reasonably assured of having some children survive to adulthood, they need to have *more* children than would be the case in a society with negligible child mortality.

Taking all this into account, if we can alleviate world poverty, this would actually reduce population growth. We'll get a smaller population,

living at a higher standard of living.

By the way, you can tell that population isn't the main cause of poverty, because there are plenty of high-population-density countries that are rich, and plenty of low-population-density countries that are poor. It's not like there's just some fixed amount of money sitting there, and we have to divide it among us, so we get smaller shares the more people there are. No, the amount of money varies with the population. In a prosperous nation, when we add new people, those people are generally productive, so the society's productivity grows roughly in proportion to the population.

If overpopulation isn't the cause of poverty, what is? Well, that's complicated, and everyone doesn't agree on it. One factor is political: Poor nations often have corrupt governments that rob the people and do a crappy job of protecting them. They may have crappy policies in place that interfere with market activity. These nations also may have general cultural features that depress economic activity, such as a tendency to distrust strangers, which interferes with economic transactions. Those problems are difficult or impossible for us to address directly.

However, the poor nations are generally growing economically, and at a faster rate than the wealthy nations, so they're on the way to catching up. We wealthy people can help this process along by providing aid for food, medical care, etc.

16.4. Effective Altruism

After Peter Singer gave his argument for the obligation of charity, some other people started a movement called the “**effective altruism**” movement. Their main idea is that we should not only give a significant amount to charity; we should also make a particularly strong effort to identify the *most effective* charities to give to, and focus our efforts on those charities. This is important because some charities are hundreds or thousands of times more effective than others.

Let's say you see a child drowning in a shallow pond. You could wade in and pull the child out. At the same time, you see another child who is feeling slightly chilly. You could go over and give her a jacket to make her feel warmer. You only have time to do one of these things. Which should you do?

Obviously, save the drowning child. If you instead decide to “save” the cold child, you're being an immoral jerk. That illustrates the intuition that givers are obligated to give to *effective* charities, ones that do a lot of good

with the money you give them, as opposed to inefficient charities that do only a small amount of good. For instance, a lot of people, after graduating and then making successful careers, come back and donate money to their university, for stuff like an endowed lecture series (paying to have someone give an academic lecture every year), or an endowed chair position (subsidizing some prestigious professor's salary). Don't do that. That's a total waste of money, when you could instead use the money to save many lives.^[106]

There are now organizations devoted to evaluating the effectiveness of various charities, including some organizations that you can give to who will re-gift the money to those charities that they find to be most efficient. Here are three important ones:

GiveWell: <https://www.givewell.org>.

These people evaluate cost-effectiveness of different charities oriented toward world poverty. You can go there and see what they think are the best causes. You can also just give money to GiveWell and have them split it among their top-rated charities. GiveWell estimates that the best charities save lives for a cost of about \$3000 per life.

Animal Charity Evaluators: <https://animalcharityevaluators.org>.

Like GiveWell, but for non-human-animal-related charities. In utilitarian terms, you can probably have a far greater impact by focusing on animal charities (animals are worse off and it's easier to reduce their suffering than to reduce human suffering).

Effective Altruism Funds: <https://app.effectivealtruism.org/funds>.

Similar to the previous two, but you can specify how you want to allocate your donation among four different areas of concern that you might have; their experts then evaluate what are the best charities in each of these areas, and they re-gift the money accordingly.

I'm telling you this just in case you were persuaded by the arguments of the previous sections *and* you happen to be among the small minority of humans who are consistently moral. In that case, you'll want to donate your money efficiently.

Remember, the purpose of giving isn't to make a sacrifice (such that as long as you've *sacrificed* enough, your obligation would be satisfied). We're not aiming at suffering. The purpose is to do some good. So direct your gifts to do the most good.

16.5. Government Policy

16.5.1. *The Argument for Social Welfare Programs*

So far, all of that was about individual ethics – what should you personally do about world poverty (or other problems)? We argued that you have an individual obligation to donate some of your excess money to do some good for others.

Now here's a political question: What should the *government* do? This is a different question because the government does not earn money in the same way that you and I do. The government, rather, acquires money by forcibly confiscating it from the rest of us. Many people believe that this implies much stricter moral constraints on what the government may legitimately do with “its” money than the moral constraints on what you or I may do with our money.

Here is an example that suggests, nevertheless, that it is appropriate for the government to fund charitable causes:

Pond Bystander: As before, you see a child drowning in a shallow pond. This time, you are unable to save the child yourself, as you are currently confined to a wheelchair. There is another bystander near the pond, who could easily wade in and save the child, yet he is doing nothing. You ask the bystander to save the child, but he objects; he doesn't want to get his clothes all wet, he says. It's clear that you won't be able to persuade him through normal means, as he is an asshole. You happen, however, to have a gun in your possession. You could point the gun at the asshole and *order* him to save the child, which you reasonably predict will result in his saving the child. Should you thus coerce the bystander to save the child?

In this case, though your use of force is of course regrettable, it seems justified. Admittedly, you may be treating the bystander as a mere means (contrary to Kantian ethics) as well as violating his right not to be coerced. But it still seems like the thing to do.

That, one might argue, is analogous to the government's situation. The government is aware of many poor people who are in need, who could easily be helped with some money. The citizens are not willing to voluntarily donate enough money to help all these people. So, just like you in the Pond Bystander case, the government resorts to coercing them to help – in this case, coercing them to donate money, which the government uses to help the poor.

16.5.2. The Charity Mugging Example

Now here's a different story:

Charity Mugging: You have a charity that you've created to help the poor. Your charity is doing good work, but you feel you are not getting enough voluntary donations. So one day, you decide to go out and start collecting "donations" by force. You go up to people on the street who look like they have money, and you rob them at gunpoint. You then funnel all this money into your charity. Is this appropriate behavior?

Most people have little trouble judging this to be impermissible. Yet it seems analogous to the government's behavior. The government also collects money from people by force, in order to fund *its* charity programs to help the poor. If it's not okay for you to do it, why would it be okay for the government?

So we have a clash of analogies. Pond Bystander suggests that government social welfare programs are permissible; Charity Mugging suggests that they are not. Which analogy is better? On the face of it, Charity Mugging is a closer analogy (it is more similar to social welfare programs). Charity Mugging, like social welfare programs, involves an ongoing program of coercion, aimed at alleviating chronic poverty. Pond Bystander, by contrast, involves an isolated act of coercion, aimed at resolving an acute emergency, a drowning child. If either of these is a fair analogy to government social welfare programs, it is the Charity Mugging case.

16.5.3. Other Problems with Government Programs

There are other problems facing government social programs in the real world. One of these is that government programs are not in fact aimed at the most needy people (let alone the most needy animals). The neediest people in the world are in the developing world, but governments in the wealthier nations have approximately zero concern for them. The reason for this is that these governments are democratic, and most voters don't care about foreigners. Indeed, many voters hate, or at least are very suspicious of, foreigners. Thus, despite that foreign aid forms a tiny portion of the budget, it's the one thing that Americans can agree on cutting (average voters also absurdly estimate foreign aid to be one of the largest items in the budget, when it is in fact under 1%).

The people helped by government programs in developed nations are generally not absolutely poor; they are merely poor *relative* to their society.

That is, they typically have their basic needs met, but government aims to modestly improve their welfare. A good deal of government aid actually goes to middle- and upper-class people. Of particular note, financial aid for college students is almost entirely a subsidy to the middle class (poor people rarely attend college, regardless of financial aid opportunities). Social Security is also a regressive redistribution program (that is, it redistributes toward the relatively wealthy), because the wealthier classes tend to both start work later and live longer than the poor; they thus pay into the program for a shorter time and draw from it for a longer time.

The Pond Bystander analogy doesn't really apply to these sorts of programs. In Pond Bystander, you point a gun at the bystander to make him rescue a drowning child. A better analogy would be pointing a gun at the bystander and demanding that he fund the child's college classes. Or pointing a gun at the bystander and demanding that he pay for an old person's retirement program or medical bills.

Another issue is that in our stories, you only ever *threaten* people with violence. Government, however, cannot merely threaten. Some people will always disobey (for example, tax evaders), and the government will then have to actually *carry out* its threats, for example, to imprison those people; otherwise, it will soon become known that the law is unenforced, and disobedience will be rampant. Now, it may well be that you are justified in threatening people in situations in which you wouldn't be justified in carrying out the threat. In Pond Bystander, it's plausible that you may threaten to shoot the bystander, to get him to save the child. But if he still refuses to save the child, you may not actually shoot him.

Finally, arguments like the Pond Bystander argument *assume* that government social programs actually help, rather than making things worse. There isn't a general agreement on whether that's true; some argue that government anti-poverty efforts are ineffective or counter-productive. One reason for this is that decades of such poverty programs in the U.S. do not seem to have resulted in a reduction in the poverty rate. This could be because these programs lure people into dependence on the state, and because they make it easier in the short run for people to do things that are self-destructive in the long-run (e.g., having out-of-wedlock births, or being unemployed).^[107]

It's *not clear* whether these things are true – there is reasonable dispute in social science about whether social programs in the developed world help

or not. (By contrast, there is little dispute about the fact that some charities really help the poor in the developed world, e.g., by stopping them from getting malaria.) This makes it harder to justify forcing everyone to contribute to these programs.

16.6. Conclusion

The case for government aid to the poor is shaky, especially the type of programs actually supported by actual governments. However, the case for an individual obligation to give to charity is very strong. If you saw some children drowning in a pond, whom you could easily save at small cost to yourself, surely you would and should save them. The fact that the children in the developing world are physically *farther away*, and that you aren't *seeing* them with your own eyes, is morally irrelevant. Those factors affect your ability to emotionally appreciate their plight, but they do not affect how important the needs of the global poor in fact are. So, just as you would save the drowning children, you should also save some of the global poor. (Or donate to other causes that are equally or more important.)

In donating to charity, you should choose charities that will do the most good for the money; there are organizations such as GiveWell, Animal Charity Evaluators, and Effective Altruism Funds to help you do this.

Thus far, I've mostly talked as if the best cause to donate to is relief of world poverty, though I have mentioned animal-related charities in passing. In fact, animal-related charities are probably much more cost-effective than human-oriented charities. However, it is even more difficult to get most people to care about other species than it is to get them to care about people in other countries. We'll talk about that in more detail in the next chapter.

17. Applied Ethics, 2: Animal Ethics

17.1. The Case for Vegetarianism

17.1.1. Where Does Our Food Come From?

Do you ever think about where your food comes from? Most human beings regularly consume products made from the bodies of other beings, especially chickens, pigs, and cows. An average person in a wealthy society consumes the equivalent of about 2,000 animals in a lifetime. Worldwide, about 74 billion animals are slaughtered for food per year, nearly ten times the human population of the Earth. [\[108\]](#)

Nearly all of those 74 billion animals are raised on **factory farms**. These are industrial operations in which large numbers of animals are raised in a small area, usually crowded together so closely, or held in such small cages, that they can barely move and they are forced to sit in their own excrement. Chickens (who are by far the majority of farm animals) commonly live in barns where the floor is covered in ammonia, because pouring ammonia on the floor is cheaper than cleaning out the excrement; hence, these animals breathe ammonia fumes all day. Farm animals are regularly subject to stress, pain, and physical mutilation. For instance, farmers regularly cut off pigs' tails and chickens' beaks (which contain sensitive tissue), without anesthetic. This probably feels about the same as having your finger cut off with a pair of hedge clippers. Farmers sometimes hold hot metal against cows' skin in order to mark them, again without anesthetic. This probably feels about the same as someone holding your hand down on a hot stove until you have third-degree burns. [\[109\]](#)

Animal welfare activists occasionally take videos inside factory farms, showing acts of abuse such as workers beating animals, throwing chickens against a wall for fun, and so on. People who watch videos of factory farm conditions commonly experience disgust and horror. No one who looks at them thinks that these conditions are humane; it is generally agreed that factory farm animals lead short, miserable lives.

The reason for this situation is that factory farms are the cheapest way of mass producing the meat and other animal products that consumers want. Consumers buy these products mainly for gustatory pleasure – i.e., we enjoy the taste – and we do not want to pay slightly more for our pleasure.

This (the desire for cheap mass production) is also why animals are regularly injected with hormones and antibiotics.

Any animal product that you see in a store or restaurant, unless it says otherwise, should be assumed to be from a factory farm.

17.1.2. The Argument from Suffering

On the face of it, all this raises an ethical issue. Is it okay that we are regularly causing suffering and death to other beings for our own gustatory pleasure? Is it okay to buy the products that come from this industry?

Many ethicists and ordinary people have concluded that this is not ethically acceptable; hence, they have become **ethical vegetarians**, people who abstain from meat for ethical reasons.^[110] (Some people are vegetarians for health reasons, but I won't discuss that.) The most common argument for ethical vegetarianism is something like this:

1. Suffering is bad.
2. It is wrong to cause an enormous amount of something bad, for the sake of relatively minor benefits for ourselves.
3. Factory farming causes an enormous amount of suffering, for the sake of relatively minor benefits for humans.
4. Therefore, factory farming is wrong.
5. If it's wrong to do something, it's wrong to pay other people to do it.
6. Buying products from factory farms is paying people for factory farming.
7. Therefore, it's wrong to buy products from factory farms.

By the way, notice that the primary conclusion here is not exactly that eating meat is wrong; it is that *buying products from factory farms* is wrong. This is what most ethical vegetarians believe (some also think that buying *any* animal product is wrong). When they say “you shouldn't eat meat”, that's because they assume that if you're eating meat, then you're buying it (or having someone buy it for you), and the meat is almost certainly from a factory farm. Hereafter, I'll take those assumptions for granted. (But we'll talk about “free range” meat and similar products later (§17.3).)

As in other cases, if you don't agree with the conclusion, then you need to identify which of the premises (1, 2, 3, 5, or 6) you disagree with (or, of course, you could just admit that you're wrong, but who likes to do that?). All of those premises seem obvious and uncontroversial in normal contexts.

The usual target is premise 1: Defenders of animal cruelty claim that only *human* suffering is bad; animal suffering isn't. They then need to

identify some *relevant difference* between humans and animals that explains why that would be so. (This follows the general principle that if there is a *moral* difference between two things, then there has to be some *descriptive* difference between them that accounts for the moral difference. This is universally accepted in ethics.) We will discuss attempts to identify the relevant difference between humans and animals in §17.2. But basically, none of them are at all plausible.

17.1.3. *Arguments by Analogy*

Here is another way to think about the basic issue. Imagine that there were an industry devoted to raising human babies. The babies are tortured for a few months, then slaughtered for food. Assume also that human baby meat was delicious. Would it be morally alright to buy this meat?

I hope you answered “no”. (If not, you might be a psychopath, in which case you can’t understand anything in this part of the book.) But why? It seems that the answer has to do with the badness of suffering and death, the wrongness of causing very bad things for the sake of minor benefits to yourself, and the wrongness of patronizing an immoral industry. Those things are also why buying meat in the actual world is wrong.

Again, that must be true unless there is some *morally relevant difference* between the two cases. Note that I say “morally relevant”; sometimes people overlook that part. I’m not asking for just *any difference*. For instance, humans are the only featherless bipeds on the Earth, but that is not morally relevant – i.e., having two legs and no feathers does not plausibly explain why one would have rights, or why one’s suffering would be bad and that of other beings would not be bad.

Here is another analogy. Most people accept the wrongness of animal cruelty in other contexts. For instance, if you see a man beating his dog just for the fun of it, chopping off its tail, or deliberately burning it, you would recognize that as wrong. But there is no morally relevant difference between dogs and pigs, chickens, or cows. Therefore, it’s also wrong to abuse animals on farms. Dogs happen to be widely accepted as pets in our society, so most people are able to empathize with them to some degree and to care about their mistreatment. But that’s not morally relevant; we could as easily have had a practice where cows were pets and dogs were used for food. (In fact, dogs are used for food in some countries, while cows are considered sacred in India.)

Another difference is that the torture of animals on factory farms is generally out of sight. Since you don't have to actually watch people doing it, you are able to ignore it. But again, that is obviously not morally relevant. Abusive behavior is not less bad merely because it is hidden from view.

17.1.4. *Animal Rights vs. Welfare*

Opponents of animal cruelty come in two varieties: **animal welfare** advocates and **animal rights** advocates. Animal welfare advocates just believe that animals have welfare – i.e., they can be harmed or benefitted – and that we should not harm them without good reason (where getting a comparatively tiny amount of pleasure out of it does not count as a good reason).^[111]

Animal rights advocates believe that animals also have *rights* – i.e., that there are deontological constraints on harming or interfering with them in certain ways, similar to the constraints that apply to treatment of human beings (see §15.1.4).^[112] They need not be exactly the same rights that humans have, though; perhaps humans have *more* rights. E.g., maybe humans and animals all have a right to life, but only humans have a right of free speech.

The animal rights position is the more extreme position. Those who merely endorse animal welfare would oppose harming animals to achieve trivial benefits; however, they would not oppose harming animals when doing so achieved *greater* benefits for other animals or people. That is, the mere animal welfare advocate would treat animals the way utilitarians treat everyone. By contrast, the animal *rights* advocate would treat animals the way that deontologists treat humans.

The basic motivation for animal welfare is simple. Other things being equal, causing harm is bad, and no one has been able to give any credible reason why harm would only be bad if it happens to *our* species.

Animal rights advocates would say that there is also no credible reason why rights would only apply to our species; hence, we should assume that other sentient species also have rights. However, this is less compelling than the animal welfare argument, since there is no general agreement among ethicists about why (if at all!) humans have rights. Among those who believe in rights, there are different accounts of what gives us rights, and some accounts depend upon our intelligence, rationality, and stuff like that. This means that there are reasonable accounts on which (most) non-

human animals would lack rights, along with some severely mentally disabled humans. There are, however, no reasonable accounts on which non-human animals or mentally disabled humans lack *interests*; hence, no plausible reason why it would be acceptable to seriously harm them for the sake of trivial benefits for ourselves.

17.2. Defenses of Meat-Eating

In this section, I'll talk about reasons that people have given for why it's okay to inflict extreme harm on members of other species for the sake of our own minor enjoyment. Some of these, by the way, might strike you as ridiculous, and you might wonder whether I'm setting up straw men here (see §4.1 for the straw man fallacy). I am not. All of these are things I have actually heard in conversation or seen in print.

Argument 1: It's okay to torture other animals for our own pleasure, because we are *intelligent*, and other animals are not.

Replies:

1. This is a non sequitur. What does intelligence have to do with the badness of pain and suffering? Is the claim that pain is only bad if you're *smart*? Why would that be?
2. Human babies are also unintelligent. Argument 1 thus implies that it would be permissible to torture babies for fun.
3. Similarly, Argument 1 implies that it would be permissible to torture mentally retarded people for fun.

Argument 2: Consumers are not responsible for the cruelty of factory farms, because the consumers are not themselves directly inflicting the pain and suffering, and they haven't specifically told the farm workers to do it either.

Reply: You don't have to cause a harm directly in order to be blameworthy. For instance, if you hire a hit man to kill someone, you will be just as blameworthy as the hit man.

You also don't have to expressly tell someone to do something wrong in order for you to be blameworthy. Suppose there's a used car dealer who obtains all his cars by murdering innocent people and stealing their cars. No one specifically told him to do this, but everyone, including you, knows that this is how he in fact gets his cars. It would be uncontroversially wrong to buy a car from this dealer. This

illustrates the principle that if it's wrong to do something, it's also wrong to pay other people for doing it.

Argument 3: How do we know that animals feel pain?

Reply: We know that farm animals feel pain because (a) they have the same kind of nerves that generate pain sensations in us, and (b) they behave exactly as if they are in pain, in circumstances that would cause pain in you. This is why no expert questions that cows, pigs, and chickens feel pain.

It is of course *logically possible* that other animals are just mindless mechanisms, and that for some unknown reason, only humans can experience anything. It is similarly logically possible that all humans other than you are just mindless mechanisms, and that for some unknown reason, only *you* can experience anything. However, it would not be rational or ethical to go around beating other people, on the theory that maybe they can't experience pain. Nor is it rational or ethical to treat animals similarly.

Argument 4: Animals eat each other, so why shouldn't we eat them?

Replies:

1. This argument presupposes that anything that an animal does, it is permissible for you to do. That's false.
2. The logic is analogous to the following: "Humans often kill each other; therefore, it's okay to kill humans." That's also wrong.
3. The animals living in factory farms have generally not, in fact, eaten any other animals. (But if they did, it would have to be because the farm workers fed the meat to them. So the animals don't deserve to be killed for having eaten other animals.)
4. Animals also don't raise other animals in factory farms, subjecting them to painful and unnatural conditions for their whole lives.

Argument 5: Rights imply obligations. Animals don't have any moral obligations; therefore, they don't have any rights.

Replies:

1. Human babies also don't have any obligations. Therefore, Argument 5 implies that babies don't have any rights. Similarly

for people who are severely mentally retarded or mentally ill.

2. Whether or not anyone has “rights”, it’s still wrong to cause enormous harm for no good reason. Argument 5 doesn’t change that.
3. The slogan “rights imply obligations” is confused. For you to have rights *does* imply that *other people* have obligations, namely, obligations to respect your rights. It does not imply that *you* have obligations; that’s just a non sequitur. That’s why there’s nothing odd about babies having rights without obligations.

Argument 6: Morality doesn’t apply to animals, because animals cannot understand morality.

Replies:

1. This is another non sequitur. There’s no reason to assume that we only have moral obligations to beings who can understand morality.
2. Argument 6 implies, again, that it would be fine to torture babies, mentally retarded people, and certain mentally ill people, since they can’t understand morality.

Argument 7: It’s okay to torture and kill animals because they don’t have souls.

Replies:

1. It is unclear what is meant by “not having a soul”. Here are two interpretations:
 1. “Animals are just mindless automata; they have no mental states.”

Reply: See reply to Argument 3 above.

2. “Animals have minds, but they don’t go to heaven after they die.”

Reply: If anything, this makes it even worse to kill them; at least humans get to go to heaven after they die. It certainly doesn't explain why it would be okay to torture them.

2. It is controversial whether souls exist at all. Suppose it turns out that they don't. Would it follow that it's perfectly alright to torture people for fun? If not, then the wrongness of torture doesn't depend on one's having a soul.

Argument 8: The Bible says that God gave us dominion over the animals.

Reply: But the Bible doesn't say that it's okay to torture and kill them for trivial reasons. It seems more likely that a benevolent deity would want us to act as responsible, benevolent stewards of the Earth and its creatures, not to indiscriminately harm and destroy for trivial reasons.

Factory farming did not exist at the time the Bible was written. Therefore, we can't infer that factory farming is acceptable merely because it wasn't condemned by the Bible. By the way, many other things are wrong that are not mentioned in the Bible. We have to exercise our conscience to identify these things.

Argument 9: If it's wrong to kill animals, then it must also be wrong to kill plants. Therefore, vegetarianism is just as bad as meat-eating!

Reply: There are two ways to understand this argument:

1. "Plants are just as sentient as farm animals, and stalks of corn are being tortured in the corn fields."

Reply: There is no reason to believe this. All mental states, as far as we know, are caused by activity in your brain. Plants have no nervous systems at all, let alone brains. They also don't exhibit any pain behavior (they don't act as if they are in pain). There also would be no evolutionary function to plant pain, since plants cannot do anything about it if you "hurt" them.

2. "Sentience doesn't matter. Only *life* has intrinsic value. All life is equally valuable, whether sentient or not."

Reply: I seriously doubt that the people giving this argument believe that. If you believe that, then you must consider it equally bad to kill a bacterium as it is to kill a human. If you believed that, you wouldn't go around just following the conventional practices of your society, killing plants and animals whenever you felt like it.

Maybe the argument is supposed to be that since life is the only thing that matters, and it's *okay* to kill some living things (like plants and bacteria), therefore it's also okay to kill animals. Notice that this argument also implies that it's fine to murder people.

Maybe the person would say, "Oh no, there are *two* things that matter: life, and intelligence." In that case, see reply to Argument 1.

Argument 10: Plant farming also kills animals! Farmers kill insects with pesticides. Even if you buy organic foods, they probably kill field mice sometimes in the process of tilling fields and harvesting vegetables with machines. Therefore, vegetarians are no better than carnivores!

Reply: Animal agriculture is worse than plant agriculture in a number of ways.

1. Factory farms confine animals in unnatural and unpleasant conditions, subjecting them to pain and suffering for their entire, brief lives before killing them. Plant farms do not do this.
2. Chickens, pigs, and cows definitely feel pain and suffering. Insects almost certainly do not, due to the absence of nociceptors (the kind of nerves that generate pain sensations in us). This is why insects can continue what they are doing even when their bodies are severely injured, they place the same weight on an injured leg as on an intact leg, etc. [\[113\]](#)
3. Animal farms require food for the animals, which comes from plant farms. The amount of food you have to feed the animals in the course of raising them is greater than the amount of food you get out of them at the end. Hence, meat production causes *more* of whatever harms are caused by plant farming, *in addition* to the harms directly inflicted on the animals in the factory farms. Thus, while it might be true that plant farming causes *some* harm, this can't be used to excuse animal farming. In general, one can't

justify some bad behavior by saying that the alternative action causes some (much smaller) amount of harm.

Argument 11: Eating meat is natural; therefore, it's okay.

Replies:

1. Many bad things are natural. For example, cancer is natural; that doesn't mean that we shouldn't try to cure it. Or, if you want examples of human behaviors, there's some reason to think that war is natural behavior for humans. Assume that's true. It obviously doesn't follow that war is okay, or that we shouldn't try to stop it.
2. Whether or not meat-eating is natural, factory farming is definitely not.

Argument 12: Farm animals wouldn't even exist if it weren't for the meat industry. Therefore, the meat industry is good for them!

Replies:

1. This argument could also be used to justify slavery. In America, after the importation of slaves was banned in 1808, slave-owners took to breeding new slaves from the existing stock of slaves in the country. The people who were created by these slave breeders would not have existed at all if not for slavery – in exactly the same way that farm animals wouldn't exist if not for the meat industry. But it obviously does not follow from this that it was okay to hold them as slaves. A similar argument could be made, by the way, if we were breeding human children for food.
2. The lives of animals on factory farms basically consist of suffering for a few months before being slaughtered. (Most of them are chickens, who live for only a few months before they're big enough to kill.) These are probably lives of negative utility.
3. In any case, we don't have a moral reason to create new beings so that they can experience brief, wretched lives. Once a being is created, however, we have a moral obligation to treat that being decently. These points are generally accepted for human beings; there's no obvious reason why they wouldn't also apply to animals.

4. Argument 12 does not explain why it would be okay to have cruel farms, rather than humane farms. So it wouldn't justify buying products from factory farms, even if there were nothing else wrong with the argument.

Argument 13: I can't make a difference anyway. The meat industry is so large that it won't respond to the actions of a single person.

Replies:

1. Again, this argument could be used to justify cannibalism. Suppose we had a very large industry that was breeding human children for food. The children are raised in crowded, disgusting conditions, mutilated, occasionally beaten, and then slaughtered for meat at the end of a brief life. Would it be alright to buy meat from human children, on the ground that "the industry is too large to respond to me"?
2. Anyway, the idea behind Argument 13 is pretty much incoherent. The person giving this argument doesn't think that the industry would keep producing exactly the same amount of meat no matter what the demand was like. Rather, they think that if demand goes down by *a lot*, then the industry will reduce production correspondingly (which is obvious), *but* if demand goes down by just *a little*, the industry won't respond *at all*. Here's why this makes no sense:

Assume that the industry would not respond at all if *fewer* than 100 people become vegetarian, but that they *would* respond if 100 people became vegetarians. That's the threshold. In that case, when 100 people become vegetarian, they'd presumably reduce production by about the amount that 100 people eat. But note that many other people have actually become vegetarian already. You don't know how many. Maybe 99 other people have become vegetarian since the last time they adjusted their production. So there's a 1/100 chance that you, by becoming vegetarian, will actually push us over the threshold where the meat industry responds by reducing production by 100 times the amount that one person eats. So it's worth doing that. (A 1% chance of

producing 100 units of benefit is just as good as a 100% chance of producing 1 unit of benefit – that’s the principle of expected utility.)

Notice that the reasoning works equally well no matter what you assume is the threshold. If you think the threshold is 1,000 instead of 100, you can just redo the reasoning with that number, concluding that you get a 1/1,000 chance of reducing meat production by the amount that 1,000 people eat. The reasoning also works equally well if you think there’s a probabilistic threshold – e.g., maybe you think that, as the number of vegetarians increases, the *probability* of the meat-industry responding by reducing production increases. That’s going to give the same result; the calculations would just be more complicated. Finally, the reasoning also works if, instead of declining, meat production is *increasing*. In that case, by eating meat, you have a chance of causing the industry to respond by *increasing* their production, which is bad, so you have reason to not buy more meat.

Argument 14: If it’s wrong for us to eat meat, then it must also be wrong for lions to eat gazelles. But this isn’t wrong.

Replies:

1. Argument 14 is analogous to the following: “If it’s wrong for us to kill people, then it must be wrong, say, for a snake to kill a person. But this isn’t wrong. Therefore, killing people is perfectly fine.”

If it’s not wrong for lions to kill gazelles, then I don’t know why it would be wrong for a snake to kill you either. But obviously it’s *not* okay for a *human* to kill you. So humans have obligations that animals don’t have. (By the way, I think that’s well-known to everyone, including the people who make Argument 14.)

2. Here are two of the differences between lions and humans:
 - a) Humans are capable of understanding morality and of controlling our own behavior to conform to moral principles. Therefore, we can be obligated to do that, and we can be blameworthy for not doing so. Lions do not have those abilities.
 - b) Humans are able to survive on a vegetarian diet. Lions are not.

Both of these points, again, are well-known. In fact, they are surely the reason why we think that it's not wrong for lions to eat gazelles in the first place. So the justification for one premise of the argument ("It's not wrong for lions to eat meat") directly refutes the other premise ("If it's wrong for us, then it's wrong for lions").

Argument 15: Okay, but surely, if ethical vegetarianism is correct, we would be obligated to stop predators from killing other animals. But we're not obligated to do that!

Reply: This is another argument where the justification for one of the premises refutes the other premise. Let's consider why someone might say that we're not obligated to stop predators from killing other animals. Here are some possible reasons:

1. Because that would result in the extinction of the predator species, which is bad.
2. Because it would disrupt the entire ecology, some prey species would multiply out of control, etc.
3. Because we in fact have no feasible way of stopping all predators from killing other animals.
4. Because in general, we're not obligated to stop harms from natural causes; we're only obligated to refrain from causing harms ourselves.

Now, notice that *none* of those reasons apply to stopping human meat consumption. If we become vegetarians, we won't go extinct; it won't disrupt the ecology (in fact, it would *reduce* the harm that we're doing to the environment); it is something we could feasibly do; and it would be stopping a harm that we ourselves are causing, not a harm from natural causes. So, even if we're not obligated to stop predators in nature, we can still be obligated to stop our own meat-eating.

Now, you might disagree with any of (a)-(d). If you disagree with all of them, though, then I have no idea why you would think that we're not obligated to stop predators from killing other animals.

Argument 16: I can imagine some circumstances in which eating meat is okay. E.g., if you're about to starve to death, and the only thing to eat is a chicken, you may eat the chicken. Or if there's a chicken that has died of natural causes. Or if you get humanely raised

chicken meat, maybe that would be okay. Therefore, it's not true that eating meat is wrong.

Reply: So you're only going to eat meat in those circumstances that you just listed, right? You're not just using this as an excuse to do whatever you feel like?

Here's an analogy. A says, "Killing other people is wrong. You shouldn't do it." B says, "No, because it's okay to kill people in self-defense. Thus, it's false that killing people is wrong." B then goes around indiscriminately murdering anyone he doesn't like. Can you spot the mistake? The mistake is that "It's permissible to kill in some very unusual circumstances" does not entail "It's permissible to kill anyone at any time."

That's like most meat-eaters. They claim that it's permissible to eat meat in some unusual circumstances, then take that as an excuse for eating meat whenever they feel like it.

The arguments that we discussed above (§§17.1.2–17.1.3) claim that it's wrong to buy products from factory farms (in the current, actual circumstances that almost all of us are in). They don't claim that it's wrong to eat meat in every logically possible circumstance. So Argument 16 does nothing to challenge them.

Argument 17: Okay, here's a theory: It's wrong to torture beings that *belong to a species* that is intelligent, whether or not the specific individual is intelligent. So it's wrong to torture babies and retarded people because normal, adult humans are smart, and they belong to the same species as the babies and retarded humans.

Replies:

1. This is *ad hoc*. The original claim was that one's *intelligence* somehow determines how bad one's pain is (Argument 1). This is refuted by the fact that it's wrong to torture babies, even though babies are dumb. So the meat-eater modifies his theory by claiming that it's the intelligence of your *species* that matters. But there is no reason for saying this other than to try to protect one's theory from refutation.

Note: Ad hocness

The term “*ad hoc*” is used to describe modifications to a theory that have no rationale other than to protect the theory from being refuted by some piece of evidence. Example: Experiments designed to detect psychic powers regularly fail to detect any. This is evidence that there aren’t any such things. Believers in psychic phenomena, however, often try to preserve their belief by claiming that the existence of doubting observers (like the scientists who are trying to test for psychic powers) interferes with the psychic powers, thus rendering them undetectable. There is no independent rationale for thinking that this would be the case; the supposition is introduced solely to protect the theory of psychic powers from being rebutted by evidence. Hence, this is said to be an *ad hoc* hypothesis, or an *ad hoc* rationalization. Scientists generally don’t accept *ad hoc* hypotheses because they could be used to defend almost any theory from almost any evidence.

2. If we’re going to start looking at the *group* that a particular creature belongs to, there is no reason to pick the *species* classification, rather than one of the indefinitely many other levels of classification. Why not say that it’s the *genus* that one belongs to that matters? Or the order? Or the race? Or the sex? Argument 17 arbitrarily picks one way of classifying beings and asserts some special moral significance to that.
3. The intrinsic badness of someone’s pain should be determined by the characteristics of the pain itself, and *maybe* other characteristics of the subject who is having the pain. Even if being smart somehow makes your pains worse, it’s obviously not the case that *other people’s* being smart makes *your* pains worse (even if the other people belong to the same species as you). Notice that Argument 17 implies that we could imagine two qualitatively identical pains, experienced by two intrinsically, qualitatively identical creatures, and yet one pain could be horrible and the other not bad at all. (This would happen because in one case the other members of the creature’s species happen to be smart, while in the other case, they were not smart. [\[114\]](#)) But this is obviously not possible.

Possible response to point 3: It's not really about the characteristics of other beings. Rather, each individual has a special property, a "species essence", which determines what species they belong to, and they would have this essence whether or not any other members of their species existed. And this species essence, in the case of humans, includes a kind of metaphysical potentiality for intelligence, even in cases where the potentiality is not realized and not practically realizable (as in the case of severely retarded people).

Replies to the possible response:

1. All observable facts about humans and animals can be explained by their ordinary physical and mental characteristics – things like the chemical composition of their bodies, their desires, their perceptions, etc. If a species essence is something else, then there is no reason for believing that any such thing exists. If, on the other hand, a species essence is just some collection of ordinary physical and mental characteristics, then the theorist should say what physical and mental characteristics are supposed to be morally relevant, instead of vaguely gesturing at something-or-other.
2. Even if such things existed, they would be irrelevant. Suppose we have these unobservable, metaphysical essences, which can give us metaphysical potentialities distinct from our actual abilities in the ordinary sense. (For instance, so that permanently mentally retarded people could have the "potential" for intelligence, even when there is no real way of making them intelligent.) So what? What would that have to do with the badness of pain?

Suppose it turns out that there *aren't* any species essences, nor are there any metaphysical potentialities. Would it then be alright to torture babies for fun? If not, then the wrongness of torture doesn't depend on species essences or metaphysical potentialities.

17.3. Other Ethical Issues

17.3.1. The Importance of Factory Farm Meat

Up till now, I have focused on human consumption of meat from factory farms, because this is the most important issue in animal ethics. It's important because (i) almost everyone in our society is doing it all the time, and (ii) the total harm caused by this practice is orders of magnitude greater than any other problem that people commonly talk about. As I mentioned at the start (§17.1.1), humans worldwide slaughter close to 74 billion animals for food every year, nearly all of them on factory farms. It is estimated that the total number of humans who have *ever* existed is about 110 billion. So in a couple of years, we kill more animals than the total number of humans who have ever existed. Since 99% of these are on factory farms, it's plausible that the total amount of pain and suffering we inflict through this practice, in just a few years, may be greater than all the pain and suffering in all of human history.

At the same time, this is, of all controversial issues, the one issue where there is the biggest disconnect between what rational reflection shows and most people's actual attitudes. Many people are entirely insensitive to the issue. Many dismiss all concern about animals without thinking. When they deign to defend their practices, people give some of the weakest, most transparent rationalizations that you can find anywhere.

That's why I have focused on the issue of factory farmed meat. But now, in the remainder of §17.3, I'm going to discuss some related issues. This includes clarifying what the arguments of §17.1 imply, as well as addressing some perhaps more difficult issues beyond factory farming.

17.3.2. Other Animal Products

Most of the discussion has been about *meat*, because this is the main animal product that people consume. But the main arguments for giving up meat also apply to dairy and eggs, since these are also produced on factory farms under extremely inhumane conditions. They also apply to buying clothing made from animal products, such as leather and wool.

Apropos of the latter, here is a line of thought that you might be tempted by: "Animal farmers are raising cows for the meat. Since they're slaughtering cows for meat *anyway*, we might as well use the hides for clothing. So, maybe it's okay to buy the leather." This is not correct. Animal farmers are not raising cows "for the meat". They are raising and slaughtering cows for the *money*. They don't care if it's money from meat or money from leather – both contribute equally to their bottom line, and both thus create the incentive to raise and slaughter cattle. The more money

is available (regardless of which animal products generate the revenue), the more cows they will raise.

17.3.3. Humane Animal Products

You've probably seen "cage free" eggs in the store, along with "free range" meat, eggs, or other animal products. Are these products ethical to buy?

Unfortunately, most of these labels mean little for animal welfare. There are few legal restrictions on what a company can call "free range", etc. "Cage free" and "free range" chickens may spend almost all their time crowded together in a giant shed, sitting in their own waste; they can also be painfully debeaked and otherwise abused. If you buy these products, you're trusting the company to monitor animal welfare on their own farms. Most companies, however, do not care at all about animal welfare, nor about honesty. They'll just write "free range" on the package if they think it'll sell more product.

There are, however, some animal welfare organizations that monitor farms and certify them as humane. You're likely to see something like this "Certified Humane" logo on some products:



which indicates that the farm treats its animals humanely, according to an animal-welfare-oriented organization (in this case, the organization is Humane Farm Animal Care; there are other organizations with other labels, but this is the one you most often see). These products are thus more ethical, or less unethical: They have greatly reduced animal cruelty, though there is still the issue of whether one should kill other creatures simply because one likes the taste of their flesh. By the way, in the chicken/egg industry, it is standard practice to kill male chicks shortly after hatching since they do not produce eggs and hence are not economically worth raising. This practice is allowed by cage-free, free range, pasture-raised, and even Certified Humane production methods.

17.3.4. Insentient Animals

Some species in the animal kingdom lack sentience – that is, they cannot experience pleasure or pain. Of particular interest are bivalves, such as clams, mussels, oysters, and scallops. These creatures have only a few ganglia, with no central nervous systems. It is thus highly improbable that they are capable of pleasure, pain, or any other experiences. (You have ganglia in your body too; if they get stimulated, but the signal never reaches your central nervous system, you will report feeling nothing.) This is of practical interest because they are commonly used for food and can supply nutrients, such as vitamin B₁₂, that are difficult or impossible to obtain from plant sources. Given their lack of sentience, there is no ethical obstacle to consuming them.

The case of insects is a little more controversial. They have (very small) brains, but they do not behave like other animals that feel pain. For instance, an insect will walk with an injured leg, resting the same amount of force on the injured leg as on an uninjured leg. Insects will sometimes carry on normal behaviors – e.g., continue eating or mating – even while their bodies are severely wounded. This is relevant for a few products that are made using insects, such as honey and silk. Strict vegans reject these products; however, the argument for giving up these products would be much weaker than the argument for giving up meat, dairy, and eggs, since it is doubtful whether insects can feel pain at all and also doubtful, even if they can, whether they would experience any suffering as a result of being “exploited” on a farm.

17.3.5. Lab Grown Meat

As of this writing, lab grown meat is in development but not widely commercially available yet. This is a kind of meat that is produced synthetically, without killing an animal. There is no cruelty, since there is no nervous system to feel any pain or suffering. This is also expected to be a more efficient way of producing meat, with less environmental impact.

It goes without saying that the arguments of §17.1 don’t apply to this product; there is no ethical reason not to buy lab-grown meat. This product is most likely what will eventually end factory farming. By the way, once our descendants have switched to eating lab grown meat, and thus their self-interest is no longer in the way, I bet almost everyone will easily see the wrongness of factory farming.

17.3.6. Animal Experimentation

Is it unethical to experiment on animals, in the way that it would be unethical to experiment on humans against their will?

That depends upon whether you take a consequentialist or a deontological approach to animal ethics. If you believe that animals have *rights*, then presumably animal experimentation (without the animal's consent, which would be nearly impossible to get) is unethical.

As discussed earlier, however (§17.1.4), a consequentialist approach is at least reasonable. On this approach, animal experimentation is ethical if and only if it produces expected benefits greater than the expected costs. If you can find a cure for cancer by experimenting on mice, then you should do so. Note, however, that the experimentation must still be done in the most humane way compatible with achieving the desired end.

It is far from obvious, by the way, that most animal experimentation even satisfies this fairly lax consequentialist criterion. Many experiments are done, for example, to test new cosmetics, when we already have plenty of good cosmetics. Even medical experiments on animals are far less useful than people imagine – most treatments that work on mice fail on humans. By the same token, it's plausible that most treatments that *would* work on humans fail on mice. We may in fact be missing out on treatments that would have saved human lives, because we insist on testing everything on mice first, and we don't pursue a treatment if it doesn't work on mice.

17.3.7. Responding to Other People's Immorality

Most human beings are horrifyingly immoral. That's been demonstrated in many ways. You can look at historical episodes where ordinary people participate in genocide, slavery, or other obvious evils. Or look at the famous psychology experiments in which two thirds of subjects are willing to electrocute an innocent person, as long as a guy in a white lab coat tells them to do it.^[115] Of course, the people you meet in ordinary life probably *seem* decent to you most of the time. They appear to care about morality, since they mostly refrain from abusing you (hopefully). In fact, though, they mostly only care about *social conformity* – the main reason your neighbors aren't abusing you is that it would be socially frowned upon. The evidence for this is that almost everyone continues to do immoral things as long as they are socially accepted, while avoiding moral things that are socially disapproved.

This is why most people refuse to reconsider their eating habits, regardless of what moral arguments they hear. It's quite common for

people, after hearing about the case against factory farming, to agree that buying from factory farms is wrong, yet admit that they're going to keep on doing it anyway. They would not do similarly immoral things that are socially disapproved, though; they just don't care about this particular immorality, because it is socially accepted.

All this explains why, if you agree with the arguments of this chapter, you should not only refrain from buying factory farm products yourself. You should also exert *social pressure* on other people around you. E.g., express serious disapproval whenever your friends buy products from factory farms. If you meet someone for a meal, you should insist on going to a vegetarian restaurant.

By the way, if you do this, you can expect other people to act resentful, and indignant, and often to insult you. This is because, again, they are horrible. Given their horribleness, their main thought when someone points out their immorality is to get angry at the other person for making them feel slightly uncomfortable. They won't blame themselves for *being* immoral; they'll blame you for making them aware of it. It's sort of like how a serial murderer would get mad at you if you tried to stop him from murdering more people. He would then blame you for being "preachy". Perhaps the murderer would then refuse to be your friend any more. If so, good riddance.

In fact, most of my readers are probably regular meat-eaters, which means that I've probably just alienated you by calling you immoral. If you're mad about that, feel free to stop reading – oh wait, you're already at the end of the book anyway.

18. Concluding Thoughts

18.1. What Was this Book Good For?

Each of the preceding chapters was a summary of a broad topic about which a ton of literature has been written. The topics I chose are the ones that are most commonly covered in Introduction to Philosophy courses, especially external-world skepticism, the existence of God, free will, and personal identity. (There's no standard curriculum in philosophy, though; each philosophy professor just decides what he or she wants to cover. It's just that some topics are especially popular.) So if you're reading this book on your own, you've probably learned pretty close to what you'd learn in an Intro course in college.

I hope you found some of the preceding ideas and arguments compelling. I hope that you now have more understanding of the world and our place in it, as a result of those ideas. But you probably still disagree with a bunch of things I said in previous chapters. That's okay; that's normal for this subject. (Contrast the fields of mathematics or science, in which if you read a textbook, you will probably agree with all or nearly all that it says.) But even if you rejected some of the substantive philosophical ideas that I presented arguments for, I hope that you still learned something less tangible.

Here's the big thing that I hope people get out of this book, or out of taking a course in philosophy: I hope people learn a little bit of *how to think like a philosopher*. (I guess I should add: like a good philosopher, not like a bad one.) That's a cognitive skill you can acquire, not a discrete set of propositions to be accepted. When I told you arguments about philosophical issues, that wasn't so that you could memorize the specific sequences of propositions in those arguments, and then be able to regurgitate them on a test. (I think some students try to read every book that way; I hope you didn't do that. If you did, go back and reread the book properly!) I told you those arguments so that you could *think them through*; so that you could acquire, through these examples, a sense of how one reasons about a philosophical issue; and hence, hopefully, so that you would yourself be able to reason a little bit better about such issues. I say "a little bit" because

learning to think like a philosopher is not something that happens quickly. It's a multi-year process. So you should continue reading and thinking.

18.2. How Good Philosophers Think

Now I can say some things about how good philosophers think. I hope you noticed these things about most of the arguments discussed in previous chapters.

First, good philosophers think about the most important and fundamental questions. Like “Where did the universe come from?”, “Are there objective facts about value?”, “How do we know about the world outside our minds?”, and “What determines a person’s identity?” During philosophical debate, it’s easy to get lost in details. It’s important always to keep in mind what the central issue is, and not to waste time debating points that don’t matter to that issue.

Second, good philosophers marshal *rational arguments*. We don’t just randomly opine or express our feelings. In case we want to say something that’s not obvious on its face, we try to find other things that are more obvious, that rationally support our point. This usually involves reflecting a lot on *why* things seem to us the way they do. (For this reason, skill at introspection is important to philosophy.)

Third, good philosophers *answer objections*. If you have a philosophical view (or any view really), and you know that a lot of smart people disagree with it, you really need to think about *why* they disagree. And I don’t mean “Because they’re jerks” or “Because they’re evil.”^[116] What you need to think about are the *best reasons* someone could have for disagreeing. If you can’t think of any, then you probably haven’t thought or read enough about the issue; you should then go look up some intelligent opponents and see what they say. And I don’t mean television pundits or celebrities on Twitter. The best defenders of a view are usually academics who have written books about it. You should then think seriously about those objections and whether they might be correct. If you don’t find them persuasive, try to figure out why. This is the part of rational thought that most human beings tend to skip.

Fourth, good philosophers are *clear*. We tend to use words with specific meanings in mind, we try to draw relevant distinctions and avoid conceptual confusions. We know what we’re saying and what other people are saying, and how different views differ from each other. (This isn’t to

deny that a lot of philosophers are unclear, by the way. Those are generally bad philosophers.)

The above points might seem obvious. It's obvious that one should address important questions, give arguments, answer objections, and be clear. But I still think it's worth saying these things, because apparently a lot of people haven't gotten these points. From casual observation, it looks to me like quite a lot of people in public discourse just assert controversial opinions without attempting to give any reasons for them. Or their "reasons" are just paraphrases of the central point that's in dispute. (This is particularly popular for politicians.) I virtually never see people seriously engaging with objections, either.

18.3. Further Reading

If you got something out of this book, then you should continue reading and learning about philosophy. There are many more fascinating ideas out there to contemplate. I'm going to recommend some books that you are likely to enjoy if you liked this one. (If you hated this book, though, then I don't have any recommendations for you. My apologies for making you suffer this far.)

First, I suggest reading anything else by me. This might sound self-serving, and that's because it is. However, it's also good advice for you: If you like one piece of writing by a given author, that strongly predicts liking other things by the same author. I suggest looking up my blog (fakenous.net), my other books (just search on Amazon), my web site (www.owl232.net), and even, if you're *really* hardcore, my academic articles (just search on PhilPapers).

A few other authors I would recommend for their clear writing and logical arguments: Bertrand Russell (especially his *The Problems of Philosophy*), John Searle (especially *Minds, Brains, and Science*), David Stove (especially *Scientific Irrationalism*), Jason Brennan (especially *Against Democracy*), Robert Nozick (especially *Anarchy, State, and Utopia*). If you liked this book, you probably won't hate any of those authors.

If you want to branch out from philosophy into economics, consider looking up David Friedman (especially *The Machinery of Freedom*) and Bryan Caplan (especially *The Myth of the Rational Voter*). If you want to learn some stuff about modern science from a philosophical perspective, you can't do better than David Albert (especially *Quantum Mechanics and*

Experience and *Time and Chance*). For some philosophical fiction, if you like both Harry Potter and rationality, try Eliezer Yudkowsky's fan fiction novel *Harry Potter and the Methods of Rationality*.

That's all for this book. Good luck, and stay rational.

Appendix: A Guide to Writing

This is a writing guide that I made for my students many years ago. It tells you broadly what a philosophy paper should be like (§§A.1–A.2), how to do research (§A.3), and a lot of mistakes to avoid (§§A.4–A.7). See how many you can avoid!

A.1. The Content of a Philosophy Paper

A philosophy paper should have the following elements:

1. *Thesis:* A philosophy paper should have an easily identifiable point, something that you're asserting. It should not (a) ramble on for a while without direction, (b) recount things that someone else said, or (c) pose a series of questions that you never answer.
2. *Arguments:* Your paper should present specific reasons for believing that thesis, not mere unsupported opinions. A good argument should be:
 - a) *Non-trivial*, i.e., something that is not immediately obvious to most readers.
 - b) *Original*, not just a repetition of something you read or heard in class.
 - c) *Plausible*: The premises, for example, should be things that would initially seem true to most people.
 - d) *Non-question-begging (non-circular)*: the reasons given for your thesis should each be statements that are significantly different from the thesis itself, and that someone might accept before having made up their mind about your thesis.
3. *Objections:* Try to think of reasons someone might give for doubting your thesis, and indicate why those reasons are ultimately not persuasive. You should address the best objections you can think of, and the ones you think people are most likely to raise.

A.2. Style

This is about the general style in which a philosophy paper should be written:

4. *Key Point: The purpose of* (non-fiction) *writing is to communicate.* It is not to make art or to impress the reader with your sophistication. Therefore ...
5. *Be forthcoming:* State your thesis explicitly, right at the beginning. Here's a good opening sentence: "In this paper, I argue that incest is praiseworthy."^[117] At the beginning of each *section* of the paper, state the point of that section.
6. *Be organized:* A paper should usually be divided into sections (unless the paper is very short and simple)—much as this document is. Each section should have a name that clearly indicates what is in it. For example, you might have:

- 1: Common views of incest
- 2: Failed arguments for the common view
- 2.1: The argument from birth defects
- 2.2: The argument from emotional harm
- 3: The virtues of incest
- 4: Objections and replies

7. *Stick to the point:* Do not raise issues that are not necessary to advancing your central thesis.
8. *Be brief:* If you have an unusually long sentence, break it into shorter sentences. After writing a paper, go over it *line by line* looking for words, sentences, or paragraphs that could be deleted without weakening your point. Examples:

Bad: The question as to whether fish can experience pain is an important one. [13 words]

Ok: Whether fish experience pain is important. [6 words]

9. *Be specific:* Do not use a vague word or phrase when a more specific one is available. The first sentence below is bad because "related" is one of the vaguest words there is; also, it doesn't say what sort of obligations are being discussed. Examples:

Bad: Rights are related to obligations.

Bad: Rights imply obligations.

Ok: If someone has the right to do A, then others have the obligation not to stop him from doing A.

10. *Use plain language:* Do not use “sophisticated” or bombastic words in place of simpler, accurate words. Doing so makes your paper harder to read, and often makes you look stupid when you misuse the word. It does not make you sound sophisticated. Examples:

Bad: I am disinclined to acquiesce to your request.

Ok: No.

Bad: I utilized a fork to ingest my comestibles.

Ok: I used a fork to eat.

11. *Give examples:* When discussing an unfamiliar concept or claim, give examples that illustrate it. Give examples for every major thesis you defend or attack. For example, see the examples used in this document.

A.3. Research

This is about doing research for a philosophy paper, which helps you to know what you’re talking about and not look silly.

12. *How much research should you do?* This depends on your professor, the level of the class, and the topic.
1. Most undergraduate papers call for little research, perhaps just the course readings, but check with your professor if you’re not sure.
 2. Undergraduate *theses* and graduate-level papers require something closer to the amount of research of a real academic article. In academic articles, there is no specific number of references you need; you just need to cite the major things that are relevant to what you’re saying. It is rare for an academic paper to have fewer than 30 references or more than 70. Of course, longer articles tend to have more.

13. *What to read?*

1. If there is a Stanford Encyclopedia of Philosophy article on your topic, read it (<https://plato.stanford.edu>). Then start looking at other things that that article cited that sound relevant to your paper. The Internet Encyclopedia of Philosophy is also frequently helpful (<https://iep.utm.edu>).
2. Look up your chosen topic on PhilPapers (<https://philpapers.org>). You will probably get a list of articles and books many times longer than you could possibly read. Look for things that (i) are relatively recent, (ii) are by famous people (if you know the famous figures in the field), (iii) are in big journals (if you can recognize those), (iv) have been cited a lot (PhilPapers will show the # of citations), and (v) sound relevant to what you're saying in your paper, based on the abstracts.
3. After reading a few articles, take note of other articles that are cited by more than one of the articles that you read. Get those other articles. You should probably continue until you've read enough items to fill up at least a page (single-spaced) with references. Notes:
 - ☐ If you cite a book, does that mean you read the whole book? No. It means that you read some portion of it that was relevant to your paper.
 - ☐ All the items in your reference list should be mentioned somewhere in the footnotes or text of the paper, and vice versa.
4. In the course of writing your paper, you should run into additional opportunities to insert footnotes – e.g., when you refer to some view someone might hold, insert a footnote citing some people who defend that view; if you cite factual evidence, insert footnotes to indicate the sources. Again, this is if you're trying to write an academic-journal-like paper.

A.4. Misused Words

The following are words commonly misused by students. Read all of them so you don't make any of these embarrassing mistakes:

14. *as such*: Do not use “as such” in place of “therefore”. (“As such” may be used only when the subject of the sentence following is the same as that of the sentence preceding.)

Bad: Clocks usually tell the time of day. As such, an appeal to a clock may be used to support a belief about the time of day.

Ok: Clocks usually tell the time of day. Therefore, an appeal to a clock may be used to support a belief about the time of day.

Ok: W is commander-in-chief of the armed forces. As such, he can order bombings of other countries. [The last sentence means: As *commander-in-chief*, he can order bombings, etc.]

15. *being that*: Never use this phrase. It is not grammatical.

Bad: Being that I just had a tofu sandwich, I am no longer hungry.

Ok: Since I just had a tofu sandwich, I am no longer hungry.

16. *it's*: “It’s” means “it is”, not “belonging to it”.

Bad: My car lost one of it’s wheels on the freeway.

Ok: My car lost one of its wheels on the freeway.

17. *their/they're/there*: The first means “belonging to them”. The second means “they are”. The third refers to a place.

Bad: Their sure that there cat is still they’re.

Ok: They’re sure that their cat is still there.

18. *reference*: “To reference” means “to cite a source”. It does not mean “to talk about”.

Bad: He should make the argument for sense-data without referencing physical objects.

Ok: He should make the argument for sense data without mentioning physical objects.

Ok: He referenced his colleague’s work.

19. *such/as such*: Do not use “such” to mean “this” or “as such” to mean “in that way”.

Bad: I believe the last step in the argument – that because *x* will most likely appear as such in the future means that *x* is as such – is a mistake.

Ok: I believe the last step in the argument – that because *x* will most likely appear a certain way in the future, it is that way – is a mistake.

20. *phenomena* is the plural of *phenomenon*.

Ok: We discovered an interesting phenomenon.

Ok: We discovered many interesting phenomena.

21. *data* is the plural of *datum*.

Bad: Russell thinks that when you look at a table, all you see is a sense data.

Ok: Russell thinks that when you look at a table, all you see is a sense datum.

Ok: Russell thinks that when you look at a table, all you see are sense data.

22. *reality*: Do not use “reality” to mean “appearance” or “belief”. Do not talk about whether reality is real or whether reality is true. “Reality” means everything that is real (everything that exists).

Bad: Whose reality is true?

Ok: Whose beliefs are true?

Bad: There are many different realities.

Ok: There are many different belief systems.

23. *true*: Do not use “true” to mean “believed”.

Bad: To the medievals, it was true that the sun went around the earth. But to us, this is not true.

Ok: The medievals believed that the sun went around the Earth, but we do not believe this.

Ok: The medievals believed that the sun went around the Earth, but that is not true.

24. *infer, imply:* Do not use “infer” to mean “imply”. To say or suggest something indirectly is to *imply* it.

Bad: Are you inferring that I had something to do with the President’s assassination?

Ok: Are you implying that I had something to do with the President’s assassination?

25. *know:* Do not use “know” to mean “believe”, and especially do not use it to mean “falsely believe”.

Bad: Back in the middle ages, everyone knew the sun went around the Earth.

Ok: Back in the middle ages, everyone thought the sun went around the Earth.

26. *refute:* “To refute” means “to prove the falsity of”. It does not mean “to deny”.

Bad: Clinton refuted charges that he had an affair with Monica.

Ok: Clinton denied that he had an affair with Monica.

27. *based off/based off of:* Never use these phrases.

Bad: Hume’s argument is based off of three premises.

Ok: Hume’s conclusion is based on three premises.

28. *reliant on:* Never say this.

Bad: Determinism is reliant on two definitions.

Ok: There are two definitions of determinism.

Bad: The state of the world at any point in time is reliant on the second before.

Ok: The state of the world at any point in time depends on the state of the world the second before.

29. *argue:* “Argue” is normally followed by “that” or “for”.

Bad: Dennett argues compatibilism.

Bad: To argue this, he uses an analogy of a chess-playing computer.

Ok: Dennett defends compatibilism.

Ok: Dennett argues that free will is compatible with determinism. To argue for this, he uses an analogy involving a chess-playing computer.

30. *begs the question:* “To beg the question” means “to give an argument in which one or more of the premises depends on the conclusion”. It does not mean “to raise the question” or “to prompt people to ask”.

Bad: Ted Honderich said that “there is no experimental evidence in a standard sense that there are any [quantum events]”, [\[118\]](#) which begs the question of what he thinks “experimental evidence in the standard sense” is.

Ok: Jon argued that we should believe the Bible because it is the word of God, and we know it is the word of God because the Bible says it is the word of God. This argument begs the question.

31. *try and:* Never say this.

Bad: I will try and explain this.

Ok: I will try to explain this.

[A.5. Punctuation & Formatting](#)

Follow these tips so you don’t look like a yokel:

32. *General formatting:* Papers for classes should generally have:

- ☐ All pages *numbered*
- ☐ A staple in the corner
- ☐ 1-inch margins
- ☐ 12 point font

- Some professors want double-spacing (not me, though; I think that just wastes paper).

33. *Indenting Quotations:* Quotations of 3 or more lines should be indented, without quotation marks, and with a blank line before and after.

Bad: Robert Nozick writes:

“Taxation of earnings from labor is on a par with forced labor. Some persons find this claim obviously true: taking the earnings of n hours labor is like taking n hours from the person; it is like forcing the person to work n hours for another’s purpose.” (Nozick 1974, 169)

Ok: Robert Nozick writes:

Taxation of earnings from labor is on a par with forced labor. Some persons find this claim obviously true: taking the earnings of n hours labor is like taking n hours from the person; it is like forcing the person to work n hours for another’s purpose. [\[119\]](#)

34. *Source citations:* Any time you say that someone held some view, cite the source, including the page number where they said that thing. Standard format for citations in a footnote:

Articles: Author, “Article Title,” *Journal Title* volume # (year): pages of the article, page where they said the thing you’re discussing. Example (note the punctuation!):

Michael Huemer and Ben Kovitz, “Causation as Simultaneous and Continuous,” *Philosophical Quarterly* 53 (2003): 556–65, p. 564.

(*Note:* In this book, I normally put commas and periods outside the quotation marks if they belong to the larger sentence and not to the material that’s being quoted. This is the British style. However, most Americans put commas and periods *inside* the quotes (just to be perverse, I guess), which is why I’ve shown it that way in the above example.)

Books: Author, *Title of Book* (City of publication: publisher, year), page where they said the thing you’re discussing. Example:

Robert Nozick, *Philosophical Explanations* (Cambridge, Mass.: Harvard University Press, 1981), pp. 117–18.

35. *Spacing:* Put spaces before opening parentheses, and after punctuation.

Bad: Nozick compares taxation to forced labor(p.169).

Ok: Nozick compares taxation to forced labor (p. 169).

36. *Punctuation & parentheses:* Punctuation goes *outside* parentheses if the parenthetical material is less than the complete sentence/clause that the punctuation is for. Otherwise, it goes inside.

Ok: Hare's definition is too narrow (it makes some physical facts "subjective"), while Adams' is too broad (it makes everything "objective").

Ok: In the weak sense, to undertake an obligation is, roughly, to *purport* to place oneself under an obligation. (The exact analysis is not important here.)

37. *Dashes:* A dash (— or –) is longer than a hyphen (-). The small dash comes with spaces on both sides; the long dash gets no spaces.

Bad: Our inalienable rights- to life, liberty, and handguns-are under constant attack by liberal sissies.

Ok: Our inalienable rights – to life, liberty, and handguns – are under constant attack by liberal sissies.

Ok: Our inalienable rights—to life, liberty, and handguns—are under constant attack by liberal sissies.

38. *Titles:* Use italics for book and journal titles; use quotes for article and short story titles.

Ok: Unger's celebrated paper, "Why There Are No People," first appeared in *Midwest Studies in Philosophy*, volume IV.

39. *Scare quotes:* Do not insert gratuitous quotation marks around words. (How would you like it if I told you that I "read" your "paper" over the week-end?)

Bad: Scientists use experiments to “prove” the “truth” of their theories.

Ok: Scientists use experiments to prove the truth of their theories.

Ok: Scientists use experiments to try to prove their theories.

A.6. Grammar

Try to avoid the following grammatical errors that will make your professor cry:

40. *Modifiers placed at the beginning of a sentence* attach to the *subject* of the sentence. The first sentence below is bad because it means that John was carrying a mouse in his mouth; it also rudely calls John “it”. The fourth is bad because it implies that it was the teacher who couldn’t master grammar.

Bad: Carrying a mouse in its mouth, John saw the cat enter the room.

Ok: Carrying a mouse in its mouth, the cat entered the room.

Ok: John saw the cat enter the room carrying a mouse in its mouth.

Bad: While unable to master grammar, the English teacher had to explain the use of adverb phrases to me again.

Ok: Since I could not master grammar, the English teacher had to explain the use of adverb phrases to me again.

41. *Parallelism:* When making a list or joining phrases that have similar functions, use parallel word structures.

Bad: We had no alcohol. We also did not have drugs.

Ok: We had neither alcohol nor drugs.

Bad: Guns are for family protection, to hunt dangerous or delicious animals, and keeping the King of England out of your face.

Ok: Guns are for protecting your family, hunting dangerous or delicious animals, and keeping the King of England out of your face.

Ok: The purpose of guns is to protect your family, hunt dangerous or delicious animals, and keep the King of England out of your face.

[\[120\]](#)

A.7. Other Bad Writing

Here are some miscellaneous errors that will cause your professor to shake his head in disbelief and think about jumping off a bridge:

42. *Misquoting:* When taking a quotation, copy down exactly what appears in the text. Do not introduce grammatical, punctuation, or spelling errors of your own. If you omit something from the text, use ellipses (...). If you need to add something to the text, put it in square brackets [like this]. For example, the following appears (exactly as written here) in a book by Bertrand Russell:

His theoretical errors, however, would not have mattered so much but for the fact that, like Tertullian and Carlyle, his chief desire was to see his enemies punished, and he cared little what happened to his friends in the process.

I might quote this as follows:

Ok: Russell writes:

[Marx's] theoretical errors ... would not have mattered so much but for the fact that ... his chief desire was to see his enemies punished, and he cared little what happened to his friends in the process. [\[121\]](#)

I insert "Marx's" in place of "His" so readers who can't see the context know whom Russell was talking about. I use square brackets to indicate that this is my insertion/substitution. I use ellipses where I omitted unnecessary words. Obviously, do not omit anything whose omission changes the meaning of the passage.

43. *Redundancy:* Omit unnecessary redundant words that you don't need.

Bad: It could be said that it is a fact about the world that clocks usually tell the time of day.

Ok: Clocks usually tell the time of day.

Bad: Testimony is not sufficient enough to defeat a perceptual belief.

Ok: Testimony is not sufficient to defeat a perceptual belief.

Bad: This sentence could possibly be phrased more concisely.

Ok: This sentence could be phrased more concisely.

44. *Passive voice:* This should usually be avoided.

Bad: That people exist has been denied by Peter Unger.

Ok: Peter Unger has denied that people exist.

45. *Repetition:* Don't repeat yourself. Also, do not say the same thing over and over again. If you're having trouble filling up enough pages, you need to think more about your topic so that you have more to say. You can also try reading more about it.

46. *Undermining your credibility:* Here are some things that make readers wonder why they're wasting their time reading your paper:

a) *Denying that you know what you're talking about*, as in "This is just my opinion", or "The conclusions defended in this paper may well be mistaken." If you have nothing definite to say about a topic then *don't write about it*. Choose a different topic.

Bad: I believe we have free will, but I don't really know anything about it.

[Then why would I care what you think?]

Bad: I am not claiming that my argument establishes the reality of free will.

[Then why did you make me waste my time reading it?]

b) *Assertions about things you are ignorant of*. For instance, if you have not read any of the literature on free will, you should not make comments about what most philosophers think about free will, and you should probably avoid saying anything about free will at all, if you can avoid it. If you have to say something about it, at least read an encyclopedia article about it (try the *Stanford Encyclopedia of Philosophy*^[122]). *Why:* Because when you discuss things you are ignorant of, it is highly likely that readers who are more knowledgeable than you will find your remarks, well, ignorant, whereupon they will distrust the rest of what you have to say.

c) *Overstated claims*. While avoiding problem (a), do not go to the opposite extreme of making overstated claims that you can't justify.

Bad: Obviously, I have conclusively refuted direct realism.

[The unlikelihood of your having done this undermines your credibility.]

Ok: The above arguments provide strong grounds for preferring representationalism over direct realism.

A.8. Recommended Reading

These sources will help you become a better writer and avoid style errors:

47. *The Elements of Style* by William Strunk Jr. and E.B. White, <https://www.amazon.com/dp/B07NPN5HTP/>. A classic little book of advice on composition, commonly used in college writing courses.
48. *Chicago Manual of Style*, <https://www.chicagomanualofstyle.org/home.html>. The best-known authority on matters of style, including punctuation, grammar, formatting of books, and so on.
49. *MLA Handbook*, <https://style.mla.org>. Another well-known style manual with slightly different style rules.
50. Paul Brians, “Common Errors in English”, <http://www.wsu.edu/~brians/errors/errors.html>. Long list of word usage errors. Also discusses some non-errors, such as splitting infinitives and ending sentences with prepositions.
51. Jim Pryor, “Guidelines on Writing a Philosophy Paper”, <http://www.jimpryor.net/teaching/guidelines/writing.html>. By another philosopher. Many philosophers like this guide, especially its advice to imagine that your reader is lazy, stupid, and mean.

Glossary

Here is a list of all the bolded terms that appeared in the book, with brief definitions, plus where they were explained in the text. These are generally important philosophical terms. I also include some abbreviations.

<i>absolute</i>	(1) Not varying from one person to another or one society to another. <i>Contrast</i> : relative. (§5.1.1) (2) <i>Of a deontological moral principle</i> : Unable to be outweighed by competing considerations. (§15.1.1)
<i>absolute deontology</i>	The view that some deontological moral principles are absolute. <i>Also called</i> : absolutism. (§15.1.1)
<i>absolute poverty</i>	The condition of not having enough resources to meet your basic needs. <i>Contrast</i> : relative poverty. (§16.1)
<i>absolute right</i>	A right that should never be violated, regardless of the consequences. <i>Contrast</i> : prima facie right. (§15.1.4)
<i>absolutism</i>	<i>In ethics</i> : The view that some deontological moral principles are absolute. <i>Also called</i> : absolute deontology. (§15.1.1)
<i>ad hoc hypothesis</i>	A hypothesis that is introduced to protect a theory from counter-evidence and has no other rationale. (§17.2)
<i>aesthetics</i>	The branch of philosophy that studies art, beauty, and stuff like that. (§1.3)
<i>affirming the consequent</i>	The fallacy of inferring from “If A then B” and “A” to “B”. (§4.1)
<i>agent causation</i>	Causation in which an effect is produced by a <i>person</i> who is acting, as opposed to by an

	event. (§11.5.1) <i>Contrast</i> : event causation.
<i>agent-neutral value</i>	The property of being simply good, as opposed to merely good <i>for</i> some particular person. (§14.5.3) <i>Contrast</i> : agent-relative value.
<i>agent-relative value</i>	The property of being good (beneficial) <i>for</i> someone or other. (§14.5.3) <i>Contrast</i> : agent-neutral value.
<i>analytic philosophy</i>	A style of philosophy that emphasizes clear expression and logical argumentation, mainly practiced in leading philosophy departments in the English-speaking world. Examples: Bertrand Russell, Gottlob Frege, G.E. Moore. <i>Contrast</i> : continental philosophy.
<i>anecdotal evidence</i>	Evidence for an inductive generalization that includes only one or a few cases. (§4.3)
<i>animal rights advocate</i>	Someone who believes that nonhuman animals have rights. (§17.1.4)
<i>animal welfare advocate</i>	Someone who believes that one should not harm animals without a good reason, whether or not they have rights. (§17.1.4)
<i>antecedent</i>	In a statement of the form “If A then B”, A is the antecedent. (§2.3) <i>Contrast</i> : consequent.
<i>anthropic principle</i>	The principle that we should only expect to observe characteristics of the universe that are compatible with the existence of observers. (§9.4.3)
<i>anti-realism</i>	In metaethics: The view that there are no objective evaluative truths. (§13.1.4)
<i>appeal to authority</i>	(1) An argument in which the opinion of some authority is given as a reason to accept the thing the authority believes. (§4.1) (2) The fallacy of appealing to authority in sense (1) when the authority figure in question lacks relevant expertise. (§4.1)

<i>argument</i>	A sequence of statements or propositions in which one (the “conclusion”) is meant to be supported by the others. (§2.5)
<i>argument by analogy</i>	An argument in which it is said that two situations are relevantly similar to each other, so what is true of the one case must be true of the other. (§2.6)
<i>argument from design</i>	An argument claiming that the universe has observable features that make it look like someone designed it for some purpose, and therefore, the universe probably had a creator. (§9.4.1)
<i>argument from evil</i>	An argument that infers, from the various bad things we see in the world, that there is no God. (§10.3)
<i>argumentum ad hominem</i>	The fallacy of arguing that some proposition is false because of irrelevant bad characteristics of one or more people who believe that proposition. (§4.1)
<i>argumentum ad ignorantiam</i>	The fallacy of arguing that P is false since we don’t know that it’s true. (§4.1)
<i>argumentum ad populum</i>	The (alleged) fallacy of arguing that a proposition is true since it is widely believed. (§4.1)
<i>assumption</i>	A proposition that one treats as true without arguing for it. (§4.3)
<i>base rate</i>	<i>Of a characteristic:</i> The frequency with which the characteristic occurs in the relevant population. (§4.3) <i>Of an event:</i> The frequency with which the event occurs in the relevant type of situation. (§4.3)
<i>base rate neglect</i>	The fallacy of ignoring the base rate of X when estimating the probability of X occurring in a particular case. (§4.3)

<i>beg the question</i>	To give an argument in which the premises contain (a paraphrase of) the conclusion, or in which one or more premises depend for their justification on the conclusion. <i>Also called:</i> giving a circular argument. (§§2.7, 4.1)
<i>BIV</i>	A brain in a vat; a brain that is being kept alive in a vat and stimulated artificially so that it experiences a simulation of normal life. (§6.2.2)
<i>BIV hypothesis</i>	The hypothesis that you are a BIV. <i>Contrast:</i> Real World hypothesis. (§§6.2.2, 6.3.4)
<i>brute fact</i>	A fact that has no explanation. (§9.3.4)
<i>burden of proof principle</i>	The principle that one who makes a positive claim (claiming that something exists or has some property) has a special obligation to supply evidence for the claim, unlike a person who merely makes negative claims (denying that something exists or has some property). (Controversial.) (§10.2)
<i>categorical imperative</i>	(1) An imperative that one has reason to follow regardless of one's desires. (§15.1.2) (2) The principle that one must always treat persons as ends-in-themselves and not as mere means. (§15.1.3) (3) The principle that one must only act according to maxims that one could will to be universal laws. (§15.1.2)
<i>category error</i>	The mistake of applying a predicate to something that it logically cannot apply to. (§2.10)
<i>causal theory of reference</i>	The theory that <i>what an idea refers to</i> is determined by <i>what caused that idea</i> . (§6.3.3)
<i>certainty skeptic</i>	A skeptic who thinks that we lack knowledge because our beliefs are not

absolutely certain. (§6.2.4) *Contrast:* justification skeptic.

cherry picking The fallacy of sifting through a body of evidence to select the pieces of evidence that support a desired conclusion, while ignoring the evidence that doesn't. (§4.3)

circular argument An argument in which the premises contain (a paraphrase of) the conclusion, or in which one or more premises depend for their justification on the conclusion. (§§2.7, 4.1) *Also called:* question-begging argument.

claim right A right that imposes an obligation on other people, either to supply some good to the right-holder or to refrain from harming or interfering with the right-holder. (§15.1.4) *Contrast:* permission right.

closure principle (1) The principle that if you're justified in believing something, then you have justification (available) for anything that that thing entails. (§8.3)

(2) The principle that if you know something, then you're in a position to know anything that that thing entails. (Not discussed.)

cogent *Of an argument:* having premises that provide at least some support for the conclusion, i.e., make the conclusion more probable. (§2.7)

compatibilism The view that the existence of free will is compatible with determinism. (§11.3.1)

complex question A question that presupposes something without actually stating it. Example: "Have you stopped beating your wife?" presupposes that the addressee has beaten his wife. (§4.1)

conceptual scheme A system of categories that we use to classify things; a way of dividing up the world into types of things. (§8.1)

<i>conclusion</i>	The statement or proposition that an argument is meant to support. (§2.5)
<i>conditional</i>	A statement/proposition of the form “If A then B” (symbolized “ $(A \rightarrow B)$ ”). (§2.3)
<i>confirmation bias</i>	The bias (widespread among humans) toward looking for evidence <i>for</i> a generalization or an existing belief but not looking for evidence against it. (§4.3)
<i>conjunct</i>	In a statement of the form “A & B”, A and B are the two conjuncts. (§2.3)
<i>conjunction</i>	A statement/proposition of the form “A and B” (symbolized “ $(A \& B)$ ”). (§2.3)
<i>consequent</i>	In a statement of the form “If A then B”, B is the consequent. (§2.3) <i>Contrast</i> : antecedent.
<i>consequentialism</i>	The view that the right action is always the action that maximizes the good (produces the most good possible). (§14.2)
<i>continental philosophy</i>	A style of philosophy that puts less emphasis on clear expression and logical argumentation, mainly practiced in France and Germany. Often advances subjectivist/relativist ideas. Examples: Martin Heidegger, Jean-Paul Sartre, Jacques Derrida. <i>Contrast</i> : analytic philosophy.
<i>contingent</i>	A statement/proposition is said to be contingent when it is neither necessary nor impossible, i.e., it could have been true and it could have been false. (§2.4)
<i>correspondence theory of truth</i>	The theory that truth is correspondence with reality. (§5.5.1)
<i>cosmological argument</i>	An argument that tries to establish the existence of God by arguing that the universe must have a cause. (§9.3.1) <i>Also called</i> : first cause argument.
<i>credulity</i>	The trait of being disposed to easily accept

	what people tell you without corroboration (often a mistake). (§4.3)
<i>cultural relativism</i>	The view that what is right or wrong depends on one's culture; the view that "right" means something like "approved by society". (§13.3.1)
<i>DDE</i>	See doctrine of double effect.
<i>deductive argument</i>	An argument in which the premises are supposed to guarantee the conclusion, in the sense that it would be contradictory for the premises to be true and the conclusion be false. (§2.6)
<i>defeasibility theory</i>	The theory that knowledge is justified, true belief with no (genuine) defeaters. (§8.4.5)
<i>defeater</i>	(1) <i>In the defeasibility theory:</i> A defeater for a proposition P is a true proposition that, when added to your beliefs, would result in your no longer being justified in believing P. (§8.4.5) (2) <i>In other contexts:</i> A defeater for a proposition P is a proposition that you have reason to believe and that diminishes or destroys your justification for P. (§8.4.5)
<i>denying the antecedent</i>	The fallacy of inferring from "If A then B" and " $\sim A$ " to " $\sim B$ ". (§4.1)
<i>deontology</i>	The view that the right action is not always the action that maximizes the good; the negation of consequentialism. (§15.1.1)
<i>descriptive</i>	<i>Of a property, statement, or proposition:</i> Not evaluative. (§13.1.1)
<i>determinism</i>	(1) The view that every event is preceded by a sufficient cause. (§11.2.1) (2) The view that the complete state of the world at any given time, plus all the laws of nature, are compatible with exactly one complete future of the world. (§11.2.1)

<i>direct realism</i>	The view that in perception, (i) we are directly aware of something in the external world, and (ii) we acquire foundational justification for some beliefs about the external-world. (§6.3.5)
<i>disjunct</i>	In a statement of the form “A or B”, A and B are the two disjuncts. (§2.3)
<i>disjunction</i>	A statement of the form “A or B” (symbolized “(A $\vee$ B)”). (§2.3)
<i>doctrine of double effect</i>	A (controversial) moral principle according to which it is easier to justify harming an innocent person as a foreseen <i>side effect</i> of one’s action than it is to justify harming them <i>intentionally</i> (either as an end or as a means to one’s end). The side-effect harm is said to be justifiable by consequentialist reasons, while the intended harm is not. (§15.1.3)
<i>dogmatism</i>	The trait (extremely common among humans) of being insufficiently willing to revise one’s beliefs in the light of evidence. (§3.2.6, 4.3)
<i>downward causation</i>	A (controversial) kind of causation in which the properties of a whole partly determine the characteristics or behavior of the parts. (§11.5.1)
<i>effective altruism</i>	Giving to charity in a way designed to produce the most good possible per dollar spent. (§16.4)
<i>emotional appeal</i>	An attempt to persuade someone of a conclusion by playing on their emotions, rather than providing objective evidence. (§4.1)
<i>epistemic circularity</i>	Using one or more belief-forming methods to test themselves. Includes using a form of inference to verify that that form of inference is cogent; asking an information source whether that source is trustworthy; using a cognitive

faculty to verify that that very faculty is reliable; using two faculties to verify each other's reliability; etc. (§7.2)

epistemology The branch of philosophy that studies knowledge and related matters, such as when beliefs are justified or rational. (§1.3)

equivocation The fallacy of using a word in two different senses during an argument but treating them as the same. (§4.1)

error theory In metaethics, the view that all moral judgments contain a mistake and hence are untrue. (§§13.1.4, 13.4.1) *Also called:* nihilism.

ethical intuitionism In metaethics, the view that there are irreducible, objective moral truths, which we can know about through intuition. (§§13.1.4, 13.6.1)

ethical naturalism In metaethics, the view that there are objective moral truths, which are reducible to descriptive facts, and which we can know about empirically. (§§13.1.4, 13.5.1)

ethical subjectivism In metaethics, the view that there are moral truths, all of which depend for their truth on the attitudes of observers towards the objects of evaluation (the things that are said to be good, bad, right or wrong). (§§13.1.4, 13.3.1)

ethical vegetarian Someone who abstains from meat for ethical reasons. (§17.1.2)

ethics The branch of philosophy that studies good, bad, right, and wrong. (§1.3)

ethnocentrism The tendency to assume that one's own culture is the best merely because it is one's own. (§5.4.4)

evaluative *Of a sentence, proposition, or property:* Entailing that something is good, bad, right, or wrong. (§13.1.1)

<i>event causation</i>	A kind of causation in which an event is caused by another event. (§11.5.1) <i>Contrast:</i> agent causation.
<i>evidential argument from evil</i>	An argument that claims that the bad things in the world constitute (non-conclusive) evidence that there is no God. (§10.3)
<i>evidentialism</i>	(1) The view that it is morally wrong to form beliefs with insufficient evidence. (§3.1.4) (2) The view that the justification for a belief depends entirely on one's evidence. (Not discussed.)
<i>expressivism</i>	In metaethics, the view that moral statements express some non-cognitive attitude (e.g., an emotion or desire), rather than a belief. (§§13.1.4, 13.2.1) <i>Also called:</i> Non-cognitivism.
<i>external world</i>	Everything that is outside one's own mind. (§6.1)
<i>external world skepticism</i>	The view that no one can know any contingent propositions about the external world. (§6.1)
<i>extrinsic</i>	<i>Of a property/predicate:</i> Depending on other objects, as opposed to a thing's own characteristics. <i>Contrast:</i> intrinsic. (§§12.3.6, 12.3.7, 12.4.2)
<i>factory farming</i>	A method of raising animals for meat or other products in which the animals are confined in small spaces, typically in inhumane conditions. (§17.1.1)
<i>false analogy</i>	An analogy that is misleading because it compares two things that are not really relevantly alike. (§4.1)
<i>false dilemma</i>	A fallacy in which one treats two alternatives as if they were the only possibilities, when in reality they are not. (§4.1)

<i>falsifiable</i>	<i>Of a theory:</i> Capable of being refuted (or at least rendered improbable) by evidence, in the sense that it makes testable predictions. (§6.3.4)
<i>fine tuning</i>	<i>Of the universe:</i> The phenomenon of having certain broad characteristics that fall within a very narrow range of possibilities that would make life possible. (§9.4.2)
<i>fine tuning argument</i>	The argument that because the universe is fine tuned, it probably had an intelligent designer. (§9.4.2)
<i>first cause argument</i>	An argument that tries to establish the existence of God by arguing that the universe must have a cause. <i>Also called:</i> cosmological argument. (§9.3.1)
<i>foundational belief</i>	A belief that has foundational justification. (§7.5.1)
<i>foundational justification</i>	Justification for a belief that does not depend upon one's having justification for any other beliefs; non-inferential justification. (§7.5.1)
<i>foundational method</i>	A belief-forming method that produces foundational beliefs. (§7.5.5)
<i>foundationalism</i>	The view that some beliefs have foundational justification, and all other justified beliefs are justified by them. (§7.5.1)
<i>free will</i>	A capacity humans are thought to have whereby they have multiple alternatives available and are able to control their own acts or decisions. (§11.1)
<i>free will defense</i>	An argument that tries to reconcile the existence of God with the existence of evil by saying that free will is very valuable, and God could not have prevented evil without interfering with our free will. (§10.4.7)
<i>Frege-Geach</i>	In metaethics, the (alleged) problem for

<i>problem</i>	expressivism/non-cognitivism that these views cannot explain the meanings of moral expressions that appear in contexts where the sentence as a whole does not give a positive or negative evaluation of anything (e.g., “If doing something is wrong, then paying someone to do it is also wrong” or “I wonder whether abortion is wrong”). (§13.2.2)
<i>fully-grounded</i>	<i>Of a belief:</i> Having no false beliefs in its evidential ancestry; not being inferred from a false belief, nor from a belief that was inferred from a false belief, etc. (§8.4.1)
<i>genetic fallacy</i>	The fallacy of confusing a thing’s origin with its current properties. (§4.1)
<i>genuine defeater</i>	<i>In the defeasibility theory:</i> The sort of defeater that stops you from having knowledge. (No general agreement on exactly how to describe them.) <i>Contrast:</i> misleading defeater.
<i>global skepticism</i>	The view that no one can know any proposition whatsoever. (ch. 7) <i>Contrast:</i> external-world skepticism.
<i>guilt by association</i>	The fallacy of rejecting an idea because of irrelevant associations with something else that’s bad, e.g., because some people who endorse the idea are bad for some other reason. (§4.1)
<i>hard determinism</i>	The view that because determinism is true, no one has free will. (§11.3.1)
<i>hasty generalization</i>	The fallacy of drawing a generalization based on insufficient evidence. (§4.1)
<i>hedonism</i>	The view that pleasure or enjoyment is the only intrinsic good. (§14.2)
<i>Humean theory of reasons</i>	The theory that all reasons for action come from desires; specifically, a person has a reason to perform an action only insofar as (they believe) the action would increase the

	probability of some outcome that the person desires. (§13.4.2)
<i>ideological “cause” judgment</i>	A judgment about what causes what in which one selects one factor to focus on, from multiple equally valid causal factors, based on ideological preferences. (§4.3)
<i>impartialism</i>	The view that one should give equal consideration to all persons, with no preference toward oneself or those close to one. (§14.2)
<i>indirect realism</i>	The view that in perception, (i) we are directly aware only of something in our own minds and indirectly aware of the external world, and (ii) we acquire foundational justification for beliefs about our own minds and only inferential justification for beliefs about the external world. (§6.3.5)
<i>inductive argument</i>	A non-deductive argument in which one draws a generalization from a number of instances of that generalization. E.g., one infers that all A's are B from many examples of A's that are B. (§2.6)
<i>inference to the best explanation</i>	A non-deductive argument in which one infers that some theory is probably true since it provides the best explanation for some set of evidence. (§2.6)
<i>inferential justification</i>	Justification for B is inferential when it depends upon one's having justification for some other belief that supports B. (§6.3.5)
<i>initial plausibility</i>	The degree to which something <i>seems</i> correct, prior to considering arguments for and against it. (§7.4)
<i>instrumental value / instrumental good</i>	Something that is good as a means; something that is valuable for the sake of something else that it helps you get. <i>Contrast:</i> intrinsic value. (§14.4.1)
<i>intentionality</i>	The property of being <i>of</i> or <i>about</i>

	something, or (purportedly) <i>representing</i> or <i>referring to</i> something. E.g., words, ideas, and pictures all have intentionality. (§6.3.3)
<i>intrinsic</i>	<i>Of a property/predicate:</i> Depending only on a thing's own characteristics, as opposed to other objects. <i>Contrast:</i> extrinsic. (§§12.3.6, 12.3.7, 12.4.2)
<i>intrinsic value / intrinsic good</i>	Something that is good as an end in itself; something that is valuable for its own sake. (§14.4.1)
<i>intuition</i>	The mental state of having something <i>seem</i> correct to you upon intellectual reflection, but prior to argument. (§§7.6.1, 13.6.1)
<i>intuitionism</i>	See ethical intuitionism.
<i>JTB analysis</i>	The theory that knowledge = justified, true belief. (§§8.2, 8.3)
<i>justification skeptic</i>	A skeptic who thinks that we lack knowledge because our beliefs are unjustified. (§6.2.4) <i>Contrast:</i> certainty skeptic.
<i>justified belief</i>	A belief that makes sense to hold, that a reasonable person would hold, that represents what a person (rationally) <i>ought</i> to think. (§6.2.4)
<i>law of excluded middle</i>	The principle that for any proposition, A, (A or ~A); that is, any proposition either is or is not the case. (§5.2.2)
<i>law of non-contradiction</i>	The principle that for any proposition, A, ~ (A & ~A); that is, there is nothing that both is and is not the case. (§5.2.1)
<i>logic</i>	The branch of philosophy that studies such things as what follows from what and what is consistent or inconsistent. (§1.3)
<i>meta-ethics</i>	The branch of ethics that studies questions about the nature of evaluative statements and judgments, e.g., how we know ethical truths,

	whether there are objective values, and what evaluative expressions mean. (§13.1.2)
<i>metaphysical possibility</i>	The broadest (most inclusive) sense of possibility, in which, roughly, any proposition that makes sense (isn't absurd) is possible (could have been true). (§2.4)
<i>metaphysics</i>	The branch of philosophy that studies broad questions about the nature of reality. (§1.3)
<i>misleading defeater</i>	<i>In the defeasibility theory:</i> The sort of defeater that does not stop you from having knowledge. (No general agreement on exactly how to describe them. Sometimes described, e.g., as a defeater that defeats by means of supporting a false proposition.) <i>Contrast:</i> genuine defeater.
<i>moderate deontology</i>	An ethical view that holds that some kinds of action are wrong even when they produce somewhat better consequences; however, they are permissible when they produce <i>vastly</i> better consequences (or avert <i>vastly</i> worse consequences). (§15.1.1) <i>Contrast:</i> consequentialism; absolute deontology.
<i>Moorean response</i>	A type of response to an argument in which one says that the argument's conclusion is so implausible that it makes more sense to reason from the denial of the conclusion to the denial of one of the premises, rather than to accept the argument as given. (§7.4)
<i>moral realism</i>	The view that there are objective evaluative truths. (It is sometimes also taken as part of moral realism that moral truths are knowable by us.) (§13.1.4)
<i>multiverse theory</i>	The theory that there exist many parallel universes, possibly an infinite number, in which all sorts of different things are happening. (§9.4.4)

<i>necessary</i>	<i>Of a proposition:</i> Could not have been false. (§2.4)
<i>negation</i>	The negation of A is the proposition that it's not the case that A (symbolized " $\sim A$ "). (§2.3)
<i>negative right</i>	A right <i>not</i> to have other people harm or interfere with you or your property in some particular way. (§15.1.4) <i>Contrast:</i> positive right.
<i>neutrality</i>	The failure to take a position on an issue. (§3.2.2) <i>Distinguish from:</i> objectivity.
<i>nihilism</i>	The view that moral evaluations are always false, i.e., nothing is good, bad, right, or wrong. (§§13.1.4, 13.4.1) <i>Also called:</i> error theory.
<i>no false lemmas</i>	A proposed condition on knowledge which says that, to know P, there must be no false beliefs in one's chain of reasoning leading to P. (§8.4.1) <i>Compare:</i> fully grounded.
<i>non sequitur</i>	A type of argument in which the premises don't at all support the conclusion; usually said when one can't figure out how the premises are supposed to have anything to do with the conclusion. (§4.1)
<i>non-cognitivism</i>	In metaethics, the view that moral statements express some non-cognitive attitude (e.g., an emotion or desire), rather than a belief. (§13.1.4, 13.2.1) <i>Also called:</i> expressivism.
<i>numerical identity</i>	The relation that every thing bears to itself and to nothing else. x is numerically identical to y when "x" and "y" are just two names for the same thing. (§12.2.3)
<i>objective</i>	(1) <i>Of a person:</i> Being disposed to consider ideas and evidence fairly, with little or no bias. (§3.2.2) (2) <i>Of a property, truth, or fact:</i> Independent of the attitudes of observers.

(§§5.1.2, 13.1.3) *Contrast*: subjective.

<i>omnibenevolent</i>	All-good; maximally good; morally perfect. Often said of God. (§9.1)
<i>omnipotent</i>	All-powerful; able to do anything that it is metaphysically possible for a being to do. Often said of God. (§9.1)
<i>omniscient</i>	All-knowing; aware of every truth and everything that exists. Often said of God. (§9.1)
<i>one-to-one relation</i>	A relation that a thing can stand in to only one thing, i.e., it's impossible to bear the relation to more than one thing. (§12.4.2)
<i>ontological argument</i>	An argument that tries to infer the existence of God from the meaning of the word "God", or an analysis of the concept of God. (§9.2.1)
<i>open question argument</i>	An argument that tries to refute reductionist definitions of "good" by citing the fact that "Is D-ness good?" is an open question (where D is any descriptive predicate), whereas "Is goodness good?" is not an open question. May also be applied to definitions of other evaluative terms. (§13.5.2)
<i>ostro-vegan</i>	Someone who consumes only plant-based foods, plus bivalves (e.g., oysters, clams, mussels, and scallops). (§17.1.2)
<i>overconfidence</i>	The mistake of having beliefs that are too strong, given the amount of evidence one has or the probability of the propositions in question. Extremely common among humans. (§4.3)
<i>oversimplification</i>	The mistake of viewing an issue in terms that are too simple, such as by overlooking more subtle or moderate positions. (§4.3)
<i>ovo-lacto vegetarian</i>	Someone who consumes only plant-based foods, plus eggs and milk products. (§17.1.2)
<i>Pascal's wager</i>	An argument that tries to show that one should believe in God since, if there is a God,

the potential cost of not believing in God is infinite, whereas if there is no God, the potential cost of believing in God is only finite. (§9.5.1)

permission right A right that entails that the right-holder may act in a particular way but that does not impose any obligation on others. *Contrast:* claim right. (§15.1.4)

persuasive definition The fallacy of trying to build a controversial conclusion into the definition of a word, e.g., defining “capitalism” as “a system of exploitation of the poor by the rich”. (§4.1)

pescetarian Someone who consumes only plant-based foods, plus fish. (§17.1.2)

p-hacking A bad form of scientific reasoning in which one tests many different hypotheses until one finds one that passes a test for statistical significance. This often results in one’s “confirming” false theories due to chance correlations. (§4.3)

phenomenal conservatism The view that it’s rational to presume that things are as they appear, until one has specific grounds for doubting that. (§7.6.1)

philosophy of mind The branch of philosophy that addresses philosophical questions about consciousness and the mind. (§1.3)

philosophy of science The branch of philosophy that addresses philosophical questions about science, such as how scientific reasoning should work and what are the philosophical implications of modern scientific theories. (§1.3)

pluralism In ethical theory, the view that there are multiple independent moral considerations that often need to be weighed against each other, rather than just one moral principle. (§15.3)

poisoning the The rhetorical strategy of trying to prevent

<i>well</i>	an opponent from making his case or responding to your argument, by warning the audience in advance not to trust your opponent. (§4.1)
<i>political philosophy</i>	The branch of philosophy that addresses philosophical questions about government and the organization of society. (§1.3)
<i>positive right</i>	A right to receive some benefit from others, e.g., “the right to health care”, “the right to an education”. (§15.1.4)
<i>post hoc ergo propter hoc</i>	The (alleged) fallacy of inferring that A causes B from the fact that A is (frequently) followed by B. (§4.1)
<i>predicate</i>	(1) A component of a proposition, namely, a property, relationship, action, or something like that, that is attributed to the subject(s). (§2.3) (2) A linguistic expression that refers to a predicate in sense (1); normally consists of a verb, possibly followed by a noun or adjective. (Not discussed.)
<i>preferentism</i>	The view that the only intrinsic good is the satisfaction of desires. (§14.2)
<i>premise</i>	A proposition used in an argument to support the conclusion. (§2.5)
<i>prima facie right</i>	A right that can in principle be outweighed by sufficiently important moral considerations, especially by large good or bad consequences. (§15.1.4) <i>Contrast</i> : absolute right.
<i>principle of sufficient reason</i>	The thesis that every contingent fact must have an explanation. (Controversial.) (§9.3.3)
<i>problem of evil</i>	The problem, for theists, of explaining why God allows bad things to exist. (§10.3)
<i>proper function analysis</i>	An analysis according to which knowledge consists of true belief formed by a properly

	functioning, reliable belief-forming mechanism aimed at the truth. (§8.4.3)
<i>proposition</i>	The sort of thing that can be true or false, but not a belief or assertion; rather, the sort of thing <i>that</i> one can believe or assert. (§2.2)
<i>qualitative identity</i>	x and y are qualitatively identical when they have the same properties. (§12.2.3)
<i>real world hypothesis</i>	The “hypothesis” that we are perceiving the real world normally. (§6.3.4) <i>Contrast:</i> BIV hypothesis.
<i>realism</i>	(1) <i>In philosophy of perception:</i> The view that there is a mind-independent world and that perception enables us to know about it. (§6.3.5) (2) <i>In metaethics:</i> The view that there are objective values. (§13.1.4)
<i>red herring</i>	An irrelevant issue that distracts from the main issue under discussion. (§4.1)
<i>reducible</i>	Capable of being explained in more fundamental terms. (§13.5.1)
<i>reductio ad absurdum</i>	A type of argument in which one claims that a theory should be rejected since it has absurd implications. (§2.6)
<i>reduction</i>	(1) A theory that explains some phenomenon (e.g., values, or the mind) in more fundamental terms. (§13.5.1) (2) The process of giving or confirming such a theory. (§13.5.1)
<i>reductionism</i>	The view that some phenomenon under consideration (e.g., values, or the mind) is reducible. (§13.5.1)
<i>relative</i>	Varying from one person to another or one society to another (as in “truth is relative”, “morality is relative”). (§§5.1.1, 13.3.1) <i>Contrast:</i> absolute.
<i>relative poverty</i>	The state of having less wealth than most

	people in your society. (§6.3.1) <i>Contrast:</i> absolute poverty.
<i>relativism</i>	(1) The view that truth is relative. (§5.1.1) (2) <i>In metaethics:</i> The view that good, bad, right, and wrong are relative. (§13.3.1)
<i>relevant alternative to P</i>	A possibility that is incompatible with P and that has to be ruled out in order for one to know that P (where P is some proposition). (§6.3.1)
<i>reliabilism</i>	An analysis according to which knowledge is true belief formed by a reliable method. (§8.4.2)
<i>right</i>	See claim right; permission right.
<i>selection effect</i>	A correlation between two properties, A and B, is said to result from a “selection effect” when the reason A and B are correlated is that the things that have A (or the A’s that were selected for study) are already more likely to have B for other reasons, not because A actually affects B or vice versa. (§4.3) <i>Contrast:</i> treatment effect.
<i>semantic externalism</i>	The view that what a person’s words mean, or what their thoughts are about, depends partly on factors external to the person’s mind. (§6.3.3) <i>Contrast:</i> semantic internalism.
<i>semantic internalism</i>	The view that what a person’s words mean, or what their thoughts are about, depends entirely on what goes on in that person’s mind; the negation of semantic externalism. (§6.3.3)
<i>sensory experiences</i>	The kind of experience you have when you see, hear, taste, touch, or smell things. Includes hallucinations and sensory illusions as well as normal perceptions. (§7.6.1) <i>Also called:</i> perceptual experiences.
<i>skepticism</i>	A philosophical view that denies that we can know some large class of things we

	normally think we know. (§6.1) <i>See also:</i> external world skepticism; global skepticism.
<i>soft determinism</i>	The view that determinism is true, yet we still have free will. (§11.3.1)
<i>soul</i>	The immaterial component of a person that has mental states and that determines one's identity. (Controversial.) (§§9.4.4, 12.3.8)
<i>sound</i>	<i>Of an argument:</i> Being valid and having all true premises. (§2.7)
<i>speculation</i>	A kind of claim that we lack decisive evidence for (where evidence would be needed to know it); roughly, a guess. Often used unpersuasively in arguing for controversial conclusions. (§4.3)
<i>steel-manning</i>	The practice of addressing the smartest, most reasonable, etc., position that conflicts with your own, as opposed to attacking only the weakest opponents. (§4.1) <i>Contrast:</i> straw man; weak man.
<i>straw man</i>	A position that your opponent in a debate does not really hold but that you attack because it is easier to rebut than their real position. Hence, the fallacy of "attacking a straw man". (§4.1)
<i>subject</i>	(1) <i>In logic:</i> The thing that a proposition is about, or the expression in a sentence that refers to that thing. (§2.3) (2) <i>In metaphysics/epistemology:</i> A thing that can have experiences or other mental states. (Not discussed.)
<i>subjective</i>	(1) <i>Of a claim:</i> Requiring an exercise of judgment to evaluate. (§4.3) (2) <i>Of a property, truth, or fact:</i> Dependent on the attitudes of observers. (§§5.1.2, 13.3.1) <i>Contrast:</i> objective.
<i>subjectivism</i>	<i>In metaethics:</i> The view that the truth of

	moral claims depends on the attitudes of observers. (§13.1.4, 13.3.1)
<i>supererogatory</i>	Praiseworthy but not required; above and beyond the call of duty. (§16.2)
<i>theodicy</i>	An attempt to explain why God would allow bad things to exist. (§10.4)
<i>tracking analysis</i>	An analysis according to which knowledge is true belief formed in such a way that if the proposition believed were false, one would not believe it, and if it were true, one would believe it. (§8.4.4)
<i>transitivity</i>	A relation, R, is said to be transitive when: “xRy” and “yRz” entails “xRz” (i.e., if x stands in the relation to y, and y stands in the relation to z, then x must also stand in the relation to z). E.g., equality is transitive; so is the greater-than relation. (§12.3.3)
<i>treatment effect</i>	Describes a correlation between two variables that exists because the one variable influences the other. (§4.3) <i>Contrast</i> : selection effect.
<i>trivial</i>	So obvious as to be not worth stating. (§5.4.1)
<i>trolley problem</i>	A hypothetical scenario (or the moral issue posed by the scenario) in which one must choose between allowing a runaway trolley to kill five people, and diverting it so that it spares the five but kills one other person. (§13.1.1)
<i>tu quoque</i>	The fallacy of trying to defend oneself from an accusation by saying that the accuser is also guilty. (§4.1)
<i>valid</i>	<i>In logic</i> : Having premises that entail the conclusion, such that it would be contradictory to suppose all of the premises true and the conclusion false. (§2.7)

<i>vegan</i>	Someone who abstains from consuming animal products. (§17.1.2)
<i>weak man</i>	A position that <i>some</i> of your opponents hold, but only the dumber ones. Often attacked because this is easier than responding to the smartest opponents; this practice is known as “weak manning”. (§4.1) <i>Compare</i> : straw man. <i>Contrast</i> : steel-manning.
<i>whataboutism</i>	The practice of responding to a criticism of <i>x</i> by raising irrelevant criticisms of something else. Often takes the form of criticizing a politician from the other side, when a politician one likes has been criticized. (§4.3)

[1] See the Philosophical Gourmet Report, <http://www.philosophicalgourmet.com/>. This is the most widely used set of rankings in philosophy.

[2] *The Truth About the World*, ed. James and Stuart Rachels. I’m hoping Stu will give me kickbacks for this plug.

[3] I’ve simplified the original myth of Theseus.

[4] Exceptions: In philosophy of mind and philosophy of science, it is common to appeal to scientific discoveries. Even so, philosophers will typically not *themselves* make any specialized observations but will simply discuss how to interpret the observations and theories of scientists. More annoying exception: Recently, some philosophers have started practicing what they call “experimental philosophy”, which usually involves taking surveys of people’s intuitions on philosophical questions.

[5] I note that all my hypothetical examples are purely fictional, and any resemblance to any actual persons, living or dead, is entirely coincidental.

[6] What I just said in the text is known as the correspondence theory of truth. There is a debate about the proper definition of “true”. The correspondence theory is the traditional view on the subject, but some people reject it, but we don’t have time to talk about that and it would just confuse you.

[7] See “The Crazy Nastyass Honey Badger”, <http://youtu.be/4r7wHMg5Yjg>.

[8] Or maybe this is a characterization of what “probable” means: Maybe the probability of a proposition should be understood as the degree to which one’s experiences and information support it. Anyway, probability and rational belief are closely connected.

[9] Evidentialism is famously defended by W.K. Clifford in “The Ethics of Belief” (1877), which is reprinted in many philosophy textbook anthologies. The following argument in the text is based on Clifford.

[10] Note that the relevant notion of risk is not a matter of the *objective chance* of a bad outcome. It’s a matter of the *epistemic probability*, for the agent, of a bad outcome. I.e., it doesn’t

matter if there are actually factors in the world that are poised to cause the bad outcome; it matters if the agent at the time has justification for thinking there might be such factors.

[11] Answers: Premises: 1, 2, and 4, or, if you want to be more detailed: 1, 2a, 2b, 4a, and 4b. Conclusion: 5. Deductive, valid, cogent, non-circular. The soundness of the argument is open to debate; you might be able to object to one or more of the premises.

[12] See Yosef Bhatti, Kasper M. Hansen, and Asmus Leth Olsen, “Political Hypocrisy: The Effect of Political Scandals on Candidate Evaluations”, *Acta Politica* 48 (2013): 408–28, <https://link.springer.com/article/10.1057/ap.2013.6>.

[13] To see this, here is an example. Rancher A and rancher B have land right next to each other, and both raise cattle. One day, rancher A decides to go out and shoot one of his cows. After he shoots it, he brings it back, and he discovers that it has rancher B’s brand on it. So it was actually B’s cow, not A’s. A goes over to apologize to B. He says ... which of the following: (i) “I shot your cow by mistake”, or (ii) “I shot your cow by accident”? Hopefully, you can see that the more correct statement is (i). This shows that “by mistake” and “by accident” are different.

[14] From the Wikipedia article on equivocation.

[15] Of course, not all law-violations are immoral. But assume that these violations were.

[16] See “Partner Abuse State of Knowledge Project Findings at-a-Glance”, [http:// www.springerpub. com/media/springer-journals/FindingsAt-a-Glance.pdf](http://www.springerpub.com/media/springer-journals/FindingsAt-a-Glance.pdf). Note: Some people dispute these results.

[17] See chapters 20–23 in *Judgment Under Uncertainty: Heuristics and Biases*, edited by Daniel Kahneman, Paul Slovic, and Amos Tversky. Or, for a popular gloss, see Larry Swedroe, “Don’t Be Overconfident in Investing”, CBS News, <https://www.cbsnews.com/news/dont-be-overconfident-in-investing/>.

[18] For more, see John Ioannidis’ terrific, famous paper, “Why Most Published Research Findings Are False”, <https://journals.plos.org/plosmedicine/article?id=10.1371/journal.pmed.0020124>.

[19] In spite of what I have said, there are some philosophers who reject the law of excluded middle, and even a few who reject the law of non-contradiction (but those who reject the law of non-contradiction *also* accept it!).

[20] I use scare quotes because this isn’t much of an argument.

[21] You can, however, find articles in *Teaching Philosophy*, an academic journal about how to teach philosophy, that discuss how to deal with the problem of “student relativism”. See Steven Satris, “Student Relativism”, https://www.pdcnet.org/teachphil/content/teachphil_1986_0009_0003_0193_0205.

[22] What about the case where you say that X is *probably* true but not certain? That’s meaningful even though it doesn’t *exclude* the possibility of ~X, right? So is that a counter-example to my principle that meaningful claims exclude alternatives? No, because the claim “X is probable” *does* exclude alternatives. It doesn’t exclude ~X, but it excludes the alternative [X is *improbable*].

[23] *Metaphysics* IV.7, 1011b25.

[24] See William James, *Pragmatism: A New Name for Some Old Ways of Thinking*.

[25] Read the quote marks as quasi-quotation marks, if you know what this means. So “P” in the text refers to the sentence that asserts proposition P, not to the letter “P”.

[26] See esp. Fred Dretske’s “The Pragmatic Dimension of Knowledge”. The duck example below is from there.

[27] This is funny because Dennett denies that people have experiences with an intrinsic, qualitative feel. See his article “Quining Qualia” or his book *Consciousness Explained*.

[28] See Hilary Putnam’s *Reason, Truth and History*.

[29] Just ignore the fact that the human body contains water. We could have devised a thought experiment involving a substance that isn’t actually found in the body.

[30] The view that the meanings of your words and thoughts are entirely determined by things going on inside your mind is known as “**semantic internalism**”, by contrast to “**semantic externalism**”, which holds that meanings are (at least partly) determined by your relation to things in the external world.

[31] See my paper, “Serious Theories and Skeptical Theories: Why You Are Probably Not a Brian in a Vat”.

[32] Note: **Inferential justification** for P is justification for P that depends upon one’s having justification for some other belief that supports P. Indirect realists think that your justification for believing things about the external world depends upon your having justification for believing stuff about the images in your own mind.

[33] For amusement, look up the song “The G.E. Moore Shift” by the 21st Century Monads (<http://youtu.be/IXdqieipJgs>). It’s about this.

[34] I put this in this cagey way, because I can’t find a lot of people asserting this definition in print. Gettier, in his famous refutation of the definition (“Is Justified True Belief Knowledge?”), cites Plato, Chisholm, and Ayer. But I don’t think Plato or Ayer held the view. Still, other epistemologists say that it was widely held, so I guess it was. The definition, however, went out of style before I was born.

[35] This is sometimes called “the **closure principle**”.

[36] Due to Bertrand Russell, though he failed to note that it refutes the JTB analysis.

[37] Or, to be a little more detailed: Valid deduction is a *conditionally reliable* process (it’s reliable when given reliable starting beliefs). Furthermore, we can stipulate that the initial belief [Jones owns a Ford] was formed reliably (however you want to characterize reliability, as long as you don’t require 100% reliability). So according to a standard form of reliabilism (ala Alvin Goldman), the belief [Jones owns a Ford or Brown is in Barcelona] counts as reliably formed.

[38] This is from Robert Nozick, in his *Philosophical Explanations*. Slight complication: Nozick later modifies the antecedent of (iii) to something like: “P were false and you used the same method to form a belief about P as the one you actually used”. A similar clause belongs in (iv).

[39] You might think (iv) is redundant with (ii). But in Nozick’s interpretation, (iv) requires that you continue to have a true belief, not just in the actual world, but in a sufficient range of worlds similar to the actual world in which P remains true.

[40] Annoyingly, there are two uses of “defeater” in epistemology. One is as I just said. The other use is this: A defeater for P is a proposition that you believe or have justification for believing, which gives you grounds for doubting P. Note that it doesn’t have to be true; also, it *actually* undermines your justification (it’s not merely that it *would* undermine your justification *if* you believed it).

[41] You might think [Tom has an identical twin] only supports that it *might* have been the twin that you saw, not that it *was* the twin. Response: [Tom has an identical twin] lowers the probability that you saw Tom by raising the probability that you saw the twin; that’s how it defeats [Tom stole the book]. Similarly, [Mrs. Grabit says Tom has a twin] lowers the probability that you saw Tom by

raising the probability that Tom has a twin and that you saw that twin, and that's how it defeats [Tom stole the book]. So these cases seem parallel.

[42] Robert Shope's *The Analysis of Knowing*.

[43] This is basically Locke's view of words and concepts. He also thought that we got our basic concepts from sensory experiences, but I won't discuss that view here.

[44] I discussed this more in "The Failure of Analysis and the Nature of Concepts" in *The Palgrave Handbook of Philosophical Methods*.

[45] Due to the monk Gaunilo in the 12th century.

[46] Read the quotes around "x" as corner quotes, if you know what that means.

[47] This argument derives from the medieval Islamic philosophers al-Ghazali and al-Kindi. In modern times, it is defended most famously by William Lane Craig (<https://www.reasonablefaith.org>). "Kalam" refers to Islamic scholastic theology, which is the context in which the argument originated. Btw, I'm discussing the Kalam argument and (in the next subsection) Clarke's argument, rather than Aquinas' arguments, because Aquinas' arguments are less natural and harder to explain.

[48] If you want to hear about it, look up Laurence Kraus' book *A Universe from Nothing*.

[49] David Albert made this point in his *NY Times* review of Kraus' book (<https://www.nytimes.com/2012/03/25/books/review/a-universe-from-nothing-by-lawrence-m-krauss.html>).

[50] See my book *Approaching Infinity* for seventeen paradoxes of the infinite.

[51] This sort of argument was advanced famously by Samuel Clarke in his *A Demonstration of the Being and Attributes of God* in 1705.

[52] *The Selfish Gene* is an excellent exposition of the theory of evolution, and one of the best popular science books ever. *The Blind Watchmaker*, also by Dawkins, is a response specifically to the Argument from Design.

[53] See Roger Penrose, *Cycles of Time*, p. 127.

[54] I recommend David Albert's excellent and fascinating *Time and Chance*.

[55] Does this mean that the creator shouldn't be called "God"? Maybe; I don't really care. But many traditions in history have believed in a god or gods, and almost none of these gods have been supposed to be triple-omni beings. So the fine tuning argument could be fairly described as an argument for "a god", if not for "God".

[56] I got this example from John Leslie's article, "Is the End of the World Nigh?"

[57] On this, see Leonard Susskind's *The Cosmic Landscape*.

[58] Actually, not just unlikely things; it could be anything with an initial probability less than 1.

[59] A googol is 10^{100} .

[60] Why the qualifier "...for a being to bring about"? If we just said that an omnipotent being can bring about any metaphysically possible event, then we'd face counter-examples like the event of "a tree falling on its own without anyone causing it" – no one could bring that about, not even God.

[61] Richard Dawkins, *River Out of Eden*, pp. 131–2.

[62] Ptolemaic astronomy held that the sun and planets orbit the Earth. The four-elements theory held that all material objects are composed of the basic elements of earth, air, fire, and water. The four-humors theory held that diseases are caused by imbalances of the body's four fluids (yellow bile, black bile, blood, and phlegm).

[63] Qualification: Maybe you can exhibit these virtues in response to *perceived* suffering, adversity, or danger, even if your perception is inaccurate. But then you'd still have to have false perceptions, which might be bad in itself.

[64] For an amazingly great exposition of the issues about the interpretation of QM, see David Albert's book, *Quantum Mechanics and Experience*.

[65] See Roger Penrose, *The Emperor's New Mind*; John Eccles and Karl Popper, *The Self and its Brain*.

[66] W.T. Stace, *Religion and the Modern Mind*.

[67] From Peter van Inwagen, *An Essay on Free Will*, p. 56.

[68] From my "Van Inwagen's Consequence Argument" in the *Phil Review* (2000).

[69] For a compelling portrayal, see the popular movie *Schindler's List*.

[70] *Epicurus: The Extant Remains*, p. 113, fragment 40.

[71] *The Freedom of the Will*, p. 115.

[72] I thank Iskra Fileva for these observations.

[73] But by the way, if you want an amazing argument on that, see my paper "Existence Is Evidence of Immortality" in *Noûs* in 2019.

[74] The evaluative also includes statements about what is beautiful, ugly, rational, vicious, justified, etc., because these all have, as part of their meaning, that a thing is good or bad in a certain respect. Sometimes it is controversial whether a statement is evaluative or descriptive. But we're not going to raise any controversial cases, because we just want to understand the basic problems about ethics. Also, btw, ethics does not study *all* evaluative statements. For instance, claims about what is beautiful fall under aesthetics, and claims about what is a just policy fall under political philosophy, rather than ethics. But again, we can ignore those things now.

[75] From Philippa Foot, "The Problem of Abortion and the Doctrine of the Double Effect".

[76] There are some variations in the way all of the above terms are used, so some people would give different definitions for them. That's because the way these terms arose was not by someone thinking about all the possible views one could have and devising a name for each one. Rather, each view was advanced by certain philosophers, working independently of each other and usually not thinking about what the other possible views were. So they might define their views in different ways that mess up the taxonomy of theories – e.g., self-described "expressivists" might give characterizations of their view that make it overlap with subjectivism. Since my priority is to help you think clearly, not to help you use words the exact way that academic philosophers do, I've given the definitions that you need to make the taxonomy clean – so that the listed possibilities are mutually exclusive and jointly exhaustive. But I'll include terminological notes in the succeeding sections, in case you run into some academic philosophy articles.

[77] Terminological note: "Non-cognitivism" is an older term. Recent philosophers like to use the word "expressivism" instead, for some unknown reason. Expressivism holds that moral statements express some non-cognitive attitude, not belief, so it's a form of non-cognitivism.

Also, some people who call themselves "non-cognitivists" or "expressivists" will claim that moral statements *can* be true or false, *even though they don't assert propositions*. These people do this by taking weird views about what "true" means that we don't need to talk about here.

[78] Terminological note: Some people use "subjectivism" for the view that ethical truths depend upon *each individual's* attitudes, so right and wrong vary from one individual to another (even within the same society). In this chapter, however, I use "subjectivism" more broadly, so it

includes cultural relativism, (at least some forms of) divine command theory, and any other view that makes morality depend on someone's attitudes.

[79] This is discussed more in my 2016 article "A Liberal Realist Answer to Debunking Skeptics" from *Philosophical Studies*. There's also more (and more famous) discussion in Steven Pinker's amazing book, *The Better Angels of Our Nature*.

[80] Usually, the argument is given by non-cognitivists. However, it has also been given by nihilists, such as J.L. Mackie in *Ethics: Inventing Right and Wrong*.

[81] "Louis CK's Justification For Meat Consumption", <http://youtu.be/r3c0THQbdDE>.

[82] That's from J.L. Mackie's *Ethics: Inventing Right and Wrong* (1977). By "queer", he meant "weird".

[83] I'm not saying it's *more* plausible than anything else. Rather, it is tied, or approximately tied, with many other statements, such as "I exist", " $3 > 2$ ", "I know I have hands", and so on.

[84] Terminological note: In 13.1.4, I described naturalism only as holding (i). I've added (ii) here because I know of no one who holds (i) without (ii). It's theoretically possible, though, so theoretically there could be a form of moral realism that's neither naturalism nor intuitionism – this would be a view that holds (i) without (ii) or vice versa. I won't discuss such positions though.

[85] These experiments were done around 1780–1800. By the way, before this stuff was discovered, hydrogen was not called "hydrogen"; it was called "inflammable air". The word "hydrogen" basically means "water maker".

[86] The most common argument is something like this: "Fetuses are people. Killing people is (normally) wrong. Therefore, abortion is (normally) wrong." They would then try to give further arguments for the first premise, "Fetuses are people." Notice, by the way, that the *moral* premise of this argument, that killing people is normally wrong, is *not* the locus of dispute.

[87] That's the only reason you're pushing the fat man. It's not because you're prejudiced against overweight people!

[88] I'm going to mention one lame theory, just in case you've heard it and swallowed it uncritically, as people sometimes do: The psychologist Joshua Greene thinks that there is no *morally relevant* difference, but that the reason people judge the cases differently is that Footbridge involves an "up close and personal" way of killing, while Trolley is more "impersonal". Some lazy journalists have reported this super-implausible theory as if it were a fact (journalists are like that). Problems: (1) If you push the fat man with a long pole, so that you're not *close* to him at the time, that does not make it intuitively moral. (2) Steering a trolley toward someone so it hits them is not any less personal than pushing someone into the path of a trolley.

[89] Terminological note: Utilitarians sometimes use "happiness" or "pleasure" in place of "enjoyment". I think "happiness" is too narrow, because in ordinary parlance, it only includes a particular positive *emotional* state, whereas the utilitarians clearly want to include sensory pleasure as well. I think the term "pleasure" may also be off, because it is possible to enjoy certain kinds of pain, such as the taste of spicy food, and I'm not sure "pleasure" applies to that. That's why I prefer "enjoyment".

[90] Sources: The Trolley case comes from Philippa Foot ("The Problem of Abortion and the Doctrine of the Double Effect"), Footbridge from Judith Jarvis Thomson ("Killing, Letting Die, and the Trolley Problem"), Organ Harvesting from James Rachels (in informal conversations), Framing from H.J. McCloskey ("An Examination of Restricted Utilitarianism"), and Electrical Accident from T.M. Scanlon (*What We Owe to Each Other*, p. 235). I don't know the source of Promise; I heard it from Robert Fogelin.

[91] Interesting variant: Suppose that someone had deliberately sabotaged the trolley in order to kill the five people on the left track. Then, if the agent pushes the fat man off the bridge, there will be one murder, but if the agent does not push, there will be *five murders*. As a benevolent observer, you would have to hope for one murder rather than five. So, even if you think a murder is more than five times worse than an accidental death, you'd still hope the fat man gets pushed.

[92] Elsewhere, however, I have shown that equality of welfare doesn't matter intrinsically. See my "Against Equality" (<http://www.ow1232.net/papers/equality.htm>) and "Against Equality and Priority" (<http://www.ow1232.net/papers/priority.pdf>; originally in *Utilitas* (2012)).

[93] This is from Robert Nozick's *Anarchy, State, and Utopia*.

[94] I guess this was Ayn Rand's view. It's also suggested by a passage in Nozick's *Anarchy, State, and Utopia* (32–3).

[95] Judith Jarvis Thomson, for example, denies that generic, agent-neutral value exists, but she is no egoist, since she accepts deontological moral constraints.

[96] See my book, *The Problem of Political Authority*.

[97] The word "deontology" derives from the Greek *deon*, meaning "duty", and *logos*, meaning something like "study". So, according to etymology, deontology is the study of duty. But that's really not a good explanation of the word's actual meaning in contemporary philosophy. After all, you could have a consequentialist theory of duty ("Your duty is to maximize the good!"), which would not be considered "deontological".

[98] Naturally, in this discussion, types of action have to be defined in ways that are independent of the benefits or harms they produce, so that an action does not cease to be of a given type merely by having better or worse consequences.

[99] A good introductory exposition appears in Onora O'Neill's "The Moral Perplexities of Famine Relief", sections 22–31.

[100] There was also a third formulation of the Categorical Imperative: "Every rational being must act as if he were by his maxims at all times a lawgiving member of the universal kingdom of ends." But people don't talk about this very much. The first two formulations are the important ones.

[101] Robert Nozick suggests this in *Anarchy, State, and Utopia* (pp. 30–32).

[102] An alternative would be to qualify the right to life: You might say that persons have a right not to be killed *in certain ways*, e.g., by an agent creating a new threat, but no right not to be killed in certain other ways. E.g., maybe you lack a right against being killed *by the diversion of a trolley that's threatening five others*.

[103] All this follows the views of the 20th-century British intuitionist, W.D. Ross (see his *The Right and the Good*).

[104] See my "A Paradox for Weak Deontology", *Utilitas* (2009), <https://philpapers.org/archive/HUEAPF.pdf>.

[105] From his famous article, "Famine, Affluence, and Morality", which is used in a lot of philosophy classes.

[106] For a contrasting viewpoint, see "Why Donate to a University?", <https://fakenous.net/?p=2037>.

[107] See Charles Murray's book *Losing Ground*. A lot of people dispute Murray's arguments, though.

[108] That figure was for 2016, based on the UN Food and Agriculture Organization's data. I talk about this in my *Dialogues on Ethical Vegetarianism*, which discusses vegetarianism at greater

length.

[109] For more detailed description, see Stuart Rachels' article, "Vegetarianism", section 1 (<http://jamesrachels.org/stuart/veg.pdf>). For a useful video, see "What Cody Saw Will Change Your Life" on YouTube (<http://youtu.be/BFO34lmAoMQ>).

[110] Some other terms: **Vegans** are people who abstain from all animal products, not just meat. **Ovo-lacto vegetarians** are vegetarians who eat eggs and milk products. **Ostro-vegans** are people who are vegan except that they eat bivalves. **Pescetarians** are people who eat fish but abstain from all other meat.

[111] The most famous animal welfare advocate in philosophy is Peter Singer, author of *Animal Liberation* (the same philosopher who devised the Shallow Pond example from §16.1).

[112] The most famous animal rights advocate in philosophy is Tom Regan, author of *The Case for Animal Rights*.

[113] See C.H. Eisemann, W.K. Jorgensen, D.J. Merritt, M.J. Rice, B.W. Cribb, P.D. Webb, and M.P. Zalucki, "Do Insects Feel Pain? — A Biological View," *Experientia* 40 (1984): 164–7.

[114] These would be "different possible worlds", as we say – i.e., we imagine different ways the world could have gone. In one, a particular being belongs to a smart species; in another possible world, an intrinsically identical being belongs to a dumb species.

[115] See Stanley Milgram's book, *Obedience to Authority*. You can also see this YouTube video: "The Milgram Experiment 1962 Full Documentary", <http://youtu.be/rdrKCilEhC0>.

[116] I'm not denying that sometimes people are jerks or evil, of course. But you should only conclude that if a thorough search fails to turn up any plausible non-jerky motivation for their statements.

[117] This example is jocular; I am not advising you to write a paper on this.

[118] This example is based on a real article. See "Determinism as True, Both Compatibilism and Incompatibilism as False, and the Real Problem" in *The Oxford Handbook of Free Will* (2002) – which makes a fair bid to be the worst academic paper ever written on free will.

[119] *Anarchy, State and Utopia*, p. 169. Btw, notice how, in this footnote, I do not need to name the author I'm quoting, because the text already indicated that it was Robert Nozick.

[120] Paraphrase of Krusty in *The Simpsons*, "The Cartridge Family".

[121] *The Basic Writings of Bertrand Russell* (New York: Simon & Schuster, 1961), p. 479.

[122] <https://plato.stanford.edu>.

Books By This Author

[The Problem of Political Authority](#)

Argues that no government has authority; explains how a non-state society could work.

[Dialogues on Ethical Vegetarianism](#)

Explains why vegetarianism is obligatory through a series of fictional dialogues between a vegetarian and a meat-eater.

[Justice Before the Law](#)

Explains how the American legal system is unjust; argues that judges, jurors, lawyers, etc., should place justice before the law.

[Ethical Intuitionism](#)

Explains how we know about objective values; refutes the four alternative theories about ethics.

[Skepticism and the Veil of Perception](#)

Defends a direct realist theory of perception and rebuts four classic skeptical arguments.

[Approaching Infinity](#)

Solves 17 paradoxes of the infinite using a new account of which infinities are possible and which are impossible.

[Paradox Lost](#)

Solves ten mind-boggling philosophical paradoxes.